



UNIVERSIDAD DE MURCIA

ESCUELA INTERNACIONAL DE DOCTORADO

**Data Analytics Approaches in IoT Based
Smart Environments**

**Análisis de Datos en Entornos Inteligentes
Basados en el Internet de las Cosas**

D^a Aurora González Vidal

2019



Department of Engineering of Information and Communications

FACULTY OF COMPUTER SCIENCE

UNIVERSITY OF MURCIA

Data Analytics Approaches in IoT based Smart Environments

Ph.D Thesis

Authored by:

AURORA GONZÁLEZ VIDAL

Supervised by:

DR. ANTONIO FERNANDO SKARMETA GÓMEZ

MURCIA, SEPTEMBER 2019



Departamento de Ingeniería de la Información y las Comunicaciones

FACULTAD DE INFORMÁTICA

UNIVERSIDAD DE MURCIA

Análisis de datos en entornos inteligentes basados en el Internet de las Cosas

Tesis Doctoral

Presentada por:

AURORA GONZÁLEZ VIDAL

Supervisada por:

DR. ANTONIO FERNANDO SKARMETA GÓMEZ

MURCIA, SEPTIEMBRE 2019

DEDICATION AND ACKNOWLEDGEMENTS

Quoting Ortega y Gasset “I am I and my circumstance; and, if I do not save it, I do not save myself”. For this reason, I would like to thank those who have shaped my environment and who have made this work possible.

First of all my thesis director, Antonio Skarmeta, thank you for giving me the opportunity to discover research and for enhancing the talent of those around you.

Antonio Maurandi, former mentor, now friend and colleague. With you I discovered the science to which I hope to dedicate, I owe you a lot of who I am nowadays. Our complicity has been a constant in the uncertainty.

To my supervisors and colleagues at the internships in ICS, Surrey and ISSNIP, Melbourne for accepting, helping and teaching me so much about research and about the rest of the world.

To my postdoc colleagues who were there at the starting point, Victoria and Fernando who inspired my work and to the present ones, Alfonso for those chats that take us out of our caves.

To my friends: Fran, José, Lidia, Leticia, Yoel, ... I love you deeply.

To my partner, Nicolás, who has shared with me the worst moment of this thesis, its conclusion, and stood firm. Our adventures will continue to take us far away.

To my parents for their support, the opportunities they gave me and their trust and to my sister for her sisterly love.

To all of you I am grateful.

DEDICATORIA Y RECONOCIMIENTOS

Estoy de acuerdo con Ortega y Gasset en que “Yo soy yo y mi circunstancia, y si no la salvo a ella no me salvo yo”. Por ello me dispongo a agradecer a aquellos que han conformado mi entorno y que han dado lugar a este trabajo.

Empezando por mi director de tesis, Antonio Skarmeta, gracias por darme la oportunidad de descubrir la investigación y potenciar el talento de los que te rodean.

Antonio Maurandi, otrora mentor, ahora amigo y compañero. Contigo descubrí la ciencia a la que espero dedicarme, te debo mucho de lo que soy hoy día. Nuestra complicidad ha sido una constante en la incertidumbre.

A mis supervisores y compañeros en las estancias del ICS en Surrey y del ISSNIP en Melbourne por aceptarme, ayudarme y enseñarme tanto sobre la investigación y sobre el resto del mundo.

A mis compañeros postdoc que estaban al principio, Victoria y Fernando que inspiraron mi trabajo, y a Alfonso por esas charlas que nos sacan de nuestras cuevas.

A mis amigos: Fran, José, Lidia, Leticia, Yoel, ... os quiero mucho.

A mi pareja, Nicolás, que ha compartido conmigo el peor momento de esta tesis, su conclusión, y se ha mantenido firme. Nuestras aventuras nos seguirán llevando lejos.

A mis padres por su apoyo, las oportunidades y la confianza que me han dado y a mi hermana por su amor de hermana.

A todos vosotros os estoy agradecida.

TABLE OF CONTENTS

	Page
List of Tables	vii
List of Figures	viii
1 Resumen	1
1.1 Motivación y Objetivos	1
1.1.1 Internet de las Cosas (IdC) y los entornos inteligentes	2
1.1.2 Análisis de datos y Big Data en entornos inteligentes	6
1.2 Resultados	10
1.3 Organización de la Tesis	14
2 Summary	17
2.1 Motivation and Goals	17
2.1.1 Internet of Things (IoT) and smart environments	18
2.1.2 Data analytics and Big Data in smart environments	21
2.2 Results	25
2.3 Organisation of the Thesis	28
3 Thesis contributions	31
3.1 Related Work	31
3.1.1 Why energy consumption prediction is useful and how has it been carried out according to literature	31
3.1.2 Time series representation	34
3.1.3 Feature selection	36
3.1.4 HVAC usage patterns	39
3.1.5 Human mobility patterns	39
3.1.6 IoT architectures and projects for smart cities and energy management .	42
3.1.7 IoT architectures and projects for behavioural analysis towards energy efficiency	44
3.1.8 Related work summary	46

TABLE OF CONTENTS

3.2	Data analysis in IoT based Smart Environments	46
3.2.1	Smart buildings data integration and statistical analysis [R1]	47
3.2.2	Data representation [R2]	50
3.2.3	Energy consumption prediction [R3]	53
3.2.4	Feature selection [R4]	61
3.2.5	HVAC patterns [R5]	65
3.2.6	Human mobility patterns at macro and micro levels [R6]	67
3.2.7	IoT-based Big Data architecture for smart cities [R7]	72
3.2.8	IoT mechanisms to provide personalized energy management and awareness services by analysing behavioural aspects related to energy efficiency [R8]	76
3.3	Lessons Learned	79
3.4	Conclusions and Future Work	81
4	Publications composing the PhD Thesis	85
4.1	BEATS: Blocks of Eigenvalues Algorithm for Time Series Segmentation	85
4.2	A methodology for Energy Multivariate Time Series Forecasting in Smart Buildings based on Feature Selection	100
4.3	Commissioning of the Controlled and Automatized Testing Facility for Human Behavior and Control (CASITA)	113
4.4	Applicability of Big Data Techniques to Smart Cities Deployments	130
4.5	An open IoT platform for the management and analysis of energy data	141
4.6	Providing Personalized Energy Management and Awareness Services for Energy Efficiency in Smart Buildings	156

LIST OF TABLES

TABLE	Page
1.1 Resultados. Ver en negrita las publicaciones que componen la tesis. El resto son nuestras publicaciones adicionales.	15
2.1 Results. In bold the publications composing the thesis. The others are our additional publications	29
3.1 Information about the buildings	48
3.2 Results obtained for each moment	58
3.3 Metrics for energy consumption forecasting	60
3.4 Proposed FS methods for energy time series forecasting	62
3.5 RMSE, MAE and CPU time(in seconds) with 10-fold cross-validation (3 repetitions) . .	63
3.6 Selected attributes with <i>MOES-RF-MAE</i> (database #1) and their ranks.	63
3.7 Evaluation on test data with RF - database #1 and TransformedDatabase (TD) . . .	64

LIST OF FIGURES

FIGURE	Page
1.1 Componentes de la Ciudad Inteligente	3
2.1 Smart City components	19
3.1 The FS flow	37
3.2 1st floor of the TTC where red labels means energy meter (left) and 2nd floor of the Chemistry Faculty (right)	48
3.3 Correlation heatmap between consumption and outdoor environmental conditions for both consumption datasets	49
3.4 Correlation heatmap between consumption and outdoor environmental conditions for both consumption datasets	49
3.5 BEATS is shown step by step with an example	52
3.6 24h predictions performed with the fitter model (blue line) and the true values (black dots) with Model 2)	57
3.7 Boxplot of the energy consumption by moments considering all data (left); and, the time series of the energy consumption by moments during January (right)	58
3.8 Models validation performance (left) and Pairwise differences between models performance (right)	59
3.9 Dual-mode RC network	60
3.10 Weekly predictions using RF and real consumption	61
3.11 Changing frequency (left) and (right)	66
3.12 Number of DTAs and average number of changes per DTA with respect to the cell size	68
3.13 System architecture. The components that are not EPRs are depicted as dashed boxes	70
3.14 Collective landmarks (left) and metrics evolution	72
3.15 Heatmaps of clusters and movement prediction between early morning and morning slots	73
3.16 Information Model	74
3.17 IoTEP workflow	75
3.18 Entropy platform architecture	77

LIST OF ACRONYMS

AI Artificial intelligence	LR Linear Regression
ANN Artificial neural network	MAE Mean Absolute Error
ARIMA Autoregressive Integrated Moving Average	MAPE Mean Average Prediction Error
BRNN Bayesian Regularized Neural Network	ML Machine Learning
CEP Complex Event Processing	MLP Multilayer Perceptron
CVRMSE Coefficient of Variation of the RMSE	MOEA Multi Objective Evolutionary Algorithms
DCT Discrete Cosine Transform	MPC Model Predictive Control
DR Demand Response	OCB Orion Context Broker
DTA Dense Transit Area	PCA Principal Component Analysis
ENORA Elitist Pareto-based MOEA for diversity reinforcement	RBF Radial Basis Function
FS Feature Selection	RC resistor-capacitor
GAUSS Gaussian Processes	RF Random Forest
HVAC Heating Ventilating and Air Conditioning	RMSE Root Mean Square Error
ICT Information and Communications Technologies	ROIs Regions of Interest
IoT Internet of Things	SVM Support Vector Machines
	TTC Technological Transfer Centre
	XGB eXtreme Gradient Boosting

RESUMEN

Este capítulo presenta la motivación y la justificación del trabajo de tesis. Establece los objetivos de la investigación y los vincula a los resultados que se exponen de manera breve y conectada, dado que ciertos objetivos y resultados surgieron de necesidades que se identificaron cuando se establecieron los objetivos.

1.1 Motivación y Objetivos

El cambio climático está perturbando las economías nacionales y afectando la vida de muchas personas en todo el mundo. Sus consecuencias están costando muy caro hoy día y, si su progresión continúa, el precio a pagar será mucho mayor en el futuro.

Los fenómenos meteorológicos son cada vez más extremos, el nivel del mar está aumentando y las emisiones de gases de efecto invernadero se encuentran en los niveles más altos de la historia. Si no se toman medidas, es probable que el calentamiento global alcance los 5°C a finales de siglo¹, lo que tendrá un enorme impacto en la vida tal y como la conocemos hoy en día.

A fin de fortalecer la respuesta mundial para prevenir el cambio climático, varios países han adoptado muchas iniciativas. En 2015, se aprobó el *Programa de Desarrollo Sostenible* para 2030 y sus objetivos de desarrollo sostenible, que constituyen un llamamiento a la acción de todos los países para promover la prosperidad y proteger al mismo tiempo el planeta². Dentro de las 17 metas, cuatro de ellas están directamente relacionadas con las metas de la tesis: la inclusión de energía limpia y asequible, ciudades y comunidades sostenibles, consumo y producción responsables y acción climática. En el Acuerdo de París de la COP21 (2016), los países

¹<https://www.consilium.europa.eu/en/policies/climate-change/>

²<https://www.un.org/sustainabledevelopment/climate-change/>

participantes acordaron trabajar para limitar el aumento de la temperatura mundial a muy por debajo de los 2 grados centígrados ³.

Europa también está dedicando un esfuerzo considerable a reducir sustancialmente sus emisiones de gases de efecto invernadero. Para 2050, como parte de los esfuerzos requeridos por los países desarrollados, la UE se propone reducir sustancialmente sus emisiones, en un 80-95 % en comparación con los niveles de 1990 ⁴.

La investigación y la innovación contribuyen de manera decisiva a la lucha contra el cambio climático y a la adaptación al mismo, y las Tecnologías de la Información y las Comunicaciones (TIC) tienen el potencial de reducir un 20% de las emisiones mundiales de CO₂ para 2030, manteniendo las emisiones a los niveles de 2015 [19]. Un informe de British Telecommunications afirma que se espera que la influencia de las TIC en la UE reduzca la huella de carbono de la UE en un 37%, manteniendo las emisiones en los niveles de 2012.

La inteligencia artificial (IA) y sus aplicaciones particulares, como el Aprendizaje Automático (*Machine Learning* en inglés (ML) están demostrando ser muy útiles para detectar las abundantes ineeficiencias de la sociedad moderna que contribuyen a la inestabilidad climática.

El trabajo de la tesis se basa en la combinación de tecnologías TIC novedosas para la recolección y gestión de los datos y su análisis a través del ML con el fin de proporcionar entornos más inteligentes que puedan hacer un uso responsable de los recursos. La aplicación de este trabajo, en un sentido amplio, contribuye a mitigar el cambio climático.

1.1.1 Internet de las Cosas (IdC) y los entornos inteligentes

Un dispositivo IdC es un objeto físico que se conecta a Internet para transferir datos. Gracias a la proliferación de los dispositivos interconectados IdC, hoy en día se están recogiendo grandes cantidades de datos. Esto permite la creación de entornos inteligentes.

Los entornos inteligentes son entornos físicos que se entrelazan invisiblemente con abundantes dispositivos IdC, es decir, sensores, actuadores, dispositivos y elementos computacionales en general, integrados sin que se aprecie en los objetos cotidianos que nos rodean, y conectados a través de una red continua. Sin embargo, no son los sensores los que hacen que un entorno sea inteligente, sino la capacidad de procesar y aprender de todos los datos que esos sensores proporcionan a través de su análisis para proporcionar servicios automáticamente.

El desarrollo y la evolución de los análisis Big Data y de las tecnologías de IdC están desempeñando un papel importante en la adopción de iniciativas de ciudades inteligentes por diversas razones. La primera razón es el crecimiento exponencial de los objetos inteligentes que pueden participar en el desarrollo de una infraestructura de IdC [20]. Cisco Internet Business Solutions Group predice que habrá 50 mil millones de dispositivos conectados para 2020 [21].

³<https://unfccc.int/process-and-meetings/the-paris-agreement/the-paris-agreement>

⁴https://ec.europa.eu/clima/citizens/eu_en

1.1. MOTIVACIÓN Y OBJETIVOS

Otras dos razones notables son el crecimiento de la población y la tendencia de urbanización que está teniendo lugar [20]. Según las Naciones Unidas, hay un total de 1.3 millones de personas que se trasladan a las ciudades cada semana, con una población urbana que crece a 6.3 mil millones, lo que representa un 68% para el año 2050⁵. Este rápido aumento de la población urbana supone un gran estrés para las infraestructuras y el medio ambiente mundial, ya que las ciudades representan más del 70% del consumo mundial de energía [22] y producen el 80% de sus emisiones de gases de efecto invernadero [23].

En este sentido, las soluciones de IdC para ciudades inteligentes ayudan a promover el desarrollo económico, a mejorar la infraestructura y el medio ambiente, y a optimizar los sistemas de transporte de manera sostenible, mejorando al mismo tiempo la calidad de vida en las ciudades. Las zonas urbanas son el laboratorio perfecto para reducir las emisiones de gases de efecto invernadero, aumentar el uso de energías renovables y mejorar la eficiencia energética. En la Fig. 2.1 se muestran algunos componentes inteligentes importantes de la ciudad.

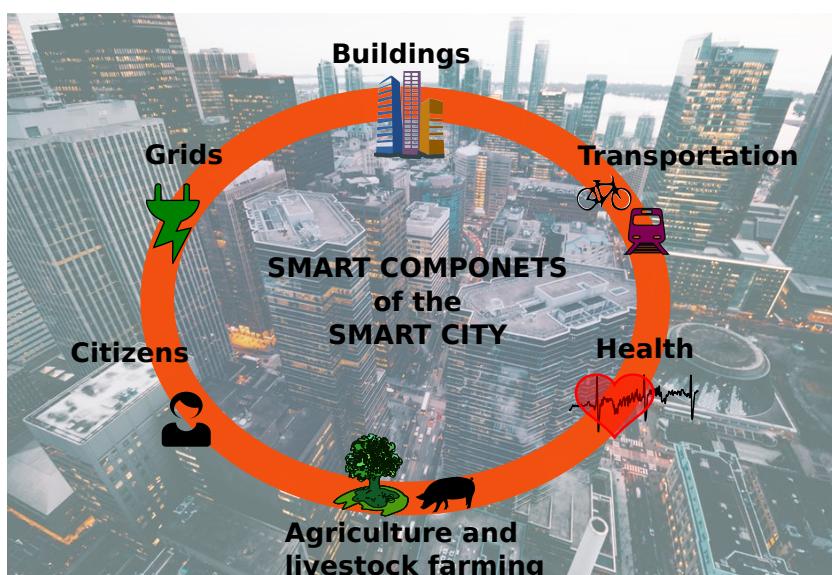


Figura 1.1: Componentes de la Ciudad Inteligente

Encontrar formas de satisfacer las necesidades energéticas de una población en crecimiento en conjunción con la creciente prosperidad económica y la escasez de recursos es un reto fundamental para lograr una sociedad sostenible. La reducción del consumo de energía y de la huella de carbono son cuestiones importantes en las ciudades inteligentes. En el desarrollo de ciudades inteligentes, la sostenibilidad se basa en la eficiencia energética y, a escala mundial, los edificios son la piedra angular de la eficiencia energética en términos de consumo de energía y emisiones de CO₂ [24].

El sector de los edificios también se ve muy afectado por la proliferación de contadores inteligentes y pantallas para el hogar. Esta tendencia parece ir en aumento si tenemos en cuenta

⁵<https://www.un.org/development/desa/en/news/population/2018-revision-of-world-urbanization-prospects.html>

CAPÍTULO 1. RESUMEN

que la Comisión Europea ha establecido que 16 Estados miembros procederán a realizar un despliegue a gran escala de contadores inteligentes para 2020 o antes[25]. Esto, junto con los nuevos avances en materia de infraestructura de datos energéticos (o en inglés *Energy Data Infrastructure* ver [5, 6]), ha creado el entorno perfecto para la creación, entre otras tecnologías, de estrategias avanzadas de retroalimentación energética para la reducción del consumo de energía en los edificios y para la educación de los ocupantes/usuarios [26], el denominado “edificio inteligente”.

Un edificio inteligente es cualquier estructura comercial, residencial o industrial en la que se han implementado procesos de automatización para controlar su funcionamiento en base a los datos recogidos por los sensores. Esto incluye tanto el ambiente interno como aparatos de aire acondicionado, iluminación, seguridad, sombreado, etc. [27] como el externo, por ejemplo el clima. Se espera que los edificios inteligentes consideren los elementos de dentro y fuera de su perímetro e interactúen con las redes eléctricas, las condiciones ambientales y los objetivos y labor de sus usuarios.

Los edificios inteligentes se consideran fundamentales para la emergencia de la ciudad inteligente. En la revista Smart Buildings Magazine, Harry G. Smeenk, vicepresidente de desarrollo de programas de la Asociación de la Industria de las Telecomunicaciones, señaló que *“el desarrollo de edificios inteligentes dará lugar a campus inteligentes, lo que fomentará comunidades inteligentes y, con el tiempo, ciudades inteligentes”. En pocas palabras, los edificios inteligentes crearán una base escalable para crear la elusiva ciudad inteligente, edificio por edificio, desde cero”*⁶.

En los países desarrollados, la energía consumida en los edificios representa entre el 20 y el 40% del consumo total de energía y es superior a la de la industria y el transporte en la UE y los EE.UU. [28, 29].

Para mitigar el cambio climático, la reducción del consumo de energía junto con el uso de fuentes de energía no fósiles es crucial. Además, la reducción del consumo de energía en los edificios debe hacerse al mismo tiempo que se garantiza la comodidad de los usuarios de los edificios y se reducen los costes para luchar contra la pobreza energética. Los análisis iniciales sugieren que la conversión de edificios en edificios inteligentes gracias a la sensorización a través del IdC, junto con el análisis de datos, puede ser una opción para resolver estos problemas.

En la encuesta de 2016 de la Continental Automated Buildings Association (CABA) denominada *Intelligent Buildings and the Impact of the Internet of Things*, se identificaron los siguientes 3 retos principales a la hora de hacer edificios más inteligentes [30]:

- Mejorar las decisiones de gasto

El hecho de que los patrones de uso de energía de los edificios a menudo no son posibles de determinar por parte de los gerentes de los edificios dificulta la identificación de las

⁶<http://www.smartbuildingsmagazine.com/features/the-smart-way-to-smart-cities-begins-with-buildings>

oportunidades adecuadas de ahorro de energía. Por lo tanto, muchas veces las medidas de ahorro de energía implementadas no mejoran la eficiencia o reducen innecesariamente la comodidad de los usuarios. Los sistemas de IdC pueden abordar este problema exponiendo datos detallados sobre el uso de la energía, permitiendo a los gestores detectar ineficiencias y crear modelos de predicción muy precisos.

- Reducir consumo energético y gasto de energía

El control de la utilización de los equipos requiere normalmente una supervisión manual. De esta manera, es complicado reducir el consumo de energía y controlar los costes. La automatización de los electrodomésticos y otros elementos de un edificio permite un mayor control de cuánto, cuándo y cómo se consume la energía.

Con el IdC, los gestores pueden observar y ajustar a distancia los sistemas de los edificios con sólo pulsar un botón, lo que facilita enormemente la reducción de costes. El ahorro de energía potencial puede mejorarse aún más con las tecnologías IdC.

- Mejora de la eficiencia operativa

La mayoría de los edificios tienen sistemas separados para el aire acondicionado, iluminación, energía, calidad del aire interior, conectividad a Internet, refrigeración, etc. Esto hace que sea muy difícil optimizar las operaciones generales del edificio. El IdC ofrece la oportunidad de integrar datos de numerosas fuentes en una única plataforma analítica. De esta manera, los gerentes pueden aplicar una estrategia holística a las operaciones de construcción. La combinación de la tecnología de IdC con los edificios inteligentes puede proporcionar un sistema de mantenimiento predictivo. Cuando los parámetros del edificio se monitorean, es más fácil detectar eventos anormales. El administrador del edificio se puede informar instantáneamente para actuar en consecuencia. De esta manera, hay menos fallos en los equipos, lo que contribuye en gran medida al ahorro de costes de los edificios inteligentes.

El último reto que destacamos se refiere a los comportamientos activos y pasivos de los ocupantes con respecto a la energía. Estos comportamientos incluyen la apertura de ventanas, el uso de electrodomésticos, el uso de persianas y sistemas de protección solar, la temperatura de consigna del aire acondicionado, la elección de la iluminación, etc [31]. Para garantizar una reducción prolongada del consumo de energía, las tecnologías de ahorro de energía deben ir acompañadas del eficiente comportamiento de los ocupantes en lo que a energía respecta [32]. Como se indica en el informe de la Agencia Europea de Medio Ambiente [33], hasta un 20 % del ahorro de energía puede lograrse a través de diferentes medidas dirigidas al comportamiento de los consumidores. Para educar a los usuarios de edificios en materia de sostenibilidad, el IdC puede contribuir con la detección de tareas específicas (dependientes del contexto a tiempo real) que los usuarios pueden llevar a cabo para mejorar la eficiencia acompañadas de un razonamiento para su interiorización.

A pesar de estas claras ventajas, muchos edificios todavía no han adoptado las tecnologías IdC.

Según [34], la escasez de infraestructuras inteligentes en los edificios implica que ningún país en Europa esté completamente preparado para la revolución inteligente. Dicho de otra forma, la falta de componentes inteligentes y conectividad entre ellos en los edificios es esencial para desentrañar las posibilidades de los edificios. Considerando los siguientes componentes: aire acondicionado, enfuches, parasoles en las ventanas y automatización de los edificios, si se mejora uno solo de estos componentes de forma aislada puede dar lugar a ahorros del 5–15%, y un sistema integrado puede conseguir un 30–50% de ahorros en edificios existentes que de otra manera resultan ineficientes [35].

Según ENERGY STAR, más de 5 millones de edificios comerciales en USA de $4600\ m^2$ o menos no contienen dispositivos inteligentes para monitorear el uso de energía, la temperatura u otros factores. Se estima que estos edificios consumen hasta un 30% más de energía de la necesaria. En todo el mundo, el número de estos edificios es mucho mayor.

1.1.2 Análisis de datos y Big Data en entornos inteligentes

La gran cantidad de datos heterogéneos que se capturan, almacenan y gestionan mediante el IdC supera las capacidades de las infraestructuras y de los motores de bases de datos tradicionales. Originalmente, las 3 Vs [36]: gran volumen, alta velocidad y gran variedad de datos se consideraron las características responsables de la aparición de las tecnologías Big Data que ayudan a resolver los problemas que superan los requisitos convencionales.

Con el tiempo, se han propuesto Vs adicionales para caracterizar al Big Data, y consideramos que las 7 Vs [37] describen mejor la complejidad del Big Data:

- Volumen: la enorme cantidad de dispositivos de IdC, incluyendo a los dispositivos portables (o *wearables* en inglés), genera enormes cantidades de datos. Los problemas que este tamaño de los datos genera son su escalabilidad, accesibilidad y capacidad de gestión.
- Velocidad: La velocidad de transferencia de datos entre la fuente y el destino.
- Variedad: Varios tipos de datos son generados: datos estructurados o no estructurados de diferentes fuentes como imagen, vídeo, texto, sensores, etc.
- Veracidad: Los datos reales que proceden del IdC casi nunca son limpios y precisos. Es necesario encontrar mecanismos para garantizar que los datos sean fiables.
- Validez: Para pasar de explorar a procesar los datos se deben validar previamente. La validez se refiere a la exactitud de los datos con respecto al uso previsto.

- Volatilidad: la retención de datos es especialmente importante en problemas de Big Data debido a su longitud. En muchos casos es crucial determinar en qué momento los datos ya no son relevantes para el análisis actual y deben dejar de almacenarse.
- Valor: El valor representa el valor de negocio que se deriva de los datos. El interés es siempre extraer el valor máximo de los datos. El valor de los datos debe superar su coste o su propiedad y gestión, incluyendo su almacenamiento.

Aunque se recogen datos en cantidades sin precedentes, menos del 1% de estos datos se están analizando [38]. Esto se debe a las complejidades derivadas de los problemas relacionados con Big Data. Existen varios desafíos en el análisis de datos reales, tales como la alta dimensionalidad, alto volumen, ruido y los *data drifts*. Los datos proporcionados por las fuentes de IdC (dispositivos sensoriales y mecanismos de detección) son multimodales y heterogéneos.

Todas estas características dificultan la ejecución y generalización de los algoritmos, por lo que hemos identificado los siguientes retos con respecto a los datos:

- Fusión de los datos procedentes de distintos sensores

La fusión de datos se define en [39] como la combinación de los datos de sensores procedentes de múltiples sensores para producir información más precisa, más completa y más fiable que no sería posible lograr a través de un solo sensor. En otras palabras, la fusión de datos es una técnica de procesamiento de datos que combina, mezcla, agrega e integra datos de varias fuentes.

Se pueden desarrollar servicios innovadores mediante la fusión de datos. En ese sentido, la fusión de datos es un reto crucial que debe abordarse. En las aplicaciones de ciudades inteligentes es esencial fusionar e interpretar los datos de forma automática e inteligente [40]. La fusión de datos y el filtrado de datos se han enumerado como dos retos principales para el IdC y sus aplicaciones, como las ciudades inteligentes [41].

- Identificación de patrones en la movilidad humana

La movilidad de las personas es especialmente importante para aplicaciones como la previsión del tráfico, la planificación urbana y el modelado epidémico. Comprender los patrones de movilidad puede ayudar a tomar decisiones basadas en datos y mejorar la calidad de vida en las ciudades inteligentes. Tradicionalmente, se utilizaban técnicas no escalables para encontrar patrones macroscópicos. Hoy en día, la incorporación de la tecnología GPS en dispositivos portátiles ha permitido recoger una gran cantidad de trazas digitales de alta resolución que permiten conocer las trayectorias espacio temporales subyacentes de las personas. Al mismo tiempo, las redes sociales han incluido capacidades basadas en la localización en sus aplicaciones. Éstas abren un sinfín de posibilidades en el análisis de los patrones de movilidad humana.

- Reducción en tiempo real de información redundante

Los algoritmos de reducción son útiles para manejar la heterogeneidad y gran volumen del Big Data reduciendo los datos a un tamaño manejable [42, 43]. Estas técnicas se aplican generalmente después de la recolección de datos [44]. Sin embargo, el almacenamiento de todos los datos complejos y de gran tamaño en bruto, redundantes, incoherentes y ruidosos que proceden de fuentes reales de IdC puede ser innecesario. La aplicación de técnicas de reducción en tiempo real puede proporcionar flujos de datos reducidos que contienen información limpia que es realmente relevante para un propósito. Por lo tanto, la aplicación de técnicas de reducción rápidas y efectivas es crucial en el desarrollo de entornos inteligentes para reducir la enorme cantidad de datos a la par que se preserva la información relevante.

- Mejora de la previsión de series temporales mediante la selección de características

Predecir valores futuros de una serie temporal es un reto al que se han enfrentado muchos investigadores durante décadas. Como en cualquier otra tarea de modelado, el preprocesamiento es un paso esencial. En particular, la selección de características, cuyo objetivo es identificar las variables de entrada más relevantes [45]. La selección de características mejora el rendimiento de las variables predictoras al eliminar variables irrelevantes, reduce los datos para acelerar el entrenamiento y aumenta la eficiencia computacional [46] y, a menudo facilita una mejor comprensión del proceso subyacente que generó los datos.

En lo que respecta a las series temporales, no solo debemos procesar las variables de entrada sino que hay más características que deben preprocesarse. Éstos son los *laggeados* y, en el caso de las series temporales multivariadas, el tamaño del conjunto de datos de entrada podría aumentar significativamente. La gestión de flujos de datos de series temporales multivariantes es necesaria para muchas aplicaciones de ciudades inteligentes, ya que los datos del IdC se recogen en múltiples ubicaciones distribuidas y periódicamente en intervalos de tiempo. Por lo tanto, es esencial el desarrollo de una metodología sistemática, automática y basada en datos para la evaluación de características, la construcción y la transformación de series temporales multivariadas que no requieran la aportación de expertos humanos.

Por lo tanto, es esencial el desarrollo de una metodología sistemática, automática y basada en datos para la evaluación de características. Esta metodología debe incluir la construcción de características y la transformación de series temporales multivariantes y no requerir la aportación de expertos humanos.

- Gobernanza de los datos para el IdC

Los datos del IdC son diferentes de los datos que las arquitecturas y plataformas típicas manejan porque son temporales, vienen en flujo y en tiempo real. Compartir y analizar la

1.1. MOTIVACIÓN Y OBJETIVOS

gran cantidad de datos que generan las nuevas tecnologías en tiempo real es clave para desarrollar las aplicaciones que soportan la automatización en escenarios inteligentes. Para hacer frente a los retos inherentes a la planificación y aplicación de soluciones complejas del IdC, necesitamos gobernar nuestros datos a través de plataformas que puedan servir para los fines de todo el proceso. Dichas plataformas también deben ser capaces de gestionar la privacidad y la seguridad de los datos a lo largo de todo el ciclo de vida: recogida de datos, calidad de los datos, almacenamiento de datos, tratamiento de datos, análisis de datos y prestación de servicios.

En resumen, el objetivo de esta tesis es explorar, analizar y aplicar formas de beneficiarse del paradigma del IdC. Este trabajo se basa en la mejora y el análisis de cada paso del proceso de análisis de datos, con el fin de proporcionar mejores servicios a los ciudadanos en entornos inteligentes, es decir, ciudades y edificios inteligentes, con especial énfasis en la eficiencia energética.

Teniendo en cuenta los retos a los que se enfrentan hoy en día tanto el análisis de datos como los edificios inteligentes, establecemos los objetivos que deben alcanzarse para que este objetivo se cumpla, lo que servirá de guía para el desarrollo de la tesis.

- O1. Identificar e integrar datos para crear conjuntos de datos relativos al consumo de energía en entornos inteligentes y determinar la naturaleza de los datos en estudio (binarios, ordinales, temporales, espaciales....). Desarrollar arquitecturas para recopilar y administrar esos conjuntos de datos.
- O2. Desarrollar técnicas de reducción de datos paralelas para las series temporales y, en particular, para los flujos del IdC, preservando sus características clave en relación con las aplicaciones Big Data.
- O3. Crear metodologías y comparar modelos de predicción del consumo de energía con varios horizontes para obtener una predicción altamente precisa y extraer patrones en el uso de la energía.
- O4. Crear características y desarrollar una metodología de reducción de características para series temporales multivariadas aplicadas a la previsión del consumo de energía.
- O5. Identificar, crear y comparar modelos para encontrar patrones en el uso de sistemas de aire acondicionado que puedan ser utilizados para acciones específicas dirigidas hacia la eficiencia energética.
- O6. Identificar patrones de movilidad humana tanto a nivel macro como microscópico utilizando datos de dispositivos portables y redes sociales.

- O7. Identificar y aplicar arquitecturas analíticas de IdC a problemas reales de ciudades inteligentes que integran todas las etapas del proceso, desde la recogida de datos hasta la prestación de servicios.
- O8. Crear mecanismos de IdC para proporcionar servicios personalizados de gestión y sensibilización en materia de energía mediante el análisis de los aspectos de comportamiento relacionados con la eficiencia energética en los edificios inteligentes.

1.2 Resultados

El cuerpo de esta tesis se incluye en varios artículos y capítulos de libros publicados. Gran parte del trabajo se basa en estudios y análisis de los datos generados por los escenarios de IdC, en particular sobre cómo utilizar los datos para la predicción de la energía consumida por los edificios. Otros estudios y publicaciones derivadas de la tesis abordan aspectos específicos relacionados con la creación de infraestructuras inteligentes y otros elementos clave para la resolución de los mencionados objetivos.

El trabajo incluye la integración de 3 conjuntos de datos recogidos en relación con 2 edificios inteligentes y su limpieza, fusión y preprocesamiento, con el fin de obtener conjuntos de datos para su análisis. El primer conjunto de datos pertenece al Centro de Transferencia Tecnológica (CTT) de la Universidad de Murcia⁷. Estos datos son las observaciones ambientales externas al edificio y el consumo total de energía del edificio del 01/12/2014 al 18/02/2018 sen intervalos de 8 horas. En total, 952 observaciones y 15 variables.

El segundo conjunto de datos pertenece a la Facultad de Química de la Universidad de Murcia y está compuesto por 5088 observaciones de 50 atributos que se miden cada hora desde el 02/02/2016 hasta el 06/09/2016.

El atributo de salida es el consumo de energía medido en KWh y hemos incluido mediciones meteorológicas de 3 fuentes diferentes que rodean el edificio, predicciones con una hora de antelación proporcionadas por un servicio web y también atributos de temporada, día de la semana y días festivos.

Por último, también se ha realizado un seguimiento del uso de los sistemas de climatización en 237 aulas de la Facultad de Química. El conjunto de datos consiste en observaciones agregadas de 12 minutos sobre la temperatura ambiente, el estado de encendido/apagado y la temperatura de consigna desde el 31/01/2015 hasta el 28/02/2017.

Estos conjuntos de datos se han creado con el propósito de investigar la interacción entre las personas y los sistemas de los edificios en relación con el consumo de energía, en un intento de extraer patrones de uso y proponer formas automáticas y eficientes para evitar el derroche de energía.

⁷www.um.es/otri/?opc=cttfuentealamo

Después de recopilar conjuntos de datos y estudiar sus características, nos dimos cuenta de la importancia de la reducción de datos y de la selección de características en entornos reales de IdC. La característica temporal de los datos procedentes de sensores (siempre vienen acompañados del momento en que se ha tomado la medición) ha sido explotada para ambos fines.

Hemos investigado métodos para la reducción de datos en entornos inteligentes, analizado sus inconvenientes y propuesto un nuevo método llamado BEATS, que cumple con los requisitos del análisis Big Data. El método propuesto se basa en la división de las series temporales (datos) en bloques que representan subconjuntos de la estructura de datos. BEATS sintetiza la información que contienen estos bloques de forma independiente, reduciendo el número de datos y conservando sus características fundamentales (perdiendo la menor cantidad de información posible). Para ello, BEATS utiliza la agregación de datos basada en matrices, la Transformada de Coseno Discreta y la caracterización de los valores propios de los datos de las series temporales. Comparamos BEATS con los algoritmos de segmentación y representación más avanzados. La mayoría de ellos asumen datos normales, no tratan los *drifts* de los datos, que son muy comunes para entornos inteligentes, y no pueden ser aplicados de manera online. BEATS está diseñado para superar estos problemas: no requiere la normalización de los datos, lo que también ayudará a preservar el valor de los datos (es decir, su magnitud), se puede aplicar de forma online mediante ventanas deslizantes y es posible calcular la distancia entre las series temporales agregadas. Para evaluar BEATS se ha utilizado en experimentos de clasificación con 6 conjuntos de datos reales. Se redujeron los datos entre un 60-70 %, mejorando significativamente el tiempo de cálculo al tiempo que se mantuvo la precisión de la clasificación. También se ha probado en técnicas de clustering donde se logró el mejor coeficiente de silueta para la mitad de los análisis, más que con ninguno de los otros métodos.

El método anterior responde a una necesidad general de los flujos de datos en el análisis de entornos inteligentes. A continuación, se ha realizado un análisis concreto del problema de la predicción del consumo de energía. Los métodos predictivos necesitan algoritmos de preprocesamiento automático que les ayuden a encontrar la mejor combinación de características para el análisis, por lo que proponemos una metodología de selección de característica multivariante que se basa en las características temporales de los datos.

La metodología se basa en *laggear* o retrasar los atributos temporales y en la configuración de una multitud de métodos diferentes de selección de características, tanto de filtro como *wrappers*, univariante y multivariante. Se han utilizado ocho métodos de selección de características para problemas de regresión y, como se esperaba, los métodos de *wrapper* han mostrado un mejor rendimiento que los métodos de filtro, y los métodos multivariantes mostraron un mejor rendimiento univariantes. Además, el Error Absoluto Medio fue mejor (EAM) que el Error Cuadrático Medio (*RMSE* en inglés) a la hora de utilizar una métrica evaluadora para los métodos de *wrapper*. Utilizando nuestra metodología, EAM se mejora en un 42.28% y RMSE en un 36.62% en comparación con no utilizar ninguna técnica de selección de características.

También se ha considerado la creación manual de características y su inclusión en el proceso descrito anteriormente. Se pueden crear variables derivadas de la relevancia retardada tales como: consumo energético a la misma hora y del mismo día pero de la semana anterior, consumo máximo en días laborables / fines de semana de la semana anterior, etc., para incluirlas en el proceso.

En esta tesis se ha realizado un gran esfuerzo para encontrar formas de predecir el consumo de energía en los edificios utilizando varios métodos, horizontes y agregaciones de los datos. De los diversos trabajos que hemos desarrollado en exclusiva para la tarea de modelado del consumo, podemos resumir los siguientes:

- Evaluación del rendimiento de los métodos Multilayer Perceptron (MLP), Bayesian Regularized Neural Network (BRNN), Support Vector Machines (SVM) with Radial Basis Function (RBF) Kernel, Gaussian Processes (GAUSS) with RBF Kernel, Random Forest (RF), eXtreme Gradient Boosting (XGB). Todos ellos entrenados y testados utilizando técnicas de validación del aprendizaje automático.
- Estudio del problema de la predicción del consumo de energía desde el punto de vista de las series temporales. Esto incluye la transformación de los datos y la comparación de algoritmos regresivos tradicionales y el nuevo algoritmo de código abierto Prophet. El modelo implementado en Prophet incorpora componentes no periódicos (utilizando una curva lineal a trozos o de crecimiento logístico), un factor de tendencia que representa los cambios periódicos y los efectos de los días festivos. Éste enmarca el problema predictivo como un ejercicio de ajuste de curvas que difiere de los modelos tradicionales utilizados para las series temporales que se basan en la dependencia temporal de los datos. En este caso hemos incluido una corrección en los datos pronosticados, mejorando la precisión del modelo.
- Uso de la corrección de las predicciones meteorológicas para mejorar el RMSE, obteniendo una mejora de un 4,54% para las predicciones de las próximas 24 horas.
- Comparativa de los diferentes modelos de datos generados (que pueden considerarse del tipo *caja negra*) entre ellos y también con los modelos tradicionales de *caja gris* para la tarea de predicción del consumo diario y semanal.
- Se ha considerado una diferenciación basada en la lógica entre situaciones que dependen del tiempo para etiquetar el comportamiento con respecto al consumo. Estas son las vacaciones y fines de semana, mañanas habituales y tardes habituales. El test no paramétrico Kruskall Wallis y las comparaciones posthoc apoyan la decisión de crear 3 diferentes modelos por día.

- Evaluación no sólo del valor puntual de RMSE, sino también de si un algoritmo de aprendizaje supera estadísticamente a los demás utilizando la prueba no paramétrica de Friedman [47] con las pruebas post-hoc correspondientes para la comparación.

Tras predecir el consumo de energía, tenemos la intención de crear medidas que reduzcan el consumo esperado para obtener un uso más eficiente de la energía. El análisis de los datos de aire acondicionado es una fuente increíble de conocimiento para hacerlo. De esta manera, hemos agregado perfiles similares de variables procedentes de los aparatos (temperatura de consigna, estado de encendido/apagado y temperatura ambiente) de acuerdo a patrones de comportamiento para poder dirigir las acciones que se deben tomar cuando se detectan ajustes de temperatura anormales y el uso de aparatos. Los resultados mostraron que los usuarios pueden ser separados en dos grupos de acuerdo con su interacción con los dispositivos: uno compuesto por aquellos que interactúan con el mando de control del aparato con frecuencia y cambian la temperatura al menos una vez a la semana y otro compuesto por aquellos que interactúan menos con los mandos.

La predicción del consumo energético de los edificios ha sido estudiada desde un punto de vista analítico, utilizando varias técnicas de preprocessamiento, horizontes y parámetros de entrada. Entendemos que existen dos escenarios principales en los que se puede predecir el consumo de energía. El primero se da cuando los modelos pueden utilizar información anterior sobre el consumo, pero solo se pueden utilizar las predicciones del resto de características o datos de entrada ya que se desea estimar el consumo a futuro, es decir, no se pueden utilizar los valores reales de las variables de entrada. El segundo se da cuando “el futuro es ahora” y queremos crear modelos de referencia para los que podamos utilizar datos de entrada reales pero sin consumo previo, ya que esto sesgaría el experimento. Dependiendo del escenario en el que nos encontremos, hemos estudiado cómo ordenar, estructurar y considerar los datos de entrada. Se ha encontrado una mejora de las predicciones al categorizar las habitaciones de acuerdo a sus patrones de uso del aire acondicionado. En ese sentido, la predicción de la movilidad humana permite a las zonas urbanas adaptar sus esfuerzos de transporte y energía a las necesidades reales de su población. Hemos desarrollado estudios preliminares basados en datos de trayectorias de dispositivos portables (*wearables* en inglés) e información geoetiquetada de redes sociales para encontrar patrones y predecir la movilidad humana.

Todos estos procedimientos analíticos que van desde el acopio y la depuración de datos hasta el análisis de los mismos y el análisis de los resultados necesitan una plataforma basada en el IdC para gestionar la interoperabilidad. La plataforma también debería permitir la integración de las técnicas óptimas de análisis de datos y aprendizaje automático para modelar relaciones contextuales y permitir la prestación de servicios. En esta tesis proponemos una arquitectura que se modela en cuatro capas: una capa de tecnologías donde se recogen los datos; una capa denominada *middleware*, donde se limpian y fusionan los datos; una capa de gestión donde se implementan las técnicas de Big Data y análisis ; y una capa de servicios donde se ofrecen

diferentes servicios que dependen del análisis previo.

Uno de los principales servicios que se han obtenido de esta tesis es la prestación de servicios personalizados de gestión y concienciación energética a los ocupantes de edificios inteligentes a través de una plataforma de IdC con el fin de aumentar la eficiencia energética. El resultado es una infraestructura que utiliza una plataforma de IdC como núcleo para administrar los datos, crear la lógica que detecta el derroche de energía, elaborar mensajes personales y cronometrados y entregar la información a través de aplicaciones móviles creadas para tal fin. Los experimentos muestran que es posible mejorar la llamada competencia de ahorro de energía, que representa el conocimiento de una persona para ahorrar energía o, en otras palabras, el potencial de un usuario para ahorrar energía utilizando las cosas que conoce. También se ha demostrado que es posible ahorrar energía a través de la retroalimentación inteligente a los usuarios del edificio.

Los resultados asociados a las contribuciones principales se presentan en la Tabla 2.1, junto al objetivo al que hacen referencia. En el capítulo 3 se explica con más detalle cómo se obtuvieron estos resultados y las principales características de las arquitecturas de IdC propuestas en esta tesis.

1.3 Organización de la Tesis

Esta tesis está organizada como un compendio de trabajos de investigación de alto impacto. Los dos primeros capítulos contienen la misma información en castellano e inglés respectivamente, y como se ha podido observar presentan tanto la motivación y justificación del trabajo como los objetivos y su vinculación con las publicaciones.

El segundo capítulo presenta la motivación y la justificación. Establece los objetivos de la investigación y los vincula a los resultados que se exponen de manera breve y conectada en el sentido de que ciertos objetivos y resultados surgieron de necesidades que se identificaron cuando se establecieron los objetivos anteriores.

El tercer capítulo es una introducción a las publicaciones donde se expone el trabajo relacionado, las brechas identificadas y los resultados, a la vez que se muestra la relación entre todas ellas. Por último, se destacan las conclusiones del trabajo.

Finalmente, el cuarto capítulo está compuesto por los 6 trabajos de investigación de alto impacto, todos ellos Q1 en el ranking de revistas científicas. Estos documentos contienen la información principal sobre los resultados presentados anteriormente. Cada uno de los trabajos de investigación va precedido de una tarjeta de presentación.

1.3. ORGANIZACIÓN DE LA TESIS

Nº	Resultado	Objetivo	Publicaciones
R1	Creación de conjuntos de datos que relacionen el tiempo atmosférico, el consumo, la ocupación y el uso de la información de los edificios. Análisis de las propiedades de éstos datos y su relación mediante análisis estadísticos.	O.1	[4, 5, 3], [7, 8]
R2	Creación de un algoritmo denominado BEATS que agrega y representa datos de series temporales en bloques de vectores de valores propios (menor dimensión). BEATS se adapta a los <i>drifts</i> de los datos reales, puede combinarse con técnicas de aprendizaje automático para su posterior análisis y está pensado para una implementación paralela, siguiendo los requisitos de Big Data.	O.2	[1]
R3	Predicción del consumo energético para varios horizontes (horario, diario, semanal) comparando modelos de caja negra y de caja gris e incluyendo comparaciones estadísticas de los resultados de los métodos más precisos.	O.3	[2, 3, 4], [7, 9, 10, 11, 12]
R4	Desarrollo de una metodología para la predicción de series temporales multivariadas de energía basada en métodos de selección de características para la regresión de series temporales que incluye métodos univariantes, multivariantes, de filtro y <i>wrappers</i> .	O.4	[2], [10]
R5	Creación de entidades de alto nivel en un edificio (grupos de usuarios / habitaciones) extrayendo perfiles de uso de los aires acondicionados usando métodos de clustering.	O.5	[5], [8]
R6	Modelado de la movilidad humana basado en áreas de tránsito denso y en datos de redes sociales con <i>Complex Event Processing</i> .	O.6	[13, 14, 15, 16]
R7	Creación de una arquitectura de Big Data basada en el IdC para proveer de servicios en las ciudades inteligentes en general que se modela en 4 capas: tecnologías, fusión, gestión y servicios; integrando funcionalidades de minería de datos en la capa de gestión. La plataforma pretende ser un paso hacia la plena adaptación del paradigma del IdC en la recuperación, gestión y análisis de datos energéticos en los edificios.	O.7	[4, 5], [17]
R8	Creación de una plataforma con mecanismos abiertos y extensibles para la gestión de datos de sensores. Combinando servicios relacionados con la energía y el análisis del comportamiento se construyen servicios de recomendación y se entregan a través de aplicaciones personalizadas a los ocupantes del edificio, lo que tiene un impacto directo en su comportamiento y, por lo tanto, aumenta la eficiencia energética.	O.8	[6], [18, 17]

Cuadro 1.1: Resultados. Ver en negrita las publicaciones que componen la tesis. El resto son nuestras publicaciones adicionales.

CHAPTER



SUMMARY

This chapter introduces the motivation and justification of the thesis. It presents the research objectives and links them to the results that are briefly explained and connected.

2.1 Motivation and Goals

Climate change is already disrupting national economies and affecting lives all around the world. Its consequences are costing dearly today and, if its progression continues the cost will be much greater in the future.

Weather events are becoming more extreme, sea levels are rising and greenhouse gas emissions are now at their highest levels in history. Without action, global warming is likely to be as much as 5°C by the end of the century¹, having a huge impact on life as we know it nowadays.

To strengthen the global response to prevent climate change, countries have adopted many initiatives. In 2015, countries adopted the 2030 Agenda for Sustainable Development and its *Sustainable Development Goals*² which are a call for action by all countries to promote prosperity while protecting the planet. Within the 17 goals, four of them are directly related to our goals: the inclusion of affordable and clean energy, sustainable cities and communities, responsible consumption and production and climate action. In the Paris Agreement at the COP21 (2016), countries agreed to work to limit global temperature rise to well below 2 degrees Celsius³.

Europe is also devoting considerable effort to cut its greenhouse gas emissions substantially. By 2050, as part of the efforts required by developed countries as a group, the EU aims to cut its

¹<https://www.consilium.europa.eu/en/policies/climate-change/>

²<https://www.un.org/sustainabledevelopment/climate-change/>

³<https://unfccc.int/process-and-meetings/the-paris-agreement/the-paris-agreement>

emissions substantially – by 80-95 % compared to the levels in 1990 ⁴.

Research and innovation make a crucial contribution to fighting and adapting to climate change, and Information and Communications Technologies (ICT) have the potential to reduce 20 % of global CO₂ emissions by 2030, holding emissions at 2015 levels [19]. A report from British Telecommunications claimed that the influence of ICTs in the EU is expected to reduce the carbon footprint of EU by 37 %, holding emissions at 2012 levels.

Artificial intelligence (AI) and its particular applications, such as Machine Learning (ML) are proving to be highly adept at spotting the many inefficiencies in modern society that contribute to climate instability.

The thesis work is based on the combination of novel ICT technologies for data collection and management and its analysis through ML in order to provide smarter environments that can make a responsible use of the resources. The application of this work contributes to mitigate climate change in a broad sense.

2.1.1 Internet of Things (IoT) and smart environments

An IoT device is a physical object that connects to the Internet to transfer data. Thanks to the proliferation of IoT devices that are interconnected, huge amounts of data are being gathered nowadays. This allows the creation of smart environments.

Smart environments are physical environments that are richly and invisibly weaved together with IoT devices; that is, sensors, actuators, gadgets, and computational elements in general, embedded seamlessly in the quotidian objects that surround us, and connected through a continuous network. However, it is not the sensors that makes an environment smart, but the ability to process and learn from all the data that they provide through its analysis in order to automatically provide services.

The development and evolution of Big Data analytics and the IoT technologies are playing a major role in the adoption of smart city initiatives for various reasons. The first reason is the exponential growth of smart objects that can participate in an IoT infrastructure [20]. Cisco Internet Business Solutions Group predicts 50 billion connected devices by 2020 [21]. Two other remarkable reasons are population growth and the urbanization trend [20]. According to the United Nations, there are a total of 1.3 million people moving into cities every week, with urban populations growing to 6.3 billion that is a 68 % by the year 2050⁵. This rapid increase in urban populations brings an intense stress on global infrastructure and environment since cities account for more than 70 % of global energy use [22] and produce 80 % of its greenhouse gas emissions [23].

In that sense, leverage IoT solutions for smart cities helps promoting economic development, upgrades infrastructure, improves environment and optimises transportation systems in a

⁴https://ec.europa.eu/clima/citizens/eu_en

⁵<https://www.un.org/development/desa/en/news/population/2018-revision-of-world-urbanization-prospects.html>

sustainable manner while improving the quality of life in the cities. Urban areas are the perfect laboratory for cutting greenhouse gas emissions, increasing the use of renewable energy and improving energy efficiency. Some important smart city components are depicted in Fig. 2.1

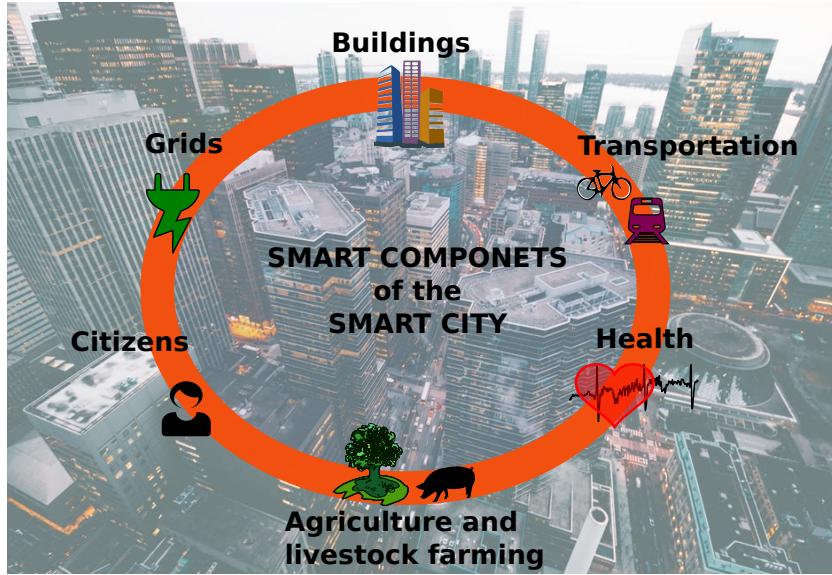


Figure 2.1: Smart City components

Finding ways to meet the energy needs of a growing population in conjunction with growing economic prosperity and resource scarcity is a fundamental challenge to achieving a sustainable society. The reduction of energy consumption and carbon footprint are important issues in smart cities. When developing smart cities, sustainability is based on energy efficiency and at a global scale, buildings are the cornerstone for energy efficiency in terms of power consumption and CO₂ emissions [24].

The building sector is also greatly affected by the proliferation of smart meters and home displays. This trend seems to be on the rise if we consider that the European Commission has established that 16 Member States will proceed with a large-scale roll-out of smart meters by 2020 or earlier [25]. This, along with new developments as regards Energy Data Infrastructure (see [5, 6]), has formed the perfect environment for the creation of, among other technologies, advanced energy feedback strategies for the reduction of energy use in buildings and for the education of building occupants/users [26], the so-called *smart building*.

A smart building is any commercial, residential or industrial structure that implements automation to control its operation based on data collected by sensors. This includes the internal environment such as Heating, Ventilation and Air Conditioning (HVAC), lighting, security, shading, etc. [27], and the external, such as the weather. Smart buildings are expected to consider elements inside and outside their perimeter and interact with electrical grids, environmental conditions, and the goals and duties of their users. Smart buildings target to improve energy efficiency, occupant comfort and environmental impact of the building as a whole.

Smart buildings are considered instrumental in bringing about the smart city. In the Smart Buildings Magazine, Harry G. Smeenk, vice president of program development at the Telecommunications Industry Association noted that “*Developing smart buildings will give rise to smart campuses, which will foster smart communities, and eventually smart cities. Simply put, smart buildings will create a scalable foundation for creating the elusive smart city, building by building, from the ground up*”⁶.

The energy consumed in buildings in developed countries comprises 20-40 % of their total energy use and it is above that of industry and transport in the EU and US [28, 29].

In order to mitigate climate change, the reduction of energy use together with the use of non-fossil energy sources is crucial. Furthermore, reducing energy consumption in buildings has to be done while ensuring buildings’ users comfort and lower costs in order to combat fuel poverty. Initial analyses suggest that the conversion of buildings into smart buildings thanks to the IoT sensorisation, together with data analytics might be an option by which to resolve these issues.

In the 2016 survey from Continental Automated Buildings Association (CABA) named *Intelligent Buildings and the Impact of the Internet of Things*, the following 3 main challenges when dealing with making buildings smarter were identified [30]

- Improving Spending Decisions

The fact that buildings’ energy usage patterns are often not possible to determine by building managers makes it difficult to identify the proper energy-saving opportunities. Therefore, many times the implemented energy-saving measures either do not improve efficiency or needlessly reduce users comfort. IoT systems can tackle this issue by exposing detailed energy use data, allowing managers to spot inefficiencies and creating highly accurate models for prediction.

- Reducing Energy Consumption and Expenditure

Controlling how equipment is used normally requires manual supervision. This way, it is complicated to reduce energy consumption and control costs. Automation of a building’s appliances and elements allows for greater control of how much, when and how energy is consumed.

With the IoT, managers can remotely observe and adjust building systems with a tap of a button, making it far easier to bring costs down. The potential energy savings can be further enhanced with IoT technology.

- Improving Operational Efficiency

Most buildings have separate systems for HVAC, lighting, power, indoor air quality, internet connectivity, refrigeration, and so forth. This makes it very difficult to optimise overall

⁶<http://www.smartbuildingsmagazine.com/features/the-smart-way-to-smart-cities-begins-with-buildings>

building operations. The IoT creates an opportunity to integrate data from numerous sources into a single analytics platform. In this way, managers can apply a holistic strategy to building operations. Pairing IoT technology with smart buildings can provide a predictive maintenance system. When the building's parameters are being monitored it is easier to detect abnormal events. The building manager can be informed instantly to act accordingly. In this way, there are fewer failures in the equipment which contributes greatly to cost savings of smart buildings.

The last challenge that we highlight relates to the active and passive behaviours of occupants with regard to energy. Those behaviours include window opening, use of appliances, solar shading and blinds, adjusting HVAC set points, lighting choice, etc. [31]. To ensure sustained reductions in energy consumption, energy-saving technologies must be accompanied by energy efficient occupant behaviour [32]. As stated in the report of the European Environment Agency [33], up to 20 % of energy savings can be achieved through different measures targeting consumer behaviour. IoT can help with the provision of specialised tasks depending on the real-time context that can lead to the education of users towards sustainability.

Despite these clear advantages, many buildings have yet to adopt IoT technology. According to [34], the lack of a smart infrastructure in buildings implies that no country in Europe is fully ready for the smart revolution. In other words, the lack of smart devices and connectivity amongst them in buildings is crucial for unlocking buildings' possibilities. Taking into account the following components: HVAC, plug load, window shading and building automation, an upgrade a single of them in an isolated way it can result in energy savings of 5–15%, and an integrated system can realise 30–50% savings in existing buildings that are otherwise inefficient [35].

2.1.2 Data analytics and Big Data in smart environments

The huge amount of heterogeneous data that is captured, stored, and managed by means of the IoT exceeds the capabilities of traditional database infrastructures and engines. Originally, the 3 Vs: [36] high volume, high velocity and high variety of data were considered the characteristics responsible for the emergence of Big Data technologies that help resolve the problems that exceed conventional requirements.

Additional V characterizations have been proposed over time, and we consider that the following 7 Vs [37] are more precise descriptors of the complexity of Big Data:

- Volume: the huge amount of IoT devices, including wearable devices, generates massive amounts of data. The concerns regarding data size are its scalability, accessibility and manageability.
- Velocity: The transfer rate of data between the source and destination.

- Variety: Several types of data: structured or unstructured data from different sources: image, video, text, sensors, etc.
- Veracity: Incoming real data from IoT are hardly ever clean and precise. It is necessary to find mechanisms to ensure that data are trustworthy.
- Validity: When the data moves from exploratory to actionable, it must be validated. Validity relates to the correctness and accuracy of data with regard to the intended usage.
- Volatility: data retention is specially important in Big Data problems due to the lengths of data. In many cases it is crucial to determine at what point data is no longer relevant to the current analysis and should not be stored.
- Value: Value represents the business value to be derived from data. The interest is always to extract maximum value from the data. Data value must exceed its cost or ownership and management, including storage.

At the same time that data are collected in unprecedented amounts, less than 1 % of these data are being analysed [38]. This is due to the complexities that problems regarding Big Data imply. There exist several challenges in the analysis of real data such as high dimensionality, high volume, noise, and data drifts. Data provided by IoT sources (sensory devices and sensing mechanisms) are multi-modal and heterogeneous.

All of the above mentioned features hinder the execution and generalization of algorithms, so we have identified the following challenges with regard to data:

- Sensor data fusion

Data fusion is defined in [39] as the combination of data from multiple sensors to produce more accurate, more complete, and more dependable information that would not be possible to achieve through a single sensor. In other words, data fusion is a processing technique that matches, merges, aggregates, and integrates data from several sources.

Innovative services can be created by the fusion of data. In that sense, data fusion is a crucial challenge that needs to be addressed. In smart cities applications it is essential to fuse, and interpret the data automatically and intelligently [40]. Data fusion and data filtering have been listed as two major challenges for the IoT and its applications, like smart cities are [41].

- Human mobility pattern identification

Human mobility is especially important for applications such as traffic forecasting, urban planning, and epidemic modeling. Understanding mobility patterns can support data-driven decisions and improve quality of life in smart cities. Traditionally, non-scalable techniques were used for finding macroscopic patterns. Nowadays, the incorporation of

GPS technology in wearable devices has made it possible to collect a large amount of high-resolution digital traces that can give insight into the underlying spatiotemporal trajectories of people. At the same time, social networks have included location-based capabilities into their applications. These open up a wealth of possibilities in the analysis of human mobility patterns.

- **Real-time redundant information reduction**

Reduction algorithms are useful to manage the heterogeneity and the big volume of Big Data by reducing data into a convenient size [42, 43]. These techniques are usually applied after data collection [44]. However, storing all the complex and large raw, redundant, inconsistent, and noisy data that come from real IoT sources might be unecessary. Applying reduction techniques in real time can provide reduced data streams containing clean information that is really relevant for a purpose. Therefore, the application of fast and effective reduction techniques is crucial in the development of smart environments to reduce the massive amount of data while relevant information is preserved.

- **Improvement of time series forecasting using Feature Selection (FS)**

Forecasting future values of a time series is a challenge that many researchers have faced for decades. As in any other modelling task, preprocessing is an essential step. In particular, FS which aims to identify the most relevant input variables [45]. FS consists of eliminating inputs that are irrelevant for the task in order to enhance predictive algorithms' performance. In that sense, FS achieves data reduction, serving for accelerating training and increasing computational efficiency [46]. Furthermore, FS can provide a better understanding of the process that generated the data.

Regarding time series, there are extra candidate features to be preprocessed. Those are the lagged values and, in the case of multivariate time series, the size of the input dataset might increase significantly. Handling multivariate time series data streams is necessary for many smart city applications since IoT data is collected from multiple, distributed locations and periodically over a time range.

Therefore, it is essential the development of a systematic, automatic, data driven methodology for feature evaluation. Such methodology should include feature construction and transformation of multivariate time series and not require input from human experts.

- **Data governance for IoT**

IoT data is different from the data that typical application architectures and platforms handle because it is temporal, on streams and real-time. Sharing and analysing the vast amount of data being generated by new technologies in real time is key in order to develop the applications that support automation in smart scenarios. To address the challenges inherent in planning and implementing complex IoT solutions, we need to govern our data

through platforms that can serve for the purposes of the whole process. Those platforms should also be capable of managing the privacy and security of the data across the entire lifecycle: data collection, data quality, data storage, data processing, data analysis and service provision.

In short, the aim of this thesis is to explore, analyse and implement ways to benefit from the IoT paradigm. This work is based on the improvement and analysis of every step in the data analysis process, leading to provide better services to citizens in smart environments, namely smart cities and smart buildings, with a special emphasis on energy efficiency.

Considering the challenges that both data analytics and smart buildings face nowadays, we set out the objectives that must be attained for this aim to be fulfilled, which will serve as a guide for the development of the thesis.

- O1. To identify and integrate data in order to create datasets relative to energy consumption in smart environments and to determine the nature of the data under study (binary, ordinal, temporal, spatial...). Develop architectures to collect and manage those datasets.
- O2. To develop parallel data reduction techniques for time series and, in particular, for IoT streams preserving their key characteristics regarding Big Data applications.
- O3. To create methodologies and compare models for energy consumption forecasting with several horizons in order to obtain highly accurate forecasting and to extract energy usage patterns.
- O4. To create features and develop a feature reduction methodology for multivariate time series applied to energy consumption forecasting.
- O5. To identify, create and compare models for finding patterns in using HVAC systems which can be used for target actions towards energy efficiency.
- O6. Identify human mobility patterns in both macro and microscopic levels using data from wearable devices and social network.
- O7. To identify and apply IoT analytic architectures to real smart city problems that integrate all steps of the process, from data collecting to service provision.
- O8. To create IoT mechanisms in order to provide personalized energy management and awareness services by analysing behavioural aspects related to energy efficiency in smart buildings.

2.2 Results

The body of this thesis is included in several published articles and book chapters. Much of the work is based on studies and analysis of the data generated by IoT scenarios, particularly on how to use the data for the prediction of the energy consumed by buildings. Other studies and publications stemming from the thesis tackle specific aspects related to the creation of smart infrastructures and other key elements for resolving the aforementioned objectives.

The work includes the integration of 3 datasets collected from 2 smart buildings and their cleaning, fusion and preprocessing in order to obtain datasets to be analysed. The first dataset belongs to the Technological Transfer Centre (TTC) at the University of Murcia⁷. These data are the environmental outdoor observations and the total energy consumption of the building from 2014-12-01 to 2016-02-18 in intervals of 8 hours. In total, 952 observations and 15 variables.

The second dataset belongs to the Faculty of Chemistry at the University of Murcia and it is composed of 5088 observations of 50 attributes that are measured hourly from 2016-02-02 until 2016-09-06. The output attribute is the *energy consumption* measured in KWh and we have included meteorological measurements from 3 different sources that surround the building, hour ahead predictions provided by a web service and also season, day of the week and holiday attributes.

Finally, we have also monitored the use of HVAC systems in 237 rooms of the Faculty of Chemistry. The dataset consist of 12-minutes aggregated observations regarding room temperature, on/off status and set point from 2015-10-31 until 2017-02-28.

These datasets were created with the purpose of investigating phenomena related to the interaction between people and buildings' systems regarding energy consumption in an attempt to extract usage patterns and propose automatic and efficient ways to avoid wasting energy.

After collecting datasets and studying the characteristics of the data, we realised the importance of both data reduction and FS in real IoT environments. Data coming from real sensors presents a temporal characteristic that has been exploited for both purposes.

We have investigated methods for data reduction in smart environments, analysed their drawbacks and proposed a novel method called BEATS, that complies with requirements of Big Data analysis. The proposed method is based on splitting time series data into blocks which represent subsets of the whole data structure. BEATS synthesizes the information that the blocks contain independently, by reducing the data points while still preserving their fundamental characteristics (loosing as little information as possible). For such purpose, BEATS uses matrix-based data aggregation, Discrete Cosine Transform (DCT) and eigenvalues characterization of the time series data. We compare BEATS with the state-of-the art segmentation and representation algorithms. Most of them assume normal data, do not handle data *drifts* —which are very common for smart environments—and they cannot be applied in a online manner. BEATS is designed to overcome those issues: it does not require normalization of the data, which will

⁷www.um.es/otri/?opc=cttfuentealamo

also help to preserve the value of the data points (i.e. magnitude of the data), can be applied in an online way using sliding windows and it is possible to compute the distance between the aggregated time series. For BEATS evaluation, it was used in 6 public real datasets. Data were reduced between 60-70 % of and computation time was significantly improved while keeping accuracy in classification. It was also tested for clustering where it achieved the best silhouette coefficient in half of the analysis, more than any of the other methods.

The prior method responds to a general necessity for data streams in smart environments analysis. Following that, a more specific analysis was carried out in the particular problem of energy consumption prediction. Predictive methods need automatic preprocessing algorithms that can help them find the best combination of features for the analysis, so we propose a multivariate FS methodology that is based on the temporal characteristics of the data. The methodology is based on lagging the temporal attributes and configuring a collection of different methods for FS, both filter and wrapper, univariate and multivariate. We applied eight different FS methods for regression and, as expected, wrapper FS methods showed better performance than filter FS methods, and multivariate FS methods showed better performance than univariate FS methods. Also, Mean Absolute Error (MAE) was better than the root mean squared error (RMSE) as metric performance in evaluators for wrapper FS methods. Using our methodology, MAE is improved by 42.28 % and RMSE by 36.62 % compared to not using any FS technique. The manual creation of features and its inclusion in the process described above has also been considered. Variables derived from the lagged-relevancy such as: energy load of the same hour and the same day but previous week, maximal load of working days / weekends during previous week and so forth can be created in order to include them in the process.

A great effort has been made in this thesis in order to find ways to predict energy consumption in buildings using several methods, horizons and aggregations of the data. From the several works that we have develop in the exclusive modelling task, we can summarise the following:

- Evaluation of the performance of Multilayer Perceptron (MLP), Bayesian Regularized Neural Network (BRNN), Support Vector Machines (SVM) with Radial Basis Function (RBF) Kernel, Gaussian Processes (GAUSS) with RBF Kernel, Random Forest (RF), eXtreme Gradient Boosting (XGB). All of them trained and tested using ML validation techniques.
- Study of the energy consumption forecasting problem from a time series point of view. This includes the transformation of the data and the comparison of traditional regressive algorithms and the novel open source library Prohpst. The implemented model on Prophet incorporates non-periodic components (using piecewise linear or logistic growth curve trend), a trend factor that represents periodic changes and holidays effects. It frames the forecasting problem as a curve-fitting exercise which differs from the traditional models used for time series that account for the temporal dependence structure in the data. In this case we have included a correction on forecasted data, improving the accuracy of the model.

- Use of weather forecasting correction for improving RMSE, obtaining a 4,54 % improvement on average for 24 hours hourly predictions.
- Comparison of the several generated data driven models (that can be considered black box) amongst them and also with traditional grey box models for the task of daily and weekly consumption prediction.
- A logic differentiation between temporal situations was considered in order to label behaviour. Those are holidays and weekends, regular mornings and regular afternoons. The non-parametric Kruskall Wallis and the posthoc pairwise comparisons support the decision of creating 3 different models per day.
- Evaluation of not only the punctual value of RMSE, but also of whether one learning algorithm out-performs statistically significantly the others using the non parametric Friedman [47] test with the corresponding post-hoc tests for comparison.

After predicting energy consumption, we intend to create measures that will reduce the expected consumption going towards a more efficient use of energy. The analysis of HVAC data is an incredible source of knowledge in order to do so. In that way, we have aggregated similar profiles of HVAC variables (setting point, on/off status and room temperature) into behaviour patterns in order to be able to direct the actions that need to be taken when detecting abnormal temperature settings and usage of HVACs. Results showed that users can be separated in two groups according with their interaction with the devices: one composed by those who often interact with the controllers and change the temperature at least once a week and another one composed by those who interact less with the controllers.

The energy consumption prediction in buildings has been studied from an analytical point of view, using several preprocessing techniques, horizons and input parameters. We understand that there are two main scenarios where energy consumption prediction takes place. The first one is when models can use available past information regarding consumption but they can just use predictions for future inputs since we want to estimate future consumption. The second one happens when “the future is now” and we want to implement baseline models for which we can use the real inputs but no prior consumption because it would bias the experiment. Depending on the scenario we are at, we have studied how to tidy, structure and consider input data. An enhancement of predictions is encountered when categorising rooms according to their HVAC usage patterns. In that direction, the prediction of human mobility allows urban areas to adapt their transport and energy efforts to the real needs of their population. We have developed preliminary studies based on trajectory data from wearable devices and geotagged information from social network in order to find patterns and predict human mobility.

All those analytic procedures that go from data collection and cleaning to the analysis of the data and analysis of results need an IoT-based platform in order to manage interoperability

aspects. The platform should also enable the integration of the optimal data analysis and machine learning techniques in order to model contextual relationships and allow service provision. In this thesis we propose an architecture that is modelled in four layers: a technologies layer where data is collected; a middleware layer where data is cleaned and fused; a management layer where Big Data techniques and analysis are implemented; and a service layer where different services that depend on the previous analysis are offered.

One of the main services that were obtained from this thesis was the provision of personalised energy management and energy awareness services to smart buildings occupants through an IoT-platform in order to increase energy efficiency. The result is a framework that uses an IoT platform as the core to administer the data, create the logic that detects energy waste, elaborates messages that are personal and timed, and deliver the information via created-for-purpose mobile apps. The experiments show that it is possible to improve the so-called energy saving competence, which represents the knowledge of a person to save energy or, in other words, the potential of a user to save energy by using things they know. It has also been proven that it is possible to save energy via intelligent feedback to building users.

The results associated with the main contributions are presented in Table 2.1, alongside the objective referred to. In Chapter 3 we explain in more detail how these results were obtained, and the principal characteristics of the IoT architectures proposed in this thesis.

2.3 Organisation of the Thesis

This thesis is organised as a compendium of high impact research papers. The first two chapters contain the same information in Spanish and English respectively and as it could be seen, they introduce both the motivation and rationale of the work and the objectives and their linkage with the publications.

The second chapter introduces the motivation and rationale. It sets the research objectives and links them to the results that are exposed in a brief and connected manner in a sense that certain objectives and results arose from necessities that were identified when other objectives were set.

The third chapter is an introduction to the research publications where the related work, the identified gaps and the results are exposed while showing the linkage between all of them. At the end it highlights the conclusions of the work.

Finally, the fourth chapter is composed by the 6 high impact research papers -all of them are ranked as Q1. Those papers contain the main information regarding the results presented above. Each of the research paper is preceded by with its presentation card.

2.3. ORGANISATION OF THE THESIS

Nb	Result	Objective	Publications
R1	Creation of datasets relating weather, consumption, occupation and usage of buildings information. Analysis of data properties and their relationship using statistical analysis.	O1	[4, 5, 3] , [7, 8]
R2	Creation of an algorithm named BEATS that aggregates and represents time series data in blocks of lower dimensional vectors of eigenvalues. BEATS adapts to drifts in real data, can be combined with machine learning techniques for further analysis and it is thought for a parallel implementation, following Big Data requirements.	O2	[1]
R3	Energy consumption forecasting for several horizons (hourly, daily, weekly) comparing black-box and grey-box models including statistical comparison of results between highly accurate methods.	O3	[2, 3, 4] , [7, 9, 10, 11, 12]
R4	Development of a methodology for energy multivariate time series forecasting based on FS methods for time series regression that includes univariate, multivariate, filter and wrapper methods.	O4	[2] , [10]
R5	Creation of higher level entities in a building (groups of users / rooms) by extracting profiles on HVAC data using clustering methods.	O5	[5] , [8]
R6	Human Mobility Modelling Based on Dense Transit Areas and Social Media with Complex Event Processing	O6	[13, 14, 15, 16]
R7	Creation of an IoT-based Big Data architecture for smart city services in general that is modelled in 4 layers: technologies, fusion, management and services integrating data mining functionalities as built-in features the management layer. The platform intends to be a stage towards the full adaptation of the IoT paradigm in the retrieval, management and analysis of energy data in buildings.	O7	[4, 5] , [17]
R8	Creation of a platform with open and extensible mechanisms for sensor data management. Combining energy and behavioural analytics and recommendation services actions are built and delivered through personalized applications to the building occupants, having a direct impact on their behaviour and, thus, increasing energy efficiency.	O8	[6] , [18, 17]

Table 2.1: Results. In bold the publications composing the thesis. The others are our additional publications

CHAPTER



THESIS CONTRIBUTIONS

This chapter is an introduction to the research publications where the related work, the identified gaps and the results are exposed while showing the linkage between all of them. At the end the conclusions of the work are highlighted.

3.1 Related Work

In this section, we make a thorough search and description of the attempts to solve the previously identified challenges.

First, we expose the concerns regarding the needs of building management with real continuous data instead of audits present in literature. We also review some machine learning approaches in which energy consumption prediction has been involved. After that, we expose the related work for the data processing problems that we faced: time series representation and FS for energy efficiency prediction. Finally, the state-of-the-art IoT architectures for smart cities, energy management and behavioural analytics for energy efficiency are presented.

At the end of each subsection, the identified gaps have been included. We have developed a series of methodologies and techniques for filling those gaps.

3.1.1 Why energy consumption prediction is useful and how has it been carried out according to literature

The application of data analytics for researching how to improve buildings' operation is widely spread in the literature. This includes, amongst others, model-based predictive controls for energy consumption and heating, Demand Response (DR), occupancy detection, and forecasting and automated fault detection. However, building managers typically rely on in-house or external

energy audits carried out yearly at best to determine suitable energy conservation measures [48]. An energy audit entails a revision of the energy efficiency of existing equipment, the operating conditions at facilities, and data collection and analysis with the aim to optimise energy usage and identify energy efficiency measures. Such activities are invasive and labour-intensive and neglect inter-building diversity and metrics that are not related to energy such as occupant satisfaction. Furthermore, several studies show that building performance decreases over time after the implementation of the findings of an energy audit [49, 50]. Therefore, the need for non-invasive energy audits over the complete life of a building motivates the data analytics research in buildings since it renders the potential for metrics that define fairer load profiles in buildings [48]. For example, [51] utilises a combination of clustering and deep learning methods, in conjunction with a weighted aggregation mechanism in order to improve load forecasting accuracy over a short period of time.

Artificial neural networks (ANNs) are able to learn the key information patterns within a multidimensional domain. These have been applied in the field of solar energy, for modeling and design of a solar steam generating plant [52], for the estimation of heating-loads of buildings [53], etc. They have also been used in HVAC systems, solar radiation, modeling and control of power-generation systems, load-forecasting and refrigeration [54]. BRNN are a type of ANN and have been used in the prediction of a series of building energy loads from an environmental input set [55]. Also, RF model has been applied in order to predict energy consumption in residential buildings [56].

Likewise, SVM have been used to predict both the total short-term electricity load and the short-term loads of individual building service systems (air conditioning, lighting, power, and other equipment) in buildings that have electricity sub-metering systems installed [57].

Another common technique for non-linear regression proposed in the literature to be applied are GAUSS with RBF Kernel [58]. It has already been used to forecast electrical load [59] or to estimate the number of occupants in a room according to data related to the room status: motion detection, CO₂ reading, sound level, ambient light and door state sensing [60].

In the reviewed literature, several time-series modelling techniques have been used for different load forecasting problems. For example, Kawashima [61] explored AutoRegressive Integrated Moving Average (ARIMA), Exponentially Weighted Moving Average (EWMA) and Linear Regression (LR) [62], to forecast cooling loads one-day in advance.

In those cases in which limited amounts of data are available and the information concerning the building architecture is partially known, grey models are suitable alternatives for the prediction of energy consumption [63]. Grey-box models use simplified physical descriptions to simulate the behaviour of a building's energy systems, and with them identify important parameters and characteristics using statistical analysis [64]. According to this nomenclature, the previously mentioned ML models are known as black-box models.

Since it was shown that resistor-capacitor (RC) networks can accurately represent the thermo-

dynamics of buildings [65], grey-box models have been used to represent the thermodynamics of buildings. Nowadays, programs such as EnergyPlus, include thermal networks in their codes [64]. This has motivated research into ideal model topologies and methodologies for these models so as to ensure that they accurately represent the responses of buildings [66]. Also, other works are now focusing on how these methods can be used for the characterization of the thermal envelope [67].

One of the objectives of load forecasting in buildings is to make short term near real-time predictions of energy demands. These forecasts can be used in planning and allocating resources to meet the demand. However, there are more research categories on data analytics-driven for buildings in which energy consumption predictions are used. Those are: (a) baseline creation, (b) Model Predictive Control (MPC) , (c) DR, and (d) occupant-centric controls.

- Baseline creation

Baselines can be used not only for performance monitoring but also for confirming and measuring energy savings derived from the implementation of energy savings actions [6].

- MPC is composed of at least two elements: (a) a forecasting algorithm and (b) an optimization algorithm to determine the optimal control sequence. Many forms of MPC have been applied for the control of HVAC equipment [68]. MPC offers the possibility of anticipating the energy needs of a building taking into account the usual requirements (e.g., comfort ranges) and the possible events that alter consumption (thanks to the prediction model), being able to optimise the building's thermal behaviour on the basis of the defined control goals. This facilitates the use of energy storage capabilities for optimising the use of renewable energy generated on-site.
- Demand Response (DR) also uses predictive models. DR energy requests, based on the current grid status, are generated from the utility or system operators and sent to the building [69]. DR can be used by building owners and operators for responding to real-time pricing [70] and also a utility company can use it to send emergency signals [71]. The objective of DR is to achieve energy cost savings through load shifting and peak load reduction strategies in response to almost real-time variations in the utility rates.
- Occupant centric control: Occupant's detection implies certain automation for lighting, however, thermal systems are not immediate and they need a certain occupancy prediction in order to improve a building's intelligence. Consumption baselines give insight into people's whereabouts, being an important element of occupant behaviour modelling. Occupancy schedules can be derived from the building's electricity consumption [72] and, at the same time, variations on consumption sometimes mean anomalous events.

In many of the reviewed studies for energy consumption forecasting, we encountered the following gaps: (i) the temporal characteristics of energy consumption and/or inputs have not

been considered, (ii) ML models have not been compared to other models that include physical parameters of the buildings and (iii) the inclusion of external predictors is not realistic or adapted to the predictive model context.

3.1.2 Time series representation

There are several approaches to represent a numeric time-dependent variable (i.e. a time series). Using basic statistics would not represent all the information that the time series contains. A classical example that supports this claim is the Anscombe's Quartet, [73] that shows how four very different datasets have identical simple statistical properties: mean, variance, correlation and regression coefficients.

In order to reduce the number of data points in a series and create a representation, segmentation methods can be used as a pre-processing step in data analytics.

Given a time series T containing n data points, segmentation is defined as the construction of a model \bar{T} , from l piecewise segments ($l < n$) such that \bar{T} closely approximates T [74].

The segmentation algorithms that aim to identify the observation where the probability distribution of a time series changes are called change-point detection algorithms. Sliding windows, bottom-up, and top-down methods are popular change-point detection based approaches. For sliding windows, each segment is grown until it exceeds an error threshold. In the bottom-up methods, the segments of data are merged until some stopping criteria is met and top-down methods partition the time series recursively until a stopping criteria is met [75].

Another way of classifying the algorithmic methods for segmentation is considering them as online and offline solutions [76]. While offline segmentation is used when the entire time series is previously given, the online segmentation deals with points that arrive at each time interval. In offline mode, the algorithm first learns how to perform a particular task and then it is used to do it automatically. After the learning phase is completed, the system cannot improve or change (unless we consider incremental learning or retraining). On the other hand, online algorithms can adapt to possible changes in the environment. Those changes are known as "drifts". Whereas top-down and bottom-up methods can only be used offline, sliding windows are applicable to both circumstances.

After segmentation, the representation of the time series based on the reduction can be regarded as an initial step that reduces the load and improves the performance of tasks such as classification and clustering. The use of such algorithms can be generally regarded in two ways:

- Representation methods: Extracting features from the whole time series or its segments and applying ML algorithms in order to classify them or compute the distance between the time series representation for clustering.
- Instanced based methods (similarities): Computing the distance matrix between the whole series and using it for clustering or classification applying a k-nearest neighbour approach

[77] by finding the most similar (in distance) time series in the training set.

We review the work made using both approaches since the ultimate goal of our time series representation is to make the time series data more compact for further processing.

- Whole series similarities: Similarity measures are used to quantify the distance between two raw time series. The list of approaches is vast and the comparison between well-known methods has lead to the conclusion that the benchmark for classification is dynamic time warping (DTW) since other techniques proposed before 2007 were found not significantly better [78].
- Intervals: For a series of length m , there are $m(m - 1)/2$ possible contiguous intervals.

Piecewise Linear Representation (PLR) [79] methods are based on the approximation of each segment in the form of straight lines and include the perceptually important points (PIP), Piecewise Aggregate Approximation (PAA) [80], and the turning point (TP) method [81].

The state-of-the-art models Time Series Forest (TSF)[82] and Learned pattern similarity (LPS)[83] generate many different random intervals and classifiers on each of them, ensembling the resulting predictions.

- Symbolic Aggregate approXimation (SAX) [84] Among all the techniques that have been used to reduce the number of points of a time series data, SAX has attracted the attention of the researchers in the field. SAX allows a time series of length n to be reduced to a string of length l ($l < n$). The algorithm has two parameters: window length w and alphabet size α , and it involves three main steps [85]:
 1. Normalization: standardizes the data in order to have a zero mean and a standard deviation of one;
 2. Piecewise Aggregation Approximation (PAA): divides the original data into the desired number of windows and calculates the average of data falling into each window; and
 3. Symbolization: discretizes the aggregated data using an alphabet set with the size represented as an integer parameter α , where $\alpha > 2$.
- Shapelets are subsequences of time series that identify with the class that the time series belongs to.
- Ensembles. COTE algorithm [86] uses a collective of ensembles of classifiers on different data transformations.

The ensembling approach in COTE is unusual because it adopts a heterogeneous ensemble rather than resampling schemes with weak learners. COTE contains classifiers constructed

in the time, frequency, change (autocorrelations), and shapelet transformation domains (35 in total) combined in alternative ensemble structures. Each classifier is assigned a weight based on the cross validation training accuracy, and new data are classified with a weighted vote.

In the reviewed literature for time series representation we can see that there are some facts that remain unsolved:

- Outliers and noise: when data are coming from sensors and physical devices usually contains noise and outliers that affect the identification of the correct parameters of the distribution.
- Data follows different distribution: some scenarios in which data does not follow a normal distribution, as assumed by some methods in the literature [84], are: radioactive decay (exponential distribution), number of cars passing through a point in a period of time (Poisson distribution), queuing models (gamma distribution), batting averages of baseball players (beta distribution).
- Fast data: two of the V's from the 7V's Big Data challenges [37] are *velocity* and *variety*. Traditionally in data mining, already collected data are processed in an offline manner using historical data. However, in IoT applications, we need to consider short-term snapshots of the data which are collected very quickly. The data are represented as streams and it can change (locally or globally) over time. Thus, we need adaptive methods that catch up with the changes and update the models online during their operation.

Taking the above-mentioned cases into account, we seek an algorithm that does not require normalization of the data. The latter will also help to preserve the value of the data points (i.e. magnitude of the data). The lack of sensitivity to magnitude in the algorithms that make assumptions about the normalized distribution and use Z-normalization makes them less efficient in analysing correlations and relatedness measures. Another requirement is the application of the algorithm in an online way and using sliding windows. Nonetheless, we have to be able to compute the distance between the aggregated time series.

3.1.3 Feature selection

An FS method is a search strategy where the performance of candidate subsets is measured with a given evaluator. A stopping criterion establishes when the FS process must finish. FS methods are typically categorized into wrapper, filter and embedded, univariate and multivariate methods. Wrapper methods [87] use a predetermined learning algorithm to determine the quality of selected features according to an evaluation metric [88]. Filter methods apply statistical measures to evaluate the set of attributes [89, 90]. Embedded methods achieve model fitting and

FS simultaneously [91]. Multivariate methods evaluate features in batches. Univariate methods evaluate each feature independently. Figure 3.1 illustrates graphically the FS flow.

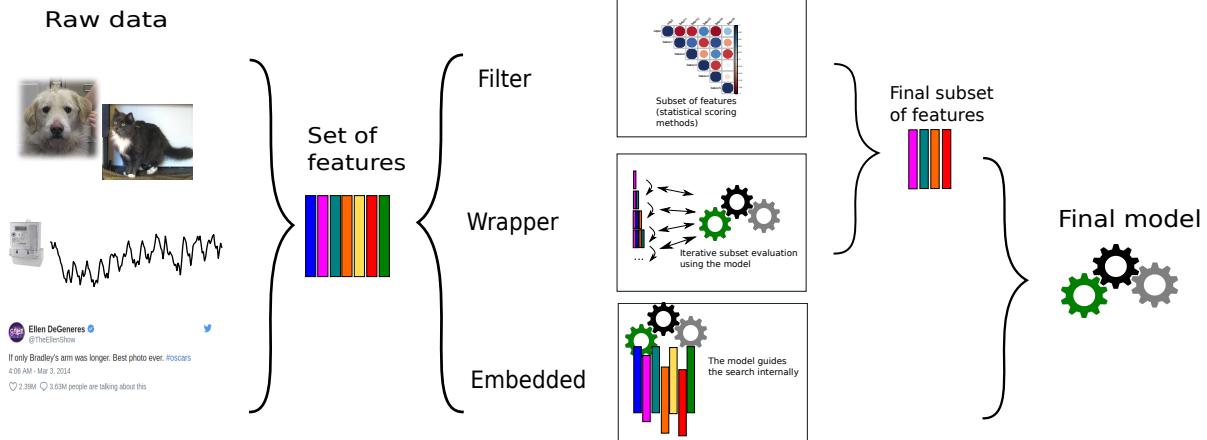


Figure 3.1: The FS flow

We have done an extensive search in order to find academic works that have carried out FS. Together with the works that address FS for energy consumption time series, we have also considered important to review FS for energy consumption when not treated as time series, and FS for time series problems in general, i.e. other approaches not specifically related to energy consumption.

The first paper that studied how the selection of subsets of features associated with building energy behaviours influences a ML model performance for energy consumption prediction used some filter methods for FS and support vector regression for forecasting [92]. A bit later, in the thesis [93], Fast Correlation-Based Filter (FCBF) is used for FS in load prediction error problems in four building areas. A meteorological dataset from several locations and also, the geographical factor are exploited by selecting variables from different locations. The baseline comparisons are done with e-SVR. According to this work, how the relationships between features change with distance motivates a greedy FS method for the electrical load forecasting. In the works [7, 4], correlation and principal components analysis (PCA) are used for FS and transformation.

FS for time series prediction has been carried out using neural networks [94]. By combining contemporaneous and lagged realisations of the independent variables and lagged dependent variables more general models of dynamic regression, autoregressive (AR) transfer functions and intervention models are constructed. It has also been done using the Granger causality discovery [95] to identify important features with effective sliding window sizes, considering the influence of lagged observations of features on the target time series.

The optimal time-windows and time lags for each variable based on feature pre-processing and sparse learning in order to configure the input dataset were searched in [96].

The forecasting of solar radiation time series is enhanced by using a train set of bootstrapped SVM in order to perform FS was done in [97]. They assure that this method is more robust than a

regular FS approach because using the later, small changes on the train set may produce a huge difference on the selected attributes. For solar radiation prediction, [98] masks the inputs as a FS step . They create their own features by defining night, sunrise, day and sunset according to the moment that their instruments perceive those. This provides certain improvements in forecast accuracy. A data-driven multi-model wind prediction methodology using a two-layer ensemble ML technique is developed in [99]. A deep FS framework is employed where four different approaches are used in order to get the input vector: PCA, Granger Causality Test, Autocorrelation and Partial Autocorrelation Analysis, and Recursive Feature Elimination. Another ensembles way of selecting features is presented in [100] and it is used for predicting the number of incoming calls for an emergency call centre in a time series manner. They use five algorithms (ReliefF, PCA, Freq. Discretization, Information Gain and K-means) that are different in nature and combine the rankings computed grouping similar approaches and computing new weights as the mean of the individual weights. After that, all variables that are ranked among the top five positions in at least three of the groups compound the selected features. In the thesis work [101] they present three case studies in which FS is a step in the model creation. They used the following methods: sequential forward/ backward selection (SFS, SBS), sequential forward/ backward floating selection (SFFS, SBFS), the n best features selection (nBest) and the best individual features.

The main data characteristics of energy time series have been specifically analysed in [102]. To explore such data from different perspectives they consider two main categories: nature (nonstationarity, nonlinearity and complexity characteristics) and pattern (cyclicity, mutability or saltation, and randomicity or noise pattern). After that, FS for electricity load forecasting was done in a time series manner using correlation and instance based methods [103]. In [104] it is presented a survey on data mining techniques for time series forecasting of electricity. The survey focuses on the characteristics of the models and their configuration. Wrapper methods, ANNs, mutual information, autocorrelation and ranking based methods are mentioned as FS techniques used in the prediction of energy consumption. Finally, the work [9] uses temperature time series together with the day of the week in order to estimate energy consumption.

To conclude, in several studies for energy consumption forecasting FS is not discussed and when it is, it does not account for the temporal characteristics of the data and it is not carried out systematically looking for the best method. Regarding the papers that focus on FS for time series prediction [94, 95], we highlight that the focus of [94] is narrowed to neural networks which is not the best for every situation since usually, neural networks are more computationally expensive and require much more data than traditional algorithms. Also, the *No Free Lunch* theorem [105] suggests that there can be no single method which perform bests on all datasets. [95] is focused on the Granger causality as FS so none of them provides a systematic comparison between the possibilities available in the FS field. There is no paper that carries out a systematic combination of univariate, multivariate, filter and wrapper methods and also checks their performance using

several predictive algorithms.

There are two main objectives that have not been considered at the same time, that is, minimise the forecasting error and also the number of variables to be used. This is an important gap that is fulfilled in this work.

3.1.4 HVAC usage patterns

Buildings occupants undertake adaptive actions in indoor environments either to change or adapt to it [106]. Within buildings, it has been seen that occupants have a substantial impact on the energy consumption [107], and it is for that reason that several studies have been carried out to understand the behaviour of building occupants [108, 109]. Also to try to reduce energy use via the change in occupant's habits [110]. Some adaptive actions that occupants can carry out for adjusting their thermal conditions are opening/closing windows and doors [106], adjusting thermostat set points [111], too hot/too cold complaint calls [112], and repositioning their window shades [113]. Research on studying these adaptive occupant behaviours has been focussing on understanding their impact on the energy performance of a building [114]. It has been shown by the literature that this change in habits can result in some cases to 20% savings [115, 116]. Given that the 50% of EU's final energy consumption is used for heating and cooling, of which 80% is used in buildings [117] HVAC are a crucial subject of study if we aim to reduce energy consumption.

The majority of studies that involve HVAC information focus on identifying anomalous behaviours in order to detect faulty equipment [118, 119]. However, the thermostat control is the main behaviour to regulate the thermal comfort, that the most energy-consuming aspect of a building. A survey of 1134 homes in England found a wide variation in thermostat settings which, in the interest of energy efficiency and sustainable development, could form the foundation of a "social norm" campaign aimed at reducing temperatures and energy use in "overheated" homes [120]. Also, the heating set point was related to outside air temperature, relative humidity and wind speed using data from 13 Danish dwellings [121].

Little research has been found in the extraction of thermostat behaviour. Two of the most relevant studies [120, 121] only consider set point for heating, leaving unstudied the cooling part of the thermostatic control. Also, the two previous studies focus on dwellings leaving unstudied non-domestic buildings. In addition to that, the study of the use of the thermostat in time seems to be under looked. This could be the result of not being able to capture such data.

In the following, we propose to associate the thermostat behaviour with energy waste.

3.1.5 Human mobility patterns

In recent years, various works have considered the processing of spatiotemporal traces for mining information about how people move [122]. These digital breadcrumbs can be collected from several sources like motion sensors [123] or smart cards [124]. GPS traces are one of the most

popular data sources in this field and social network data has been used for unravelling the goals of this movement [125]. Trajectory pattern mining examples are frequent item mining, trajectory clustering, and graph-based trajectory mining [126]. More novel approaches use historical routes to generate probabilistic models. Examples of these novel approaches are listed below.

Location Prediction

Location prediction is based on the assumption that people follow daily routines and, thus, have only a set of frequently visited locations [127]. This makes regular trips quite predictable due to their high level of repetition [128].

There are two main trends for personal location prediction: (i) using a geometrical approach so as to predict a path in the Euclidean space [129] by applying a mathematical function to the current location and velocity of the target person and (ii) pattern matching solutions that compare the route in progress with a set of mobility models. In that sense, Bayesian networks [130] and hidden Markov models [131] have been some of the applied solutions given high-resolution mobility datasets (e.g., those comprising GPS-based traces).

Regarding social media datasets, the probabilistic model W^4 [132] follows a Bayesian-network approach that forecasts the next location and activity of a user by also taking into account temporal factors. Mobility features selection is also studied [133].

Social Media for Human Mobility

It is possible to classify social media for human mobility under three different categories.

Firstly, processing geotagged tweets in order to create spatial regions depending on their usage (e.g., leisure, home, and work) using visual analytics [134], clustering algorithms [135] or classifiers [136]. Secondly, automatically extracting events (e.g., live shows, earthquakes [137], traffic jams or anomalies, etc.). For instance, [138, 139] creates smart social agendas that can be updated in real time. In a road-traffic monitoring scenario, several works make use of social media data in order to either detect or semantically enrich traffic anomalies. Correlating tweets [140] and using official traffic-management institutions' Twitter accounts [141] in order to detect road-traffic incidents are some examples. The third category includes social media as a new data source for detecting the movement of different kinds of people among places such as dynamic labelling by semantical enrichment of spatiotemporal trajectories [142] and the statement of a worldwide mobility report based geotagged Twitter data [143].

Mobile Crowd Sensing (MCS) Based Mobility Mining Solutions

An important line of research makes use of MCS for mapping activities by collecting the spatiotemporal traces of contributors [144, 145] by composing collaborative maps comprising road networks, bike and hike routes.

Also, works on real time traffic monitoring propose the distributed architectures to keep track or predict road traffic congestions within an area by means of the mobility reports generated by the on-board units of vehicles [146, 147]. As a result, now we can find solutions that combine static and vehicle-mounted and smartphone sensors to detect road traffic in an area [148, 149].

Regions of Interest (ROIs) Detection

Given a collection of spatiotemporal traces, different types of clustering algorithms can be applied to the spatiotemporal traces to uncover the target ROIs. Density-based clustering has been a prominent solution [150]. Distance between GPS points or density connectivity in a two-dimensional Cartesian space are used as features. A different approach makes use of frequency map based spatial-temporal clustering methods [151].

Complex Event Processing (CEP)

CEP is an evolution of the former publish/subscribe model that deals with more complex subscription patterns, so it can be considered a recent technology [152]. Despite CEP's widespread usage, there exists a scarcity of CEP solutions that use spatiotemporal data since only a few works actually propose practical CEP applications [153]. The GPS-based solution [154] formulates a framework to timely detect spatiotemporal relationships between moving entities.

The reviewed studies frequently suffer from some of the following limitations:

- Reliance on GPS traces datasets even though GPS feed is one of the most battery-draining sensors of a mobile device.
- Either they centre on extracting general mobility information related to a particular urban area or use models for every single user as W^4 [132]. In that sense, detecting personal mobility models that also count with the crowd-dynamics could be of interest to come up with personalised but also informed location services.
- The whole available dataset, which sometimes might not be available, is required for preprocessing.
- Moreover, most works using social networks do not fully unlock its potential since they only use the spatiotemporal aspects of the data (check-in posts) but do not consider textual details.
- The anticipation of peoples' activities and locations using geotagged social media documents is scarce.
- The mobility knowledge extracted by the aforementioned solutions focuses on road traffic features so it depends on the road network topology. We plan to capture human dynamics from a wider perspective.
- Only theoretical solutions to define formal event-based information models and architectures for social media processing have been put forward using CEP earlier [155].
- Detecting ROIs related to the movement of people instead of where people tend to remain stopped is missing in the literature.

3.1.6 IoT architectures and projects for smart cities and energy management

Due to the importance of the building sector in energy consumption, it becomes a foremost task to achieve meaningful energy savings that will reduce this energy use in reality. Despite the fact that IoT technologies have been widely used for the realization of the smart building concept, the simple sensorization of buildings is not enough to make a housing stock that consumes fewer energy resources a reality. IoT is also required to properly process, manage and, above all, analyse the energy-related data that would help to develop final energy-aware services targeting the energy efficiency goal.

An overview of both the management of energy data and the implementation of IoT platforms is put forward.

During the last years, some initiatives within the cloud computing domain have been made to intelligently manage energy data of buildings. In that sense, Big Data energy management models have been created ranging from the collection and preprocessing of data to its further analysis and the final exposition to services [156].

From a practical perspective, the Dynamic Demand Response platform [157] makes use of public and private clouds combined with infrastructure and platform as a service for data storage. This platform was extended with Cryptonite, a repository to store sensitive Smart Grid data [158]. Then, different classes of data-driven forecasting models were generated on top of the whole platform with the purpose of carrying out energy prediction among others.

ElasticStream also provides a prototype solution for energy data management and analysis. In this case, the mechanism transfers energy data to a cloud platform for further analysis on the basis of rate changes in the input data streams [159].

The MultiAgent System (MAS) named SAVES (Sustainable multiAgent systems for optimizing Variable objectives including Energy and Satisfaction) defined in [160] is used in [161] regarding actual occupant preferences and schedules, actual energy consumption and loss data measured from a real testbed building at the University of Southern California in order to predict energy consumption at different levels (frequency of prediction and device aggregation). Other works provide energy data management solutions without focusing on analytic aspects. This is the case of the Virtual SCADA architecture for cloud computing (VS-Cloud) that encompasses Cloud Computing for energy data storage [162]. VS-Cloud mainly focuses on the orchestration of components in Smart Grids and the save storage of sensitive data executed actions, incidents or alarms. Therefore, its domain of application is more related to risk management. Similarly, the work in [163] proposes an automation platform for energy monitoring. However, such a platform does not provide any particular feature to support energy data analytics as it focuses more on the definition of control strategies for energy saving.

When it comes to the development of IoT solutions, most of them are just vertical silos that do not support interoperability and with inappropriate models of governance. For that reason,

some architectures and platforms have been developed to lower those barriers.

IoT-A, the IoT Architecture (EU project from 2009 to 2012)¹ defines an Architecture Reference Model (ARM) which ensures the interoperability also scalability requirements and the security and privacy in its design, which are so often neglected. This solution rests upon the creation of an architecture reference model together with an initial set of key building blocks, principles and instructions for the technical design of the protocols, interfaces and algorithms suitable for any IoT system.

Webinos² creates an Open Source Platform and software components for the Future Internet in the form of web runtime extensions, to enable web applications and services to be used and shared over a large amount of connected devices in a consistent and secure way.

Buttler³ platform is a set of enablers and services that supply means for building context-aware applications on smart things. It provides generic APIs to access resources provided by IoT devices and other services such as security, localization, behaviour prediction and context management. Those services enhance the user experience and security. The BUTLER Platform is oriented for IoT devices and applications and provides homogeneous access to the underlying networks.

FI-LAB⁴ conforms live instances of FIWARE⁵ architecture and generic enablers, for free experimentation. FIWARE is an open Core Platform of the Future Internet, introducing an innovative infrastructure for cost-effective creation and distribution of digital services, providing security guarantees. FI-LAB forms a meeting point between sponsors and application developers.

Nowadays, several analytics software packages are available for buildings. SkySpark⁶ is well-known and it mainly runs personal rules that depend on the data collected in a building and identifies non-obvious operational problems. The ability to use artificial intelligence, instead of writing custom programming, to extract knowledge in operational data should be exploited.

Some European projects dedicated to combining IoT platforms and energy management are:

- SINFONIA⁷ created a set of measures that include optimisation of the electricity grid and solutions for district heating and cooling in order to set up a large-scale, integrated and scalable energy solutions in mid-sized European cities. SINFONIA allowed the cooperation between cities that belong to the same climate zone with the goal of reducing energy needs to meet people's activity requirements, and the consequent CO₂ emissions in order to guarantee a reliable and progressive transition toward low carbon cities.
- CityPulse⁸ provided reliable knowledge extraction techniques that were used to create new

¹<https://cordis.europa.eu/project/rcn/95713/factsheet/es>

²<https://cordis.europa.eu/project/rcn/95713/factsheet/es>

³<https://cordis.europa.eu/project/rcn/101349/factsheet/en>

⁴<https://account.lab.fiware.org/>

⁵<https://www.fiware.org/developers/>

⁶<https://skyfoundry.com/>

⁷<https://cordis.europa.eu/project/rcn/197825/factsheet/en>

⁸<https://cordis.europa.eu/project/rcn/109806/factsheet/en>

smart city applications. This was done by developing, building and testing a framework for real-time IoT stream processing and large-scale data analytics.

- e-balance⁹ goal was the integraton of customers into the smart-grids in order to tackle present and future environmental problems with ICT based solutions, new business models and citizens' behaviour in real world conditions. It focused on investigating the economic and social aspects of the energy efficiency, integrating ICT for decentralized power management providing more autonomy for delivering local decisions including intelligent power generation, consumption.

As above mentioned, several multi-purpose IoT platforms already provide generic solutions to manage IoT data. However, there is a lack of platforms in this field focusing on (1) the household energy domain and (2) providing support for data analytics.

Some energy models only provide a theoretical approach [156]. Also, the aforementioned initiatives do not constitute holistic energy data management and analysis solutions. The platforms do not include explicit features that are necessary for all energy monitoring scenarios like data volatility monitoring and outliers detection to ease the deployment of data mining algorithms and other services over of the stored data.

Another neglected feature by existing IoT platforms is the support of built-in data mining features able to generate new useful knowledge from the collected and stored data [164]. In real IoT deployments, this processing and analysis task has been frequently done by third-party services. However, integrating certain data mining functionalities as built-in features of platforms would provide a great benefit in a wide range of domains, for example quick statistics, easy to generate digests or sanity checks. In that sense, only a few IoT platforms actually include native data analytics features. As a matter of fact, SensorCloud¹⁰ enables a simple interface for common operations like smoothing, filtering and interpolation whereas GroveStreams¹¹ provides some real-time data analytics mechanisms. However, none of them supports sensor heterogeneity nor follows an open source approach.

3.1.7 IoT architectures and projects for behavioural analysis towards energy efficiency

In order to realize a sustainable energy transition, human behaviour should be considered regarding educational aspects (raising awareness of the benefits) and technologies understanding [165]. Whilst new technologies and materials are available, certain initiatives are required in order to encourage individuals to use them properly.

⁹<https://cordis.europa.eu/project/rcn/109806/factsheet/en>

¹⁰<https://sensorcloud.com/>

¹¹<https://grovestreams.com/>

There is a scarcity of academic works that evaluate how behaviour-change interventions affect energy efficiency. [166]. Organisations have initiated interventions based on holistic programmes including gamification and rewardings for real-live changes [167, 168], comparative feedback and competitions [169], various feedback types toward energy savings [170], prompts, peer-education and dashboards [171].

Several studies try to find how and why the installation of smart metres changed households' electricity use [172]

Some programmes combine these tactics with installing new technological features such as a system that determines activity in homes by integrating motion, door, lighting and temperature sensors which allows adapting the services depending on the behaviour patterns [173]. Also, modelling residents' behaviour by studying presence, lighting and window status sensors can decrease energy consumption by up to one third [174].

There are many European projects that are targeting the integration of technological advances for incentivising behavioural change towards energy efficiency. Some of them are:

ChArGED¹² (CleAnweb Gamified Energy Disaggregation) created a framework that encourages the achievement of energy efficiency and the reduction of wasted energy in public buildings. The framework set up low-cost devices to improve energy disaggregation at several level. Energy waste is targeted by a gamified application that feeds personalized real-time recommendations to individuals. The game is designed in a way that helps users to understand how their actions affect the environment.

enCOMPASS¹³ (Collaborative Recommendations and Adaptive Control for Personalised Energy Saving) developed and validated several digital tools that make energy consumption data accessible and understandable for all stakeholders: from residents to building managers and ICT-providers. Those tools were integrated and provide visualisation of energy-related data, energy-saving recommendations that depend on the context, intelligent control and adaptive gamified incentives.

TRIBE¹⁴ (TRaIning Behaviours towards Energy efficiency: Play it!) is based on the development of a serious game implemented through social networks (for information and experience exchange) for the energy sector in which building users adopt energy efficient attitudes. It includes a simulation engine and real time collection data from the ICTs installed in the pilots, enabling a dynamic interaction building-consumer and moving towards a change in players' behaviour. In addition, some tools and guidelines were set up to be used by users and owners of public buildings, including: (1) an initial energy audit and diagnosis, (2) the creation of a virtual pilot similar to the real buildings, (3) an energy efficiency deployment plan based on ICT, (4) a funding plan (5) a user engagement campaign for the detected behaviour change challenges.

¹²<http://www.charged-project.eu/>

¹³<http://www.encompass-project.eu/>

¹⁴<http://tribe-h2020.eu/>

As can be seen in the literature, there has not been a work that combines the methods, analytics and properties developed for designing an ecosystem that targets at improving energy efficiency through consumers' understanding, engagement and behavioural changes.

Regarding data, they focus either on survey data [172] or on monitored data. The combination of this two ways of data acquisition is not carried out. However, surveys can be of great help at the beginning of an energy efficiency campaign, when data is not available in order to make a profile of the users that can later on be transformed due to the sensed data and performance evaluation of users.

Regarding methodology, they focus either on the competition factor [169] or in the provision of feedback depending on the specific users' internal values [26], which is not updated or verified regarding its behaviour towards the tools.

Within the little number of available platforms, they do not include the possibility of incorporating algorithms. This is a very limiting characteristic since keeps the platforms isolated from the fast evolution of analytic advances.

A unifying platform that gathers all those characteristics is needed in order to achieve good results regarding energy literacy and consequently efficient behaviour towards energy.

3.1.8 Related work summary

Every subsection in this related work ends with a reflection on the missing parts detected on the reviewed literature of the subject. Briefly this includes the application of machine learning techniques for energy consumption forecasting using external predictors, the comparison between black-box and grey-box models and the lack of a complete methodology for multivariate feature selection for energy consumption prediction. Regarding pattern extraction, we found that the set point of HVACs are not fully exploited and human mobility pattern extraction can be improved using social networks data. The collected data for these scenarios is now huge and its volume implies low performance. We found that methods for data reduction do not take into account IoT and Big Data characteristics: variability, volume, velocity, etc. Looking at the platforms, there are no specific ones that address energy management and provide support for data analytics at the same time. The majority of the attempts to do so do not take into account the human behaviour factor, that is also key for the achievement of realistic energy efficiency.

In the following, we expose how our work tackles each of the previously stated gaps, ordered by result as presented in Table 2.1.

3.2 Data analysis in IoT based Smart Environments

In urban environments, there is a wide variety of data sources. A wealth of sensors are distributed around cities, in both indoor and outdoor spaces. This situation has brought new analytics mechanisms and tools that provide insight into the data which allows building powerful systems

and applications in an efficient and collaborative way [175]. Mobile sensors are playing a major role in the development of applications too since they are embedded in our lives as smartphones, smart cards, wearable technology and, in the case of vehicles, on-board sensors. Urban dynamic patterns can be detected and studied thanks to the information that these sensors provide. In the following, we summarise the results of the dissertation. Every subsection refers to a result that is presented in Table 2.1. In the same table, we can see the publications' references to each of the results.

3.2.1 Smart buildings data integration and statistical analysis [R1]

In order to create ML models that forecast energy consumption, it is crucial to understand the relationship between energy consumption and other attributes. Possible influencers or input attributes for energy consumption prediction are weather variables and occupation patterns that refer to the day of the week, the hour of the day, the kind of day, etc. In order to do so, the integration of 3 smart buildings datasets from 2 different buildings formed by collected data from several sources is carried out. Table 3.1 contains a summary of the buildings' information. The first dataset belongs to the TTC Fuente Alamo¹⁵ of the University of Murcia (see Fig. 3.2 left). This data consist of the environmental outdoor observations and the total energy consumption of the building from 1st December, 2014 to 18th February, 2016 in intervals of 8 hours and the origin of the consumption (HVAC, lighting or other electrical equipment) is unknown. In total, 952 observations and 15 variables. Outdoor environmental measures are acquired from The Research Institute of Agriculture and Food Development of Murcia (IMIDA)¹⁶ that provides real time records of weather from several stations across the region of Murcia. The following variables are included in the dataset: temperature (mean, min and max) (°C), humidity (mean, min and max) (%), radiation (mean and max) (w/m^2), wind speed (mean and max) (m/s^2), wind direction (mean) (degrees), precipitation (mm), dew point (°C) and vapour pressure deficit (kPa). Station CA91 with latitude 37.699033 and longitude 1.238044.

The second dataset belongs to the Faculty of Chemistry at the University of Murcia (see Fig. 3.2 right) and it is composed of 5088 observations of 50 attributes that are measured hourly from 2016-02-02 until 2016-09-06. The output attribute is the energy consumption measured in KWh and we have included meteorological measurements from 2 different meteorological stations from IMIDA that are close to the building (MO12 lat:38.007031 long: 1.302564 and MU62 lat: 37.940067 long: 1.134719), hour ahead predictions and other weather related variables provided by Weather Underground¹⁷ web service and also season, day of the week and holiday attributes are included.

Weather Underground is a web service that through its API provides the following real

¹⁵<https://www.um.es/web/otri/contenido/ctt>

¹⁶<http://www.imida.es/>

¹⁷<https://www.wunderground.com/>

Name & Country	Envelope area (m^2)	Orientation	Coordinates	Cons. year
TTC Fuente Alamo. Spain	3323	South-West	Latitude: 37.724383 Longitude: 1.093324	2004
Chemistry faculty. Spain	1500	South-West	Latitude: 38.020939 Longitude: -1.169722	1944

Table 3.1: Information about the buildings

values: temperature ($^{\circ}$ C), apparent temperature ($^{\circ}$ C), dew point ($^{\circ}$ C), humidity (%), wind speed (m/s), mean sea level pressure (mbar), visibility (km) and precipitations in last hour (mm). We also use *one-hour predictions* for the first six previous attributes, together with *probability of precipitations (%)*, *sky cover (%)* and *wind direction* (degrees).

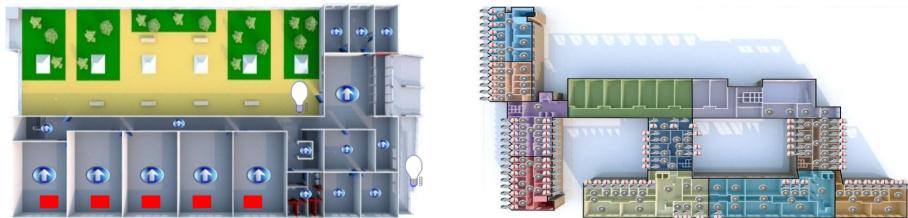


Figure 3.2: 1st floor of the TTC where red labels means energy meter (left) and 2nd floor of the Chemistry Faculty (right)

Finally, we have also monitored the use of HVAC systems in 237 Chemistry Faculty rooms. The dataset consist of 12-minutes aggregated observations regarding room temperature, on/off status and set point from 2015-10-31 until 2017-02-28.

Fig. 3.3 shows the pairwise correlations between all variables involved in the energy consumption forecasting problem. A blue circle indicates positive correlation and a red circle indicates negative correlation. The radius of the circle indicates the magnitude of the correlation. In TTC (left) and focusing on the first row, we see that energy consumption correlates significantly ($\alpha = 0.05$) and positively (blue circle) with temperature, radiation, wind speed variables, vapour pressure deficit and dew point, and negatively (red circle) with wind direction and humidity variables. This means that we can use safely these variables as inputs of the energy consumption model of our reference building, because they all have a clear impact on the energy consumption except precipitations (crossed out because they are not significant) [7]. Since we collected more attributes for the Chemistry faculty dataset at the right side of Fig. 3.3 we have shown only the relationship between energy consumption and the rest of the variables. In this case, we should exclude from the analysis all variables related to precipitations, dewpoint and sky visibility.

In this preprocessing stage, the study of correlations can be complemented con PCA analysis. PCA is a method that reduces the dimensionality of the data by creating new uncorrelated variables that successively maximize variance. Those new variables are found by solving an

3.2. DATA ANALYSIS IN IOT BASED SMART ENVIRONMENTS

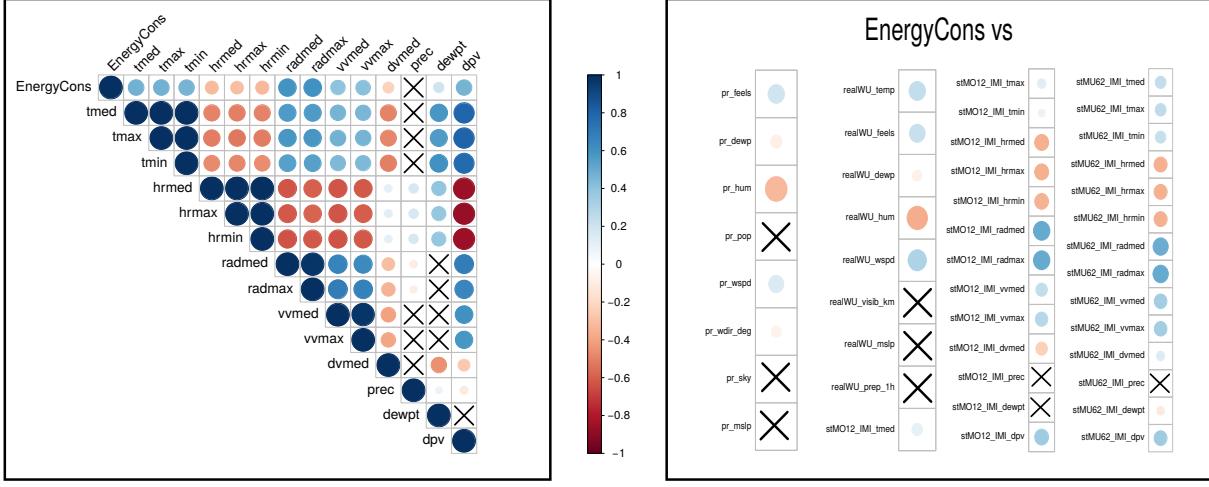


Figure 3.3: Correlation heatmap between consumption and outdoor environmental conditions for both consumption datasets

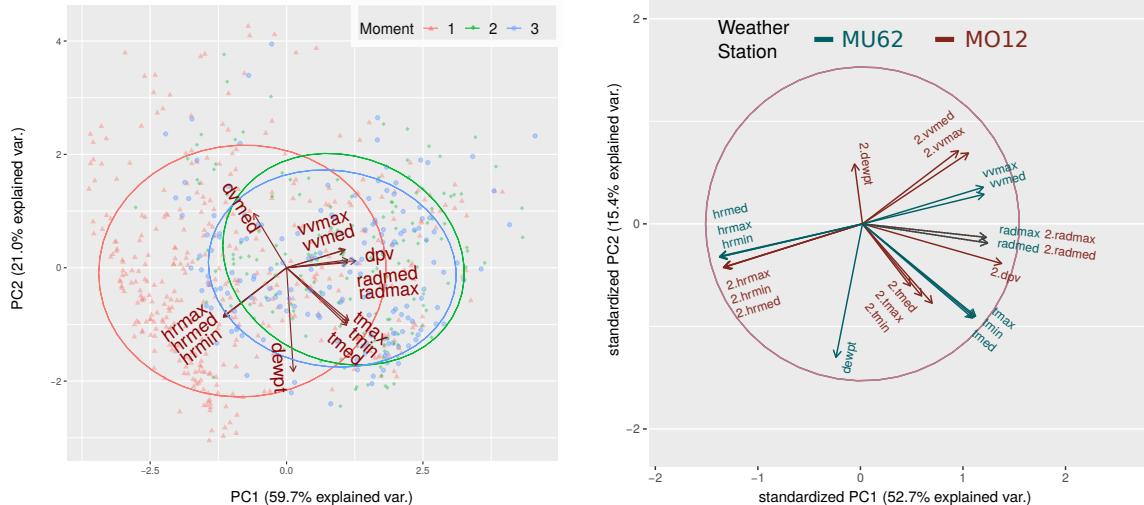


Figure 3.4: Correlation heatmap between consumption and outdoor environmental conditions for both consumption datasets

eigenvalue/eigenvector problem [176]. It is especially important to apply it when inputs might be correlated, as it is our case since dependent variables can hinder ML models' accuracy.

PCA provides intuitive visualizations for exploring relationships in data since subjects are projected into the new dimensions, which are linear combinations of the initial features. In both cases of Fig. 3.4 around 60 % of the variance can be explained with 2 dimensions. The correlation circle information gives a similar result than Fig. 3.3 and also the first dimension in both cases mainly reflects radiation. There are certain parallelisms between PCAs of both datasets, however, something to be noted at the right images is that dew points from different sources are negatively correlated and there exist differences between both sources for wind velocity.

3.2.2 Data representation [R2]

Due to the huge volumes of data that are provided in smart environments, it is of special interest to investigate methods for data reduction. We attempt to aggregate and represent large volumes of data in an efficient and higher-granularity form in order to create sequences of patterns and data segments that occur in large-scale IoT data streams. The contribution of our approach is to do such representation on-the-fly since usually data treatment has to be done very quickly, adapting to unpredictable changes in the data or even without prior knowledge.

Our proposed method is based on splitting time series data into blocks. These blocks can be either overlapping or non-overlapping and they represent subsets of the whole data structure. The method synthesizes independently the information that the blocks contain. It reduces the data points while still preserving their fundamental characteristics (losing as little information as possible). We propose a novel technique using matrix-based data aggregation, DCT and eigenvalues characterization of the time series data. The algorithm is called Blocks of Eigenvalues Algorithm for Time series Segmentation (BEATS).

Transforms, in particular, integral transforms are used to reduce the complexity in mathematical problems. In order to decorrelate the time features and reveal the hidden structure of the time series, they are transformed from the time domain into other domains.

DCT uses cosines obtained from the discretization of the kernel of the Fourier Transform. DCT transfers the series to the frequency domain. Among the four different cosine transformations classified by Wang [177], the second one (i.e. DCT-II) is regarded as one of the best tools in digital signal processing [178] (time series can be regarded as a particular case of signals). Due to its mathematical properties such as unitarity, scaling in time, shift in time, the difference property, and the convolution property, DCT-II is asymptotically equivalent to the KLT where under certain (and general) conditions KLT is an optimal but impractical tool to represent a given random function in the mean square error (MSE) sense. KLT is said to be an optimal transform because:

- It completely decorrelates the signal in the transform domain;
- It minimizes the MSE in bandwidth reduction or data compression;
- It contains the most variance (energy) in the fewest number of transform coefficients; and
- It minimizes the total representation entropy of the sequence.

The details of the proof of the above statements can be found in [178]. Understanding the properties of the DCT, we use it to transform our time series data.

We apply the transformation essentially by using the compression of a stream of square 8x8 blocks, taking reference from the standards in image compression [179] where DCT is widely used (e.g. JPEG). Since 8 is a power of 2, it will ease the performance of the algorithm.

As an illustration, we provide an example in Fig. 3.5. We have divided the time series as blocks of 64 observations that are shown using a dashed red line. If we arrange the first block

into a squared matrix M , we can visualize that the information is spread through the matrix as a heatmap. Intuitively, each 8×8 block includes 64 observations of a discrete signal which is a function of a two-dimensional (2D) space. The DCT decomposes this signal into 64 orthogonal basis signals. Each DCT coefficient contains one of the 64 unique *spatial frequencies* which comprise the *spectrum* of the input series. The DCT coefficient values can be regarded as the relative amount of the spatial frequencies contained in the 64 observations [179].

Let M be the 8×8 input matrix. Then, the transformed matrix is computed as $D = UMU^T$, where U is an 8×8 DCT matrix. U coefficients for the $n \times n$ case are computed as shown in Eq. 3.1:

$$(3.1) \quad U_{ij} = \begin{cases} \frac{\sqrt{2}}{2} & i, j = 1 \\ \cos\left(\frac{\pi}{n}(i-1)(j-\frac{1}{2})\right) & i, j > 1 \end{cases}$$

After applying DCT, the information is accumulated in its upper-left part.

Each of the 64 entries of the matrix D is quantized by point-wise division of the matrices D and Z , where the elements of the quantization matrix Z are integer values ranging from 1 to 255.

Quantization is the process of reducing the number of bits needed to store an integer value by reducing the precision of the integer. Given a matrix of DCT coefficients, we can divide them by their corresponding quantizer step size and round it up depending on its magnitude, normally 2 decimals. If the maximum of the DCT matrix is small, the number of decimals is selected by the operation $\lfloor \log_{10} \max \rfloor - 4$, where $\lfloor \log_{10} \max \rfloor$ returns the position of the first significant figure of the maximum number in the transformed matrix D . This step is used to remove the high frequencies or to discard information which is not very significant in large-scale observations.

The selected matrix Z is the standard quantization matrix for DCT [180]. After the quantization process, a large number of zeroes appears in the bottom-right position of the matrix $Q = \frac{D}{Z}$, i.e. it is a sparse matrix.

We extract the 4×4 upper-left matrix that contains the information of our 64 raw data and compute the eigenvalues.

Using BEATS so far we have significantly reduced the number of points of our time series from 64 to 4 but we have also converted its components into complex numbers. These complex numbers (eigenvalues vector) represent the original block in a lower dimension. This eigenvalues vector is used in BEATS to represent the segments and hence, it is the potential input for the ML models. However, it is not always possible to feed ML algorithms with complex numbers and the eigenvalues could be complex numbers. To solve this problem, we compute the modulus of the eigenvalues and remove the repeated ones (they are presented in pairs so the information would be repeated).

In case that there are no complex numbers in the output of BEATS, we will conserve the first three values, since the latter values are sorted in descending order. This means that we have represented the original 64 observations as three values. In our example, the final representation (modulus of the eigenvalues) consists of 0.1860, 0.0246, 0.0085.

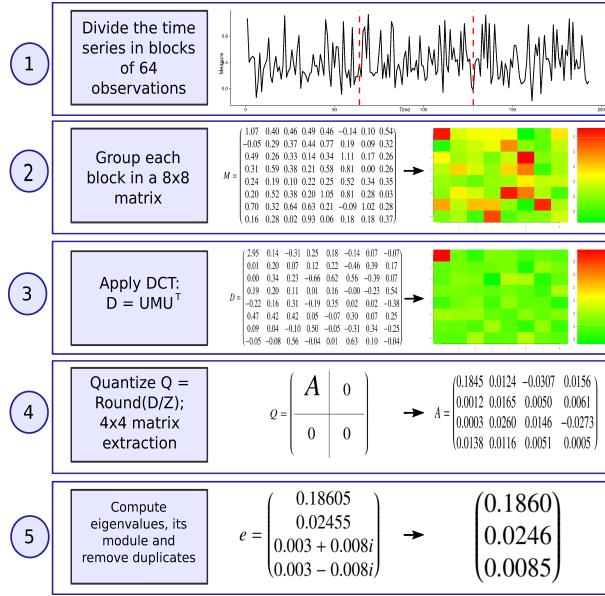


Figure 3.5: BEATS is shown step by step with an example

The BEATS process is summarized in Fig. 3.5.

Big Data BEATS implementation

In contrast to the traditional analysis procedure where data are first stored and then processed in order to deploy models, the major potential of the data generated by IoT is accomplished by the realisation of continuous analytics that allows making decisions in real time.

There are three types of data processing: Batch Processing, Stream Processing and Hybrid Processing.

Batch processing operates over a group of transactions collected over a period of time and reports results only when all computations are done, whereas stream processing produces incremental results as soon as they are ready [181].

Regarding the available Big Data Tools, we have considered Hadoop and Spark Big Data frameworks. Hadoop was designed for batch processing. All data are loaded into HDFS and then MapReduce starts a batch job to process that data. If the data changes the job needs to be run again. It is step by step processing that can be paused or interrupted, but not changed.

Apache Spark allows performing analytical tasks on distributed computing clusters. Spark's real-time data processing capability provides substantial lead over Hadoop's MapReduce and it is essential for online time series segmentation and representation.

The Spark abstraction for a continuous stream of data are called a Discretized Stream or DStream . A DStream is a micro-batch of Resilient Distributed Datasets, RDDs. That means, a DStream is represented as a sequence of RDDs. RDDs are distributed collections that can be operated in parallel by arbitrary functions and by transformations over a sliding window of data (windowed computations).

BEATS adapted to Spark technology

For the online implementation of BEATS we have decided to use pyspark, the Spark Python API that exposes the Spark programming model to Python.

There are many works proposing online time series processing but few of them that have implemented it. In [182] is highlighted that MapReduce is not the appropriate technology for rolling window time series prediction and proposes an index pool data structure.

Pyspark allows us to use the Spark Streaming functionalities that are needed in order to implement BEATS online. BEATS algorithm can be separately applied to windows of the data. Therefore we associate the data received within one window to one RDD, that can be processed in a parallel way.

A suitable type of RDDs for our implementation is key/value pairs. In detail, the key is an identifier of the time series (e.g., sensor name) and the value is the sequence of values of our time series that fall in the window. That way the blocks are exposed to operations that give the possibility to act on each key in parallel or regroup data across the network.

The transformations that we use are:

- Window: used for creating sliding window of time over the incoming data.
- GroupByKey: grouping the incoming values of the sliding window by key (for example, same sensor data).
- Map: The Map function applied in parallel to every pair (key, value), where the key is the time series, values are a vector and the function depends on what has to be done.

3.2.3 Energy consumption prediction [R3]

Energy consumption prediction is a complex task that depends on the context for which the forecast is needed. We have studied several short-term horizons: hourly, 3 moments a day, daily and weekly. For each of them there exists the possibility of predict instances as independent subjects adding temporal characteristics as extra variables (day of the week, hour of the day, etc.) or use ARIMA and temporal models. Also, the possibility to use real or predicted inputs depends on the application:

- Campaign evaluation: it is possible to use real weather and occupation information, but close prior consumption information would bias the experiment. If a model depends on the consumption of the previous day, on the second day of the campaign this model will be biased because previous day consumption is altered by the campaign measures.
- Resources optimisation: it can use prior energy consumption, but no real information regarding the future so occupation, weather and other inputs must be also estimated.

3.2.3.1 The models

Support Vector Regressor (SVR) works in a similar fashion to Support Vector Machine (SVM). Whereas SVM is a classification technique, SVR fits the optimal curve out of which the training data do not deviate more than a small number ϵ . More specifically, during classification the samples that are close to the margin are penalized even if they are correctly classified, whereas in the regression method an acceptable deviation margin of the samples from the prediction curve is set. The free hyperparameter of this model is C , the penalty parameter of the error term. C is the weight of how much the samples inside the margin contribute to the overall error.

Random Forest (RF) is an ensemble learning method in which a group of weak models are combined to form a more powerful model. In RF, multiple regression/classification trees are grown from random with replacement samples. Each tree provides its own prediction and all results are averaged. For each node, m_{try} (hyperparameter) variables are selected at random out of the number of inputs. The best split in these m_{try} is then used to split the node [183].

eXtreme Gradient Boosting (XGB) is built on the principles of gradient boosting and is designed for speed and performance (extreme). XGB generates a prediction by means of an ensemble of weak prediction models that, in our case, are decision trees. The concept is to sequentially build the model by fitting a weak prediction model on the weighted training data set, in which the higher weights are assigned to samples that were previously difficult to predict. The free hyperparameters that are adjusted in this model are the maximum depth limit of the number of nodes in the tree, the minimum number of samples required to split an internal node and the learning rate by which the contribution of each tree is shrunk.

The **Artifical Neural Networks** (ANN) here used are Multilayer Perceptron (MLP) and Bayesian Regularized Neural Networks (BRNN). ANN consist of three layers: input (contains independent variables), hidden, and output layers. The hidden layer contains activation functions and it calculates the weights of the variables in order to explore the effects of predictors. In the output layer, results are presented with an estimation error. The error values are then propagated backwards, starting from the output, until each neuron has an associated error value which roughly represents its contribution to the original output. However, because of the big number of connections, overfitting may occur. Regularization techniques with the backpropagation training are used in order to have a smoother network that is less likely to overfit BRNNs are more robust than standard back-propagation nets.

GAUSS with RBF Kernel (GAUSS). A Gaussian process is a random process where any point x is assigned a random variable $f(x)$ and where the joint distribution of a finite number of these variables is itself Gaussian, that is: $p(f|X) = \mathcal{N}(f|\mu, \mathcal{K})$, where \mathcal{K} is a positive definite kernel function. The kernel function returns a measure of similarity between two points that also encodes how similar its realisations should be. If points x_i and x_j are considered to be similar by the kernel the function values at these points, $f(x_i)$ and $f(x_j)$, can be expected to be similar too. Here, we have used the squared exponential kernel, also known as RBF kernel:

$k(x_i, x_j) = \sigma_f^2 \exp(-\frac{1}{2l^2}(x_i - x_j)^T(x_i - x_j))$. The length parameter l controls the smoothness of the function and σ_f the vertical variation. [184]

Prophet is a modular regression model with interpretable parameters that can be adjusted with domain knowledge about the time series [185]. Prophet conducts an automatic optimisation procedure for forecasting time-series data by fitting a non-linear additive model and generates uncertainty intervals. Three main model components compose the model: trend, seasonality, and holidays: $y(t) = g(t) + s(t) + h(t) + \epsilon_t$, where

- $g(t)$: represents non-periodic components using piecewise linear model with automatic change point selection or logistic growth curve trend.
- $s(t)$: trend factor that represents periodic changes. Time series often have multi-period seasonality as a result of the human behaviours they represent. This part relies on Fourier series to provide a flexible model of periodic effects.
- $h(t)$: effects of holidays (a list provided by the user). Holidays often do not follow a periodic pattern, so their effects are not well modeled by a smooth cycle.
- ϵ_t : error which will be assumed to follow a normal distribution.

3.2.3.2 Model assessment energy consumption

The RMSE and MAE [186] are two of the most common metrics used to measure accuracy for continuous variables and they are appropriate for model comparisons because they express average model prediction error in the units of the variable of interest as can be seen by their definition in the following equations:

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y}_i)^2}, \quad MAE = \frac{\sum_{i=1}^n |y_i - \bar{y}_i|}{n}$$

where y_i is the real consumption, \bar{y}_i is the predicted consumption and n is the number of observations.

However, in order to compare energy consumption prediction within works that do not use the same dataset or the same values of energy to be predicted it is not useful to compare such metrics whose output depends on the magnitude of the output data.

For that reason, we complement the information with the coefficient of variance of the RMSE. The Coefficient of Variation of the RMSE (CVRMSE) is a non-dimensional measure calculated by dividing the RMSE of the predicted energy consumption by the mean value of the actual energy consumption. For example, a CVRMSE value of 5% would indicate that the mean variation in actual energy consumption not explained by the prediction model is 5% of the mean value of the actual energy consumption [187]. CVRMSE has often been used in energy prediction studies [188]. Similarly, the Mean Average Prediction Error (MAPE) metric has been used in a wide

number of electricity prediction studies [189, 190]. It expresses the average absolute error as a percentage. They are calculated as follows:

$$CVRMSE = \frac{\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y}_i)^2}}{\bar{y}} \times 100, \quad MAPE = \frac{1}{n} \sum_{i=1}^n \left| \frac{y_i - \bar{y}_i}{y_i} \right| \times 100.$$

3.2.3.3 Hourly predictions

We were able to control an office based on an occupation schedule in order to test how the inclusion of external weather predictions would influence accuracy on energy consumption. In this case, we computed each hour's consumption prediction in a horizon of 24h.

Data was collected from 10 June to 14 August 2017. In this period, up to 4 workers were working in a normal schedule from 9:00 to 17:00. Equipment was turned on and off and the settings of the conditioning system were modified remotely.

We have tested the Prophet software, which was previously used in other disciplines and found that it could be excellent for experiments that require the prediction of pertinent variables. Our contribution has also been to prove that this package is an ideal “soft” addition to the infrastructure.

There was a positive correlation between the predicted external temperature (1 day before) and the measured external temperature $r = 0.9$, $p\text{-value} < 0.01$ so it is feasible to anticipate and predict the real outdoor conditions. A significant regression equation was found ($F(1,1582) = 6285$, $p\text{-value} < 0.01$, $R^2 = 0.8$). The real temperature is equal to $8.59 + 0.66 \times (\text{predicted}) \text{ } ^\circ\text{C}$.

Energy consumption in buildings has several characteristics appropriate for the Prophet algorithm and thus should perform well for energy prediction: strong multiple human-scale seasonality (such as day of the week and the time of year), important holidays that occur at irregular intervals that are known in advance and a certain random component.

We evaluated two scenarios in order to asses the inclusion of temperature forecasts.

- Model 1: Previous energy consumption, previous occupation and future occupation with a known pattern and schedule (RMSE = 286.73 KWh, CVRMSE = 10.2 %).
- Model 2: Model 1 + outdoor temperature values with corrected temperature predictions (RMSE = 268.56 KWh, CVRMSE = 9.5 %).

Out of working times, consumption stays always the same. For that reason, we have computed metrics for the working hours and we can see that CVRMSE is better using Model 2 than Model 1, justifying the inclusion of weather forecasts on the modelling.

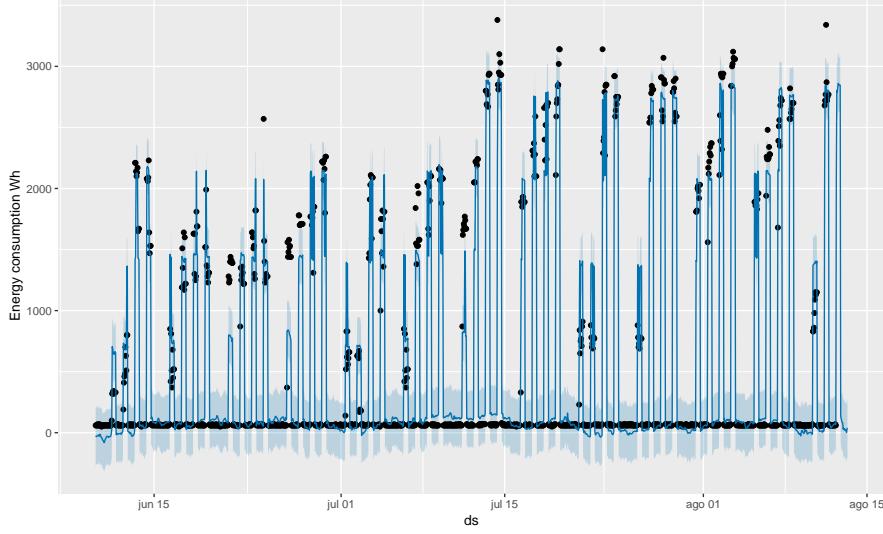


Figure 3.6: 24h predictions performed with the fitter model (blue line) and the true values (black dots) with Model 2)

3.2.3.4 Moment of the day predictions

We have displayed an outline based on basic and logic usability estimations of the building and we have included their Consumption Range (CR) and Consumption Mean (CM):

- Moment 1: holidays, weekends, nights (22h-6h). CR = [3.578, 14.1] KWh, CM = 7.904 KWh
- Moment 2: regular mornings (6h-14h). CR = [26.01, 86.19] KWh, CM = 54.27 KWh.
- Moment 3: regular afternoons (14h-22h). CR = [6.357, 53.290], CMs = 31.48 KWh.

The visual differences that are noticeable at Fig. 3.7 were confirmed with a Kruskall Wallis H[191] test. There is a significant difference between groups ($H(2) = 547.7$, p-value < 0.01). An analysis of the differences by pairs performing the post-hoc Wilcoxon test [191], determines that it is possible to divide data in those moments. This reasoning leads us to suggest three different models corresponding to the just mentioned partitions.

We considered 8 different observations for each environmental input (one per hour). Also, we created two new variables for every attribute by taking their mean and median. Just to clarify the considered inputs, for situation 1 and, for example, temperature, we will have 11 attributes: temperature at 6 AM, at 5 AM, ... at 22 PM, mean of temperature (from 6AM to 22PM) and median of temperature.

After training the models using several combinations of inputs we achieved the best results using the day of the week, month, season, mean temperature and mean humidity with RF algorithm for moment 1 (mtry = 4, RMSE = 1 KWh) and moment 3 (mtry = 2, RMSE = 3.87 KWh) and BRNN for moment 2 (number of neurons = 2, RMSE = 7.08 KWh) as can be seen in Table 3.2. All these values represent between a 12.09% and a 12.86% of error (CVRMSE).

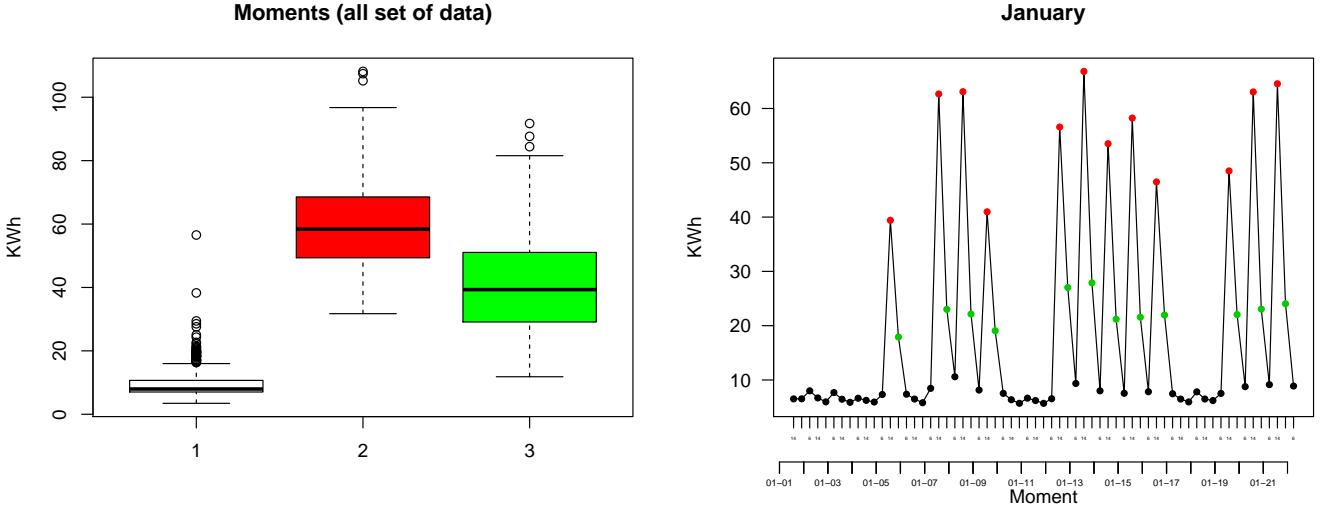


Figure 3.7: Boxplot of the energy consumption by moments considering all data (left); and, the time series of the energy consumption by moments during January (right)

Moment	Technique	Best Parameter	RMSE (KWh)	CVRMSE (%)	R^2
1	Gauss	$\sigma = 0.1$	1.1	13.43	0.57
	MLP	size = 34	1.1	13.46	0.55
	SVR	cost = 4	1.09	13.26	0.58
	BRNN	neurons = 3	1.1	13.47	0.55
	RF	mtry = 4	1	12.18	0.65
2	Gauss	$\sigma = 0.1$	7.76	14.1	0.67
	MLP	size = 37	1.56	15	0.68
	SVR	cost = 1	4.26	13.4	0.71
	BRNN	neurons = 2	7.08	12.86	0.75
	RF	mtry = 2	7.48	13.59	0.72
3	Gauss	$\sigma = 0.1$	4.5	14.07	0.67
	MLP	size = 37	4.81	15.03	0.69
	SVR	cost = 1	4.20	13.14	0.73
	BRNN	neurons = 5	4.31	13.45	0.73
	RF	mtry = 2	3.87	12.09	0.76

Table 3.2: Results obtained for each moment

Having trained and tested 5 different models, it is necessary to find statistical evidence that the selected one outperforms better not just in a punctual way. Our 10-fold cross-validation with 5 times repetition strategy generates a set of 50 measurements for each model. In Figure 3.8 (left) it is displayed the model's performance for situation 3, where red crosses show the median of RMSE for each model. For every situation, the Friedman test, that is the nonparametric alternative for repeated measures (within subjects) Analysis Of Variance (ANOVA) is significant ($p\text{-value} < 0.05$), and looking for corrected pairwise differences, we find that RF is the only one

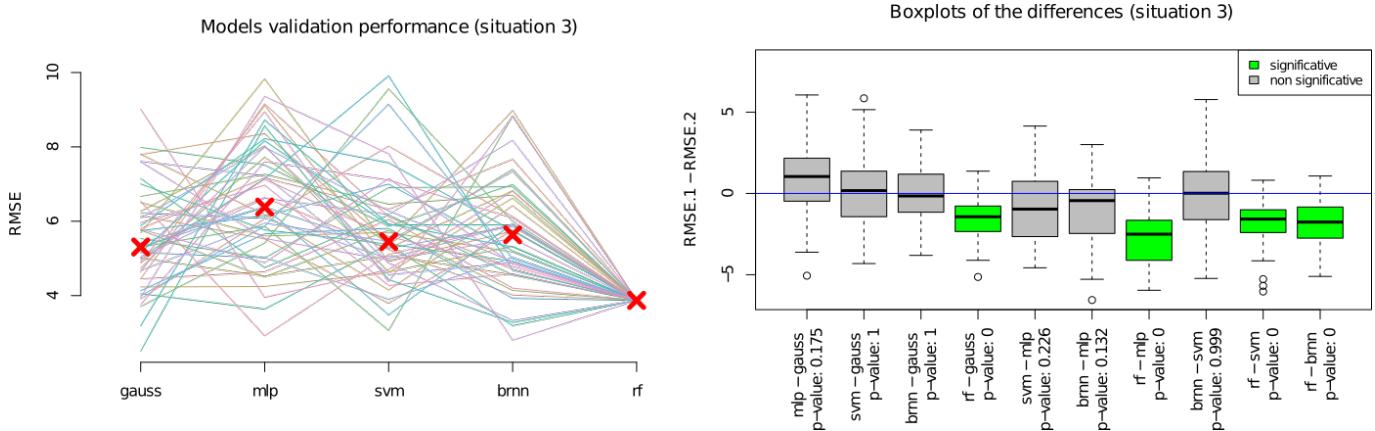


Figure 3.8: Models validation performance (left) and Pairwise differences between models performance (right)

that differs from the others. Between the other models' performances, there aren't significant differences as can be observed in Figure 3.8 (right), where boxplots of the differences are coloured in green when the differences are significantly different from zero.

3.2.3.5 Daily and weekly predictions

Our interest lies in the weekly quantification of energy use. However, daily dynamics are useful since there are patterns that can be found depending on the day of the week. Our model predicts daily energy consumption and then computes the metrics in an aggregated manner, so the global quantification takes place on a weekly basis.

The data that is used in order to build and train our baseline corresponds to 1 year's worth of data from a whole building, from February 2016 to February 2017 and we will compare a grey-box and black-box methodologies.

In order to make use of grey-box models, the set of outputs and inputs have to be defined together with the topology of the system. The most common mathematical representation of lumped parameter models is the state-space representation. The general form for time-invariant models can be written as shown in Eq. 3.2

$$(3.2) \quad \begin{cases} x'(t) = Ax(t) + Bu(t) \\ y(t) = Cx(t) + Du(t) \end{cases}$$

where x is a vector concerning the states of the model, in our case the temperatures, x' is the derivative (rate of change) of the states, A is a characteristic matrix of the model, B defines the effect of the inputs in the model, and u are the inputs, in our case the outside temperature and gains in electric. In this formulation, y represents the variables that are measured, in our case electricity, C is the identity matrix; and D is zero in all cases for this work. Using this

		Models					
		SVR	RF	XGB	TWT	Gauss	Grey
Daily	CVRMSE	12.4	9	11	14.9	17.45	33.57
	MAPE	7.2	6	7.3	12.3	15.01	43.02
Weekly	CVRMSE	6.4	5	6.2	11.1	16.8	19.53
	MAPE	5.2	4.5	5.5	9.4	12.3	15.48

Table 3.3: Metrics for energy consumption forecasting

formulation, every time a solution has to be evaluated, the built-in GNU Octave function **lsim** was used.

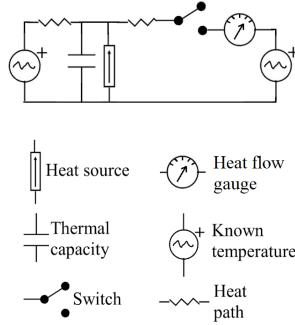


Figure 3.9: Dual-mode RC network

We have considered a dual-mode RC-network as the one shown in Fig. 3.9 and previously introduced by Ramallo-González on [192]. These grey-box models have been largely used in the past for building energy simulation. The reader is referred to [66], [193] and [192] for quantitative evidence of the accuracy of this kind of models.

Our black-box methodology is highly versatile with respect to the input data since it allows the addition of variables with minimal effort. We create the method in a constructive manner by relating the 24 temperature values of each day with the energy consumption of the building.

The subject building has several features that are typical of educational buildings: the load on weekends is substantially lower than that of weekdays and there might be also differences among weekdays. In these terms, we used ANOVA in order to determine whether there were differences between the consumption on the different days of the week ($p\text{-value} = 0.001 < 0.05$). After carrying out a **post-hoc** test we concluded that Fridays could be considered to behave differently to the other days of the week, which could be owing to lower occupation. Having attained this knowledge, we considered it necessary to add a dichotomous variable that indicates the kind of day of the week. Weekend and holiday consumption is estimated using the mean of previous weekends and holidays.

The algorithms that were found to be relevant for use within our black-box methodology are: SVR, RF and XGB.

The prediction metrics are summarized in Table 3.3. The first three methods: SVR, RF, XGB

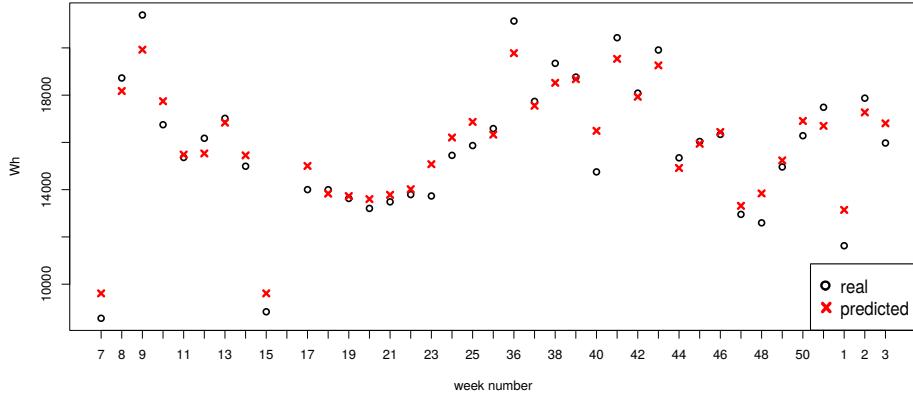


Figure 3.10: Weekly predictions using RF and real consumption

belong to the group of black-box models and are blind to the physics of the problem. We also have the TWT, the GAUSS and the Grey-box model, the last of which contains information regarding the physical phenomenon of the model topology. As can be seen, they return the best results when compared to the Gaussian method, which is applied in a more traditional manner, i.e., by relating the instantaneous consumption measurement with the instantaneous inputs measurements and also with our grey-box model approach.

Of the three black-box methods, RF is that which stands out since it attained a CVRMSE of 9 and 5 % and a MAPE of 6, and 4.5 % for the daily and weekly predictions respectively. We have plotted the weekly consumption in Fig. 3.10 .

3.2.4 Feature selection [R4]

In our framework for energy prediction we have devised a methodology that consists of: database transformation, FS, regression, decision making and forecasting. Next, each step is described separately.

After transforming the database, it is necessary to define the search strategy and the evaluator. We have configured a multitude of different methods for FS, both filter and wrapper, univariate and multivariate. For multivariate wrapper FS methods, we evaluate attribute sets by using a learning scheme with cross-validation and a performance measure. For univariate wrapper FS methods, we evaluate the worth of an attribute by using a user-specified classifier, cross-validation and a performance evaluation measure. The FS and classification processes have been executed in batch mode.

We considered for this research five wrapper and three filter FS methods, five multivariate and three univariate as it is shown in Table 3.4.

3.2.4.1 Strategies

As multivariate FS methods, we used the probabilistic strategy MultiObjectiveEvolutionary-Search [194] (two objectives: performance metric and attribute subset cardinality), and the

Database #Id.	Type of FS method	Search strategy	Evaluator	Acronym
#1	Wrapper Multivariate	MultiObjectiveEvolutionarySearch	RF (MAE)	MOES-RF-MAE
#2	Wrapper Multivariate	MultiObjectiveEvolutionarySearch	RF (RMSE)	MOES-RF-RMSE
#3	Wrapper Multivariate	MultiObjectiveEvolutionarySearch	kNN (RMSE)	MOES-kNN-RMSE
#4	Wrapper Multivariate	MultiObjectiveEvolutionarySearch	LR (MAE)	MOES-LR-MAE
#5	Wrapper Univariate	Ranker	RF (RMSE)	RANKER-RF-RMSE
#6	Filter Multivariate	GreedyStepwise	ConsistencySubsetEval	GS-CFSSE
#7	Filter Univariate	Ranker	ReliefFAttributeEval	RANKER-RFAE
#8	Filter Univariate	Ranker	PrincipalComponents	RANKER-PCA

Table 3.4: Proposed FS methods for energy time series forecasting

deterministic strategy GreedyStepwise [195]. The two most popular Multi Objective Evolutionary Algorithms (MOEA) are Elitist Pareto-based MOEA for diversity reinforcement (ENORA) [196], on which our team is intensively working over the last decade and NSGA-II (elitist non-dominated sorting genetic algorithm) [197]. In [198] is statistically shown that ENORA performs better than NSGA-II in terms of hypervolume [199, 200] for regression tasks, for which we have decided to use ENORA in this work. GreedyStepwise performs a greedy forward or backward search stopping when the addition or deletion of any of the remaining attributes results in a decrease in evaluation, thus, it has no backtracking capability.

For univariate FS, we used Ranker method [201] that ranks attributes by their individual evaluations is used.

3.2.4.2 Evaluators

For the wrapper methods, we used RF, k-Nearest Neighbors (kNN) and LR, with the metrics RMSE and MAE in order to test the features subsets. Those methods offer a good compromise between performance and computational time.

We considered the multivariate filter evaluator ConsistencySubsetEval [202], which scores a subset of features as a whole, by projecting the training instances according to the attribute subset. As univariate filter evaluators we used ReliefFAttributeEval [203] and PrincipalComponents [204]. *ReliefFAttributeEval* evaluates the value of an attribute by repeatedly sampling an instance and examining the value of the given attribute for the nearest instance of the same and different class. *PrincipalComponents* performs dimensionality reduction by choosing enough eigenvectors to account for some percentage of the variance in the original data (default 95%), and then transforms back to the original space.

3.2.4.3 Regression

After FS we obtained 8 features subsets. The used predictive algorithms, evaluated with cross-validation, were: RF, kNN, LR, SVRs [205] , GAUSS. Table 3.5, shows the evaluation in 10-fold cross-validation, 3 repetitions (a total of 30 models with each regression algorithm in each database), for the metrics *RMSE*, *MAE*. We have used a corrected paired t-test in order to test whether results are statistically significantly different to #1 . A mark * denotes that it is worse,

3.2. DATA ANALYSIS IN IOT BASED SMART ENVIRONMENTS

		#1	#2	#3	#4	#5	#6	#7	#8	TD
RMSE	RF	12.6685	12.9133	13.3814 *	14.4111 *	15.4203 *	18.7996 *	13.9174 *	36.7834 *	21.3455 v
	kNN	17.7612	20.8112 *	17.2423	25.0447 *	25.8680 *	25.7792 *	22.7562 *	37.8315 *	29.5936 *
	LR	19.3960	18.3234 v	18.5017 v	18.1808 v	18.6416	22.0898 *	18.1092 v	53.6083 *	17.7597 v
	SVR	20.0636	18.8337 v	18.9237 v	18.6714 v	18.9697 v	23.0016 *	18.8051 v	55.5770 *	18.2458 v
	GAUSS	21.9231	21.7133	24.7083 *	19.4114 v	18.8832 v	22.1160	18.3440 v	54.6361 *	17.8482 v
MAE	RF	5.8264	6.0012 *	6.2785 *	6.8621 *	7.5242 *	9.0191 *	6.4675 *	23.3071 *	12.6164 *
	kNN	8.8796	10.0150 *	8.6797	13.1307 *	13.3098 *	12.7038 *	11.0927 *	17.4372 *	14.3159 *
	LR	11.2708	10.1276 v	10.2363 v	9.7287 v	10.1738 v	13.2454 *	10.5091 v	38.3269 *	10.6126 v
	SVR	10.0410	9.0122 v	9.0806 v	8.9702 v	9.0669 v	11.5835 *	8.9962 v	36.7292 *	8.9314 v
	GAUSS	15.3332	15.1402 v	17.9369 *	12.0857 v	10.6294 v	13.3943 v	11.0307 v	38.8031 *	10.9028 v
CPU time	RF	0.9474	1.0349 *	0.7792 v	1.3432 *	0.9714	0.5708 v	0.7802 v	1.6078 *	3.0609 *
	kNN	0.0005	0.0005	0.0000	0.0000	0.0005	0.0000	0.0016	0.0000	0.0005
	LR	0.0042	0.0109	0.0026	0.0125 *	0.0089	0.0063	0.0115	0.0057	4.2172 *
	SVR	31.9255	29.0380 v	26.4307 v	62.5958 *	87.1901 *	75.5521 *	141.1615 *	9.0995 v	1626.4151 *
	GAUSS	115.4714	115.3620	115.4219	115.5542	110.7714 v	110.4302 v	110.5990 v	110.6505 v	114.0536

Table 3.5: RMSE, MAE and CPU time(in seconds) with 10-fold cross-validation (3 repetitions)

Input attribute	Rank	Importance
Lag_energy-1	1	7.398
Lag_stMO12_IMI_radmax+0	2	1.337
holiday	3	0.367
Lag_energy-3	4	0.357
ArtificialTimeIndex	5	0.302
Lag_stMO12_IMI_radmed-3	6	0.273
Lag_pr_feels-2	7	0.248
Lag_pr_temp-2	8	0.172

Table 3.6: Selected attributes with *MOES-RF-MAE* (database #1) and their ranks.

a mark *v* denotes a statistically better result, and no mark denotes no statistically meaningful difference.

3.2.4.4 Decision making

Looking at Table 3.5 we see that the best results have been obtained with the FS method *MOES-RF-MAE* (database #1) when *RandomForest* is used as regression algorithm, which shows statistically significant differences with respect to the rest and *MOES-RF-MAE* is also superior, with statistically significant differences except for the FS method *MOES-RF-RMSE*. With respect to the UserCPU_Time_training performance metrics, its results are acceptable in comparison to the rest of the methods. We can then choose the FS method *MOES-RF-MAE* and the database #1 for the final forecasting process.

Table 3.6 shows the selected attributes with *MOES-RF-MAE* and their rank and importance for each of the datasets. An attribute is evaluated by measuring the impact of leaving it out from the full set.

	1 step		2 steps		3 steps		Average	
	#1	TD	#1	TD	#1	TD	#1	TD
MAE	10.994	26.758	20.465	34.777	32.749	49.7	21.403	37.079
RMSE	16.051	36.556	28.768	45.079	44.834	59.821	29.884	47.152
# Instances	1526		1525		1524		-	

Table 3.7: Evaluation on test data with RF - database #1 and TransformedDatabase (TD)

3.2.4.5 Forecasting

Finally, we analyse the prediction ability of the forecaster obtained with the selected attributes. First, we train the model on the data, and then it is applied to make a forecast at each time point by stepping through the data. These predictions are collected and summarized, using MAE and RMSE metrics.

Table 3.7 with the databases #1 and *TransformedDatabase* (with all lagged variables and all overlay variables), on test data (30%). The reduced database #1 improves the 1,2,3-steps-ahead predictions using the database without performing FS (TD). Using the averages of the steps-ahead predictions we see that with our methodology MAE is improved by 42.28% and RMSE by 36.62%, evaluated on test data.

3.2.4.6 Comparison with other methods proposed in the literature

For current and future comparisons with further research, the hourly *CVRMSE* was 20 % and we have also averaged it per day obtaining a daily *CVRMSE* = 11 % for the 1-step case.

Multivariate ARIMA: we have used the traditional time series method *ARIMA* with exogenous regressors [206]. Results are much worst than using our ML oriented approach. Using our selected features, mean *MAE* is 119 and mean *RMSE* is 126. This results are way worse than ours but still better than using all variables with *ARIMA*, for which *MAE* increases between 35 and 55 KWh and *RMSE* increases between 37 and 58 Kwh.

3.2.4.7 Analysis of results and discussion

As expected, wrapper show better performance than filter FS methods, and multivariate show better performance than univariate FS methods. Multivariate methods can identify interaction amongst features simultaneously, specially wrapper-based FS methods [207]. To reduce the computational time of multivariate wrapper FS methods (NP-hard), deterministic search strategies, such as *GreedyStepwise*, can be used but hidden and basic interactions could be missed due to the way the search space is traversed [208]. Probabilistic search techniques, such as *MultiObjectiveEvolutionarySearch*, can overcome these difficulties by allowing to generate new subsets in different locations of the search space guided by a metaheuristic. In the thesis, we propose to use a multivariate wrapper FS method where the search strategy is based on multi-objective

evolutionary computation, thus intrinsically overcoming the problem of interactions between features.

For wrapper FS methods, the RF evaluator has proven more effective than kNN and LR based evaluators. SVR and GAUSS are discarded as evaluators for wrapper methods because of their excessive computational time. Run time of RF is acceptable, and this method is not very sensitive to the variation of its parameters.

MAE has shown better behaviour than *RMSE* as metric performance in evaluators for wrapper FS methods since *MOES-RF-MAE* (database #1) produces better results than the method *MOES-RF-RMSE* (database #2) (see Table 3.5) when evaluated on cross-validation with *RandomForest* using the *RMSE* metric (12.6685 vs. 12.9133, an improvement of 1.9%). This improvement can also be observed in Table 3.5 when both databases are evaluated with the *MAE* metric (5.8264 vs. 6.0012, an improvement of 2.91% in this case).

3.2.5 HVAC patterns [R5]

Each terminal unit has a remote controller that facilitates the interaction of the user with the conditioning system. The user can turn on or off the room unit at their will, but cannot program the operation based on a timer. Also, the user can control the set point temperature. This means that the user can change at any time the thermostat control of the unit to any value between 16 and 29-degree C at their will. Each room has also a wall-mounted screen that shows the temperature of the room captured by the machine, the set point (thermostat) temperature and the fan operation mode. Every 12 minutes the following data was gathered: room temperature (°C), on/off status, set point (°C).

The intention is to create virtual areas comprising several building space areas, finding patterns in the HVAC use and consequently in the energy-related use and defining these virtual areas according to such information to optimise the content of information.

To do so, we aggregate each attribute per energy device daily. We can represent each device as a time series and with this, it is possible to fit a model or find a clustering algorithm that groups every attribute of the time series finding some distinctions between them

- Interaction frequency to turn it on/off
- Interaction frequency to change the set point
- Daily hours of operation (how many hours the machine is on)
- Average and standard deviation of the daily set point preferences

We have defined two ways to find patterns: based on the interaction of people with the machines in order to change the set point and based on the temperature preferences.

Fig. 3.11 (left) shows the histogram of interactions of the users with their controllers normalised by hours of use. Due to the skewness of the data, we then applied a logarithmic transformation.

We wanted to investigate the yearly fluctuations of data that could be found in thermostat values because on the ASHRAE Standard 55 for thermal environmental conditions and on [209] is stated that preferences differ throughout the year. In order to smooth the data yet imposing the yearly periodicity we fit a sinusoidal function with a fixed wavelength of 365 days. The equation was:

$$T_{set}(time) = a + b \sin\left(\frac{2\pi time}{356} + \lambda\right),$$

where a is the constant term, b is the yearly swing, and λ is the phase or lag.

We plotted the kernel density functions of the set point changes per hour, grouping the curves by people who used the machines for more than the hours given in the legend of Fig. 3.11 (right), showing a bimodal nature. As a result, two groups of users can be defined: one that interacts with the controller often (a change every week) and also consumes less and another that tends to not to interact with the set point (1 or 2 changes in the whole period) and are higher consumers.

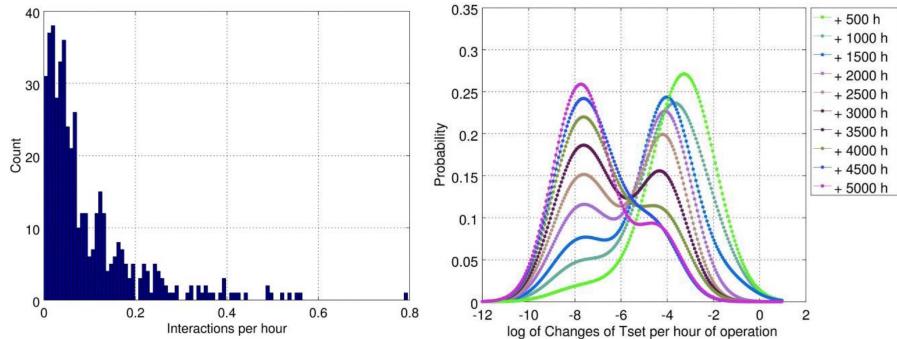


Figure 3.11: Changing frequency (left) and (right)

According to ASHRAE, the bounds of comfort for summer and winter are 25-27, 21-24 respectively. To evaluate the number of hours that the systems were pushed outside the comfort ranges, and how much, the integral of the area defined by the curve representing the set point temperature and the upper or lower bound were calculated. This provided with an indicator of overheating or over-cooling on Kelvin-hour, Kh a measurement well recognized on the Building Physics community.

It was seen that users tend to tolerate high temperatures much more than cold temperatures and use values of the thermostat for cooling that are close to the upper bound of the ASHRAE comfort range. The mean of overheating was 292.7 Kh, which is substantial. However, in the case of overcooling the figures were more prominent: a mean of 866.4 Kh. Those users that interact with the machines less often have registered the larger overcooling and no particular relationship between the number of interactions with the controls and overheating was found.

We also propose another way to classify the rooms based on the clustering of the raw set point time series. In this case, we used several algorithms: Hierarchical clustering, longitudinal K-means, DBSCAN and Spectral clustering.

Once every device is assigned to a cluster or virtual area, the mean of the elements of each cluster is computed in order to get an average measurement. Finally, each generated cluster is stored as an instance of a virtual energy area. Those virtual sensors were useful in order to:

- Send specific actions by checking the cluster that the user/room belongs to,
- Improve energy consumption prediction aggregating by cluster,
- Detecting outlier data which implies device failures [210],
- Reduce the monitored number of HVAC in order to obtain similar results in the analysis.

3.2.6 Human mobility patterns at macro and micro levels [R6]

Wearable devices are equipped with sensors like GPS that allow capturing a large amount of high-resolution digital traces which are instrumental for the mobility mining discipline, focusing on giving insight into the spatiotemporal trajectories of people. Looking at a macro level we try to detect regions with a high density of human transit. At the same time, many social network sites have included location-based capabilities into their smartphone's applications so most of the data belonging to those sites can be geotagged. This can be used to recognise human mobility models and patterns at both micro and macro levels.

3.2.6.1 Dense Transit Areas identification with GPS (macro)

We propose to characterise the flow of people inside regions using the online aggregation of the spatiotemporal traces. Our mechanism allows discovering Dense Transit Areas (DTA) that represent a spatial region of a city that is visited by a large number of citizens' routes in real time. Such monitoring aims to detect relevant changes of human mobility within regions that can be signs of events of interest, like unplanned demonstrations or serious traffic problems. The present system supports two different modes of execution, DTA discovery (generating areas) and DTA monitoring (controlling if areas remain as such).

DTA discovery

Routes are composed by the GPS sensor tuples (x, y, t) where (x, y) are latitude-longitude coordinates at instant t . The route is delivered to the central server and also stored in the local personal routes repository that keeps the last routes covered by the user within the mobile client. Distinguishing between low speed (walking) and high speed (vehicle) routes.

The spatial region under study is divided into squared cells and subcells. The route density is calculated in each of them considering: length of the cell, average speed and spatial length of the route and the coming and outgoing side of the route. Two areas are merged if the transit

information they represent is strongly related (similar speed and direction). Once a consistent set of DTAs is generated (when the ratio between new DTAs and the number of routes is below a threshold) DTAs monitoring starts.

DTA monitoring

This phase focuses on controlling the evolution of mobility features of the DTAs to early detect potential mobility shifts using a reliable subset of participants. For each person, the probability of visiting a DTA within a timestamp is calculated in the case that they have initiated a route nearby. For each period, the subset of users providing the best coverage of the detected DTAs is then used. For them, the ongoing route's sequence is stored. The similarity between the current and the historical state of each DTA is computed and if either the speed or direction similarity is below a predefined threshold during consecutive sampling periods then the algorithm infers that the human dynamics inside the DTA under review have changed so we go back to the DTA discovery phase.

Evaluation

In order to test our system, we have used the GeoLife dataset (GL) [211], a public collection of human trajectories produced by 178 users carrying different GPS feeds for over three years in Beijing city (China).

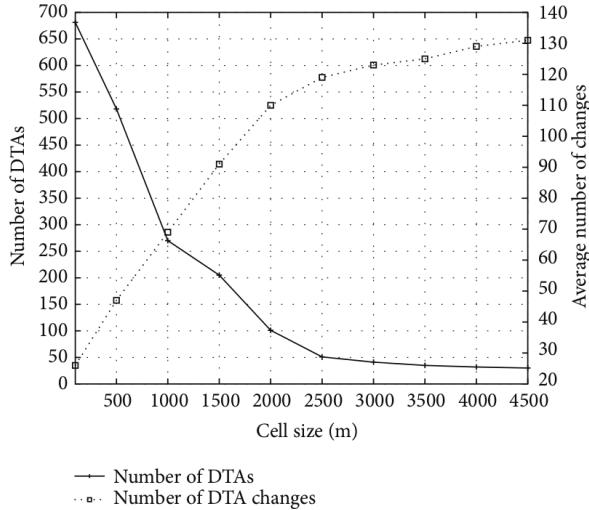


Figure 3.12: Number of DTAs and average number of changes per DTA with respect to the cell size

In Fig. 3.12 we can see that setting a small cell size generates a large number of DTAs and when increasing cell size DTAs become more sensitive to changes. We also analysed the accuracy of the approach by comparing the capability of the system to detect variations of speed and directional features with respect to the DSSIM function [212] -that measures the dissimilarity between two spatiotemporal trajectories during a time period- and the event-based mechanism proposed in [213] to detect abnormally high or low speeds of moving objects in real time. The

results achieved by our proposal vary according to the DSSIM function. Our system had a higher precision for large DSSIM values that indicate very evident differences between current and historical routes. When DSSIM ranged between 0.8 and 1.0, our proposal achieved acceptable precision results, about 0.8. However, for DSSIM values below 0.6 ours system's precision decays. This is because certain differences of the routes do not imply changing their incoming or outgoing sides in the DTA which are the features used by the proposal to detect transit changes. An strength of our method is that it des not need the whole sequence of timestamped locations to operate whereas the others do.

3.2.6.2 Human routes identification with social media (macro and micro)

The previous approach does not take into account the activity level of the users within each detected area. Here we propose two ways for predicting the movement of the population within the city on an online fashion using the social network Twitter¹⁸. Those are a graph-based and a cluster-based approach.

Graph-based

A user's graph is generated and updated on the basis of his geotagged documents ("tweets") gathered on composing a hierarchy that represents the mobility flows of a large spatial region. The graphs include frequent locations ("landmarks") and frequent topics used for predicting the next meaningful location, or landmark, that will be visited by a person. The system architecture is compound by a client-side that runs on the mobile device of each user for detecting personal landmarks and a server-side responsible for composing collective ones.

For preprocessing, CEP is used. A CEP system is useful for the timely detection of situations of special significance that cannot be directly handled by humans. It consists of a palette of asynchronously interconnected Event Processing Rules (EPRs), defined by expert knowledge. Our defined CEP events are *tweet*: a raw tweet wit textual content and metadata (user nick, timestamp, and geotagging) as attributes; *tweet with topic*: includes the most probable topics that it refers to using a bag of words; *landmark*: indicates if a tweet has been written inside any personal or collective landmark; and *route*: represents a completed route as a sequence of personal or collective landmark. Retweets, URL links, mentions to other users, and stop words are deleted.

Route composition

The landmark comprising each new tweet with topic event including spatial region and its frequently associated activities is detected. Personal and collective landmarks are spatial regions with a high density of tweets related to one or more activities. As a result, a slightly modified version of the online landmark discovery algorithm (LDM) [214] has been applied to *tweets* locations for landmark detection.

Graph Generation

¹⁸<https://twitter.com/>

Personal and collective mobility graphs using the completed routes' sequences are generated on the fly. Both graphs encode the statistical information from the routes as a directed multigraph where each vertex represents a unique landmark. The personal mobility graph is updated when a route is completed. If all elements (landmark{activity}) are already vertices of the graph and there is a route identifier whose edges connect these elements in the same order, its identifier is extracted. Otherwise, a new identifier is generated. Then, the frequency attributes of each edge associated with this identifier is incremented or created. We perform a similar task also with the collective routes and create two graphs (working days and weekends).

Location prediction

Each time a new landmark event is appended, the route is delivered to the Local Predictor Maker (LPM) as Fig. 3.13 shows. LPM focuses on forecasting the next landmark (spatial region and associated topic) covered by the ongoing route. The algorithm detects the historical routes that best fit the ongoing route. This detection is done by searching the maximum set of edges that connect the visited landmarks in the same order. If target user statistics do not provide a good prediction then the algorithm makes use of the collective statistics.

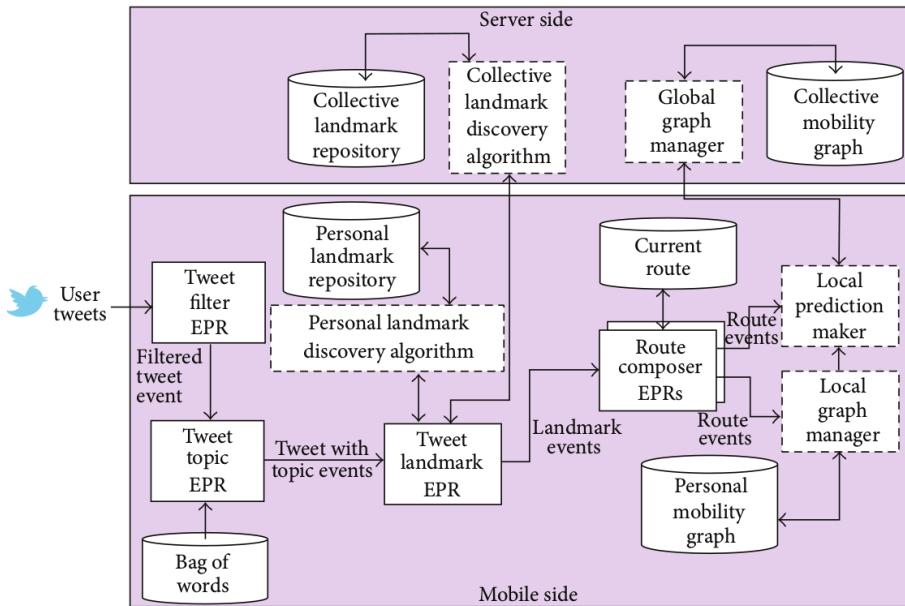


Figure 3.13: System architecture. The components that are not EPRs are depicted as dashed boxes

Evaluation metrics

The prediction rate (PR) and the prediction error (PE) metrics are used. PR counts the number of routes for which at least one landmark is provided as a prediction when a new element is appended (coverage). PE is the average of all distance deviations across each prediction (deviation from the actual next landmark $real_l$). Since each landmark may be associated with an activity, the distance between the predicted ($pred_a$) and the real activity ($real_a$) must be also

considered so we have used the Semantic-Hierarchical Similarity (SHS) [214].

$$PR = \frac{\# \text{routes with prediction}}{\# \text{routes}}, \quad PE = w_1 \times (1 - \frac{dist(real_l, pred_l)}{dmax}) + w_2 \times SHS(real_a, pred_a)$$

where $dmax$ is the maximum distance between two points in the dataset's spatial region and w_1, w_2 are adjustable coefficients.

Cluster-based

We cluster social-media documents with the fuzzy c-means (FCM) algorithm for different predefined time slots. Those are: early morning (0am-8am), morning (8am-12am), evening (12am-4pm), late (4pm-9pm), night (9pm-0am). For such clustering, the spatiotemporal features of the documents are considered. On the basis of these clusters, we define different profiles to predict the movement of the population within the city. Furthermore, the levels of social media activity have been measured for each cluster. As we will see, a correlation exists among the detected clusters, their associated activity and the prediction levels of its related movement.

For preprocessing, we aggregate tweets that have been posted in similar locations and time in order to avoid the disturbance of the real prediction with the same user's tweets which do not represent a real movement in space-time dimension.

Clustering

The FCM clustering algorithm [215] is applied to each of the five time-slots. The result is a membership matrix between all tweets and all clusters. For each time-slot, the cluster with the highest membership is selected to be the representative one for the user in that timeslot. As a result, five pairs of centroid-time slots are obtained which represent the usual movement of the user during the day across the time slots. Users are classified depending on the average posts per day, users are classified into three levels: low, medium, high. The activity level of each cluster is measured in order to enrich the information about the users and discover the kind of users in each cluster.

Movement prediction between clusters and time slots

For this task, the percentage of users in each time slot is calculated using the representative cluster for each time slot and user. Then, using the information between the pairs of clusters in consecutive time slots (e.g. from early morning to morning) we obtain the percentage of users that flow from one cluster to another. This information, combined with the activity level for each cluster provides a global and precise vision of the behaviour of the population studied in the total area. Using the information between the pairs of clusters in consecutive time slots (e.g. from early morning to morning) we obtain the percentage of users that flow from one cluster to another.

Experiments and results

In order to evaluate the proposals we used the Twitter Crawling API targeting Madrid city during 82 days using 181581 tweets from 41008 users.

Figure 3.14 (left) shows the collective landmarks generated by the system. We can see that the higher concentration of collective landmarks corresponds to spatial areas with a high density

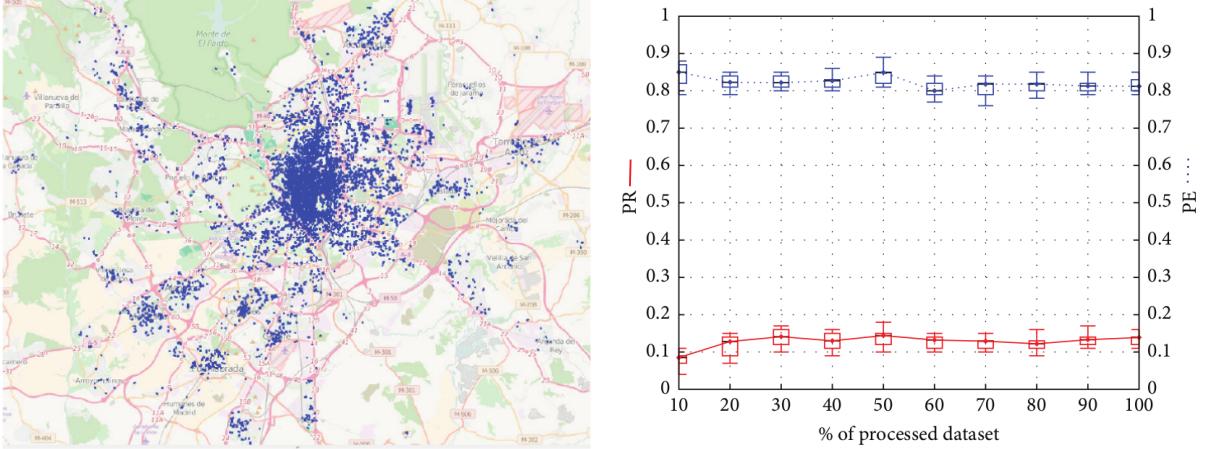


Figure 3.14: Collective landmarks (left) and metrics evolution

of human movement like the city centres. The evolution of the PR and PE of the system can be seen in Fig. 3.14 (right). There we appreciate that the PE remained more or less flat with no significant variance. This way, when only 10% of the dataset was processed our proposal was capable of achieving around 0.85 PE, thanks to the spatial distribution of the tweets. Users tend to post tweets located relatively close among them, limiting the locations that should be considered by the system and, thus, it is more likely to correctly predict the next location of a user.

Fig. 3.15 depicts the heat map of resulting digital traces of the datasets in the early morning and morning time slots and also the flow of people within time slots. Unsurprisingly, remarkable density of tweets exists in the centre of the city whereas the tweets in the suburbs are more spread.

3.2.7 IoT-based Big Data architecture for smart cities [R7]

The previously described methods need a common environment to interact with. In that sense, it is necessary to create an IoT-based platform to share the large volumes of heterogeneous information and to manage all interoperability aspects and enable the integration of the ML techniques above described.

The IoT platform is compliant with the FIWARE architecture, a key initiative of the Future Internet Public–Private Partnership (PPP) to create a well-aligned set of open enablers to receive, process, contextualize and publish IoT data from and for smart cities including from city-wide information to dwelling specific data¹⁹. In particular, the Orion Context Broker (OCB)²⁰ and the COMET²¹ modules are used in order to store in a NoSQL repository the historical data, that are

¹⁹<https://www.fiware.org/>

²⁰<https://fiware-orion.readthedocs.io/en/master/>

²¹<https://fiware-sth-comet.readthedocs.io/en/latest/>

3.2. DATA ANALYSIS IN IOT BASED SMART ENVIRONMENTS

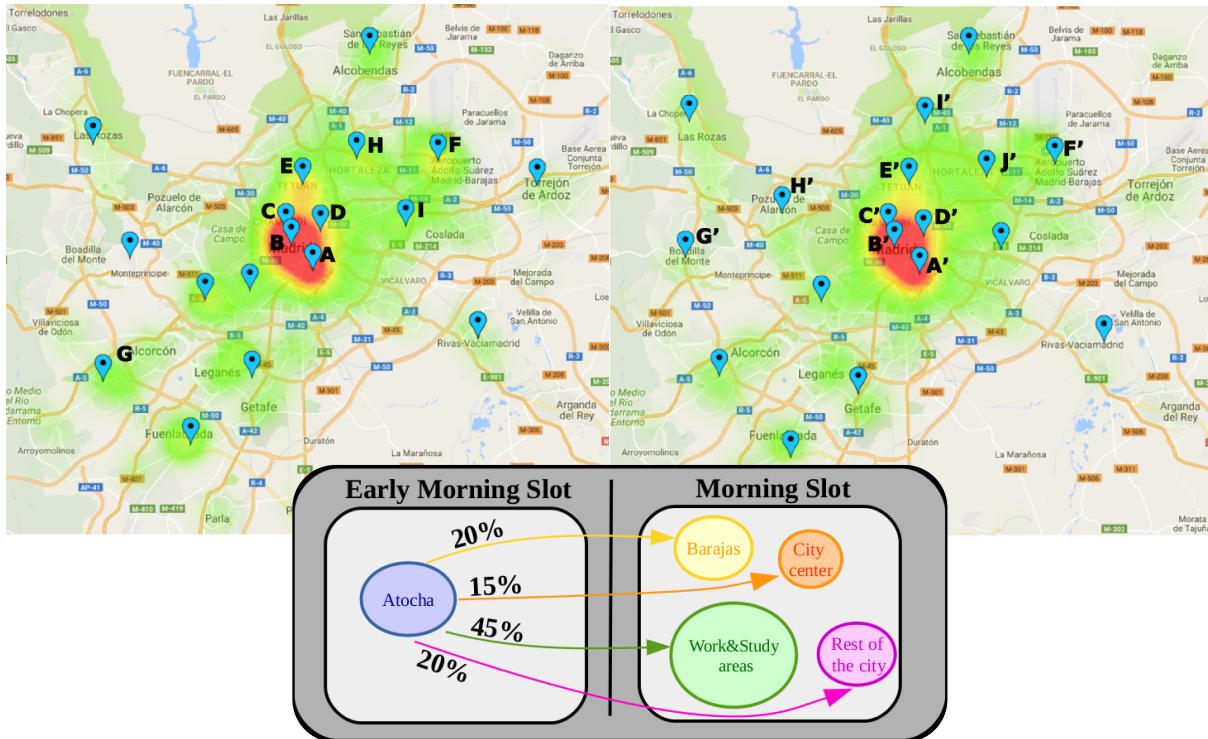


Figure 3.15: Heatmaps of clusters and movement prediction between early morning and morning slots

the measurements from the several data sources.

We started by defining a model compliant with the NGSI information model accepted in the FIWARE ecosystem that follows an entity–attribute approach. Each entity represents real or virtual elements of interest and has a type that allows defining type-based hierarchies. In this way, an entity has its own defined attributes and the inherited ones from its ancestors. Among its components there are three key entity groups related to the energy ecosystem of a building by means of NGSI entities:

- *Building entity*: the operational (opening hours, building use, etc.) and architectonic (fabrics, orientation, etc.) details of the building are the attributes of this entity.
- The *Spacial region* entity serves to link buildings with similar energy usage patterns because they are located in similar geographic regions.
- The *Building space area* entity represents the inner structure of a building (e.g., classrooms, corridors, etc.).

Introducing data related to the previous entities facilitates the transfer of information between platforms and Building Information Modeling (BIM).

- Entities referring to the energy-related sensors: building sensor, power meter, and HVAC entities. Each entity includes the set of attributes monitored by the corresponding sensor along with other metadata (e.g., location, timestamp). The clean version of these entities refers to the filtered data.
- Entities that represent sensors outside the building that may provide useful information. As Fig. 3.16 shows, this is defined by means of the *external sensor* and *weather conditions* entities.

Finally, only the entities in gray in Fig. 3.16 have instances stored in ORION and COMET as we will see later.

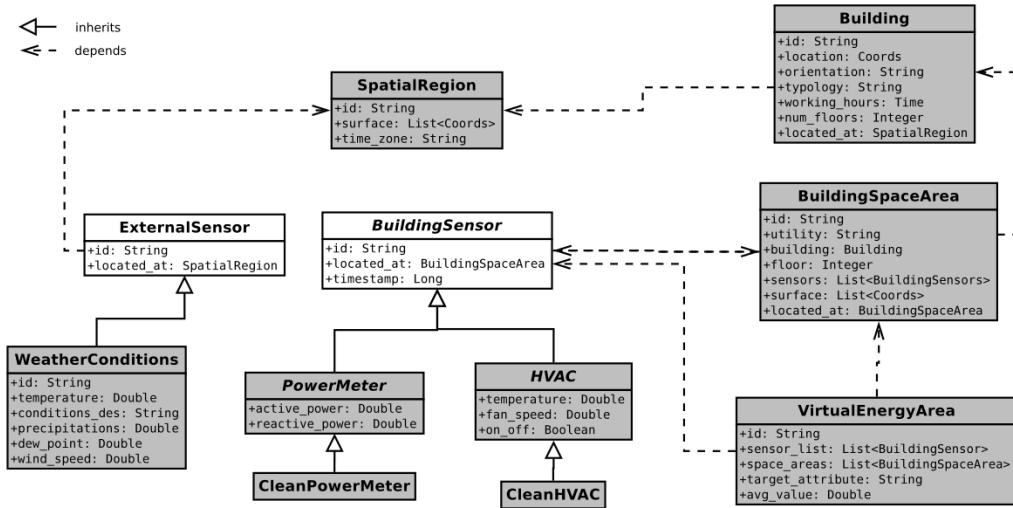


Fig. 2. IoTEP information model.

Figure 3.16: Information Model

The proposed platform that was created for creating applications and services for smart cities, and that especially covers the household energy domain and the provision of support for data analytics: sensorisation, homogenisation and storage, analytics and services is composed by several layers.

- Sensorisation layer

This layer is in charge of connecting physical devices or actuators that provide data to the platform. Then it maps the collected data to the NGSI entities of the information model using the FIWARE IoT Agent enabler²² and sends the mapped information to the next layer.

- Homogenisation and storage layer

²²<https://fiware-tutorials.readthedocs.io/en/latest/iot-agent/index.html>

3.2. DATA ANALYSIS IN IOT BASED SMART ENVIRONMENTS

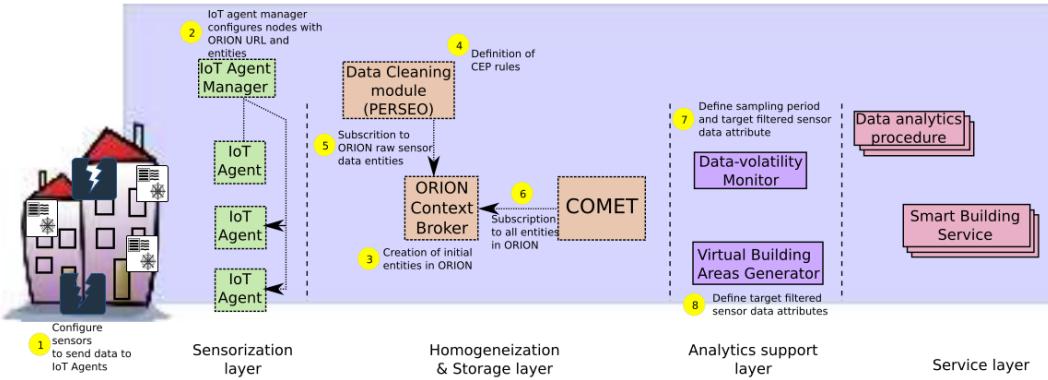


Figure 3.17: IoTEP workflow

This layer addresses the heterogeneity of the incoming data and contains real time data cleaning stage which ensures the quality of the data collected. Firstly, OCB implements a publish-subscribe store providing data access and the IoT Agents in the sensorization layer update the sensor entities' attributes in real time with the new readings from the devices. Secondly, the COMET enabler supports the access to historical time series data and incorporates an ad-hoc API to retrieve raw historical sensor data along with several built-in simple aggregation functions.

- **Analytics support layer**

The third layer embraces all the functionalities of the platform to provide support for data mining services that can run on top of the platform. In particular, we have included predictive ML algorithms (including the algorithms described in 3.2.3), an energy data volatility detector and a virtual entities generator (including the algorithms described in 3.2.5). The data volatility monitor detects either abnormal energy consumption related to building spaces or an abnormal temperature setting related to HVACs. An alarm is triggered when the current averaged value of an attribute differs substantially from a recent historic rate of change.

- **Service layer**

This layer serves as an interface between the IoTEP and the user, that could be anything from a building services manager to the back end of a smartphone application.

Smart-building services can be nested at this level of the IoTEP platform, and will allow features such as advanced HVAC predictive control, home automation, fuel poverty evaluation, sick building syndrome diagnostics, risk situations for vulnerable people (as in heat waves), smart tariff strategies, manage emergencies, saving energy and many others including results visualisation. These actions can either involve managers or be automatically set.

3.2.8 IoT mechanisms to provide personalized energy management and awareness services by analysing behavioural aspects related to energy efficiency [R8]

The combination of IoT technologies, data modelling, management and fusion, Big Data analytics, and personalized recommendation mechanisms has resulted in an open and extensible architectural approach able to exploit in a homogeneous, efficient and scalable way the vast amount of energy, environmental, and behavioural data collected in energy efficiency campaigns and lead to the design of energy management and awareness services targeted to the occupants' lifestyles. It is named ENTROPY platform.

Two actors are needed to start an energy efficiency campaign : managers and end users. Potential campaign managers are: administrators, energy efficiency experts, data scientists and behavioural scientists. They are responsible for setting up sensor data monitoring, data analysis, and personalized recommendation delivery processes, that will depend on the kind of building and the kind of users engaged. They are able to define the set of buildings along with their division in subareas and their characteristics (surface, working hours, location, etc.). Next, sensors per area are assigned and queries are designed. This data is used as input for data mining and analysis processes (specifying which algorithm, the required input and the desired output variable). The campaign administrator should be also in charge of undertaking corrective actions in general.

End users may consist of citizens, students, academic personnel, employees, etc. As a starting point, they get a profile by means of a questionnaire regarding its personality, work engagement, energy conservation habits, and game interaction preferences. Through the ENTROPY mobile applications and serious games they get information regarding energy consumption and environmental parameters in the areas that they have activities at and receive personalized recommendations and requests for action. At the end of an energy efficiency campaign, they fill in an evaluation questionnaire, targeting at measuring the perception of behavioural change, as well as any changes with regards to their gaming profile.

A high-level view of the ENTROPY energy-aware IT ecosystem architectural approach is provided at Fig. 3.18.

The IoT management and data aggregation layer is responsible for IoT nodes registration, sensor activation, management and data aggregation and cleaning functionalities at the edge part of the infrastructure.

The data representation and fusion layer represents data based on a set of defined semantic models [216], supports a set of data fusion mechanisms over active data streams and then stores them in the Big Data repository based on MongoDB. Upon the activation of a new sensor data stream, the manager denotes the mapping between the monitored sensor metric with the relevant parameter in the semantic model, supporting the unified access to the collected data.

The Energy Semantic Model is similar to the one described in 3.2.7 where entities related to

3.2. DATA ANALYSIS IN IOT BASED SMART ENVIRONMENTS

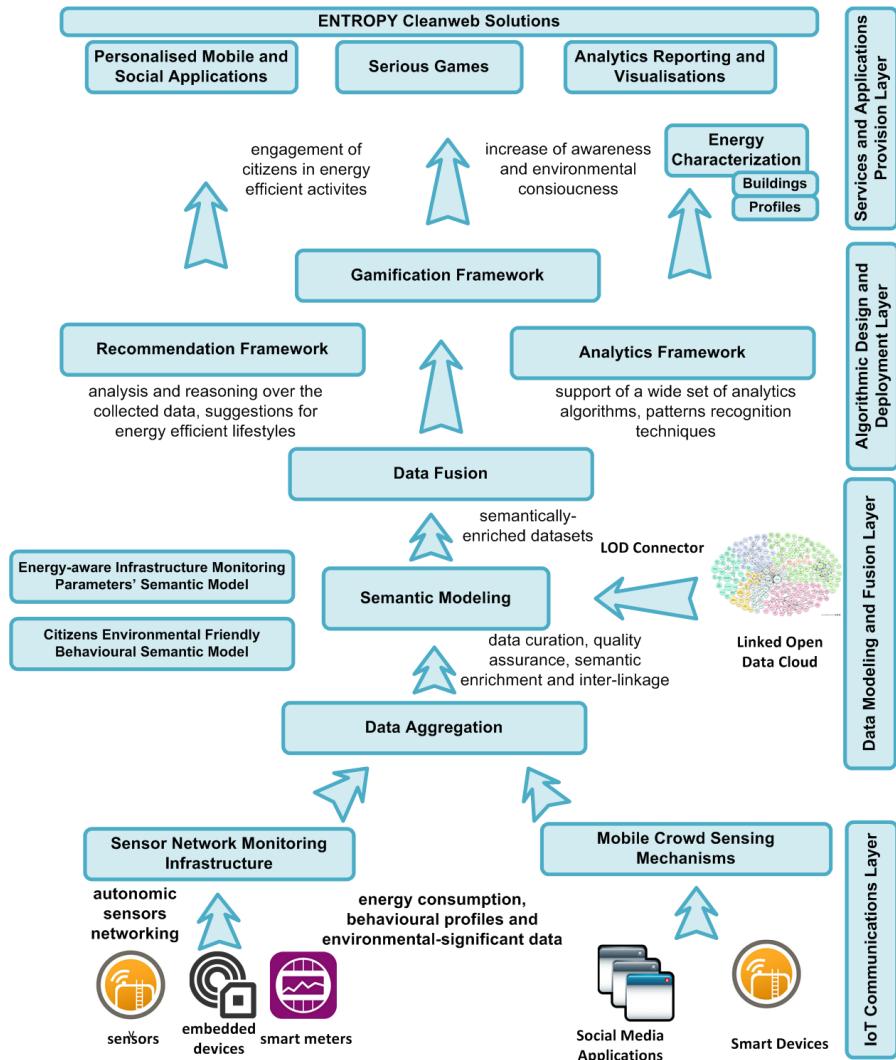


Figure 3.18: Entropy platform architecture

building areas, building spaces (windows, doors, etc.), equipment and outside sensors are defined and related to each other.

The Behavioural Semantic Model facilitates the categorization of users and the provision of personalized content and recommendations for achieving behavioural change. The main concepts regard the Agent and the Recommendation. An Agent can be a Person or a Group to whom personalized recommendations can be sent. It has several characteristics: personality traits such as extroversion and agreeableness; work engagement characterized by vigour, dedication, and absorption; gaming preferences such as socializer, free spirit, achiever and prize preferences such as rewards, badges, points, etc. Recommendations can have the form of a Message, a Quiz/Challenge, an Action whose result contributes to the elimination of a certain energy waste cause or a Question that leads to the collection of crowd-sensing feedback (e.g., comfort level).

The smart energy management services layer is responsible for providing advanced analytics

and recommendations to end users, as well as incorporating learning techniques for continuously exploiting the produced output by each service. These mechanisms work in a complementary fashion since the produced output from an analysis process can trigger the provision of a new recommendation.

A rule consists of a condition element which connects a context change with specific target user group criterion. When a rule is fired, the recommendation engine selects the set of target users based on their filters (location, responsive at the proposed actions through the personalized recommendations, etc) and creates a personalized recommendation for each of them by using the defined recommendation template. A produced recommendation contains the target user, the related content, the measurement attributes that are involved in the creation of the recommendation, the possible reward for the completion of the recommendation, as well as the validation method for it.

This layer also includes the support of a set of Big Data mining and analysis techniques towards the extraction of energy and behavioural analytics. An analysis process is based on the selection of an analysis template and the selection of the queries to be executed for providing the input datasets (training and/or evaluation datasets). Each analysis template represents a specific algorithm and provides the user with the flexibility to adjust the relevant configuration parameters. A set of initial algorithms are considered, however, the overall implementation facilitates the incremental addition of further analysis mechanisms.

The interconnection of the platform components with the analysis toolkits is based on the OpenCPU system for embedded scientific computing that provides a reliable and interoperable HTTP API for data analysis based on the R Project for Statistical Computing [217]. In the case of large-scale data processing and the need for a Big Data analysis framework, the Apache Spark engine is used, where the analysis process is realized in a set of worker nodes, each one of which is hosting an Apache Spark OpenCPU Executor²³.

The end user applications layer is responsible for the design of personalized mobile applications and web-based serious games able to take advantage of the set of services provided by the lower layers.

Some indicators for the evaluation of the different aspects of the platform are: energy savings and users behavioural change.

- Indicators of energy savings: savings at users' level, savings at areas' level, savings at buildings' level and savings extrapolation.
- Indicators of users' interactions: changes on behaviour based on data, self-awareness of change, changes in participants' perceived norm, changes on participants' personal values.
- Indicators of users' behavioural change: indicators of users' improvements on energy literacy, results from surveys and results from games.

²³<https://github.com/onetaqbeyond/opencpu-spark-executor>

There is a total of four buildings in the three pilots in which the platform is tested.

- A positive behaviour change was reported by the participants as they perceived it. Overall, the ENTROPY intervention led to an improvement on all the behavioural parameters we recorded, and especially on the participants' self-determination to conserve energy at work (+61.36%) and, also notably on the strength of their energy-saving competence (+14.35%), negative attitude towards saving energy at work (-13.77%), energy-saving at work as a habit (+13.60%), and the intention to save energy at work (+11.23%)
- Too much interaction does not lead to better results.

The best results were obtained in UMU, where the users' interaction with the apps was mid-level. In POLO, on the other hand, where the average interaction with the apps was by far the highest (compared to UMU and HESSO), the strength of behaviour change was not the highest. It was higher than in HESSO, where minimum interaction took place, but lower than in UMU, where the interaction was a fair amount but not as much as in POLO. A fair amount of interaction with the apps was the optimum remedy to effect behaviour change. Too little led to lower behavioural change, whereas too much interaction from a point on was in lines with "bombarding" the users with content from the apps, and overall led to less behaviour change.

The campaigns used a series of automatic triggers that allowed to identify specific behaviours that were not optimal in terms of energy use. Based on these triggers, the platform was able to tackle behaviours that resulted in energy waste. Also and in parallel with this more focused actions, energy literacy was improved in the participants as means of background-permanent improvement of the behaviour in terms of energy use. As the sensors continued providing information during and after the campaigns, it was possible to identify how effective the campaigns were. In addition to that, data from the interaction of the users was captured. This was of great use, as one can consider that engagement and effectiveness of energy behaviour advice can only work together, and engagement is driven by a variety of factors, some of which escape from the scope of this project.

Through several campaigns we obtained a 19.7% of savings in heating, 12.3% in cooling and 16.16 % in electricity in average for the 3 pilots.

3.3 Lessons Learned

The main driver of this work is real data. Data coming from sensors and Internet of Things devices capture the real dynamics of the environment and it has been shown that its analysis can provide the improvement of services and the creation of new ones.

In this thesis, we have developed a set of analytical tools that intend to provide cities with intelligence by means of data analysis, including Machine Learning techniques.

Realistic scenarios in which energy consumption prediction is useful are provided. Also, the development of methodologies for the proper use of algorithms depending on the forecasting needs, improving accuracy with respect to previous works. The best algorithm in accuracy and time was RF in the majority of the scenarios. We have also studied ways to preprocess inputs, including a thorough Feature Selection methodology that reduces computational time. We found multiple objective optimisation using MAE to be the most suitable tool for Feature Selection in energy consumption forecasting scenarios.

One of the main consumers in buildings are Heating, Ventilating and Air Conditioning systems. We were able to characterise people behaviour towards their use, finding two groups depending on how often they change their set point. Those that interact little with the controller use the equipment for more hours in the year. In contrast, those users that interact with the machines less often registered the larger overcooling. We were also able to define virtual areas depending on the temperature preferences using DBSCAN clustering. These results may serve well when defining actions towards users regarding controllers and thermal preferences.

Another very important factor in energy efficiency is human mobility. Our efforts on human mobility focused on finding Dense Transit Areas with GPS at a macro level and identifying routes and predicting location using social networks at both micro and macro level, using collective behaviour in personal models. These efforts are mainly applicable for intelligent public transports and traffic forecasting, event detection, and urban planning in general but can also be considered as a first step towards the linking of human mobility, occupation and efficiency in buildings.

The deployed platforms allowe to fuse data and support the integration of data mining procedures for the provision of final services for energy data mining. Data pre-processing, clustering and forecasting are integrated in the platform. This enables the development of more sophisticated energy-aware services. For example, through the ENTROPY mobile applications and serious games, users get information regarding energy consumption and environmental parameters in the areas that they have activities at and receive personalized recommendations and requests for action. Those are big steps towards a more efficient energy-literate society. Together with algorithms, we have studied human patterns and also we have provided educational tools that make possible the achievement of better services. This work favours not only the emergence of smart cities but also smarter citizens. Automation processes are of great interest for some scenarios, however, they can lead to misunderstanding of the surrounding environment and can also lead to lack of reaction under failures. At the end, combining behavioural analytic with technological advances and ML we aim to use our resources sustainably, specifically energy.

Finally, even though the majority of the thesis results are applications, the fundamental properties of Internet of Things data have also been investigated. A method for time series data representation named BEATS was developed. Data is segmented and represented in order to extract their key characteristics in lower-dimensionality. Time series are segmented and reduced at high rates when using overlapping windows. The independence between blocks that our

algorithm provides is one of its most important features, that it can be applied in a distributed, online manner. BEATS also presents other qualities such as adapting to drifts and low latency.

3.4 Conclusions and Future Work

The work on this thesis is focused on two problems: real data analysis and its use for improving energy efficiency. It was our goal to find and fill some of the gaps that prevent us from using the precious information hidden in data. The combination of fine data management with ML serves to extract knowledge that can be used for many problems in smart cities, especially improving and creating services regarding efficiency in buildings.

Using statistical analysis and FS techniques we have found that the most important variables for the problem of energy consumption prediction were temperature, radiation, occupation and previous values of consumption. In the collection of data, we see that radiation predictions are not always available so its use is restricted. Additionally, occupation is not always available and it is difficult to predict it accurately, so we proposed to analyze the patterns by differentiating between working hours and days and non-working days. The results show that there is an important difference between mornings and afternoons and also between all days of the week and Fridays in the studied buildings. Regarding occupation, we have also studied mobility patterns using wearable devices and social media. The extracted patterns could further be applied in the estimation of buildings occupation and its relation with consumption.

We have analysed several scenarios for energy consumption prediction and we have found that RF is an outstanding method in many of them, very appropriate due to its easy parallelisation. In order to obtain a horizon of predictions, multivariate time series methods also work well.

Results also show well-defined usage patterns of HVAC, one of the main contributors for energy consumption in buildings. The findings support the fact that there are two kinds of HVAC users: those who interact a lot with the set point and those who do not. The former appear to be higher consumers. These findings have been used to design strategies for energy consumption reduction.

We observed that the data that have been collected from real scenarios have several characteristics that pose challenges: their volume and their temporal nature. In that sense, we have also developed also a segmentation and representation algorithm called BEATS that transforms the data so that it provides similar amounts of information in a more compact manner. We have proved its effectiveness in classification and clustering problems using real data.

Finally, we developed IoT architectures that integrate all the steps developed in this thesis from the collection to the analysis of the data and even the provision of personalised services. Those services were designed for the improvement of energy efficiency in smart buildings targeting behavioural change and have proven to be useful for reducing energy consumption and improving energy literacy.

The following are some future research options that could be developed using the results of this thesis as a starting point both on the field of energy and data analysis:

- Going further than smart buildings, towards smart grids: integrating the created IoT platforms and predictive models into the novel smart grid scenarios

The connection of the monitoring and management IoT platforms to the plethora of energy-demanding devices, energy generating local systems and energy exchange platforms between all kinds of buildings, prosumers and companies needs to be accomplished. This targets ultimately the automated effective orchestration and guides the actuation and decision-making at all levels of the energy system: Transmission and Distribution System Operators (TSOs, DSOs), Energy Services Companies (ESCOs), prosumers and consumers.

- Cross-building knowledge transfer by using ML time series prediction techniques

All methods aiming to predict energy consumption that we have studied require labelled data, such as historical data. Such labels and datasets are not always available and it takes a long time and effort to collect, clean and manage them. At this point, data are employed in these forecasting methods in a non-adaptable way since they are completely static and thus without considering future events or changes which can occur in the network and the infrastructure of the buildings.

It is not always economically feasible or possible time-wise to develop an IoT infrastructure in all buildings. In that sense, the need of unsupervised methods for energy consumption prediction is evident, especially after showing what it is possible to do with this information.

Reducing the total time required for the analytic procedure is the key to scaling the deployment of energy efficiency projects in general, and reducing overall costs [218]. It is for this reason that a transfer learning approach should be considered in future studies in order to reduce the quantity of data that needs to be collected to create a reasonable building model.

- Finding further scenarios and ways to apply BEATS

The development of the time series representation method BEATS is an important achievement of this work that can be further explored in the following ways:

- Adding 3-dimensional (3D) data to the possible inputs and studying the modification of BEATS by substituting Discrete Cosine Transform by its 3D version.
- Considering multi-sensor data. So far, BEATS is applied on each sensor in order to represent data in lower dimensions.
- Studying the possibility to apply BEATS for dimensionality reduction by using the obtained eigenvectors to project the data in a similar fashion than to Principal Component Analysis (PCA).

3.4. CONCLUSIONS AND FUTURE WORK

- Connecting mobility results with energy consumption

The investigation of the effects of urban mobility on energy consumption of specific urban areas is another line that can be derived from this thesis. That is, developing a deeper understanding of whether a similar spatial dependency exists in human mobility as an indicator for urban consumption of energy.

- Extend the context in which to apply the developed methods

The application of the forecasting, feature selection and time series representation methods that have been developed in this thesis can be of great interest in further contexts than smart buildings and energy. Also, the developed platforms can serve for the connection between analysis and services in other areas such as smart agriculture, water management, etc.

For example, smart agriculture scenarios are emerging and for the discovery of useful trends and patterns, there is a need to work on large sets of data obtained across multiple farms. After collecting more data and measurement about the production: soil quality, irrigation levels, weather, presence of insects and pests, its fusion using our IoT platform and the analysis through techniques here developed can serve of great help for the realisation of a smarter agriculture and livestock farming.

PUBLICATIONS COMPOSING THE PHD THESIS

4.1 BEATS: Blocks of Eigenvalues Algorithm for Time Series Segmentation

Title	BEATS: Blocks of Eigenvalues Algorithm for Time Series Segmentation
Authors	Aurora González-Vidal, Payam Barnaghi, and Antonio F. Skarmeta
Type	Journal
Journal	IEEE Transactions on Knowledge and Data Engineering
Impact factor (2018)	3.857
Rank	Q1
Publisher	IEEE
Volume	30
Issue	11
Pages	2051-2064
Year	2018
Month	March
ISNN	1041-4347 (Print), 1558-2191 (Electronic)
DOI	10.1109/TKDE.2018.2817229
URL	https://ieeexplore.ieee.org/document/8319952/
State	Published
Author's contribution	The PhD student, Aurora González Vidal, is the main author of the paper

BEATS: Blocks of Eigenvalues Algorithm for Time Series Segmentation

Aurora González-Vidal[✉], Payam Barnaghi[✉], Senior Member, IEEE,
and Antonio F. Skarmeta[✉], Member, IEEE

Abstract—The massive collection of data via emerging technologies like the Internet of Things (IoT) requires finding optimal ways to reduce the observations in the time series analysis domain. The IoT time series require aggregation methods that can preserve and represent the key characteristics of the data. In this paper, we propose a segmentation algorithm that adapts to unannounced mutations of the data (i.e., data drifts). The algorithm splits the data streams into blocks and groups them in square matrices, computes the Discrete Cosine Transform (DCT), and quantizes them. The key information is contained in the upper-left part of the resulting matrix. We extract this sub-matrix, compute the modulus of its eigenvalues, and remove duplicates. The algorithm, called BEATS, is designed to tackle dynamic IoT streams, whose distribution changes over time. We implement experiments with six datasets combining real, synthetic, real-world data, and data with drifts. Compared to other segmentation methods like Symbolic Aggregate approXimation (SAX), BEATS shows significant improvements. Trying it with classification and clustering algorithms it provides efficient results. BEATS is an effective mechanism to work with dynamic and multi-variate data, making it suitable for IoT data sources. The datasets, code of the algorithm and the analysis results can be accessed publicly at: <https://github.com/auroragonzalez/BEATS>.

Index Terms—BEATS, SAX, data analytics, data aggregation, segmentation, DCT, smart cities

1 INTRODUCTION

LESS than 1 percent of the data that are nowadays captured, stored, and managed by means of the Internet of Things (IoT) and Big Data technologies is being analysed [1]. There exist several challenges in the analysis of data such as high dimensionality, high volume, noise, and data drifts. Data provided by IoT sources (sensory devices and sensing mechanisms) are multi-modal and heterogeneous. Since all of the above mentioned features hinder the execution and generalization of the algorithms, many higher-level representations or abstractions of the raw data have been proposed to address these challenges.

In this paper, we attempt to aggregate and represent large volumes of data in efficient and higher-granularity form. The latter is an attempt to create sequences of patterns and data segments that occur in large-scale IoT data streams. The contribution of our approach is to do such representation on-the-fly since usually data treatment has to be done very quickly, adapting to unpredictable changes in the data or even without prior knowledge.

A use case where large and dynamic datasets are present is smart cities. Data aggregation and pattern representation enables us to find underlying patterns, providing further understanding of *the city data*. Big Data analytics, machine learning and statistical techniques are used to predict, classify and extract information that empowers machines with decision-making capabilities.

IoT data is usually related to physical objects and their surrounding environment. Normally, IoT data is collected together with a timestamp. The collection of several points spaced in time, having a temporal order is known as time series data. Time series can be analysed using various techniques such as clustering, classification and regression (as inputs of models) in the fields of data mining, machine learning, signal processing, communication engineering, and statistics.

Our proposed method is based on splitting time series data into blocks. These blocks can be either overlapping or non-overlapping and they represent subsets of the whole data structure. The method synthesizes independently the information that the blocks contain. It reduces the data points while still preserving their fundamental characteristics (losing as little information as possible). We propose a novel technique using matrix-based data aggregation, Discrete Cosine Transform (DCT) and eigenvalues characterization of the time series data. The algorithm is called Blocks of Eigenvalues Algorithm for Time series Segmentation (BEATS). We compare BEATS with the state-of-the art segmentation and representation algorithms. We also compare and evaluate the approaches in two of the most common machine learning tasks, classification and clustering, by comparing metrics between each of the transformed datasets. We also present a

- A. González-Vidal and A. F. Skarmeta are with the Department of Information and Communications Engineering, University of Murcia, Murcia 30100, Spain. E-mail: {aurora.gonzalez2, skarmeta}@um.es.
- P. Barnaghi is with the Institute for Communication Systems, University of Surrey, Guildford, Surrey GU2 7XH, United Kingdom. E-mail: p.barnaghi@surrey.ac.uk.

Manuscript received 21 June 2017; revised 21 Jan. 2018; accepted 9 Mar. 2018. Date of publication 19 Mar. 2018; date of current version 4 Oct. 2018. (Corresponding author: Aurora González-Vidal.)

Recommended for acceptance by E. Terzi.

For information on obtaining reprints of this article, please send e-mail to: reprints@ieee.org, and reference the Digital Object Identifier below. Digital Object Identifier no. 10.1109/TKDE.2018.2817229

use case that is related to smart cities showing the suitability of BEATS for real time data stream analysis. This is shown by explaining how to apply it within a Big Data framework.

The remainder of the paper is organized as follows: Section 2 describes the related work. Section 3 motivates the need of a new approach. Section 4 details the algorithm and briefly explains the mathematical background of the work. Section 5 includes the evaluations in several scenarios using different datasets and a use-case related to smart cities. Section 6 discusses the results of the experiments and Section 7 concludes the paper and describes the future work.

2 RELATED WORK

There are several approaches to represent a numeric time-dependent variable (i.e., a time series). The most basic one is to compute the mean and standard deviation among other statistical measures (e.g., variance, mode). Using those statistics it is not possible to represent all the information that the time series contains. A classical example that supports this claim is the Anscombe's Quartet, [2] that shows how four very different datasets have identical simple statistical properties: mean, variance, correlation and regression coefficients.

In order to reduce the number of data points in a series and create a representation, segmentation methods can be used as a pre-processing step in data analytics.

Definition 1 (Segmentation). *Given a time series T containing n data points, segmentation is defined as the construction of a model \bar{T} , from l piecewise segments ($l < n$) such that \bar{T} closely approximates T [3].*

The segmentation algorithms that aim to identify the observation where the probability distribution of a time series changes are called change-point detection algorithms. Sliding windows, bottom-up, and top-down methods are popular change-point detection based approaches. For sliding windows, each segment is grown until it exceeds an error threshold. The next block starts with the new data point not included in the newly approximated segment and so on. In the bottom-up methods, the segments of data are merged until some stopping criteria is met and top-down methods partition the time series recursively until a stopping criteria is met [4].

Another way of classifying the algorithmic methods for segmentation is considering them as online and offline solutions [5]. While offline segmentation is used when the entire time series is previously given, the online segmentation deals with points that arrive at each time interval. In offline mode, the algorithm first learns how to perform a particular task and then it is used to do it automatically. After the learning phase is completed, the system cannot improve or change (unless we consider incremental learning or retraining). On the other hand, online algorithms can adapt to possible changes in the environment. Those changes are known as "drifts". Whereas top-down and bottom-up methods can only be used offline, sliding windows are applicable to both circumstances.

After segmentation, the representation of the time series based on the reduction can be regarded as an initial step that reduces the load and improves the performance of

tasks such as classification and clustering. The use of such algorithms can be generally regarded in two ways:

- Representation methods: Extracting features from the whole time series or its segments and applying machine learning algorithms (Support Vector Machines, Random Forest, etc) in order to classify them or compute the distance between the time series representation for clustering.
- Instanced based methods (similarities): Computing the distance matrix between the whole series and using it for clustering or classification applying a k-nearest neighbour approach [6] by finding the most similar (in distance) time series in the training set.

BEATS is based on the first perspective since as stated in *Bagnall et al The greatest improvement can be found through choice of data transformation, rather than classification algorithm* [7]. However, we review the work made using both approaches since the ultimate goal of our time series representation is to make the time series data more aggregated and better represented for further processing.

2.1 Whole Series Similarities

Similarity measures are used to quantify the distance between two raw time series. The list of approaches is vast and the comparison between well-known methods has lead to the conclusion that the benchmark for classification is dynamic time warping (DTW) since other techniques proposed before 2007 were found not significantly better [8].

Similar results have been stated in [9] when comparing DTW with more recent distance measures as: Weighted DTW [10], Time warp edit (TWE) [11] and Move-split-merge (MSM) [12] together with a slight accuracy improvement (1 percent) when using Complexity invariant distance (CID) [13] and Derivative transform distance (DTD_C) [14].

When computation time is not a problem, the best approach is to use a combination of nearest neighbour (NN) classifiers that use whole series elastic distance measures in the time domain and with first order derivatives: Elastic ensemble (EE) [15]. However, if a single measure is required a choice between DTW and MSM is recommended, with MSM preferred because of its overall performance.

In the clustering domain, the number of evaluated similarity distances is even higher, due to the nature of the problem. An extensive description of similarity measures can be found in [16]. DTW and CID are also used in clustering the raw time series [17], [18].

2.2 Intervals

Various algorithms focus on deriving features from intervals of each series. For a series of length m , there are $m(m - 1)/2$ possible contiguous intervals.

Piecewise Linear Representation (PLR) [19] methods are based on the approximation of each segment in the form of straight lines and include the perceptually important points (PIP), Piecewise Aggregate Approximation (PAA) [20], and the turning point (TP) method [21].

The state-of-the-art models Time Series Forest (TSF) [22] and Learned pattern similarity (LPS) [23] generate many different random intervals and classifiers on each of them, ensembling the resulting predictions.

TSF trains several trees in a random forest fashion but each tree uses as data input the $3\sqrt{m}$ statistics features (mean, standard deviation and slope) of the \sqrt{m} randomly selected intervals.

LPS can be regarded as an approximation of an autocorrelation function. For each series, they generate a random number l of series by randomly selecting a fixed number w of elements of the primitive one. A column of the generated $l * n \times w$ matrix is chosen as the class and a regression tree is built (autocorrelation part). After that, for every series the number of rows of the matrix (originated by the raw series) that reside in each leaf node is counted. Concatenating these counts the final representation of the series is formed. Then, a 1-NN classifier is applied to process the time series data.

2.3 Symbolic Aggregate Approximation (SAX)

Among all the techniques that have been used to reduce the number of points of a time series data, SAX has specially attracted the attention of the researchers in the field. SAX has been used to asses different problems such as finding time series discords [24], finding motifs in a database of shapes [25], and to compress data before finding abnormal deviations [26] and it has repeatedly been enhanced [27], [28], [29].

SAX allows a time series of length n to be reduced to a string of length l ($l < n$). The algorithm has two parameters: window length w and alphabet size α , and it involves three main steps [30]:

- Normalization: standardizes the data in order to have a zero mean and a standard deviation of one;
- Piecewise Aggregation Approximation (PAA): divides the original data into the desired number of windows and calculates the average of data falling into each window; and
- Symbolization: discretizes the aggregated data using an alphabet set with the size represented as an integer parameter α , where $\alpha > 2$.

As normalized time series data assumes a Gaussian distribution for the data, the discretization phase allows to obtain a symbolic representation of the data by mapping the PAA coefficients to a set of equiprobable breakpoints that are produced according to the alphabet size α . The breakpoints determine equal-sized areas under the Gaussian curve [31] in which each area is assigned to an alphabet character.

Since SAX representation does not consider the segment trends, different segments with similar average values may be mapped to the same symbols. Among the multiple enhancements done to SAX (see related work section of [28] and [29]) we highlight the following works:

- Extended SAX (ESAX) [27]: adds maximum and minimum along with the original SAX representation.
- SAX Trend Distance (SAX_{TD}) [28]: defines the trend distance quantitatively by using the starting and ending point of the segment and improved the original SAX distance with the weighted trend distance.
- SAX with Standard Deviation (SAX_{SD}) [29]: adds the standard deviation of the segment to its SAX representation.

The Vector Space Model (VSM) is combined with SAX in [32] in order to discover and rank time series patterns by

their importance to the class. Similarly to shapelets, SAX-VSM looks for time series subsequences which are characteristic representatives of a class. The algorithm converts all training time series into bags of SAX words and uses *tf-idf* weighting and cosine similarity in order to rank by importance the subsequences of SAX words according to the classes.

2.4 Shapelets

Shapelets are subsequences of time series that identify with the class that the time series belongs to.

The Fast shapelets (FS) [33] algorithm discretises and approximates shapelets using SAX. The dimensionality of the SAX dictionary is reduced through masking randomly selected letters (random projection).

Learned shapelets (LS) [34] optimizes a classification loss in order to learn shapelets whose minimal euclidean distances to the time series are used as features for a logistic regression model. An improvement of such model is the use of DTW instead of euclidean distance [35].

The Fused Lasso Generalized eigenvector method (FLAG) [36] is a combination of the state-of-the-art feature extraction technique of generalized eigenvector with the fused LASSO that reformulates the shapelet discovery task as a numerical optimization problem instead of a combinatorial search.

Finally, we take into consideration the clustering algorithm k-shape [37], a centroid-based clustering algorithm that can preserve the shapes of time-series sequences. They capture the shape-based similarity by using a normalized version of the cross-correlations measure and claims to be the only scalable method that significantly outperforms k-means.

2.5 Ensembles

So far we have reviewed how data transformation techniques are applied to different algorithms in order to improve their accuracy and to reduce the computation time. COTE algorithm [38] uses a collective of ensembles of classifiers on different data transformations.

The ensembling approach in COTE is unusual because it adopts a heterogeneous ensemble rather than resampling schemes with weak learners. COTE contains classifiers constructed in the time, frequency, change (autocorrelations), and shapelet transformation domains (35 in total) combined in alternative ensemble structures. Each classifier is assigned a weight based on the cross validation training accuracy, and new data are classified with a weighted vote.

The results of evaluations in COTE show that the simple collective formed by including all classifiers in one ensemble is significantly more accurate than any of its components.

3 MOTIVATION AND CONTRIBUTIONS

As it can be seen among the segmentation techniques that we referenced in section 2, we have mentioned not only the representation techniques but also how the whole classification and clustering procedure is performed by combining representation with machine learning algorithms. We intended to show that our representation method is an efficient alternative segmentation method to be employed in time series data processing.

One commonality of the several studies that we have reviewed is that most of the existing algorithms use normalization that re-scales the data.

However, there are few studies that do not apply re-scaling and normalization. BEATS uses a non-normalized algorithm for constructing the segment representation.

The concept *drift* appears when a model built in the past is no longer fully applicable to the current data. Concept drift is due to a change in the data distribution according to a single feature, to a combination of features or in the class boundaries, since the underlying source generating the data is not stationary.

The potential changes in the data might happen in:

- The prior probability $P(y_i)$;
- The conditional probability $P(x|y_i)$;
- The posterior probability $P(y_i|x)$; and
- A combination of the above.

Where x is the predicted data and y_i is the observed data.

These changes can cause two kinds of concept drift: real and virtual [39].

If only the data distribution changes without any effect on the output, i.e., changes in $P(y_i)$ and/or $P(x|y_i)$ that does not affect $P(y_i|x)$, it is called virtual drift.

When the output, i.e., $P(y_i|x)$, also changes it is called real concept drift.

In the IoT domain and especially in smart city data analysis, we are interested in the second type of drift which will be referred as *data drift* in this paper [40]. Some examples where a data drift may occur in smart cities are related to the replacement of sensors (different calibration), sensor wear and tear [41] or drastic changes to the topics of discussion in social media used for crowdsensing [42].

There are several existing methods and solution addressing the concept drift for supervised learning [41], and some recent ones also for unsupervised learning [40]. However, we focus on the initial step of the analysis (i.e., pre-processing). We claim that not only the model has to be adaptive but also the way that we segment the inputs has to take into account the dynamics of the data and be able to efficiently deal with the changes in the structure of the data.

A considerable challenge in segmentation is to find a common way to represent the data. This is due to the variety of ways to formulate the problem in terms of defining the key parameters (number of segments, segmentation starting point, length of segments, error function, user-specified threshold, etc.).

The first step in SAX algorithm is assuming that for a particular problem that we deal with, the data follows a normal distribution or at least we have a sufficiently large number of samples in order to say that the distribution of the data is approximately normal, appealing to the central limit theorem [43]. Nevertheless, this is a strong assumption because there are many scenarios in which this might not be the case; for example:

- Outliers and noise: data from physical devices usually contains noise and outliers that affect the identification of the correct parameters of the distribution.
- Data follows different distribution.

- Fast data: two of the V's from the 7V's Big Data challenges [44] are *velocity* and *variety*. Traditionally in data mining, batch data is processed in an offline manner using historical data. However, in IoT applications we need to consider short-term snapshots of the data which are collected very quickly. Thus, we need adaptive methods that catch up with the changes during their operation.

All mentioned algorithms lack of at least one of such 3 problems too. We have developed an algorithm that does not require normalization of the data. The latter will also help to preserve the value of the data points (i.e., magnitude of the data). The lack of sensitivity to magnitude in the algorithms that make assumptions about the normalized distribution and use Z-normalization makes them less efficient in analysing correlation and regression. Another requirement is the application of the algorithm in an online way and using sliding windows. Nonetheless, we have to be able to compute the distance between the aggregated time series. Considering these requirements we have designed the BEATS algorithm.

4 BEATS PRESENTATION

This section describes our proposed algorithm and discusses its mathematical and analytical background. We present BEATS and show the effect of each step of the algorithm in a block of data.

4.1 BEATS Construction

Transforms, in particular integral transforms, are used to reduce the complexity in mathematical problems. In order to decorrelate the time features and reveal the hidden structure of the time series, they are transformed from the time domain into other domains. Well-known transformations are the Fourier Transform, which decomposes a signal into its frequency components, and the Karhunen-Loeve Transform (KLT) which decorrelates a signal sequence.

Discrete Cosine Transform (DCT) is similar to Discrete Fourier Transform (DFT) but uses cosines obtained from the discretization of the kernel of the Fourier Transform. DCT transfers the series to the frequency domain. Among the four different cosine transformations classified by Wang [45], the second one (i.e., DCT-II) is regarded as one of the best tools in digital signal processing [46] (times series can be regarded as a particular case of signals). Due to its mathematical properties such as unitarity, scaling in time, shift in time, the difference property, and the convolution property, DCT-II is asymptotically equivalent to the KLT where under certain (and general) conditions KLT is an optimal but impractical tool to represent a given random function in the mean square error sense (MSE). KLT is said to be an optimal transform because:

- It completely decorrelates the signal in the transform domain;
- It minimizes the MSE in bandwidth reduction or data compression;
- It contains the most variance (energy) in the fewest number of transform coefficients; and
- It minimizes the total representation entropy of the sequence.

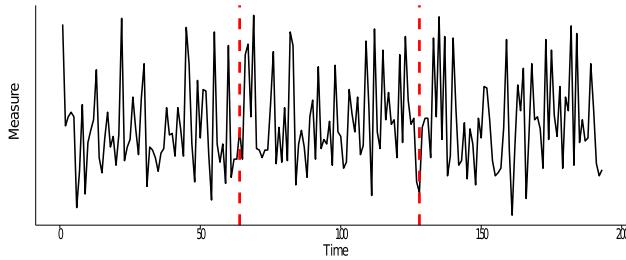


Fig. 1. An example of a time series divided into blocks of 64 observations.

The details of the proof of the above statements can be found in [46]. Understanding the properties of the DCT, we use it to transform our time series data.

We apply the transformation essentially by using the compression of a stream of square 8×8 blocks, taking reference from the standards in image compression [47] where DCT is widely used (e.g., JPEG). Since 8 is a power of 2, it will ease the performance of the algorithm.

As an illustration, we provide an example. We have divided the time series shown in Fig. 1 as blocks of 64 observations that are shown using a dashed red line. If we arrange the first block row-wise into a squared matrix M , we can visualize that the information is spread through the matrix as the heatmap shown in Fig. 2.

It should be noted that while our raw time series data is represented in value/time, a 2D transformation is applied to the data. This is based on the assumption that in each block, the neighbour values of a selected observation m_{ij} (eg. $m_{i-1,j}, m_{i,j-1}, m_{i-1,j-1}$) are correlated. In time series with very rapid changes in the data, small block sizes will be more suitable and if the changes are not very rapid size block can be larger. In this paper, we use a common 8×8 block size for our description.

Intuitively, each 8×8 block includes 64 observations of a discrete signal which is a function of a two-dimensional (2D) space. The DCT decomposes this signal into 64 orthogonal basis signals. Each DCT coefficient contains one of the 64 unique *spatial frequencies* which comprise the *spectrum* of the input series. The DCT coefficient values can be regarded as the relative amount of the spatial frequencies contained in the 64 observations [47].

Let M be the 8×8 input matrix. Then, the transformed matrix is computed as $D = UM^T$, where U is an 8×8 DCT matrix. U coefficients for the $n \times n$ case are computed as shown in Eq. 1:

$$U_{ij} = \begin{cases} \frac{\sqrt{2}}{2} & i, j = 1 \\ \cos\left(\frac{\pi}{n}(i-1)(j-\frac{1}{2})\right) & i, j > 1 \end{cases} \quad (1)$$

The formula of Eq. (1) is obtained using Eq. (5) (Appendix 8). Finally, we multiply the first term by $\frac{1}{\sqrt{2}}$ in order to make the DCT-II matrix orthogonal. After applying DCT, the information is accumulated in its upper-left part, as it is shown in the heatmap in Fig. 3.

Each of the 64 entries of the matrix D is quantized by pointwise division of the matrices D and Z , where the elements of the quantization matrix Z are integer values ranging from 1 to 255.

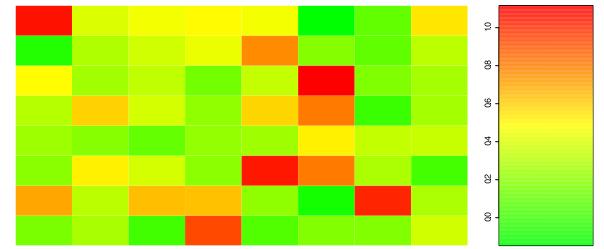


Fig. 2. The heatmap of the matrix obtained from the first block of time series data.

Quantization is the process of reducing the number of bits needed to store an integer value by reducing the precision of the integer. Given a matrix of DCT coefficients, we can divide them by their corresponding quantizer step size and round it up depending on its magnitude, normally 2 decimals. If the maximum of the DCT matrix is small, the number of decimals is selected by the operation $\lfloor \log_{10} \max \rfloor - 4$, where $\lfloor \log_{10} \max \rfloor$ returns the position of the first significant figure of the maximum number in the transformed matrix D . This step is used to remove the high frequencies or to discard information which is not very significant in large-scale observations.

The selected matrix Z is the standard quantization matrix for DCT [48].

After the quantization process, a large number of zeroes appears in the bottom-right position of the matrix $Q = \frac{D}{Z}$, i.e., it is a sparse matrix.

We extract the 4×4 upper-left matrix that contains the information of our 64 raw data and compute the eigenvalues, which in our case are: $0.18605, 0.02455, 0.00275 + 0.00843i, 0.00275 - 0.00843i$.

Using BEATS so far we have significantly reduced the number of points of our time series from 64 to 4 but we have also converted its components into complex numbers. These complex numbers (eigenvalues vector) represent the original block in a lower dimension. This eigenvalues vector is used in BEATS to represent the segments and hence, it is the potential input for the machine learning models. However, it is not always possible to feed machine learning algorithms with complex numbers and the eigenvalues could be complex numbers. To solve this problem, we compute the modulus of the eigenvalues and remove the repeated ones (they are presented in pairs so the information would be repeated).

In case that there are no complex numbers in the output of BEATS, we will conserve the first three values, since the latter values are sorted in a descending order. This means that we have represented the original 64 observations as

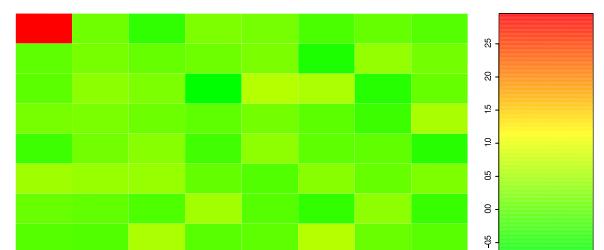


Fig. 3. The heatmap of the DCT matrix.

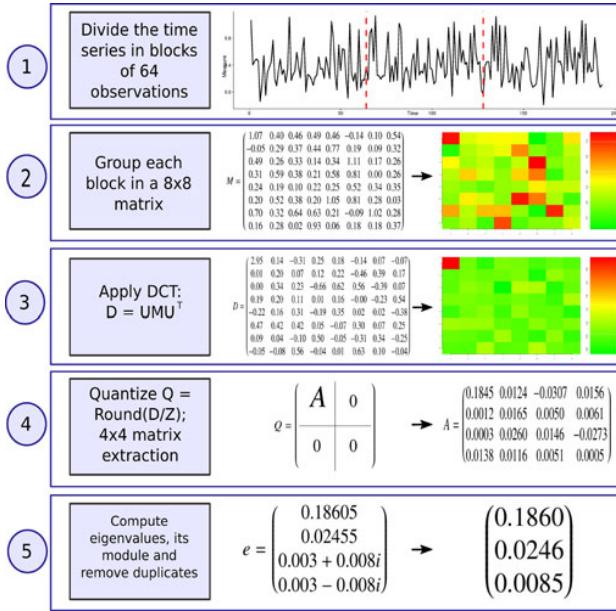


Fig. 4. BEATS is shown step by step with an example.

three values. In our example, the final representation (modulus of the eigenvalues) consists of 0.1860, 0.0246, 0.0085.

The BEATS process is summarized in Fig. 4.

We also consider the relevance of the direct computation of the eigenvalues of the 8×8 matrix M in order to assure that the DCT and its quantization contribute to the aggregation of the information. We refer to this method throughout the paper as Eigen.

4.2 Complexity Analysis of BEATS

The time complexity is represented as a function of the input time series size (n). Regarding the different steps of BEATS, the processes that have a key impact on the run time are DCT, which is a double matrix multiplication, i.e., $O(n^3)$; pointwise matrix division for the quantization, i.e., $O(n^2)$ and eigenvalue computation, i.e., $O(\beta^3)$, where n is the size of the matrix block (square root of the amount of data that compounds each block), and $\beta (\leq n)$ is the size of the extracted matrix from which we compute the eigenvalues. Although we have set the values to $n = 8$ and $\beta = 4$, we compute the complexity in general terms.

So far, the dominant task regarding the complexity is the DCT function. For about the past 40 years, many fast algorithms have been reported to enhance the computation of discrete cosine transforms [49]. In order to improve the efficiency of the algorithm, we have implemented a popular way of computing the DCT of our N -points time series. We use a $2N$ -points Fast Fourier Transform (FFT). This has reduced the complexity to $O(n^2 \log(n))$ [50].

Hence, for each block we have a complexity of $O(n^2 \log(n) + \beta^3)$. Let N be the size of our time series data; if we do not use sliding windows, we will apply the algorithm $\frac{N}{n \times n}$ times, so the complexity is $\frac{N}{n \times n} O(n^2 \log(n) + \beta^3)$. As we can see, the complexity of the algorithm grows linearly depending on the number of blocks where we have to apply the computations.

By applying multiple processing architectures, the complexity problem nowadays can also depend on how efficiently

we can parallelize the processing load. Parallelising the BEATS algorithm is very simple since the computations are *block dependent* and no information out of the block is required for each individual calculation. This makes the process ideal to be done using graphics processing units (GPUs), and thereby minimising the latency of the computation.

5 EXPERIMENTAL EVALUATION

We perform two data mining processes: classification and clustering. Following our approach the data is going to be transformed by the two methods: BEATS and Eigen, summarized as follows:

- BEATS: 8×8 matrix blocks of the data, discrete cosine transformation, and quantization of each of the matrices, reduction to a 4×4 matrix, removal of the duplicated modulus of the complex eigenvalues and selection of the first three values.
- Eigen: 8×8 matrix blocks of the data, computation of the eigenvalues of the matrices, removal of the duplicated modulus of the complex eigenvalues, and selection of the first three values.

Having introduced several algorithms in Section 2, we compare BEATS and Eigen with common existing state-of-the-art methods that show an improvement in comparison with the primitive ones.

The algorithms' code has been accessed from the authors' public repositories when available. When not, R software and Python have been used in order to program them.

We perform each of the techniques using several datasets in order to analyse the type of problems that our algorithm performs better than other methods. It is possible to use sliding windows for our method. In the experiment, we consider a slide of 8 observations. The evaluations also include a cross validation step in order to find their parameters.

A smart cities use case where we cluster traffic data is also presented. The intention is to see how BEATS is suitable for different scenarios including online smart cities applications.

5.1 Datasets

We give a short explanation of the datasets that are used to evaluate the algorithm. Four of the datasets are obtained from the UCR Time Series Classification Archive [51]: ArrowHeads, Coffee, FordA, Lightning7 and ProximalPhalanxOutlineAgeGroup. For each dataset we use, when provided, the train sample in order to find the hyperparameters of the model and then, we test their classification performance with the test set. For clustering we use only the training set. When the split is not provided, which is the case in one of the datasets (the randomly generated by us), we use 75 percent of the samples for the training set and 25 percent of the samples for testing.

The datasets that are used in the experiments are briefly described below.

Arrow Heads (Real and Without Drifts). The Arrow Heads dataset¹ contains 211 series having 192 observations classified into three different classes. The arrowhead data consists

1. http://www.cs.ucr.edu/~eamonn/time_series_data/

of outlines of the images of arrowheads [52]. The shapes of the projectile points are converted into a time series using the angle-based method and they are classified based on shape distinctions such as the presence and location of a notch in the arrow. The classification of projectile points is an important topic in anthropology. According to our method, we reduced the dataset to 72 observations.

Lightning7 (Real and Long). We use the Lightning7 dataset that gathers data related to transient electromagnetic events associated with the lightning natural phenomenon. Data is gathered with a satellite with a sample rate of 800 microseconds and a transformation is applied in order to produce series of length 637.

The classes of interest are related to the way that the lightning is produced.²

Initially, each measurement (time series) carries 320 variables. Using our method, we have reduced the dataset to 96 variables.

Random LHS Generator Lift (Synthetic and with Drifts). A dataset with data drifts is also used in our experiments. In this case, we have evaluated the algorithms with the data generated by using the code from the Repository³ described in [53], which was first used in [40]. The drift is introduced both by shifting the centroids in randomized intervals and by changing the data distribution function used to randomly draw the data from the centroids that are selected through Latin Hypercube Sampling (LHS). This dataset is created for smart cities data analysis and allows to create sample datasets that simulate dynamic and multi-variate data streams in a smart environment. The data generator is developed in the context of the CityPulse smart city project.⁴

The number of centroids is set to ten and we generated 300 series that follow three different distributions (triangular, Gaussian and exponential). Initially, each set (time series) carries 192 variables. Using our method, we reduced the dataset to 51 variables.

Coffee (Real-World Data). The Coffee dataset¹ contains 56 series having 286 observations classified into two different classes. The Coffee data consists of the series generated by the Fourier transform infrared spectroscopy of two species of coffee: Arabica and Robusta. Originally, such method intended to serve as an alternative to wet chemical methods for authentication and quantification of coffee products [54]. Using BEATS, we reduced the dataset to 57 observations which represent the patterns that occur in the dataset. This can be used for further analysis and classification of coffee types.

FordA (Real-World Data). The FordA dataset¹ contains 4921 series having 500 observations each classified into two different classes. The data was generated on the context of a classification competition. The problem is to diagnose whether a certain symptom exists in a automotive subsystem using the engine noise as a measurement. Both training and test data set were collected in typical operating conditions, with minimal noise contamination. Using BEATS, we reduced the dataset to 100 observations. The BEATS observations are

more resilient to noise and provide an efficient way to discover and extract patterns from real-world raw data.

ProximalPhalanxOutlineAgeGroup (Real-World Data from Images). The ProximalPhalanxOutlineAgeGroup dataset¹ contains 605 series having 80 observations each classified into three different classes. The dataset was created [55] for testing the efficacy of hand and bone outline detection and whether these outlines could be helpful in bone age prediction. The problem involves using the outline of one of the phalanges of the hand in order to predict whether the subject is one of three age groups. Using BEATS, we reduced the dataset to 9 observations per subject. This observations provide a reduced feature set that ease the analysis tasks.

5.2 Classification

Classification of time series analysis is a classic problem consisting of building a model based on labelled time series data and using the model to predict the label of unlabelled time series samples.

The applications of this technique are widely extended in many areas, ranging from epilepsy diagnosis based on time series recorded by electroencephalography devices (electrical activity generated by brain structures over the scalp) [56] to uncovering customers' behavior in the telecommunication industry [57], and predicting traffic patterns in a smart city environment.

After transforming our data using BEATS and Eigen, we followed the general data modelling process proposed in [58] to classify the series: standardization, splitting the dataset into training and test sets, choosing the model, selecting the best hyperparameters of each model using 10-fold cross validation on the training set and checking the accuracy of the model using the test set. With respect to the methodology followed in [58], we improve the way of looking for the hyperparameters of the algorithms using the python package optunia since it contains various optimizers for hyperparameter tuning.

Among other options like grid search, random search and genetic algorithms, we have chosen particle swarm implementation since it is shown to surpass the performance of other solutions [59].

The models that we use to combine with BEATS and Eigen are the widely known Random Forest (RF) and Support Vector Machines (SVM) with Radial Basis Function Kernel.

Whereas Random Forest deals with *small n large p-problems*, high-order interactions and correlated predictor variables, SVMs are more effective for relatively small datasets with fewer outliers. Generally speaking, Random Forests may require more data. Both of the algorithm show better performance when combined with SVM.

The tuning of SVM has been done without deciding the kernel in advance. That means, the kernel (linear, polynomial or RBF) is considered as an hyperparameter.

According to the discussion in Section 2, we compare our method with:

- Original time series (i.e., raw data): DTW with 1-NN classification since, after many trials, it is still the benchmark of comparison for distance based classification. Having a complexity of $O(n^2)$ that under

2. <http://www.timeseriesclassification.com/description.php?Dataset=Lightning7>

3. https://github.com/auroragonzalez/BEATS/tree/master/data/random_LHS_generator_drift

4. <http://www.ict-citypulse.eu>

TABLE 1
Accuracy of Each Method Using As Inputs Each of the Segmented Time Series

Model \ dataset	Arrow Heads	Lightning7	Random Generator	Coffee	Ford A	Proximal
Model						
BEATS-SVM	0.81	0.7	0.75	1	0.75	0.85
Eigen-SVM	0.79	0.72	0.73	1	0.74	0.8
DTW-1NN	0.67	0.75	0.71	0.87	0.66	0.81
SAX-VSM	0.68	0.59	0.52	0.96	0.09*	0.75
TSF	0.73	0.75	0.75	0.97	0.75	0.85
FLAG	0.57	0.76	0.67	1	0.73	0.64
COTE	0.78	0.8	0.7	1	0.75	0.83

*The bag of words generated by a wide majority of the test subjects is not related to the ones generated by the train step. This implies that their TF*IDF weights are not computed and it is not possible to compute the cosine similarity. In consequence, the method is not valid for many of the cases, producing the reported bad results.

certain circumstances [60] could be reduced to $O(n)$ using lower bounds such as LB_{Keogh} or $LB_{Improved}$ [61].

- Intervals: We choose TSF in order to make the comparison since it is more modern and quicker than the rest. Its complexity is $O(t * m * n * logn)$, where t = number of trees and m = number of splits or segments.
- Symbolic approximations: In the classification task, we use SAX-VSM. The complexity is linear: $O(n)$.
- Shapelets: FLAG is the newest, the quickest and claims to be better than its predecessors. Its complexity is $O(n^3)$.
- Ensembles: COTE. It is an ensemble of dozens of core classifiers many of which having a quadratic, cubic or even bi-quadratic complexity. It is the most computationally expensive in this list.

The results are shown in Table 1. It is important to mention that not only accuracy results but also the time that it takes the algorithm to run both training and test phases including input transformation, has improved. This runtime is shown in Fig. 5, where a logarithmic transformation is applied to the data in order to improve visibility.

We have depicted both metrics: accuracy and running time in a plot that summarises the results over all the datasets. Both metrics have been scaled per dataset and we have computed the average performance per model that is represented by the bigger points in the plot.

In order to make a more consistent analysis of the results, we have generated 100 Random LHS Generator Lift datasets and the model accuracy of the models using violin plots (see Fig. 6), which together with the regular statistics that

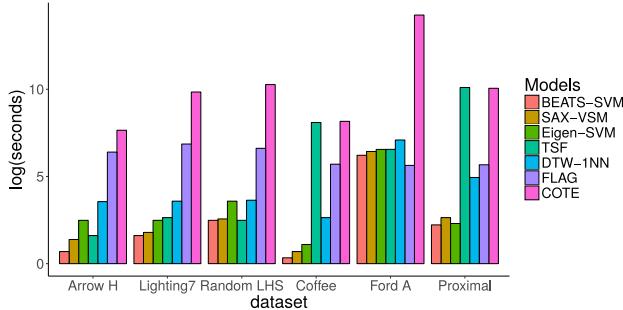


Fig. 5. Running time (log(sec)) and programming language of the algorithms.

boxplot provide they show the probability density of the data at different values of accuracies. While the differences between BEATS-SVM, TSF and COTE are not statistically significant (p-value = 0.7 > 0.05), BEATS-SVM is very quick in comparison to COTE and that BEATS is also more versatile than the rest since it can be combined with any classification algorithms.

5.3 Clustering

Clustering is used to identify the structure of an unlabeled dataset by organising the data into homogeneous groups where within-group-object similarity is minimized and between-group-object dissimilarity is maximized. The process is done without consulting known class labels. Clustering is an unsupervised machine learning method. In particular, time series clustering partitions time series data into groups based on similarity or distance; so that time series data in the same cluster are similar.

Clustering has tackled tasks such as the assignment of genes with similar expression trajectories to the same group [62]. The creation of profiles of the trips carried out by tram users [63] or the acquisition of energy consumption predictions by clustering houses [64] are among examples of using clustering methods.

After transforming our data using BEATS and Eigen, we applied the connectivity based algorithm *hierarchical agglomerative clustering* and the centroid based algorithm *k-means* to

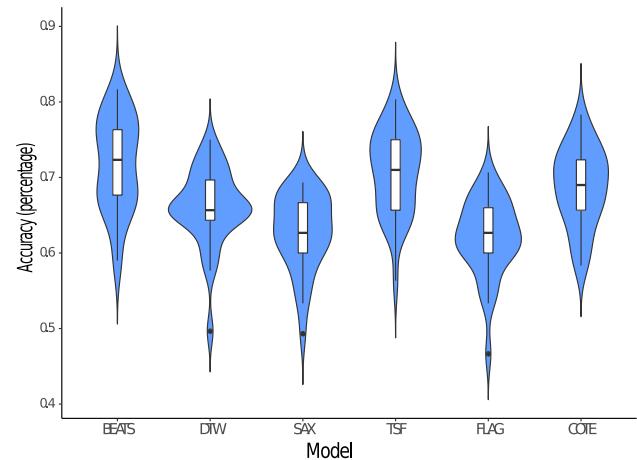


Fig. 6. Classification accuracy on the 100 randomly generated datasets.

TABLE 2
Silhouette Coefficient of Each Method Using As Inputs Each of the Segmented Time Series

Model \ dataset	Arrow Heads	Lightning7	Random Generator	Coffee	Ford A	Proximal
BEATS-HC	0.6	0.25	0.45	0.25	0.46	0.4
Eigen-HC	0.58	0.31	0.25	0.26	0.36	0.38
DTW	0.33	0.21	0.44	0.21	0.12	0.31
SAX_{SD}- HC	0.53	0.06	0.19	0.13	0	0.33
k-shape	0.44	0.19	0.05	0.43	0.38	0.5

cluster the time series datasets. In the hierarchical clustering, the selected agglomerative method is *complete linkage*, meaning that the distance between two clusters is the maximum distance between their individual components (in each time series). Hierarchical clustering seems to be a better partner for both of them.

The dissimilarity matrix contains the distances between the pairs of time series. We use the cosine dissimilarity for the rest of the segmentations (BEATS and Eigen). The cosine dissimilarity is calculated as one minus the cosine of the included angle between elements of the time series (see Eq. (2))

$$\text{dissimilarity} = 1 - \frac{\mathbf{XY}}{\|\mathbf{X}\| \|\mathbf{Y}\|} = 1 - \frac{\sum_{i=1}^n X_i Y_i}{\sqrt{\sum_{i=1}^n X_i^2} \sqrt{\sum_{i=1}^n Y_i^2}}. \quad (2)$$

Finally, for both methods we have used a fixed number of clusters. As we were aware of the classification groups (our data is labeled), we applied the algorithms setting apriori the number of clusters k and used the silhouette coefficient as a metric for measuring the cluster quality.

The silhouette coefficient is an internal measure that combines the measurement of cohesion and separation. Cluster cohesion measures how closely related the objects in a cluster are. Cluster separation measures how well separated the clusters are from each other. The silhouette coefficient for a subject i is defined as

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}, \quad (3)$$

where $a(i)$ is the average distance between i and each of the points of the assigned cluster and $b(i)$ is the average distance between i and each of the points of the next best cluster. This value can be used to compare the quality of different cluster results.

From the definition it is clear that $s(i) \in [-1, 1]$. Meanwhile a silhouette coefficient value closer to 1 means that the clustering is good; a value close to -1 represents less efficiency in the

categorization for the clusters. When it is close to 0, it means that the point is in the border between two clusters.

According to the discussion in Section 2 we will analyse:

- Original time series: DTW distance using the tight lower bound of [61], that makes it faster.
- Symbolic approximations: We have taken the most modern improvement that SAX has experienced: SAX_{SD}. The MINDIST function that returns the minimum distance between the original time series of two words [65] is enhanced with the distance between the standard deviation of each segment.
- Shapelets: k-shape is the model chosen in this direction.

The results of the clustering experiments done in the training sets are shown in Table 2. The run time of the algorithms is shown in Fig. 7. In this case, all the algorithms have been coded using the same programming language so we consider that the graph is enough in order to estimate the different algorithms complexity regarding time.

5.4 Big Data Use Case: Traffic in Smart Cities

In this section we apply BEATS in a smart cities related use-case: traffic data clustering, done in an online and distributed way.

5.4.1 BEATS Implementation for Big Data

In contrast to the traditional analysis procedure where data is first stored and then processed in order to deploy models, the major potential of the data generated by IoT is accomplished by the realization of continuous analytics that allow to make decisions in real time.

There are three types of data processing: Batch Processing, Stream Processing and Hybrid Processing.

Batch processing operates over a group of transactions collected over a period of time and reports results only when all computations are done, whereas stream processing produces incremental results as soon as they are ready [66].

Regarding the available Big Data Tools, we have considered Hadoop⁵ and Spark⁶ Big Data frameworks. Hadoop was designed for batch processing. All data is loaded into HDFS and then MapReduce starts a batch job to process that data. If the data changes the job needs to be ran again. It is step by step processing that can be paused or interrupted, but not changed.

Apache Spark allows to perform analytical tasks on distributed computing clusters. Sparks real-time data

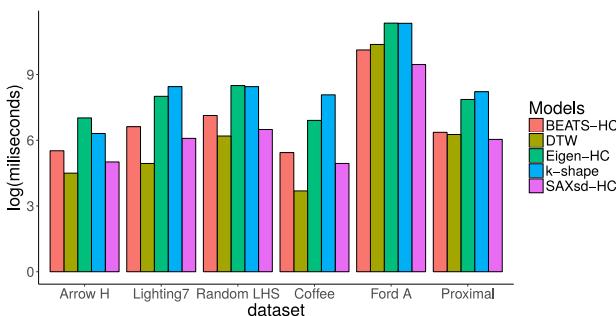


Fig. 7. Running time (log(milisec)) of the clustering algorithms.

5. <http://hadoop.apache.org/>

6. <https://spark.apache.org/>

processing capability provides substantial lead over Hadoop's MapReduce and it is essential for online time series segmentation and representation.

The Spark abstraction for a continuous stream of data is called a Discretized Stream or DStream. A DStream is a micro-batch of Resilient Distributed Datasets, RDDs. That means, a DStream is represented as a sequence of RDDs. RDDs are distributed collections that can be operated in parallel by arbitrary functions and by transformations over a sliding window of data (windowed computations).

5.4.2 BEATS Adapted to Spark Technology

For the online implementation of BEATS we have decided to use pyspark, the Spark Python API that exposes the Spark programming model to Python.

There are many works proposing online time series processing but few of them that have implemented it. In [67] is highlighted that MapReduce is not the appropriate technology for rolling window time series prediction and proposes a index pool data structure.

Pyspark allows us to use the Spark Streaming functionalities that are needed in order to implement BEATS online. In Section 3 we have seen that BEATS algorithm can be separately applied to windows of the data. Therefore we associate the data received within one window to one RDD, that can be processed in a parallel way.

A suitable type of RDDs for our implementation is key / value pairs. In detail, the key is an identifier of the time series (e.g., sensor name) and the value is the sequence of values of our time series that fall in the window. That way the blocks are exposed to operations that give the possibility to act on each key in parallel or regroup data across the network.

The transformations that we use are:

- Window: use for creating sliding window of time over the incoming data.
- GroupByKey: grouping the incoming values of the sliding window by key (for example, same sensor data).
- Map: The Map function applied in parallel to every pair (key, value), where the key is the time series, values are a vector and the function depends on what has to be done.

5.4.3 The Applied Scenario

We use one of the real-world datasets obtained from the collection of datasets of vehicle traffic in the City of Aarhus in Denmark for a period of 6 months.⁷ The dataset is provided in the context of the CityPulse smart city project.

The selected dataset gathers 16971 samples of data from sensors situated in lamp posts covering an area around 2345m.⁸ The variables considered for the analysis are: flow (numbers of cars between two points) and average speed. Each variable is a time series.

In order to simulate an online application we consider that the BEATS segmentation is carried out on hourly based data.

7. <http://iot.ee.surrey.ac.uk:8080/datasets.html#traffic>

8. http://iot.ee.surrey.ac.uk:8080/datasets/traffic/traffic_june_sep/index.html

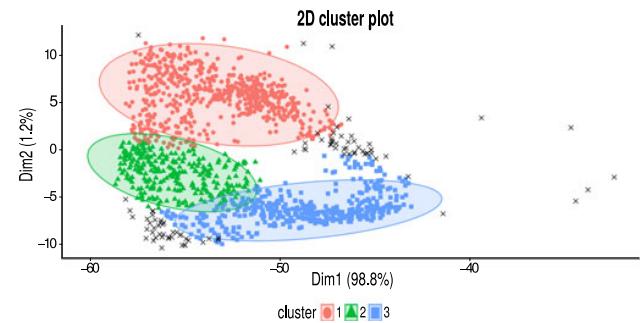


Fig. 8. Plot of the DBSCAN clusters using traffic data.

To achieve this, since the data is collected every 5 mins., a sliding window of size 12 is selected. The goal of the clustering is to determine the status of the road in terms of the traffic flow and occurrences. For every window of 128 observations (64 for each variable) BEATS obtains three flow related representatives and three speed related representatives.

Each observation of the final input dataset for the clustering model represents one window of the raw data. The final dataset has 6 variables and 1409 samples. This means a reduction of around 75 percent of data.

The data is gathered by anonymously collecting Wi-Fi and Bluetooth signals transmitted by travelers' smartphones or in-vehicle systems. This infrastructure provides noisy data in cases such as stopped vehicles in traffic jam, buses with a lot of passengers.

In order to tackle the presence of outliers and noise, the selected clustering technique is density-based spatial clustering (DBSCAN). DBSCAN groups points that are closely packed together. Points that do not fit into any of the main groups because they lie in low-density regions are marked as outliers. The hyper-parameters of DBSCAN are minimum number of points required to form a dense region (MinP) and ϵ in order to find the ϵ -neighborhood of each point. We set that clusters contain at least a 20 percent of the data and $\epsilon = 4.014$. Using such configuration, we obtain 3 different clusters and a 8 percent of data that cannot be classified in any of the previous, i.e., outliers. The description of the clusters, including the number of points n that belong to each of the clusters and the mean μ and standard deviation sd for both flow and speed is:

- Cluster 1 ($n = 618$): High flow ($\mu = 30.97, sd = 12.66$) and medium speed ($\mu = 102.5, sd = 10.2$);
- Cluster 2 ($n = 271$): Medium flow ($\mu = 15.97, sd = 8.4$) and high speed ($\mu = 110, sd = 9.21$); and
- Cluster 3 ($n = 432$): Low flow ($\mu = 6.1, sd = 5.56$) and low/medium speed ($\mu = 97.8, sd = 14.3$).

In order to represent the data in lower dimension, we select the first two principal components of the data using Principal Components Analysis (PCA). The obtained clusters are shown in Fig. 8. Crosses in black colour represent the noise data. We have also projected the clusters in the three flow related components of BEATS, so that clusters can be visualized in a 3D form as presented in Fig. 9.

Regarding this application, we can conclude that clustering methods applied to the segments generated by BEATS are able to characterise the status of the roads by grouping the values in an effective form.

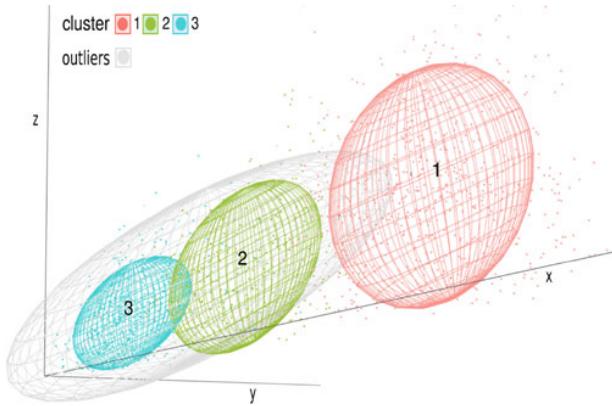


Fig. 9. 3D plot of the DBSCAN clusters using traffic data.

Using a computer with an Intel i5 Processor, 8GB RAM Memory, Ubuntu 16.04 operative system and the statistical software R 3.4.3 [68], the running time of DBSCAN using BEATS segmented data is 0.25 seconds. However, to run the DBSCAN with raw data it takes around 35 seconds. The later confirms again the suitability of BEATS in current IoT scenarios.

6 DISCUSSION

As we have described in the paper, the randomness and predictability of a real-world time series changes over time due to several factors.

The existing solutions for pattern creation and abstraction in time-series data often work based on statistical measures (which have limited representation and granularity), symbolic methods such as SAX (which assumes that the data is normally distributed and requires normalization of the data), or signal processing and stream processing methods such as wavelet or Fourier transforms (which act as filters and can extract features from the data but do not provide a pattern representation/abstraction).

Our proposed model combines a series of methods to create a window based abstraction of time series data and uses a frequency domain function combined with characteristic value measures that represents the overall direction of the dataframe (i.e., an n-dimensional matrix constructed during our windowing/slicing process) as a vector.

BEATS is an algorithm that process data streams whose randomness and predictability varies depending on the segment of data. The proposed algorithm is useful specially in applications such as smart cities where results of the segmentation and processing algorithms are used in order to make fast decisions regarding traffic, energy, light regulation, etc. This can be made by combining various sensory data and other historical data. In general terms, the intention is to predict and manage what is occurring in order to provide informed or automated decisions for repetitive tasks that can be handled by machines. BEATS offers a powerful solution to aggregate and represent large-scale streaming data in a quick and adaptable way. It uses blocks of eigenvalues in a much lower-dimensionality (with a high aggregation rate) which preserves the main information and characteristics of the data. Since BEATS uses eigenvalues, it provides a homogeneous way to represent multi-modal and heterogeneous

streaming data. In other words, all different types of numerical streaming data are transformed into vectors of eigenvalues that, in principal, preserve and represent the magnitude and overall direction of the data in a lower-dimensionality space. This not only allows to compare and combine different blocks of data from various data streams, but also provides a unified way to represent the blocks of data as patterns in the form of eigenvalues.

In this paper, we mainly target a key step after collection of the data: aggregation. Aggregation of data becomes a very significant task in order to extract the key characteristics of the data in lower-dimensionality. We segment the time series and make a reduction for each time series at a rate of 60 ~ 70 percent when using overlapping windows. The independence between blocks that our algorithm provides is one of its most important features. BEATS also presents other qualities such as adapting to drifts and low latency.

BEATS reduces the data by using the eigenvalues of a submatrix of the DCT transformation. These eigenvalues represent the key-characteristics of the data.

The evaluation is performed using classification and clustering, two of the classical machine learning tasks using several types of datasets. The inputs of the models are the different representations introduced in the paper: BEATS and Eigen together with raw data for the other models.

Classification is measured by accuracy. This allows us to perform a test for equality of proportions, that is a χ^2 test of independence in order to assure that the differences between accuracies are statistically significant.

For the Arrow Heads dataset we find that BEATS combined with SVM outperforms all the algorithms. However, the differences between COTE and BEATS are not statistically significant ($\chi^2(1) = 0.37$, p-value = 0.54 > 0.05). On the other hand, the difference between TSF and BEATS are statistically significant ($\chi^2(1) = 4.8$, p-value = 0.04 < 0.05).

In the case of Lightning7, there are several models that outperform BEATS. The winning one is COTE. Nonetheless, COTE is very complicated, time demanding and computationally expensive. The rest only overperforms BEATS by 6 percent at most.

In the case of Random LHS Generator Lift, TSF and BEATS perform similarly.

In the Coffee dataset, we observe that several approaches (including BEATS) achieve a 100 percent accuracy on classification.

In FordA, BEATS, TSF and COTE perform similarly. However, BEATS is the quickest amongst them.

Finally, in the Proximal dataset TSF and BEATS perform similarly in terms of accuracy. However, BEATS is again quicker.

Even though COTE and TSF are strong rivals to BEATS, it should be noted that the computation time and simplicity of BEATS makes it useful to use in rapid analysis having still good results. Also, due to its nature is very adaptable and easy to combine with any other classification algorithm different than SVM.

The clustering experiment is evaluated by comparing the hundredths of the silhouette coefficients, where each hundredth is going to be counted as a point in the below description.

BEATS is 7 points above SAX_{SD} for the Arrow Heads dataset, 1 point above DTW in the Random LHS Generator Lift set and 8 points above k-shape in Ford A. Being the most computationally expensive of all the clustering algorithms under study, as it can be seen in Fig. 7, k-shape outperforms BEATS in two datasets: Coffee and Proximal.

It can be said that in clustering, BEATS behaves better when we are using long datasets since it outperforms every algorithms in both metrics: silhouette coefficient and running time in the biggest dataset: *FordA*.

Finally, by applying DBSCAN to cluster traffic data, we noticed that BEATS performs efficiently since the clusters represent different situations of the use-case in terms of traffic flow and speed.

7 CONCLUSIONS AND FUTURE WORK

In this paper, we introduce a novel algorithm called BEATS, which aggregates and represents time series data in blocks of lower-dimensional vectors of eigenvalues. BEATS is not sample dependent so it adapts to data drifts in the underlying data streams.

The BEATS abstractions can be combined with various machine learning models to discover patterns, identify correlations (within or between data streams), extract insights and identify activities from the data. In this paper, we have used several datasets and have shown several use cases that demonstrate how the BEATS abstractions can be used for clustering, analysis and grouping the activities and patterns in time-series data.

Compared to existing segmentation methods, BEATS shows significant improvements in representing datasets with drifts. When combined with classification and clustering methods, we have shown that it can obtain competitive results compared with other state-of-the-art but more complex and time consuming methods.

For the BEATS algorithm evaluation we have fixed the length of the segments at 64; so the only parameter to take into consideration was the slide of the window, that we have kept constantly equal to 8, so the blocks of transformed data intersect. Nevertheless, the optimization of the sliding window is an open issue to be addressed in future work.

For the clustering tasks, it is important to take into account that the definition of similarity is subjective. The similarity depends on the domain of application.

By using BEATS, we are able to restructure the streaming data in a 2D way and then transform it into the frequency domain using DCT. The algorithm finds a smaller sequence that contains the key information of the initial representative. This aggregation provides an opportunity to eliminate repetitive content and similarities that can be found in the sequence of data.

The eigenvalues vectors are a homogeneous representation of the data streams in BEATS that allow us to go one step further in understanding of the sequences and patterns that can be considered as the data structure of a data series in an application domain (e.g., smart cities).

Its applications can be extended to several other domains and various patterns/activity monitoring and detection methods. The future work will focus on applying 3D cosine transform and adaptive block size estimation.

APPENDIX A

Definition A.1 (Integral transform). *The integral transform of the function $f(t)$ with respect to the kernel $K(t, s)$ is*

$$F(t) = \int_{-\infty}^{\infty} K(t, s)f(s)ds, \quad (4)$$

if the integral exists.

The kernel of the Fourier Transformation is $K(t, s) = e^{-its}$, and, in particular for the cosine fourier transformation $K(t, \omega) = \cos(t, \omega)$. If we discretize the kernel we can reach that $K_c(j, k) = \cos\left(\frac{jk\pi}{N}\right)$, where N is an integer.

Definition A.2. (Discrete Cosine Transformation (DCT) - II). *DCT is a linear and invertible function*

$$f : \mathbb{R}^n \rightarrow \mathbb{R}^n$$

where \mathbb{R} denotes the set of real numbers or, equivalently, on a $n \times n$ matrix, defined by:

$$f_j = \sum_{k=0}^{n-1} \cos\left(\frac{\pi}{n} j \left(k + \frac{1}{2}\right)\right) \text{ where } j = 0, 1, \dots, n-1 \quad (5)$$

ACKNOWLEDGMENTS

This work has been partially funded by MINECO grant BES-2015-071956, PERSEIDES TIN2017-86885-R project, and ERDF funds, by the European Comission through the H2020-ENTROPY-649849 EU Project, and the H2020 FIESTA Project under grant agreement no. CNECT-ICT-643943.

REFERENCES

- [1] J. Gantz and D. Reinsel, "The digital universe in 2020: Big data, bigger digital shadows, and biggest growth in the far east," *IDC iView: IDC Analyze Future*, vol. 2007, pp. 1-16, 2012.
- [2] D. Abbott, *Applied Predictive Analytics: Principles and Techniques for the Professional Data Analyst*. Hoboken, NJ, USA: Wiley, 2014.
- [3] E. J. Keogh and M. J. Pazzani, "An enhanced representation of time series which allows fast and accurate classification, clustering and relevance feedback," in *Proc. 4th Int. Conf. Knowl. Discovery Data Mining*, vol. 98, pp. 239-243, 1998.
- [4] E. Keogh, S. Chu, D. Hart, and M. Pazzani, "Segmenting time series: A survey and novel approach," *Data Mining Time Series Databases*, vol. 57, pp. 1-22, 2004.
- [5] H. Aksoy, A. Gedikli, N. E. Unal, and A. Kehagias, "Fast segmentation algorithms for long hydro meteorological time series," *Hydrological Processes*, vol. 22, no. 23, pp. 4600-4608, 2008.
- [6] X. Xi, E. Keogh, C. Shelton, L. Wei, and C. A. Ratanamahatana, "Fast time series classification using numerosity reduction," in *Proc. 23rd Int. Conf. Mach. Learn.*, 2006, pp. 1033-1040.
- [7] A. Bagnall, L. Davis, J. Hills, and J. Lines, "Transformation based ensembles for time series classification," in *Proc. SIAM Int. Conf. Data Mining*, 2012, pp. 307-318.
- [8] X. Wang, A. Mueen, H. Ding, G. Trajcevski, P. Scheuermann, and E. Keogh, "Experimental comparison of representation methods and distance measures for time series data," *Data Mining Knowl. Discovery*, vol. 26, pp. 1-35, 2013.
- [9] A. Bagnall, J. Lines, A. Bostrom, J. Large, and E. Keogh, "The great time series classification bake off: A review and experimental evaluation of recent algorithmic advances," *Data Mining Knowl. Discovery*, Springer, vol. 31, no. 3, pp. 606-660, 2017.
- [10] Y.-S. Jeong, M. K. Jeong, and O. A. Omatamu, "Weighted dynamic time warping for time series classification," *Pattern Recognit.*, vol. 44, no. 9, pp. 2231-2240, 2011.
- [11] P.-F. Marteau, "Time warp edit distance with stiffness adjustment for time series matching," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 31, no. 2, pp. 306-318, Feb. 2009.

- [12] A. Stefan, V. Athitsos, and G. Das, "The move-split-merge metric for time series," *IEEE Trans. Knowl. Data Eng.*, vol. 25, no. 6, pp. 1425–1438, Jun. 2013.
- [13] G. E. Batista, X. Wang, and E. J. Keogh, "A complexity-invariant distance measure for time series," in *Proc. SIAM Int. Conf. Data Mining*, 2011, pp. 699–710.
- [14] T. Gorecki and M. Łuczak, "Non-isometric transforms in time series classification using DTW," *Knowl.-Based Syst.*, vol. 61, pp. 98–108, 2014.
- [15] J. Lines and A. Bagnall, "Time series classification with ensembles of elastic distance measures," *Data Mining Knowl. Discovery*, vol. 29, no. 3, pp. 565–592, 2015.
- [16] S. Aghabozorgi, A. S. Shirkhorshidi, and T. Y. Wah, "Time-series clustering—decade review," *Inf. Syst.*, vol. 53, pp. 16–38, 2015.
- [17] V. Hautamaki, P. Nykänen, and P. Franti, "Time-series clustering by approximate prototypes," in *Proc. 19th Int. Conf. Pattern Recognit*, 2008, pp. 1–4.
- [18] G. E. Batista, E. J. Keogh, O. M. Tataw, and V. M. De Souza, "Cid: An efficient complexity-invariant distance for time series," *Data Mining Knowl. Discovery*, vol. 28, no. 3, pp. 634–669, 2014.
- [19] Y. Zhu, D. Wu, and S. Li, "A piecewise linear representation method of time series based on feature points," in *Proc. Int. Conf. Knowl.-Based Intell. Inf. Eng. Syst.*, 2007, pp. 1066–1072.
- [20] E. J. Keogh and M. J. Pazzani, "A simple dimensionality reduction technique for fast similarity search in large time series databases," in *Proc. Pacific-Asia Conf. Knowledge Discovery Data Mining*, 2000, pp. 122–133.
- [21] N. T. Nguyen, B. Trawiński, R. Katarzyniak, and G.-S. Jo, *Advanced Methods for Computational Collective Intelligence*. Berlin, Germany: Springer, 2012, vol. 457.
- [22] H. Deng, G. Runger, E. Tuv, and M. Vladimir, "A time series forest for classification and feature extraction," *Inf. Sci.*, vol. 239, pp. 142–153, 2013.
- [23] M. G. Baydogan and G. Runger, "Time series representation and similarity based on local autopatterns," *Data Mining Knowl. Discovery*, vol. 30, no. 2, pp. 476–509, 2016.
- [24] E. Keogh, J. Lin, and A. Fu, "Hot sax: Efficiently finding the most unusual time series subsequence," in *Proc. 5th IEEE Int. Conf. Data Mining*, 2005, pp. 8–pp.
- [25] X. Xi, E. J. Keogh, L. Wei, and A. Mafra-Neto, "Finding motifs in a database of shapes," in *Proc. Int. Conf. Data Mining*, 2007, pp. 249–260.
- [26] C. D. Stylios and V. Kreinovich, "Symbolic aggregate approximation (SAX) under interval uncertainty," in *Proc. Annu. Conf. North Amer. Fuzzy Inf. Process. Soc. held Jointly 5th World Conf. Soft Comput.* IEEE, 2015, pp. 1–7.
- [27] B. Lkhagvaa, Y. Suzuki, and K. Kawagoe, "Extended SAX: Extension of symbolic aggregate approximation for financial time series data representation," in *Proc. Data Eng. Workshop*, vol. 4A-18, 2006.
- [28] Y. Sun, J. Li, J. Liu, B. Sun, and C. Chow, "An improvement of symbolic aggregate approximation distance measure for time series," *Neurocomputing*, vol. 138, pp. 189–198, 2014.
- [29] C. T. Zan and H. Yamana, "An improved symbolic aggregate approximation distance measure based on its statistical features," in *Proc. 18th Int. Conf. Inf. Integr. Web-Based Appl. Serv.*, 2016, pp. 72–80.
- [30] J. Lin, E. Keogh, L. Wei, and S. Lonardi, "Experiencing sax: A novel symbolic representation of time series," *Data Mining Knowl. Discovery*, vol. 15, no. 2, pp. 107–144, 2007.
- [31] S. Kolozali, D. Puschmann, M. Bermudez-Edo, and P. Barnaghi, "On the effect of adaptive and non-adaptive analysis of time-series sensory data," *IEEE Internet Things J.*, vol. 3, no. 6, pp. 1084–1098, 2016.
- [32] P. Senin and S. Malinchik, "Sax-VSM: Interpretable time series classification using sax and vector space model," in *Proc. IEEE 13th Int. Conf. Data Mining*, 2013, pp. 1175–1180.
- [33] T. Rakthanmanon and E. Keogh, "Fast shapelets: A scalable algorithm for discovering time series shapelets," in *Proc. SIAM Int. Conf. Data Mining*, 2013, pp. 668–676.
- [34] J. Grabocka, N. Schilling, M. Wistuba, and L. Schmidt-Thieme, "Learning time-series shapelets," in *Proc. 20th ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2014, pp. 392–401.
- [35] M. Shah, J. Grabocka, N. Schilling, M. Wistuba, and L. Schmidt-Thieme, "Learning DTW-shapelets for time-series classification," in *Proc. 3rd IKDD Conf. Data Sci.*, 2016, Art. no. 3.
- [36] L. Hou, J. T. Kwok, and J. M. Zurada, "Efficient learning of time-series shapelets," in *Proc. 13th AAAI Conf. Artif. Intell.*, 2016, pp. 1209–1215.
- [37] J. Paparrizos and L. Gravano, "k-shape: Efficient and accurate clustering of time series," in *Proc. ACM SIGMOD Int. Conf. Manag. Data*, 2015, pp. 1855–1870.
- [38] A. Bagnall, J. Lines, J. Hills, and A. Bostrom, "Time-series classification with COTE: The collective of transformation-based ensembles," *IEEE Trans. Knowl. Data Eng.*, vol. 27, no. 9, pp. 2522–2535, Sep. 2015.
- [39] M. Sayed-Mouchaweh, *Learning from Data Streams in Dynamic Environments*. Berlin, Germany: Springer, 2016.
- [40] D. Puschmann, P. Barnaghi, and R. Tafazolli, "Adaptive clustering for dynamic iot data streams," *IEEE Internet Things J.*, vol. 4, no. 1, pp. 64–74, Feb. 2017.
- [41] J. Gama, I. Ziobraité, A. Bifet, M. Pechenizkiy, and A. Bouchachia, "A survey on concept drift adaptation," *ACM Comput. Surveys* [Online]. Available: <http://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=6703726>, vol. 46, no. 4, 2014, Art. no. 44.
- [42] C. Lifna and M. Vijayalakshmi, "Identifying concept-drift in twitter streams," *Procedia Comput. Sci.*, vol. 45, pp. 86–94, 2015.
- [43] C. M. Grinstead and J. L. Snell, *Introduction to Probability*. Providence, RI, USA: American Mathematical Society, 2012.
- [44] M. Ali-ud-din Khan, M. F. Uddin, and N. Gupta, "Seven v's of big data understanding big data to extract value," in *Proc. Zone 1 Conf. Amer. Soc. Eng. Educ.*, 2014, pp. 1–5.
- [45] Z. Wang, "Fast algorithms for the discrete W transform and for the discrete fourier transform," *IEEE Trans. Acoust. Speech Signal Process.*, vol. 32, no. 4, pp. 803–816, Aug. 1984.
- [46] K. R. Rao and P. Yip, *Discrete Cosine Transform: Algorithms, Advantages, Applications*. Cambridge, MA, USA: Academic Press, 2014.
- [47] G. K. Wallace, "The jpeg still picture compression standard," *IEEE Trans. Consum. Electron.*, vol. 38, no. 1, pp. xviii–xxxiv, Feb. 1992.
- [48] A. C. Bovik, *The Essential Guide to Image Processing*. Cambridge, MA, USA: Academic Press, 2009.
- [49] G. Bi and Y. Zeng, *Transforms and Fast Algorithms for Signal Analysis and Representations*. Berlin, Germany: Springer, 2004.
- [50] J. Makhoul, "A fast cosine transform in one and two dimensions," *IEEE Trans. Acoust. Speech Signal Process.*, vol. 28, no. 1, pp. 27–34, Feb. 1980.
- [51] Y. Chen, E. Keogh, B. Hu, N. Begum, A. Bagnall, A. Mueen, and G. Batista, "The UCR time series classification archive," Jul. 2015, [Online]. Available: www.cs.ucr.edu/~eamonn/time_series_data/
- [52] L. Ye and E. Keogh, "Time series shapelets: A new primitive for data mining," in *Proc. 15th ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2009, pp. 947–956.
- [53] D. Puschmann, "Random lhs generator drift," 2016. [Online]. Available: <https://github.com/UniSurreyIoT/random-LHS-generator-drift>
- [54] R. Briandet, E. K. Kemsley, and R. H. Wilson, "Discrimination of arabica and robusta in instant coffee by fourier transform infrared spectroscopy and chemometrics," *J. Agricultural Food Chemistry*, vol. 44, no. 1, pp. 170–174, 1996.
- [55] L. M. Davis, "Predictive modelling of bone ageing," Ph.D. dissertation, University of East Anglia, Norwich, Norfolk, 2013.
- [56] A. Anguera, J. Barreiro, J. Lara, and D. Lizcano, "Applying data mining techniques to medical time series: An empirical case study in electroencephalography and stabilometry," *Comput. Structural Biotechnology J.*, vol. 14, pp. 185–199, 2016.
- [57] Y. Yang, Q. Yang, W. Lu, J. Pan, R. Pan, C. Lu, L. Li, and Z. Qin, "Preprocessing time series data for classification with application to crm," in *Proc. Australasian Joint Conf. Artifi. Intell.*, 2005, pp. 133–142.
- [58] A. González-Vidal, V. Moreno-Cano, F. Terroso-Sáenz, and A. F. Skarmeta, "Towards energy efficiency smart buildings models based on intelligent data analytics," *Procedia Comput. Sci.*, vol. 83, pp. 994–999, 2016.
- [59] S.-W. Lin, K.-C. Ying, S.-C. Chen, and Z.-J. Lee, "Particle swarm optimization for parameter determination and feature selection of support vector machines," *Expert Syst. Appl.*, vol. 35, no. 4, pp. 1817–1824, 2008.
- [60] C. A. Ratanamahatana and E. Keogh, "Three myths about dynamic time warping data mining," in *Proc. SIAM Int. Conf. Data Mining*, 2005, pp. 506–510.
- [61] D. Lemire, "Faster retrieval with a two-pass dynamic-time-warping lower bound," *Pattern Recognit.*, vol. 42, no. 9, pp. 2169–2180, 2009.
- [62] Z. Bar-Joseph, A. Gitter, and I. Simon, "Studying and modelling dynamic biological processes using time-series gene expression data," *Nature Rev. Genetics*, vol. 13, no. 8, pp. 552–564, 2012.

- [63] M. V. Moreno, F. Terroso-Saenz, A. Gonzalez, M. Valdes-Vela, A. F. Skarmeta, M. A. Zamora-Izquierdo, and V. Chang, "Applicability of big data techniques to smart cities deployments," *IEEE Trans. Ind. Inform.*, vol. 13, no. 2, pp. 800–809, Apr. 2017.
- [64] C. Costa and M. Y. Santos, "Improving cities sustainability through the use of data mining in a context of big city data," in *Proc. Int. Conf. Data Mining Knowl. Eng.*, 2015, vol. 1, pp. 320–325.
- [65] J. Lin, E. Keogh, S. Lonardi, and B. Chiu, "A symbolic representation of time series, with implications for streaming algorithms," in *Proc. 8th ACM SIGMOD Workshop Res. Issues Data Mining Knowledge Discovery*, 2003, pp. 2–11.
- [66] M. Hirzel, H. Andrade, B. Gedik, G. Jacques-Silva, R. Khandekar, V. Kumar, M. Mendell, H. Nasgaard, S. Schneider, R. Soulé, et al., "Ibm streams processing language: Analyzing big data in motion," *IBM J. Res. Develop.*, vol. 57, no. 3/4, pp. 7–1, 2013.
- [67] L. Li, F. Noorian, D. J. Moss, and P. H. Leong, "Rolling window time series prediction using MapReduce," in *Proc. IEEE 15th Int. Conf. Inf. Reuse Integr.*, 2014, pp. 757–764.
- [68] R Core Team, *R: A Language and Environment for Statistical Computing*, R Foundation for Statistical Computing, Vienna, Austria, 2016. [Online]. Available: <https://www.R-project.org/>



Aurora González Vidal received the graduated degree in mathematics from the University of Murcia, in 2014. In 2015, she received a fellowship to work in the Statistical Division of the Research Support Services, where she specialized in statistics and data analysis. Since 2015, she has been working toward the PhD degree in computer science, focusing her research on data analytics for energy efficiency and studied for a master's degree in Big Data. She was a visiting PhD student at the University of Surrey, where she worked on the study of segmentation of time series.



Payam Barnaghi is a reader in machine intelligence in the Institute for Communication Systems Research with the University of Surrey. He was the coordinator of the EU FP7 CityPulse project on smart cities. His research interests include machine learning, the Internet of Things, the Semantic Web, adaptive algorithms, and information search and retrieval. He is a senior member of the IEEE and a fellow of the Higher Education Academy.



Antonio F. Skarmeta received the BS (Hons.) degree in computer science from the University of Murcia, the MS degree in computer science from the University of Granada, Spain, and the PhD degree in computer science from the University of Murcia. He is a full professor with the Department of Information and Communications Engineering, University of Murcia. He is involved in numerous projects, both European and National. Research interests include mobile communications, artificial intelligence, and home automation. He is a member of the IEEE.

▷ For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.

4.2 A methodology for Energy Multivariate Time Series Forecasting in Smart Buildings based on Feature Selection

Title	A methodology for Energy Multivariate Time Series Forecasting in Smart Buildings based on Feature Selection
Authors	Aurora González-Vidal, Fernando Jiménez and Antonio Skarmeta-Gómez
Type	Journal
Journal	Energy and Buildings
Impact factor (2018)	4.495
Rank	Q1
Publisher	Elsevier
Volume	196
Pages	71-82
Year	2019
Month	August
ISNN	0378-7788
DOI	doi: 10.1016/j.enbuild.2019.05.021
URL	https://www.sciencedirect.com/science/article/pii/S0378778818338775
State	In Press, Accepted Manuscript
Author's contribution	The PhD student, Aurora González Vidal, is the main author of the paper



A methodology for energy multivariate time series forecasting in smart buildings based on feature selection



Aurora González-Vidal*, Fernando Jiménez, Antonio F. Gómez-Skarmeta

Department of Information and Communication Engineering, Faculty of Informatics, University of Murcia, 30100, Spain

ARTICLE INFO

Article history:

Received 25 December 2018

Revised 22 April 2019

Accepted 10 May 2019

Available online 11 May 2019

Keywords:

Feature selection

Energy efficiency

Time series

Smart buildings

Smart cities

ABSTRACT

The massive collection of data via emerging technologies like the Internet of Things (IoT) requires finding optimal ways to reduce the created features that have a potential impact on the information that can be extracted through the machine learning process. The mining of knowledge related to a concept is done on the basis of the features of data. The process of finding the best combination of features is called feature selection. In this paper we deal with multivariate time-dependent series of data points for energy forecasting in smart buildings. We propose a methodology to transform the time-dependent database into a structure that standard machine learning algorithms can process, and then, apply different types of feature selection methods for regression tasks. We used Weka for the tasks of database transformation, feature selection, regression, statistical test and forecasting. The proposed methodology improves MAE by 59.97% and RMSE by 40.75%, evaluated on training data, and it improves MAE by 42.28% and RMSE by 36.62%, evaluated on test data, on average for 1-step-ahead, 2-step-ahead and 3-step-ahead when compared to not applying any feature selection methodology.

© 2019 Elsevier B.V. All rights reserved.

1. Introduction

Energy efficiency is the goal to optimise the amount of energy required to provide products and services. Energy consumption is increasing with the growing population and intensified in highly populated parts of cities [1]. Energy efficiency is in the interest of everyone, from individuals to governments, since it yields economical savings, reduces greenhouse gas emissions and alleviates energy poverty [2]. In order to achieve energy efficiency, smart grids, open data platforms and networked transport systems are proliferating for managing and monitoring resources automatically. This provides the emergence of smart cities, which thanks to the collection of data using sensors that are interconnected through the internet (*Internet of Things*) allow the extraction of insights that are necessary in order to provide better services to the citizens that also include energy efficiency.

The huge amounts of data that are collected via the IoT are consequently analysed in order to extract the knowledge necessary for achieving energy efficiency. For example, the main source of data in smart grids is the *Advanced Metering Infrastructure* (AMI) which deploys a large number of smart meters at the end-user side. The amounts of AMI data grow very quickly. If data is collected ev-

ery 15 mins by 1 million metering devices, the total records reach 35.04 billion and the volume of meter reading data surge up to 2920 Tb [3,4]. An actual example of this is the *Electricity Load Diagrams* dataset from the *UCI Machine Learning Repository* [5] that contains 140,256 attributes.

However, in order to realise such analysis it is desirable to reduce the dimensionality of the data for easing the models performance. In order to do so there exist several approaches such as *segmentation* and *representation of attributes* [6] or *feature selection* [7]. We are going to focus on feature selection since it has shown its effectiveness in many applications by building simpler and more comprehensive models, improving learning performance, and preparing clean, understandable data [8].

In this work, we use *time series data* from the Chemistry Faculty of the University of Murcia to generate energy consumption forecasts [9,10]. *Time series forecasting* differs from typical machine learning applications where each data point is an independent example of the concept to be learned, and the ordering of data points within a dataset does not matter. For this reason, standard machine learning methods should not be used directly to analyze time series data. In this paper, we propose a methodology to, firstly, transform the time series into a form that standard machine learning algorithms can process, and then, systematically apply a set of feature selection methods for regression that includes *univariate*, *multivariate*, *filter* and *wrapper* methods [11]. Time series data is transformed by removing the temporal ordering of individual

* Corresponding author.

E-mail address: aurora.gonzalez2@um.es (A. González-Vidal).

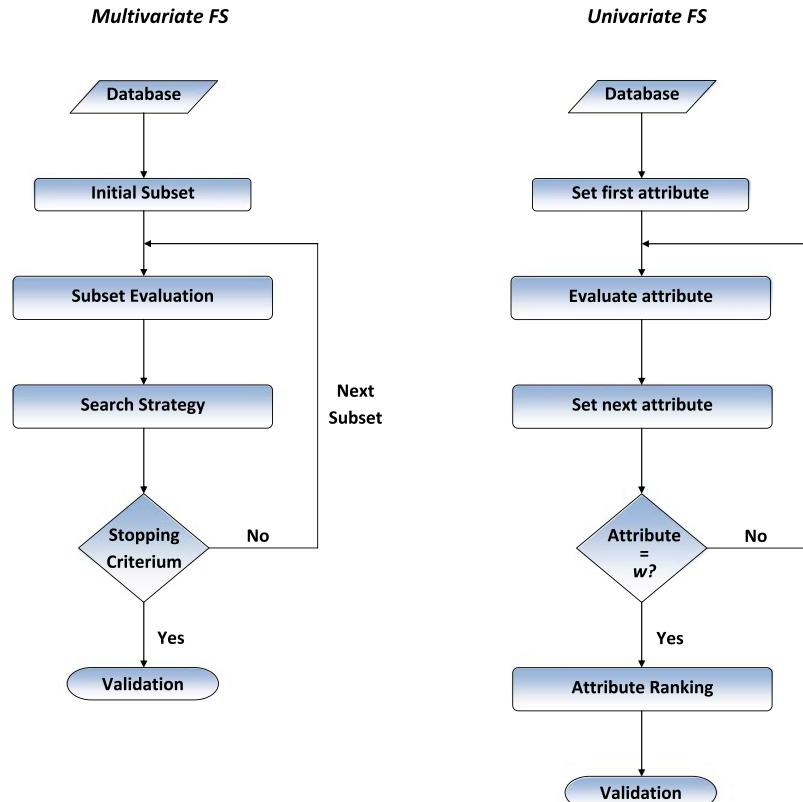


Fig. 1. General schemes for multivariate and univariate feature selection.

input examples and adding a set of delays to the input which are called *lagged attributes* and provide the temporal information. The methodology also allows dealing with *intervention attributes*, which are to be considered external to the data transformation and closed-loop forecasting processes. This approach to time series forecasting is more powerful and more flexible than classical statistics techniques such as ARMA and ARIMA [12]. Feature selection methods are applied for the selection of both lagged and intervention attributes. *Random Forest*, *instance-based learning* and *linear regression* algorithms are used for regression with the different reduced databases. Finally, the best reduced database together with the best regression algorithm are used for the predictions *1-step-ahead*, *2-step-ahead* and *3-step-ahead* evaluated in training data and test data, and the results are compared with the predictions obtained with the original database. The experiments have been carried out using the *Waikato Environment for Knowledge Analysis (Weka)* [13].

With this background, the paper has been organized as follows: **Section 2** describes the background of the paper; **Section 3** proposes a methodology for the energy efficiency analysis in smart buildings based on feature selection; **Section 4** analyzes and discusses the results; **Section 5** introduces some other methods used for the same purpose in the literature, and finally **Section 6** depicts the main conclusions and future work.

2. Background

This section defines the concept of feature selection and their categorization, shows some important related works in literature,

emphasizes the contributions of the paper, and describes the dataset used for experiments.

2.1. Feature selection

Feature Selection (FS) is defined in [7] as the process of eliminating features from the database that are irrelevant to the task to be performed. FS facilitates data understanding, reduces the measurement and storage requirements, the computational process time, and the size of a dataset, so that model learning becomes an easier process. An FS method is basically a *search strategy* where the performance of candidate subsets is measured with a given *evaluator*. The search space for candidate subsets has cardinality $O(2^w)$, where w is the number of features. A *stopping criterion* establishes when the FS process must finish. It can be defined as a control procedure that ensures that no further addition or deletion of features produces a better subset, or it can be as simple as a counter of iterations. FS methods are typically categorized into *wrapper*, *filter* and *embedded*, *univariate* and *multivariate* methods. *Wrapper methods* [14] use a predetermined learning algorithm to determine the quality of selected features according to an evaluation metric [15]. *Filter methods* apply statistical measures to evaluate the set of attributes [16–18]. *Embedded methods* achieve model fitting and FS simultaneously [19]. *Multivariate methods* evaluate features in batches. *Univariate methods* evaluate each feature independently. **Fig. 1** shows general schemes for multivariate and univariate FS.

2.2. Related work

We have carried out an extensive search in order to find other academic works that have solved a similar problem than ours. To-

gether with the works that address FS for energy consumption time series, we have also considered important to review FS for energy consumption when not treated as time series, and FS for time series problems in general, i.e. other approaches not specifically related to energy consumption.

The first paper that studied how the selection of subsets of features associated with building energy behaviours influences a machine learning model performance for energy consumption prediction used some filter methods for FS and support vector regression for forecasting [20]. A bit later, in the thesis [21], *Fast Correlation-Based Filter (FCBF)* is used for FS in load prediction error problems in four building areas. A meteorological dataset from several locations and also, the geographical factor are exploited by selecting variables from different locations. The baseline comparisons are done with e-SVR. According to this work, how the relationships between features change with distance motivates a greedy FS method for the electrical load forecasting. In the works [22,23], *correlation and principal components analysis (PCA)* are used for FS and transformation.

Feature selection for time series prediction has been carried out using neural networks [24]. By combining contemporaneous and lagged realisations of the independent variables and lagged dependent variables more general models of dynamic regression, autoregressive (AR) transfer functions and intervention models are constructed. It has also been done using the Granger causality discovery [25] to identify important features with effective sliding window sizes, considering the influence of lagged observations of features on the target time series.

Other studies have searched for the optimal time-windows and time lags for each variable based on feature pre-processing and sparse learning in order to configure the input dataset [26].

In other works, the forecasting of solar radiation time series is enhanced by using a train set of bootstrapped Support Vector Machines in order to perform FS [27]. They assure that this method is more robust than a regular FS approach because using the later, small changes on the train set may produce a huge difference on the selected attributes. Other studies related to solar radiation prediction mask the inputs as a FS step [28]. They create their own features by defining night, sunrise, day and sunset according to the moment that their instruments perceive those. This provides certain improvements on forecast accuracy. A data-driven multi-model wind prediction methodology using a two-layer ensemble machine learning technique is developed in [29]. A deep FS framework is employed where four different approaches are used in order to get the input vector: *PCA*, *Granger Causality Test*, *Autocorrelation and Partial Autocorrelation Analysis*, and *Recursive Feature Elimination*. Another ensembles way of selecting features is presented in [30] and it is used for predicting the amount of incoming calls for an emergency call center in a time series manner. They use five algorithms (*ReliefF*, *PCA*, *Freq. Discretization*, *Information Gain* and *K-means*) that are different in nature and combine the rankings computed grouping similar approaches and computing new weights as the mean of the individual weights. After that, all variables that are ranked among the top five positions in at least three of the groups compound the selected features. In the thesis work [31] they present three case studies in which FS is a step in the model creation. They used the following methods: *sequential forward/ backward selection (SFS, SBS)*, *sequential forward/ backward floating selection (SFFS, SBFS)*, the *n best features selection (nBest)* and the *best individual features*.

The main data characteristics of energy time series have been specifically analysed in [32]. To explore such data from different perspectives they consider two main categories: nature (nonstationarity, nonlinearity and complexity characteristics) and pattern (cyclicity, mutability or saltation, and randomicity or noise pattern). After that, FS for electricity load forecasting was done in a

time series manner using correlation and instance based methods [33]. In [34] it is presented a survey on data mining techniques for time series forecasting of electricity. The survey focuses on the characteristics of the models and their configuration. *Wrapper methods*, *Artificial Neural Networks*, *mutual information*, *autocorrelation* and *ranking based methods* are mentioned as FS techniques used in the prediction of energy consumption. Finally, the work [2] uses temperature time series together with day of the week in order to estimate energy consumption.

2.3. Contributions of the work

Regarding the papers that also focus on feature selection for time series prediction [24,25], we highlight the aspects that make our work standout the previous. The focus of Crone and Kourentzes [24] is narrowed to neural networks which it is not the best for every situation since usually, neural networks are more computationally expensive and require much more data than traditional algorithms. Also, the *No Free Lunch theorem* [35] suggests that there can be no single method which perform bests on all datasets. [25] is focused on the Granger causality as feature selection so none of them provide a systematic comparison between the possibilities available in the feature selection field. We have carried out such comparisons by combining univariate, multivariate, filter and wrapper methods and also we have checked the performance of the several databases obtained in a plethora of prediction algorithms.

Additionally, we have followed a multi-objective evolutionary search strategy which is more advance than the other procedures allowing to minimise *Root-Mean-Square Error (RMSE)* and *Mean Absolute Error (MAE)* and also the number of variables. In addition, in this work we use a multi-objective evolutionary search strategy, which simultaneously minimizes the error - *Root-Mean-Square Error (RMSE)* or *Mean Absolute Error (MAE)* - and minimizes the number of attributes, unlike the single-objective search strategies that only minimize the error. Evolutionary techniques are metaheuristics for global search, unlike other deterministic search strategies that tend to fall in local optima. We have measured the feature selection effectivity using both metrics *RMSE* and *MAE*. We have obtained that minimizing *MAE* provides better results in the posterior prediction phase. In the feature selection process, the methodology also allows dealing with intervention attributes, which are to be considered external to the data transformation.

Finally, to the best of our knowledge this is the first time that a multivariate time series feature selection methodology in proposed for predicting energy consumption in smart buildings.

2.4. Energy efficiency dataset

The reference building in which the energy consumption forecasting has been carried out is the Chemistry Faculty of the University of Murcia, which is a building used as a pilot for the H2020 ENTROPY project (Grant Agreement No 649849).¹

The dataset is composed of 5088 observations of 50 attributes that are measured hourly from 2016-02-02 00:00:00 until 2016-09-06 23:00:00, where time-stamps from 2016-02-05 00:00:00 until 2016-05-07 23:00:00 are missing data. Table 1 shows the number, name and sources of the dataset attributes. The output attribute is the *energy consumption* measured in KWh. Attributes *datetime* ("yyyy-MM-dd HH:mm:ss"), *season* (1–4), *day of the week* (1–7), and *holiday* (0,1) have been extracted from the date's observation. We have used meteorological data gathered from several sources and stations with the purpose to select the attributes

¹ <http://entropy-project.eu>.

Table 1

Attributes and data sources of the energy consumption dataset used in this paper.

Number	Name	Data source
1–8	realWU_temp, realWU_feels, realWU_dewp, realWU_hum, realWU_wspd, realWU_visib_km, realWU_mslp, realWU_prep_1h	Weather Underground
9–17	pr_temp, pr_feels, pr_dewp, pr_hum, pr_pop, pr_wspd, pr_wdir_deg, pr_sky, pr_mslp	Weather Underground
18–33	stMO12_IMI_tmmed, stMO12_IMI_tmin, stMO12_IMI_hrmed, stMO12_IMI_hrmax, stMO12_IMI_hrmin, stMO12_IMI_radmed, stMO12_IMI_radmax, stMO12_IMI_vvmed, stMO12_IMI_vvmax, stMO12_IMI_dvmed, stMO12_IMI_prec, stMO12_IMI_dewpt, stMO12_IMI_dpvt, stMU62_IMI_tmmed, stMU62_IMI_tmax, stMU62_IMI_tmin, stMU62_IMI_hrmed, stMU62_IMI_hrmax, stMU62_IMI_hrmin, stMU62_IMI_radmed, stMU62_IMI_radmax, stMU62_IMI_vvmed, stMU62_IMI_vvmax, stMU62_IMI_dvmed, stMU62_IMI_prec, stMU62_IMI_dewpt, stMU62_IMI_dpvt	IMIDA MO12
34–45	energy	IMIDA MU62
46	season, day_of_the_week, holiday, datetime	Output attribute Date's observations
47–50		

from the most explanatory source according to our feature extraction analysis.

Weather Underground² is a web service that through its API provides the following real values: *temperature* ($^{\circ}\text{C}$), *apparent temperature* ($^{\circ}\text{C}$), *humidity* (%), *wind speed* (m/s), *mean sea level pressure* (mbar), *visibility* (km) and *precipitations in last hour* (mm). We also use *one-hour predictions* for the first six previous attributes, together with *probability of precipitations* (%), *sky cover* (%) and *wind direction* (degrees).

IMIDA³ (The Research Institute of Agriculture and Food Development of Murcia) provides real time records of weather. We have selected two weather stations regarding proximity to the building: MO12 and MU62 and from each of them we have collected the following variables: *temperature* (mean, minimum and maximum) ($^{\circ}\text{C}$), *humidity* (mean, minimum and maximum) (%), *radiation* (mean and maximum) (w/m^2), *wind speed* (mean and maximum) (m/s^2), *wind direction* (mean) (degrees), *precipitation* (mm), *dew point* ($^{\circ}\text{C}$) and *vapour pressure deficit* (kPa).

3. A methodology for energy multivariate time series forecasting based on feature selection

We have followed the methodology shown in the Fig. 2 to perform the energy time series forecasting. The following six steps have been systematically applied: database transformation, feature selection, regression, statistical tests, decision making and forecasting. Next, each step is described separately, and some of the names of the Weka classes and methods that are required throughout the process are indicated.

3.1. Database transformation

The first step of our methodology is to transform the database by creating lagged versions of variables for use in the time series problem. For this, the following steps are carried out:

1. Set an *artificial time-stamp* with start value 1. We use an artificial time index for convenience. In this way, no instances are inserted in the training data for the missing time-stamps.
2. Set the attributes to *lag*. The system can jointly model multiple attributes to lag simultaneously in order to capture dependencies between them. Because of this, modelling several series simultaneously can give different results for each series than modelling them individually. The rest of the attributes (*non lagged* attributes) are considered as *intervention* attributes (also called *overlay data*). We set attributes 1 to 46 as lagged attributes. Attributes 47, 48 and 49 are intervention attributes.

² <https://www.wunderground.com/>.

³ <http://www.imida.es/>.

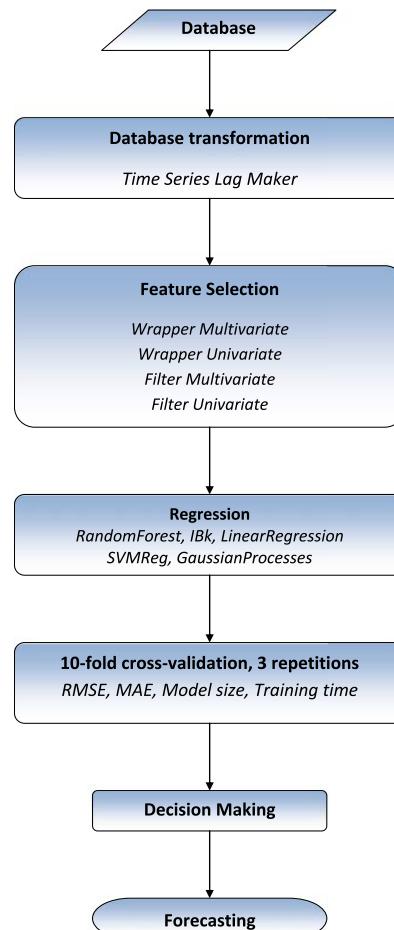


Fig. 2. Methodology for feature selection for energy time series forecasting.

3. Set the minimum previous time step to create a lagged field. We set to 0 the *minimum lag length* to create. A value of 0 means that a lagged variable will be created that holds target values at time 0.
4. Set the maximum previous time step to create a lagged variable. We set to 3 the *maximum lag length* to create. A value of 3 means that a lagged variable will be created that holds target values at time -3. All time periods between the minimum and maximum lag will be turned into lagged variables. In this

Table 2
Proposed feature selection methods for energy time series forecasting.

Database #Id.	Type of FS method	Name	Search strategy	Evaluator
#1	Wrapper Multivariate	MOES-RF-MAE	MultiObjectiveEvolutionarySearch	RandomForest (MAE)
#2	Wrapper Multivariate	MOES-RF-RMSE	MultiObjectiveEvolutionarySearch	RandomForest (RMSE)
#3	Wrapper Multivariate	MOES-IBk-RMSE	MultiObjectiveEvolutionarySearch	IBk (RMSE)
#4	Wrapper Multivariate	MOES-LR-MAE	MultiObjectiveEvolutionarySearch	LinearRegression (MAE)
#5	Wrapper Univariate	RANKER-RF-RMSE	Ranker	RandomForest (RMSE)
#6	Filter Multivariate	GS-CFSSE	GreedyStepwise	CfsSubsetEval
#7	Filter Univariate	RANKER-RFAE	Ranker	ReliefFAttributeEval
#8	Filter Univariate	RANKER-PCA	Ranker	PrincipalComponents

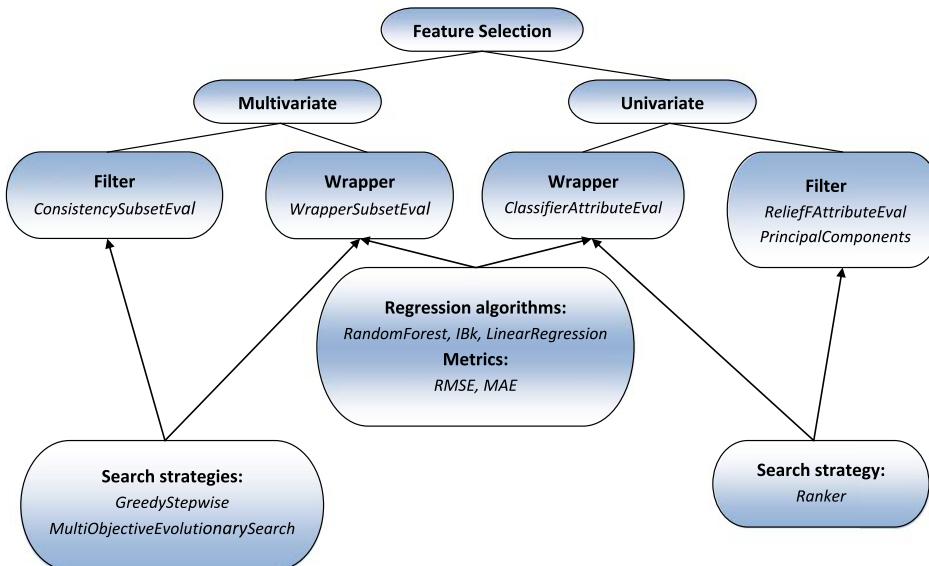


Fig. 3. Organization chart of the proposed feature selection methods for energy time series forecasting.

way, for example the variable *energy* will be transformed into 4 lagged variables *Lag_energy+0* (equivalent to the variable *energy*), *Lag_energy-1*, *Lag_energy-2* and *Lag_energy-3*.

5. Perform database transformation. A database of 189 attributes has been generated with the transformation.
6. Save transformed database with the name *TransformedDatabaseAux*. This auxiliary transformed database will be used later in the forecasting phase.
7. Remove *datetime* attribute. When using an artificial time index, the attribute *ArtificialTimeIndex* is added to the database, so the attribute *datetime* must be removed.
8. Save the final transformed database with the name *TransformedDatabase*. The final number of attributes of the transformed database is 188.

We use the class *weka.classifiers.timeseries.core.TSLagMaker* for this task. Data transformation can be done from the plugin tab in Weka's graphical "Explorer" user interface, or and using the API through a Java program.

3.2. Feature selection

Once the task of transforming the database is done, the next step is to apply FS on the *TransformedDatabase2* database. In Weka, FS is implemented with the class *weka.attributeSelection.AttributeSelection* through two components: the *search strategy* (*weka.attributeSelection.ASSearch* abstract class) and the *evaluator* (*weka.attributeSelection.ASEvaluation* abstract class). This allows users and programmers to configure a multitude of dif-

ferent methods for FS, both filter and wrapper, univariate and multivariate. Evaluators with names ending in *SubsetEval* configure multivariate methods, whereas those with names ending in *AttributeEval* configure univariate methods. For multivariate wrapper FS methods, the *weka.attributeSelection.WrapperSubsetEval* class evaluates attribute sets by using a learning scheme with cross-validation and a performance measure. For univariate wrapper FS methods, the *weka.attributeSelection.ClassifierAttributeEval* class evaluates the worth of an attribute by using a user-specified classifier, cross-validation and a performance evaluation measure to use for selecting attributes. Since the FS and classification processes must be executed in batch mode, Weka offers the class *weka.classifiers.meta.AttributeSelectedClassifier* which is a meta-classifier where dimensionality of data is reduced by attribute selection before being passed on to a learning algorithm. Table 16 summarizes the packages and classes for FS in Weka used in this paper.

We applied eight different FS methods for regression shown in Table 2 and graphically in Fig. 3. In Table 2, Database #Id denotes the identifier of the reduced database generated with each FS method. Each FS method is the result of a specific choice among the search strategy and the evaluator. We considered for this research five wrapper FS methods and three filter FS methods. Among them, five FS methods are multivariate and three FS methods are univariate. Table 14 shows the parameters used for each FS method. Next we show the search strategies and evaluators considered in this paper.

3.2.1. Search strategies

As multivariate FS methods, we use a *probabilistic search strategy* and a *deterministic search strategy*. *MultiObjectiveEvolutionarySearch* [36] is the probabilistic strategy, and *GreedyStepwise* [37] is the deterministic strategy. *MultiObjectiveEvolutionarySearch* use multi-objective evolutionary computation where two objectives are optimized: the first one is a performance metric or statistical measure chosen by user with the evaluator, while the second one is the attribute subset cardinality, and it is to be minimized. The final output is given by the non-dominated solutions in the last population having the best fitness score for the first objective. *MultiObjectiveEvolutionarySearch* class has two multi-objective evolutionary algorithms implemented, *ENORA* and *NSGA-II*. *ENORA* is our MOEA, on which we are intensively working over the last decade. We have applied *ENORA* to constrained real-parameter optimization [38], fuzzy optimization [39], fuzzy classification [40], feature selection for classification [41] and feature selection for regression [42]. In this paper, we apply it to feature selection for regression in times series forecasting. *NSGA-II* algorithm has been designed by Deb et al. and has been proved to be a very powerful and fast algorithm in multi-objective optimization contexts of all kinds. In [42] is statistically tested that *ENORA* performs better than *NSGA-II* in terms of *hypervolume* [43,44] for regression tasks, for which we have decided to use *ENORA* in this work. *GreedyStepwise* performs a greedy forward or backward search through the space of attribute subsets, stopping when the addition (forward direction) or deletion (backward direction) of any of the remaining attributes results in a decrease in evaluation, thus, it has no backtracking capability.

For univariate FS methods, *Ranker* method [45] is required. *Ranker* method ranks attributes by their individual evaluations. A threshold, or the number of attributes to retain, allows reducing the attribute set.

3.2.2. Evaluators

We considered the multivariate filter evaluator *ConsistencySubsetEval* [46]. *ConsistencySubsetEval* scores a subset of features as a whole, by projecting the training instances according to the attribute subset, and considering the consistency of class values in the obtained instance sets. As far as univariate filter evaluators are concerned, *RelieffAttributeEval* [47] and *PrincipalComponents* [48] were considered. *RelieffAttributeEval* evaluates the worth of an attribute by repeatedly sampling an instance and considering the value of the given attribute for the nearest instance of the same and different class. Can operate on both discrete and continuous class data. *PrincipalComponents* performs a principal components analysis and transformation of the data. Dimensionality reduction is accomplished by choosing enough eigenvectors to account for some percentage of the variance in the original data (default 95%). Attribute noise can be filtered by transforming to the principal components space, eliminating some of the worst eigenvectors, and then transforming back to the original space.

We use the wrapper *WrapperSubsetEval* [14] for multivariate FS methods and *ClassifierAttributeEval* [49] for univariate FS methods in conjunction with the predictors *RandomForest* [50], *IBk* [51] and *LinearRegression* [52], and with the metrics *RMSE* and *MAE* [53]. *RandomForest* is an *ensemble learning* method which constructs a forest of random trees with controlled variance, for classification or regression purposes. *IBk* is a simple instance-based learner that uses the class of the nearest k training instances for the class of the test instances and it is also valid for regression. *LinearRegression* uses the *Akaike* criterion for model selection, and is able to deal with weighted instances. Note that not all regression algorithms can be used as evaluators in wrapper FS methods due to their high computational time. *RandomForest*, *IBk* and *LinearRegression* are learning algorithms that offer a good compromise between

performance and computational time so they are suitable as evaluators in wrapper FS methods.

3.3. Regression

Once FS was made, the next step was to perform regression with the reduced and *TransformedDatabase2* databases using different regression algorithms. We considered *RandomForest*, *IBk* and *LinearRegression* since these algorithms were used as evaluators in the wrapper FS methods. Additionally we used *Support Vector Machines* [54] and *Gaussian Processes* [55], which are widely used for time series forecasting [56], concretely the *Weka* implementations *SMOreg* and *GaussianProcesses*. *SMOreg* [57] implements the support vector machine for regression. The parameters can be learned using various algorithms, being *RegSMOImproved* the most popular algorithm. *GaussianProcesses* implements Gaussian processes for regression without hyperparameter-tuning. To make choosing an appropriate noise level easier, this implementation applies normalization/standardization to the target attribute as well as the other attributes. Both *SMOreg* and *GaussianProcesses* can use *PolyKernel*, *PrecomputedKernelMatrixKernel*, *Puk*, *RBFKernel* or *StringKernel*. Table 15 shows the parameters used for the regression methods. Tables 3 and 4 show the evaluation in full training set for the *RMSE* and *MAE* metrics respectively.

3.4. Statistical test

In order to detect over-fitting and prediction ability, the regression models have also been evaluated with cross-validation. Tables 5–8 show the evaluation in 10-fold cross-validation, 3 repetitions (a total of 30 models with each regression algorithm in each database), for the metrics *RMSE*, *MAE*, *Serialized_Model_Size* and *User_Time_training*⁴ respectively. The result of the experiment has been analysed through a *paired t-test (corrected)*, with 0.05 significance, being #1 the test base. For each result, a mark * denotes that the result is statistically worse than the test base; similarly, a mark v denotes a statistically better result, and no mark denotes no statistically meaningful difference.

3.5. Decision making

Looking at Tables 5 to 8 we can make a decision for choosing the best reduced database and, therefore, the best FS method. The best results have been obtained with the FS method *MOES-RF-MAE* (database #1) when *RandomForest* is used as regression algorithm, which show statistically significant differences with respect to the rest of the analysed FS methods for the *MAE* performance metric. For *RMSE* performance metric, FS method *MOES-RF-MAE* is also superior to the rest of FS methods, with statistically significant differences except for the FS method *MOES-RF-RMSE*. With respect to the *Serialized_Model_Size* and *UserCPU_Time_training* performance metrics, the results of the FS method *MOES-RF-MAE* by using *RandomForest* are acceptable in comparison to the rest of the methods. We can then choose the FS method *MOES-RF-MAE* and the database #1 for the final forecasting process.

Table 9 shows the selected attributes with *MOES-RF-MAE*. Table 9 shows the selected attributes and their ranks and importances for each of the datasets. The rank and importance of the attributes has been obtained through a univariate wrapper feature selection method, where the search strategy is the *ranker* method, and the evaluator is *ClassifierAttributeEval* with classifier = *RandomForest* (with default parameters), *evaluationMeasure* = *MAE*, and

⁴ Intel (R) Core (TM) i5-4460 @ 3.20 GHz 3.20 GHz RAM 8.00 GB Operating Systems 64 bits, processor × 64.

Table 3
RMSE with full training set.

	#1	#2	#3	#4	#5	#6	#7	#8	TransformedDatabase
<i>RandomForest</i>	5.0930	5.0286	5.2923	5.5701	5.9543	7.2083	5.3704	13.5680	7.8809
<i>IBk</i>	3.0201	3.5134	2.4937	1.3826	2.7927	1.5093	1.4044	0.0000	2.1045
<i>LinearRegression</i>	19.5455	18.4759	18.7110	18.3092	18.7878	22.1723	18.2264	53.5429	17.2416
<i>SMOref</i>	20.2988	19.1824	19.2648	19.0136	19.3193	23.1186	19.1566	55.5580	18.4857
<i>GaussianProcesses</i>	21.9302	21.7321	24.7275	19.4525	18.9750	22.1774	18.4000	54.5592	17.4686

Table 4
MAE with full training set.

	#1	#2	#3	#4	#5	#6	#7	#8	TransformedDatabase
<i>RandomForest</i>	2.5667	2.6015	2.7341	2.8990	3.1639	3.7101	2.7528	8.5778	4.7050
<i>IBk</i>	0.0730	0.0824	0.0465	0.0271	0.0559	0.0231	0.0284	0.0000	0.0419
<i>LinearRegression</i>	11.2387	10.0955	10.2126	9.6797	10.1295	13.2144	10.4735	38.2477	10.1297
<i>SMOref</i>	10.0226	8.9893	9.0665	8.9401	9.0363	11.5677	8.9673	36.6050	8.7132
<i>GaussianProcesses</i>	15.2061	15.0741	17.8833	11.9989	10.5405	13.3437	10.9358	38.7237	10.5050

Table 5
RMSE with 10-fold cross-validation (3 repetitions).

	#1	#2	#3	#4	#5	#6	#7	#8	TransformedDatabase
<i>RandomForest</i>	12.6685	12.9133	13.3814*	14.4111*	15.4203*	18.7996*	13.9174*	36.7834*	21.3455 v
<i>IBk</i>	17.7612	20.8112*	17.2423	25.0447*	25.8680*	25.7792*	22.7562*	37.8315*	29.5936*
<i>LinearRegression</i>	19.3960	18.3234 v	18.5017 v	18.1808 v	18.6416	22.0898*	18.1092 v	53.6083*	17.7597 v
<i>SMOref</i>	20.0636	18.8337 v	18.9237 v	18.6714 v	18.9697 v	23.0016*	18.8051 v	55.5770*	18.2458 v
<i>GaussianProcesses</i>	21.9231	21.7133	24.7083*	19.4114 v	18.8832 v	22.1160	18.3440 v	54.6361*	17.8482 v

Table 6
MAE with 10-fold cross-validation (3 repetitions).

	#1	#2	#3	#4	#5	#6	#7	#8	TransformedDatabase
<i>RandomForest</i>	5.8264	6.0012*	6.2785*	6.8621*	7.5242*	9.0191*	6.4675*	23.3071*	12.6164*
<i>IBk</i>	8.8796	10.0150*	8.6797	13.1307*	13.3098*	12.7038*	11.0927*	17.4372*	14.3159*
<i>LinearRegression</i>	11.2708	10.1276 v	10.2363 v	9.7287 v	10.1738 v	13.2454*	10.5091 v	38.3269*	10.6126 v
<i>SMOref</i>	10.0410	9.0122 v	9.0806 v	8.9702 v	9.0669 v	11.5835*	8.9962 v	36.7292*	8.9314 v
<i>GaussianProcesses</i>	15.3332	15.1402 v	17.9369*	12.0857 v	10.6294 v	13.3943 v	11.0307 v	38.8031*	10.9028 v

Table 7
Serialized_Model_Size ($\times 10^6$ bytes) with 10-fold cross-validation (3 repetitions).

	#1	#2	#3	#4	#5	#6	#7	#8	TransformedDatabase
<i>RandomForest</i>	11.9955	11.9149 v	13.4332*	12.6169*	14.7657*	16.3795*	14.9757*	20.3033*	16.7962*
<i>IBk</i>	0.5064	0.5796*	0.4330 v	0.8735*	0.5798*	0.4329 v	0.5799*	0.5804*	0.7081*
<i>LinearRegression</i>	0.1278	0.1274 v	0.1275 v	0.1285*	0.1285*	0.1275 v	0.1285*	0.1278*	0.1604*
<i>SMOref</i>	0.1734	0.7654*	0.6609 v	0.1060*	0.1028*	0.8805*	1.1013*	0.7658*	7.6196*
<i>GaussianProcesses</i>	168.4590	168.4900*	168.3855 v	168.7841*	168.7528*	168.6051*	168.8259*	168.4904*	175.3427*

Table 8
UserCPU_Time_training (seconds) with 10-fold cross-validation (3 repetitions).

	#1	#2	#3	#4	#5	#6	#7	#8	TransformedDatabase
<i>RandomForest</i>	0.9474	1.0349*	0.7792 v	1.3432*	0.9714	0.5708 v	0.7802 v	1.6078*	3.0609*
<i>IBk</i>	0.0005	0.0005	0.0000	0.0000	0.0005	0.0000	0.0016	0.0000	0.0005
<i>LinearRegression</i>	0.0042	0.0109	0.0026	0.0125*	0.0089	0.0063	0.0115	0.0057	4.2172*
<i>SMOref</i>	31.9255	29.0380 v	26.4307 v	62.5958*	87.1901*	75.5521*	141.1615*	9.0995 v	1626.4151*
<i>GaussianProcesses</i>	115.4714	115.3620	115.4219	115.5542	110.7714 v	110.4302 v	110.5990 v	110.6505 v	114.0536

`leaveOneAttributeOut = true`. An attribute is evaluated by measuring the impact of leaving it out from the full set.

3.6. Forecasting

Finally, in this section we analyse the prediction ability of the forecaster obtained with the selected attributes. We use the class `weka.classifiers.timeseries.WekaForecaster` for this task. Forecasting can be done from the plugin tab in Weka's graphical "Explorer" user interface, or using the API through a Java program. When an evaluation is performed, firstly the forecaster is trained on the

data, and then it is applied to make a forecast at each time point (in order) by stepping through the data. These predictions are collected and summarized, using MAE and RMSE metrics, for each future time step predicted. We use in this paper three time units to forecasts, i.e. all the 1-step-ahead, 2-steps-ahead and 3-steps-ahead predictions are collected and summarized. This allows us to see, to a certain degree, how forecasts further out in time compare to those closer in time.

Tables 10 and 12 show the evaluation of the forecaster, with the database #1, on training data (70%) and test data (30%) respectively. The last 500 training data and the first 500 test data of these

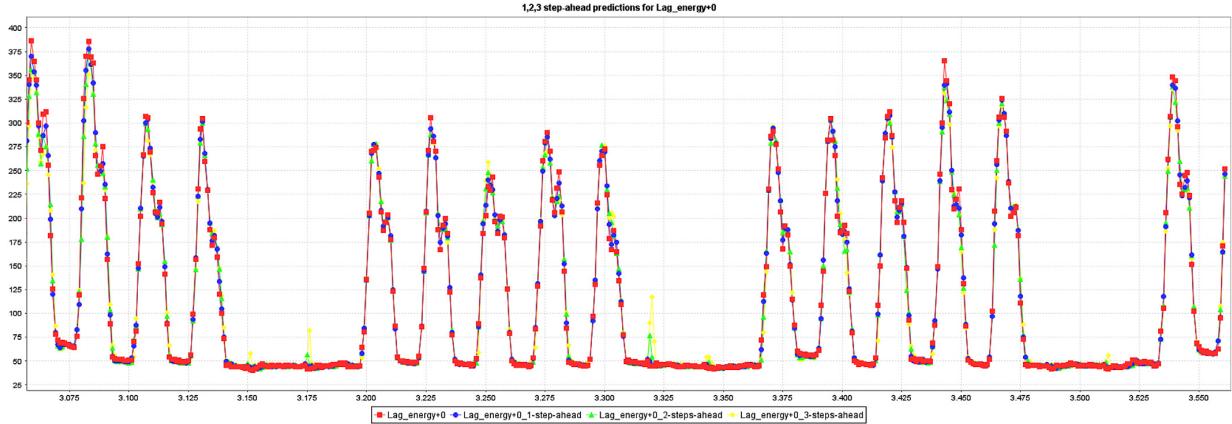


Fig. 4. 1,2,3-step-ahead predictions for Lag-energy+0 evaluated on the last 500 training data with *RandomForest* - database #1.

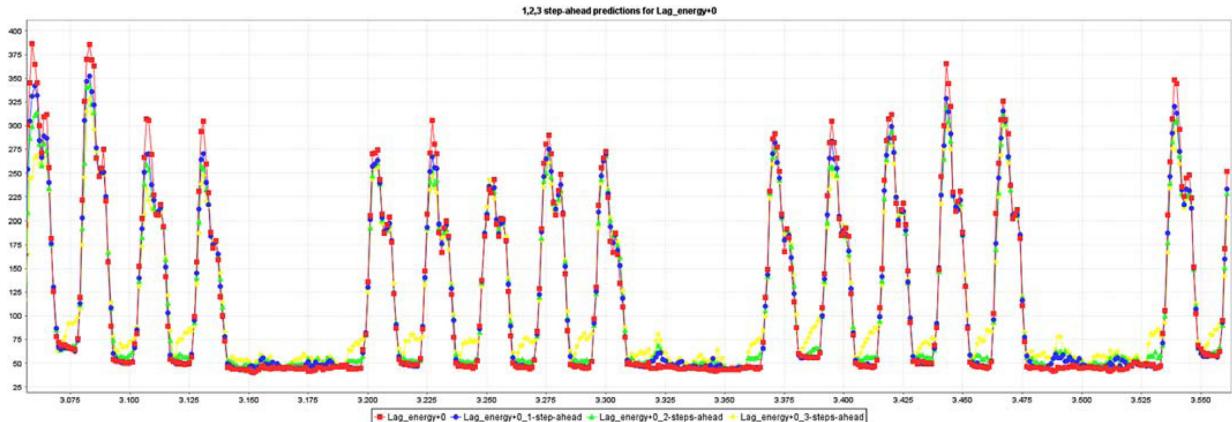


Fig. 5. 1,2,3-step-ahead predictions for Lag-energy+0 evaluated on the last 500 training data with *RandomForest* - *TransformedDatabase*.

Table 9
Selected attributes with MOES-RF-MAE (database #1) and their ranks.

Input attribute	Rank	Importance
Lag_energy-1	1	7.398
Lag_stMO12_IML_radmax+0	2	1.337
holiday	3	0.367
Lag_energy-3	4	0.357
ArtificialTimeIndex	5	0.302
Lag_stMO12_IML_radmed-3	6	0.273
Lag_pr_feels-2	7	0.248
Lag_pr_temp-2	8	0.172

Table 10
Evaluation on training data (3562 instances) with *RandomForest* - database #1.

	1-step-ahead	2-steps-ahead	3-steps-ahead	Average
Number of instances	3559	3558	3557	—
MAE	2.6684	4.3897	5.8962	4.3181
RMSE	5.3008	9.4352	13.0256	9.2539

evaluations are also shown graphically in Figs. 4 and 6 respectively. To verify if the FS process has been effective both for the reduction of the complexity of the model and for the increase of its predictive capacity, the forecasting process has also been carried out on the database *TransformedDatabase* (with all lagged variables and all overlay variables). Tables 11 and 13 show the evaluation of the

forecaster with the database *TransformedDatabase*, and Figs. 5 and 7 show graphically the evaluation of the last 500 training data and the first 500 test data respectively.

4. Analysis of results and discussion

When observing the results of the experiments carried out using the proposed methodology, the following statements can be derived:

A. Regarding the FS process:

- As expected, wrapper FS methods show better performance than filter FS methods, and multivariate FS methods show better performance than univariate FS methods. Multivariate methods can identify interaction amongst features simultaneously, specially wrapper-based FS methods [58]. To make it possible, multivariate methods evaluate the relevance of sets of features to determine which are the bests according to certain performance measure for a given task. However, multivariate wrapper feature selection methods present a high computation costs, since the number of possible subsets of feature is very high (2^w , being w the number of features) making the problem of finding the best subsets to be NP-Hard. To reduce the computational time, some deterministic search strategies, such as *GreedyStepwise*, can be used. The main disadvantage of these deterministic search techniques is that hidden and basic

Table 11Evaluation on training data (3562 instances) with *RandomForest - TransformedDatabase*.

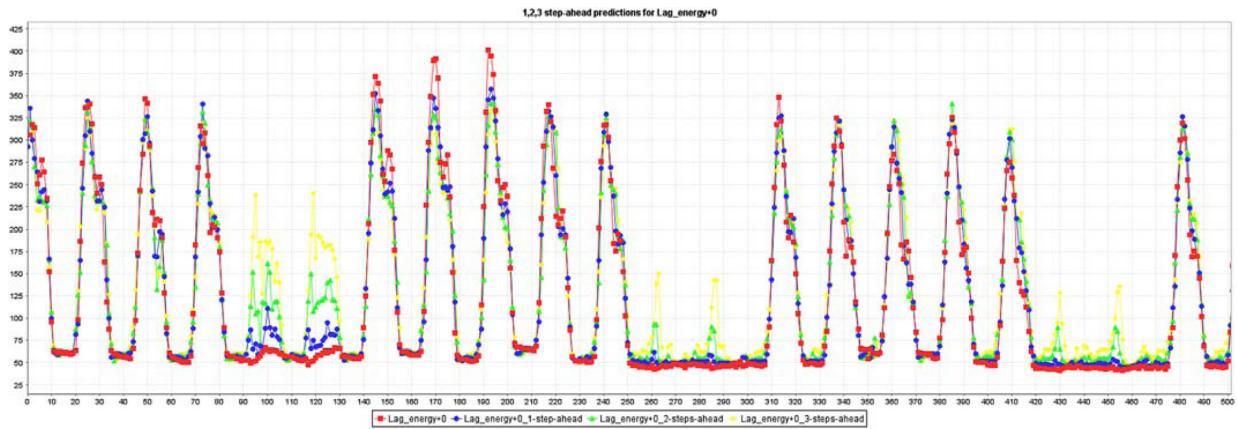
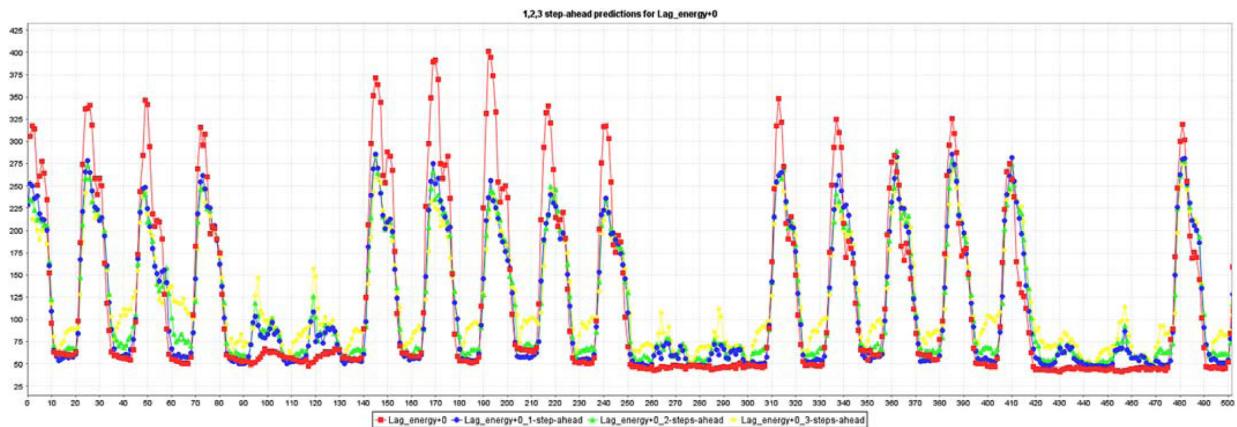
	1-step-ahead	2-steps-ahead	3-steps-ahead	Average
Number of instances	3559	3558	3557	–
MAE	4.4041	9.6858	18.2695	10.7865
RMSE	7.5987	14.0861	25.1676	15.6175

Table 12Evaluation on test data (1526 instances) with *RandomForest - database #1*.

	1-step-ahead	2-steps-ahead	3-steps-ahead	Average
Number of instances	1526	1525	1524	–
MAE	10.9941	20.4655	32.7499	21.4032
RMSE	16.0509	28.7680	44.8343	29.8844

Table 13Evaluation on test data (1526 instances) with *RandomForest - TransformedDatabase*.

	1-step-ahead	2-steps-ahead	3-steps-ahead	Average
Number of instances	1526	1525	1524	–
MAE	26.7583	34.7768	49.7004	37.0785
RMSE	36.5563	45.0787	59.8209	47.1520

Fig. 6. 1,2,3-step-ahead predictions for Lag-energy+0 evaluated on first 500 test data with *RandomForest - database #1*.Fig. 7. 1,2,3-step-ahead predictions for Lag-energy+0 evaluated on first 500 test data with *RandomForest - TransformedDatabase*.

interactions could be missed due to the way the search space is traversed [59]. Probabilistic search techniques, such as *MultiObjectiveEvolutionarySearch*, can overcome this difficulties by allowing to generate new subsets in different locations of the search space guided by a metaheuristic. In this paper, we propose to use a multivariate wrapper feature selection method where the search strategy is based on multi-objective evolutionary computation, thus intrinsically overcoming the problem of interactions between features.

- For wrapper FS methods, the *RandomForest* evaluator has proven more effective than *IBk* and *LinearRegression* based evaluators. *SMOreg* and *GaussianProcesses* are discarded as evaluators for wrapper methods because of their excessive computational time. Run time of *RandomForest* is acceptable for wrapper FS methods setting the number of iterations to 10 (-I 10), and this method is not very sensitive to the variation of its parameters. However, *RandomForest* generates regression models larger than *IBk*, *LinearRegression* and *SMOreg*.
- *IBk* is very prone to over-fitting. Although in the evaluation on full training data the best results have been obtained with *IBk*, these results become poor when the evaluation is done on cross-validation, which indicates that *IBk* over-fits the regression models.
- *LinearRegression*, *SMOreg* and *GaussianProcesses* are not prone to over-fitting, but it has not been efficient for this problem.
- *MAE* has shown better behaviour than *RMSE* as metric performance in evaluators for wrapper FS methods. This can be seen in Table 5: the FS method *MOES-RF-MAE* (database #1) produces better results than the method *MOES-RF-RMSE* (database #2) when evaluated on cross-validation with *RandomForest* using the *RMSE* metric (12.6685 vs. 12.9133, an improvement of 1.9%). This improvement can also be observed in Table 6 when both databases are evaluated with the *MAE* metric (5.8264 vs. 6.0012, an improvement of 2.91% in this case).

B. Regarding the forecasting process:

Tables 10 to 13 show how 1,2,3-steps-ahead predictions using the reduced database #1 improve the 1,2,3-steps-ahead predictions using the database without performing feature selection. Using the averages of the 1,2,3-steps-ahead predictions (shown also in Tables 10 to 13) we can calculate the percentage differences between the average predictions by doing feature selection and without doing so. With our methodology, *MAE* is improved by 59.97% and *RMSE* by 40.75%, evaluated on training data, and *MAE* is improved by 42.28% and *RMSE* by 36.62%, evaluated on test data.

5. Comparison with other methods proposed in literature

The metrics *RMSE* and *MAE* are two of the most common metrics used to measure accuracy for continuous variables and they are appropriate for model comparisons because they express average model prediction error in units of the variable of interest. However, in order to compare energy consumption prediction within several papers that do not use the same dataset or the same values of energy to be predicted it is not useful to compare such metrics whose magnitude depend on the range of the output data.

For that reason, we choose the coefficient of variance of the *RMSE*. *CVRMSE* is a non-dimensional measure calculated by dividing the *RMSE* of the predicted energy consumption by the mean value of the actual energy consumption. For example, a *CVRMSE* value of 5% would indicate that the mean variation in actual energy consumption not explained by the prediction model is 5% of the mean value of the actual energy consumption [60].

In the work with similar objectives [22], the preprocessing is carried out through correlation and *Principal Components Analysis* [48] and each day is divided in three moments alluding to occupation: morning, afternoon and night. That way, 3 different models are trained and the results are the following: *Random Forest* is selected at night and in the afternoon providing a *RMSE* of 1 and 3.87 *KWh* and *Bayesian Regularized Neural Networks* [61] is selected for the morning with *RMSE* = 7.08 *KWh*. In that sense, we could say that our FS approach overcomes this method in general. In the work [2], the temperature time series together with day of the week are used in order to estimate energy consumption. Results show again *Random Forest* as the outstanding model and the daily *CVRMSE* = 9%.

For current and future comparisons with further research, we obtained an hourly *CVRMSE* = 20% and we have also averaged it per day obtaining a daily *CVRMSE* = 11% for the 1-step case.

Although there are other methodologies for time series forecasting, such as *Wavelet Transform* and fusions (with *Artificial Neural Networks*, *Support Vector Machines*, etc. [62,63]), in this paper we have compared our proposal with *ARIMA*, which is one of the most used for time series prediction methodologies used in the literature.

Multivariate ARIMA: we have used the traditional time series method *ARIMA* with exogenous regressors [64]. Results are much worst than using our machine learning oriented approach. Using our selected features, mean *MAE* is 119 and mean *RMSE* is 126. This results are way worst than ours but still better than using all variables with *ARIMA*: *MAE* increases between 35 and 55 *KWh* and *RMSE* increases between 37 and 58 *Kwh*.

6. Conclusions

In this work we have proposed a methodology for energy multivariate time series forecasting. The methodology is based on, firstly, database transformation into a form that standard machine learning algorithms can process, and then, systematically apply a set of feature selection methods for regression. The methodology deals with both lagged and intervention variables, unlike other works in the literature where only lagged variables are treated or the time series problem is univariate. The results of the experiments carried out show that the proposed methodology effectively reduces both the complexity of the forecast model and their *RMSE* and *MAE* in 1,2,3-steps-ahead predictions. The results of our methodology improve those obtained with other works reported in the literature, as well as those obtained with the *marima* package for multivariate time series forecasting.

Conflict of interest

None.

Acknowledgments

This work has been sponsored by MINECO through the PER-SEIDES project (ref. TIN2017-86885-R) and grant BES-2015-071956 and by the European Comission through the H2020 IoTcrawler (contract 779852) EU Project. This work was also partially funded by the Spanish Ministry of Science, Innovation and Universities under the SITUS project (Ref: RTI2018-094832-B-I00), and by the European Fund for Regional Development (EFRD, FEDER).

Table 14
Parameters of the proposed feature selection methods for energy time series forecasting.

Database #Id.	Parameters
#1	-E "weka.attributeSelection.WrapperSubsetEval -B weka.classifiers.trees.RandomForest -F 5 -T 0.01 -R 1 -E DEFAULT - -P 100 -I 10 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S1" -S "weka.attributeSelection.MultiObjectiveEvolutionarySearch -generations 500 -population-size 100 -seed 1 -a 0 " -E "weka.attributeSelection.WrapperSubsetEval -B weka.classifiers.trees.RandomForest -F 5 -T 0.01 -R 1 -E MAE - -P 100 -I 10 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S1" -S "weka.attributeSelection.MultiObjectiveEvolutionarySearch -generations 500 -population-size 100 -seed 1 -a 0 " #2 #3 #4 #5 #6 #7 #8
#2	-E "weka.attributeSelection.WrapperSubsetEval -B weka.classifiers.trees.RandomForest -F 5 -T 0.01 -R 1 -E MAE - -P 100 -I 10 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S1" -S "weka.attributeSelection.MultiObjectiveEvolutionarySearch -generations 500 -population-size 100 -seed 1 -a 0 " -E "weka.attributeSelection.WrapperSubsetEval -B weka.classifiers.lazy.IBk -F 5 -T 0.01 -R 1 -E DEFAULT - -K 1 -W 0 -A "weka.core.neighboursearch.LinearNNSearch -A "weka.core.EuclideanDistance -R first-last"" -S "weka.attributeSelection.MultiObjectiveEvolutionarySearch -generations 500 -population-size 100 -seed 1 -a 0 " -E "weka.attributeSelection.WrapperSubsetEval -B weka.classifiers.functions.LinearRegression -F 5 -T 0.01 -R 1 -E MAE - -S 0 -R 1.0E-8 -num-decimal-places 4" -S "weka.attributeSelection.MultiObjectiveEvolutionarySearch -generations 500 -population-size 100 -seed 1 -a 0 " -E "weka.attributeSelection.ClassifierAttributeEval -execution-slots 1 -B weka.classifiers.trees.RandomForest -F 5 -T 0.01 -R 1 -E DEFAULT - -P 100 -I 10 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S1" -S "weka.attributeSelection.Ranker -T -1.8E308 -N10" -E "weka.attributeSelection.CfsSubsetEval -P 1 -E1" -S "weka.attributeSelection.GreedyStepwise -T -1.8E308 -N -1 -num-slots 1" -E "weka.attributeSelection.ReliefFAttributeEval -M -1 -D 1 -K10" -S "weka.attributeSelection.Ranker -T -1.8E308 -N10" -E "weka.attributeSelection.PrincipalComponents -R 0.95 -A5" -S "weka.attributeSelection.Ranker -T -1.8E308 -N10"
#3	
#4	
#5	
#6	
#7	
#8	

Table 15
Parameters of the regression methods.

Name	Parameters
RandomForest	-P 100 -I 100 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S 1
IBk	-K 1 -W 0 -A "weka.core.neighboursearch.LinearNNSearch -A "weka.core.EuclideanDistance -R first-last""
LinearRegression	-S 0 -R 1.0E-8 -num-decimal-places 4
SMOreg	-C 1.0 -N 0 -I "weka.classifiers.functions.supportVector.RegSMOImproved -T 0.001 -V -P 1.0E-12 -L 0.001 -W1" -K "weka.classifiers.functions.supportVector.PolyKernel -E 1.0 -C250007"
GaussianProcesses	-L 1.0 -N 0 -K "weka.classifiers.functions.supportVector.PolyKernel -E 1.0 -C250007" -S 1

Table 16
Packages and classes for feature selection in Weka used in this paper.

Name	Description
weka.classifiers.timeseries.core.TSLagMaker	Class for creating lagged versions of target variable(s) for use in time series forecasting
weka.attributeSelection	Package for feature selection
weka.attributeSelection.AttributeSelection	Class for feature selection
weka.attributeSelection.ASearch	Abstract class for search strategy
weka.attributeSelection.ASEvaluation	Abstract class for evaluation
weka.classifiers.AbstractClassifier	Abstract classifier
weka.classifiers.SingleClassifierEnhancer	Abstract utility class, extends AbstractClassifier
weka.classifiers.meta.AttributeSelectedClassifier	Meta-classifier for feature selection + classification/regression, extends SingleClassifierEnhancer
weka.attributeSelection.GreedyStepwise	Class for greedy stepwise search strategy, extends ASSearch
weka.attributeSelection.MultiObjectiveEvolutionarySearch	Class for multi-objective evolutionary search strategy, extends ASSearch
weka.attributeSelection.PSOSearch	Class for particle swarm optimization search strategy, extends ASSearch
weka.attributeSelection.Ranker	Class to rank attributes in univariate feature selection methods, extends ASSearch
weka.attributeSelection.WrapperSubsetEval	Class for multivariate wrapper feature selection methods, extends ASEvaluation
weka.attributeSelection.ConsistencySubsetEval	Class for multivariate filter feature selection methods, extends ASEvaluation
weka.attributeSelection.ClassifierAttributeEval	Class for univariate wrapper feature selection methods, extends ASEvaluation
weka.attributeSelection.ReliefFAttributeEval	Class for univariate filter feature selection methods, extends ASEvaluation
weka.attributeSelection.PrincipalComponents	Class for univariate filter feature selection methods, extends ASEvaluation
weka.classifiers.trees.RandomForest	Class for constructing a forest of random trees, extends weka.classifiers.meta.Bagging
weka.classifiers.lazy.IBk	Class that implements an instance-based learning algorithm, extends weka.classifiers.Classifier
weka.classifiers.functions.LinearRegression	Class for using linear regression for prediction, extends weka.classifiers.AbstractClassifier
weka.classifiers.timeseries.WekaForecaster	Class that implements time series forecasting using a Weka regression scheme

References

- [1] M. Akcin, A. Kaygusuz, A. Karabiber, S. Alagoz, B.B. Alagoz, C. Keles, Opportunities for energy efficiency in smart cities, in: Smart Grid Congress and Fair (ICSG), 2016 4th International Istanbul, IEEE, 2016, pp. 1–5.
- [2] A. González-Vidal, A.P. Ramallo-González, F. Terroso-Sáenz, A. Skarmeta, Data driven modeling for energy consumption prediction in smart buildings, in: Big Data (Big Data), 2017 IEEE International Conference on, IEEE, 2017, pp. 4562–4569.
- [3] K. Zhou, C. Fu, S. Yang, Big data driven smart energy management: from big data to big insights, *Renewable Sustainable Energy Rev.* 56 (2016) 215–225, doi:[10.1016/j.rser.2015.11.050](https://doi.org/10.1016/j.rser.2015.11.050).
- [4] Lavastorm, Big data, analytics, and energy consumption, 2014.
- [5] D. Dua, C. Graff, Uci machine learning repository, 2017.
- [6] A. Gonzalez-Vidal, P. Barnaghi, A.F. Skarmeta, Beats: blocks of eigenvalues algorithm for time series segmentation, *IEEE Trans. Knowl. Data Eng.* (2018).
- [7] H. Liu, H. Motoda, *Feature Selection for Knowledge Discovery and Data Mining*, Kluwer Academic Publishers, Norwell, MA, USA, 1998.
- [8] Y. Li, T. Li, H. Liu, Recent advances in feature selection and its applications, *Knowl. Inf. Syst.* 53 (3) (2017) 551–577, doi:[10.1007/s10115-017-1059-8](https://doi.org/10.1007/s10115-017-1059-8).
- [9] L. Dannecker, *Energy Time Series Forecasting: Efficient and Accurate Forecasting of Evolving Time Series from the Energy Domain*, 1st ed., Springer Vieweg, 2015.
- [10] C. Deb, F. Zhang, J. Yang, S.E. Lee, K.W. Shah, A review on time series forecasting techniques for building energy consumption, *Renewable Sustainable Energy Rev.* 74 (2017) 902–924, doi:[10.1016/j.rser.2017.02.085](https://doi.org/10.1016/j.rser.2017.02.085).
- [11] I. Guyon, A. Elisseeff, An introduction to variable and feature selection, *J. Mach. Learn. Res.* 3 (2003) 1157–1182.
- [12] R. Adhikari, R.K. Agrawal, An introductory study on time series modeling and forecasting, *CoRR abs/1302.6613* (2013).
- [13] I.H. Witten, E. Frank, M.A. Hall, *Data Mining: Practical Machine Learning Tools and Techniques* (Third Edition), The Morgan Kaufmann Series in Data Management Systems, Morgan Kaufmann, Boston, 2011.
- [14] R. Kohavi, G.H. John, Wrappers for feature subset selection, *Artif. intell.* 97 (1–2) (1997) 273–324.
- [15] N. Japkowicz, M. Shah, *Evaluating Learning Algorithms: A Classification Perspective*, Cambridge University Press, New York, NY, USA, 2011.
- [16] A.G. Karegowda, A.S. Manjunath, M.A. Jayaram, Comparative study of attribute selection using gain ratio and correlation based feature selection, *Int. J. Inf. Technol. Knowl. Manage.* 2 (2) (2010) 271–277.
- [17] M.A. Hall, Correlation-based Feature Selection for Machine Learning, Technical Report, University of Waikato, 1999.
- [18] A. Ahmad, L. Dey, A feature selection technique for classificatory analysis, *Pattern Recognition Letters* 26 (1) (2005) 43–56.
- [19] S. Salzberg, C4.5: Programs for machine learning by J. Ross Quinlan, Morgan Kaufmann Publishers, Inc., 1993, *Mach. Learn.* 16 (3) (1994) 235–240, doi:[10.1007/BF00993309](https://doi.org/10.1007/BF00993309).
- [20] H.-X. Zhao, F. Magoulès, Feature selection for predicting building energy consumption based on statistical learning method, *J. Algorithms Comput. Technol.* 6 (1) (2012) 59–77.
- [21] O. Utterbäck, Feature selection methods with applications in electrical load forecasting, Master's Theses in Mathematical Sciences (2017).
- [22] A. González-Vidal, V. Moreno-Cano, F. Terroso-Sáenz, A.F. Skarmeta, Towards energy efficiency smart buildings models based on intelligent data analytics, *Procedia Comput. Sci.* 83 (2016) 994–999.
- [23] M.V. Moreno, F. Terroso-Sáenz, A. González-Vidal, M. Valdés-Vela, A.F. Skarmeta, M.A. Zamora, V. Chang, Applicability of big data techniques to smart cities deployments, *IEEE Trans. Ind. Inf.* 13 (2) (2017) 800–809.
- [24] S.F. Crone, N. Kourentzes, Feature selection for time series prediction—a combined filter and wrapper approach for neural networks, *Neurocomputing* 73 (10–12) (2010) 1923–1936.
- [25] Y. Sun, J. Li, J. Liu, C. Chow, B. Sun, R. Wang, Using causal discovery for feature selection in multivariate numerical time series, *Mach. Learn.* 101 (1–3) (2015) 377–395.
- [26] S. Hido, T. Morimura, Temporal feature selection for time-series prediction, in: 2012 21st International Conference on Pattern Recognition (ICPR 2012), IEEE, 2012, pp. 3557–3560.
- [27] O. García-Hinde, V. Gómez-Verdejo, M. Martínez-Ramón, C. Casanova-Mateo, J. Sanz-Justo, S. Jiménez-Fernández, S. Salcedo-Sanz, Feature selection in solar radiation prediction using bootstrapped svrs, in: Evolutionary Computation (CEC), 2016 IEEE Congress on, IEEE, 2016, pp. 3638–3645.
- [28] D. O'Leary, J. Kubby, Feature selection and ann solar power prediction, *J. Renewable Energy* 2017 (2017).
- [29] C. Feng, M. Cui, B.-M. Hodge, J. Zhang, A data-driven multi-model methodology with deep feature selection for short-term wind forecasting, *Appl. Energy* 190 (2017) 1245–1257.
- [30] R.G. Pajares, J.M. Benítez, G.S. Palmero, Feature selection for time series forecasting: a case study, in: Hybrid Intelligent Systems, 2008. HIS'08. Eighth International Conference on, IEEE, 2008, pp. 555–560.
- [31] E. Ferreira, Model selection in time series machine learning applications, PhD Thesis, University of Oulu Graduate School, University of Oul (2015).
- [32] L. Tang, C. Wang, S. Wang, Energy time series data analysis based on a novel integrated data characteristic testing approach, *Procedia Comput. Sci.* 17 (2013) 759–769.
- [33] I. Koprinska, M. Rana, V.G. Agelidis, Correlation and instance based feature selection for electricity load forecasting, *Knowl.-Based Syst.* 82 (2015) 29–40.
- [34] F. Martínez-Alvarez, A. Troncoso, G. Asencio-Cortés, J.C. Riquelme, A survey on data mining techniques applied to electricity-related time series forecasting, *Energies* 8 (11) (2015) 13162–13193.
- [35] C. Goutte, Note on free lunches and cross-validation, *Neural Comput.* 9 (6) (1997) 1245–1249.
- [36] F. Jiménez, G. Sánchez, J. García, G. Sciacivico, L. Miralles, Multi-objective evolutionary feature selection for online sales forecasting, *Neurocomputing* (2016).
- [37] S. Russell, P. Norvig, *Artificial Intelligence: A Modern Approach*, 2nd ed., Prentice-Hall, 2003.
- [38] F. Jiménez, A. Gómez-Skarmeta, G. Sánchez, K. Deb, An evolutionary algorithm for constrained multi-objective optimization, in: Proceedings of the Evolutionary Computation on 2002. CEC '02. Proceedings of the 2002 Congress, in: CEC '02, vol. 2, IEEE Computer Society, Washington, DC, USA, 2002, pp. 1133–1138.
- [39] F. Jiménez, G. Sánchez, P. Vasant, A multi-objective evolutionary approach for fuzzy optimization in production planning, *J. Intell. Fuzzy Syst.* 25 (2) (2013) 441–455.
- [40] F. Jiménez, G. Sánchez, J.M. Juárez, Multi-objective evolutionary algorithms for fuzzy classification in survival prediction, *Artif. intell. Med.* 60 (3) (2014) 197–219.
- [41] F. Jiménez, E. Marzano, G. Sánchez, G. Sciacivico, N. Vitacolonna, Attribute selection via multi-objective evolutionary computation applied to multi-skill contact center data classification, in: Proc. of the IEEE Symposium on Computational Intelligence in Big Data (IEEE CIBD 15), IEEE, 2015, pp. 488–495.
- [42] F. Jiménez, G. Sánchez, J. García, G. Sciacivico, L. Miralles, Multi-objective evolutionary feature selection for online sales forecasting, *Neurocomputing* 234 (2017) 75–92.
- [43] E. Zitzler, K. Deb, L. Thiele, Comparison of multiobjective evolutionary algorithms: empirical results, *Evol. Comput.* 8 (2) (2000) 173–195.
- [44] E. Zitzler, L. Thiele, M. Laumanns, C. Fonseca, V. Grunert da Fonseca, Performance assessment of multiobjective optimizers: an analysis and review, *IEEE Trans. Evol. Comput.* 7 (2002) 117–132.
- [45] J. Novakovic, Toward optimal feature selection using ranking methods and classification algorithms, *Yugoslav J. Oper. Res.* 21 (1) (2016).
- [46] H. Liu, R. Setiono, A probabilistic approach to feature selection - a filter solution, in: Proceedings of the 13th International Conference on Machine Learning (ICML), vol. 96, 1996, pp. 319–327.
- [47] K. Kira, L.A. Rendell, A practical approach to feature selection, in: Proceedings of the Ninth International Workshop on Machine Learning, in: ML92, Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1992, pp. 249–256.
- [48] H. Abdi, L.J. Williams, Principal component analysis, *Wiley Interdiscip. Rev. 2* (4) (2010) 433–459, doi:[10.1002/wics.101](https://doi.org/10.1002/wics.101).
- [49] R. Schafer, Accurate and efficient general-purpose boilerplate detection for crawled web corpora, *Lang. Resour. Eval.* (2016), doi:[10.1007/s10579-016-9359-2](https://doi.org/10.1007/s10579-016-9359-2). Online first
- [50] L. Breiman, Random forests, *Mach. Learn.* 45 (1) (2001) 5–32.
- [51] D.W. Aha, D. Kibler, M.K. Albert, Instance-based learning algorithms, *Mach. Learn.* 6 (1) (1991) 37–66, doi:[10.1023/A:1022689900470](https://doi.org/10.1023/A:1022689900470).
- [52] X. Yan, *Linear Regression Analysis: Theory and Computing*, World Scientific Publishing Company Pte Limited, 2009.
- [53] C.J. Willmott, K. Matsuurra, Advantages of the mean absolute error (MAE) over the root mean square error (RMSE) in assessing average model performance, *Clim. Res.* 30 (1) (2005) 79–82. Cited By (since 1996) 149
- [54] C. Cortes, V. Vapnik, Support-vector networks, *Mach. Learn.* 20 (3) (1995) 273–297.
- [55] C.E. Rasmussen, C.K.I. Williams, *Gaussian Processes for Machine Learning (Adaptive Computation and Machine Learning)*, The MIT Press, 2005.
- [56] G. Rubio, H. Pomares, L.J. Herrera, I. Rojas, Kernel methods applied to time series forecasting, in: F. Sandoval, A. Prieto, J. Cabestany, M. Graña (Eds.), *Computational and Ambient Intelligence*, Springer Berlin Heidelberg, Berlin, Heidelberg, 2007, pp. 782–789.
- [57] S.K. Shevade, S.S. Keerthi, C. Bhattacharyya, K.R.K. Murthy, Improvements to the SMO algorithm for SVM regression, *IEEE Trans. Neural Netw.* 11 (5) (2000) 1188–1193, doi:[10.1109/72.870050](https://doi.org/10.1109/72.870050).
- [58] L. Chuang, C. Ke, C. Yang, A hybrid both filter and wrapper feature selection method for microarray classification, *CoRR abs/1612.08669* (2016).
- [59] L.C. Molina, L. Belanche, Á. Nebot, Feature selection algorithms: a survey and experimental evaluation, in: Proceedings of the 2002 IEEE International Conference on Data Mining (ICDM 2002), 9–12 December 2002, Maebashi City, Japan, 2002, pp. 306–313.
- [60] T.A. Reddy, N.F. Saman, D.E. Claridge, J.S. Haberl, W.D. Turner, A.T. Chalifoux, Baseline methodology for facility-level monthly energy use-part 1: Theoretical aspects, in: *ASHRAE transactions, ASHRAE*, 1997, pp. 336–347.
- [61] B. F., W. D., Bayesian Regularization of Neural Networks, vol. vol. 458, Livingstone D.J. (eds) Artificial Neural Networks. Methods in Molecular Biology, Humana Press, pp. 23–42.
- [62] M. K., M.C. Lineesh, C.J. John, Wavelet neural networks for nonlinear time series analysis, *Appl. Math. Sci.* 4 (2010) 2485–2495.
- [63] O. Kisi, M. Cimen, Precipitation forecasting by using wavelet-support vector machine conjunction model, *Eng. Appl. Artif. Intell.* 25 (4) (2012) 783–792.
- [64] H. Spliid, Multivariate ARIMA and ARIMA-X Analysis, CRAN, 2017. License GPL-2, Version 2.2, RoxygenNote 5.0.1.

4.3. COMMISSIONING OF THE CONTROLLED AND AUTOMATIZED TESTING FACILITY
FOR HUMAN BEHAVIOR AND CONTROL (CASITA)

4.3 Commissioning of the Controlled and Automatized Testing Facility for Human Behavior and Control (CASITA)

Title	Commissioning of the Controlled and Automatized Testing Facility for Human Behavior and Control (CASITA)
Authors	Ignacio Rodríguez-Rodríguez, Aurora González-Vidal, Alfonso P. Ramallo-González and Miguel Ángel Zamora
Type	Journal
Journal	Sensors
Impact factor (2018)	3.031
Rank	Q1
Publisher	MDPI
Volume	18
Issue	9
Year	2018
Month	September
ISSN	1424-8220
DOI	doi:10.3390/s18092829
URL	https://www.ncbi.nlm.nih.gov/pubmed/30150573
State	Published
Author's contribution	The PhD student, Aurora González Vidal, contributed to the methodology, software, analysis and research and writing the publication

Article

Commissioning of the Controlled and Automatized Testing Facility for Human Behavior and Control (CASITA)

Ignacio Rodríguez-Rodríguez * , Aurora González Vidal , Alfonso P. Ramallo González  and Miguel Ángel Zamora 

Departamento de Ingeniería de la Información y las Comunicaciones, Facultad de Informática, Universidad de Murcia, 30100 Murcia, Spain; aurora.gonzalez2@um.es (A.G.V.); alfonop.ramallo@um.es (A.P.R.G.); mzamora@um.es (M.A.Z.)

* Correspondence: ignacio.rodriguez1@um.es

Received: 19 July 2018; Accepted: 24 August 2018; Published: 27 August 2018



Abstract: Human behavior is one of the most challenging aspects in the understanding of building physics. The need to evaluate it requires controlled environments and facilities in which researchers can test their methods. In this paper, we present the commissioning of the Controlled and Automatized Testing Facility for Human Behavior (CASITA). This is a controlled space emulation of an office or flat, with more than 20 environmental sensors, 5 electrical meters, and 10 actuators. Our contribution shown in this paper is the development of an infrastructure-Artificial Intelligence (AI) model pair that is perfectly integrated for the study of a variety of human energy use aspects. This facility will help to perform studies about human behavior in a controlled space. To verify this, we have tested this emulation for 60 days, in which equipment was turned on and off, the settings of the conditioning system were modified remotely, and lighting operation was similar to that in real behaviors. This period of commissioning generated 74.4 GB of raw data including high-frequency measurements. This work has shown that CASITA performs beyond expectations and that sensors and actuators could enable research on a variety of disciplines related to building physics and human behavior. Also, we have tested the PROPHET software, which was previously used in other disciplines and found that it could be an excellent complement to CASITA for experiments that require the prediction of several pertinent variables in a given study. Our contribution has also been to proof that this package is an ideal “soft” addition to the infrastructure. A case study forecasting energy consumption has been performed, concluding that the facility and the software PROPHET have a great potential for research and an outstanding accuracy.

Keywords: modelling; energy; buildings

1. Introduction

Energy is one of the most important topics of study worldwide. Most governments have implemented initiatives that aim at more energy-efficient societies because of an urgent need to decelerate (1) energy resources exhaustion and (2) greenhouse gas emissions. Buildings are responsible for up to 40% of the carbon emissions in developed countries [1]; it has been seen that their energy use can be reduced substantially not only via renovation of their thermal envelope [2,3] but also via the modification of users’ behavior [4].

This opportunity to reduce energy use via changes in behavior has come at the same time as a technological revolution. There now exist sophisticated information management systems to control the different working points of building infrastructures. These systems have already been proved to

be effective solutions to the problem of high energy consumption associated with comfort spaces in buildings, see for example [4–6].

Energy consumption is a complex phenomenon in which many aspects play a role; only a comprehensive way of studying it can fully cover its social, economic, and behavioral aspects. The effectiveness of technological solutions for modifying human behavior seems to vary depending on the study [7,8]. It is for this reason that more experimental research on human behavior is needed. This is a topic that was being studied as early as 1960 when Newton et al. [9] outlined the difficulties of understanding human behavior in buildings. Today, there have been advances in the understanding of human behavior, such as the work of [10,11]. However, more testing is needed to continue improving this field of research. A research facility that can serve as a testing arena for this kind of experimentation with the control of all aspects of functioning systems in a building is highly valuable in this field.

The analysis of energy efficiency in built environments has received growing attention in the last decade [12–14]. One possible method to lower energy use could be to generate a management system to tackle this challenge. A home automation system based on the internet of things (IoT) can monitor and control intelligently the different infrastructures involved in a building's energy consumption, while being able to provide comfort, security and communication, energy efficiency, and promote water, electricity, and fuel conservation. Hence, the research community is not only interested in the understanding and modeling of human behavior, but also on the developing and testing of control strategies for building automation based on IoT.

With respect to the advances in sensing and control infrastructure, the growth on information and communication technologies (ICT) offer an even greater potential in the near future [15], and has opened a door for considering homes as environments with many more devices (such as sensors controllers or actuators). A facility that serves to understand the interactions between humans and buildings will need to have all those components to perform valid research.

The internet of things represents a radical evolution of the current internet to a network of interconnected devices that not only harvest information from the environment (sensing), but also allow interacting, managing, and storing easily any kind of data [16–18]. Following an IoT approach, new home automation systems could allow fulfilling the requirements posed by the social changes and new trends in our way of life, facilitating the design of more human, personal, multifunctional, and flexible homes. This change seems to be coming soon as the European Commission has established that 16 of the European Union (EU) member states will implement a large-scale smart-meter rollout by 2020 [19].

The efficiency and accuracy of any home automation system is possible as far as good predictions can be achieved by developing models about the building status. Ergo, different variables have to be taken into account regarding their impact on the energy consumption of buildings, while attempting to consider them in an integral vision [20]. Making a suitable selection and analysis of them is not obvious. Not only do environmental parameters such as humidity and temperature have to be studied, others like human behavior, weather forecast, insulating materials, or thermal inertia should be also considered in order to obtain patterns that will make it possible to anticipate changes in order to avoid declines in comfortable conditions and rises in energy consumption.

With this purpose, the available data about a building and its context have to be interpreted to obtain valuable knowledge. Statistical and novel methods of data analysis allow researchers to establish correlations between variables and to generate performance models of a building, which can be used to ensure efficient responses by the automation system. Thus, in the context of data science, many new and more powerful technologies are bringing alternatives, or even breakthroughs, in the prediction of building energy consumption associated with thermal comfort [21].

The facilities we present for the Controlled and Automatized Testing Facility for Human Behavior (CASITA) have an IoT-based home automation system installed and operational, where experiments can be done in order to test human behavior and IoT solutions. It is located at the Technology Transfer Center of the University of Murcia, Murcia, Spain. This test lab has numerous sensors, actuators, and controllers providing data, which are able to be used to generate accurate models in order to

predict energy consumption and many other variables related to building physics. In addition to this, we have coupled the software PROPHET as the soft component of the functioning of the infrastructure for variable forecasting and completion. In this work, two models of energy consumption forecasting will be presented and discussed.

For the commissioning of this infrastructure, a model of the energy consumption based on the novel PROPHET package has been developed within mathematical software R. It measures several variables and evaluates variables that are beneficial for weather forecasting, thereby filling the future time series of outdoor conditions to validate the infrastructure.

All the steps proposed in this paper describe how preliminary testing on the research facility was performed, which can be used to design efficient management systems for saving energy that are fully scalable and that can be applied with the same goal in other buildings with similar sensing and actuation levels. With this paper we contribute to the development of a facility that is pioneer according to the knowledge of the authors, as it sums up the IoT and hardware infrastructure to a soft facet consistent on algorithms of prediction included on the PROPHET library.

This paper is structured as follows: Section 2 describes the infrastructure. Section 3 describes the commissioning and a pilot study to verify the validity of the data and the analysis methods available in CASITA. Section 4 shows conclusions and further work, followed by the references.

2. The Controlled and Automatized Testing Facility for Human Behaviour (CASITA)

Currently, a smart building can be equipped with information and communication technology (ICT) systems, as can be seen in [22], where a sensor network is deployed in a house. Another example is shown in [23].

Although CASITA has been used before in other studies [6,24–29], the commissioning and description of the research facility had not been published yet. This paper aims to provide the necessary documentation to close this gap. CASITA (see Figure 1) is a case of a smart space with a wide deployment of sensors and devices integrated as if it was a home/office automation system.

In this highly sensed habitat, data referring to human behavior and to outdoor and indoor environmental parameters are collected.

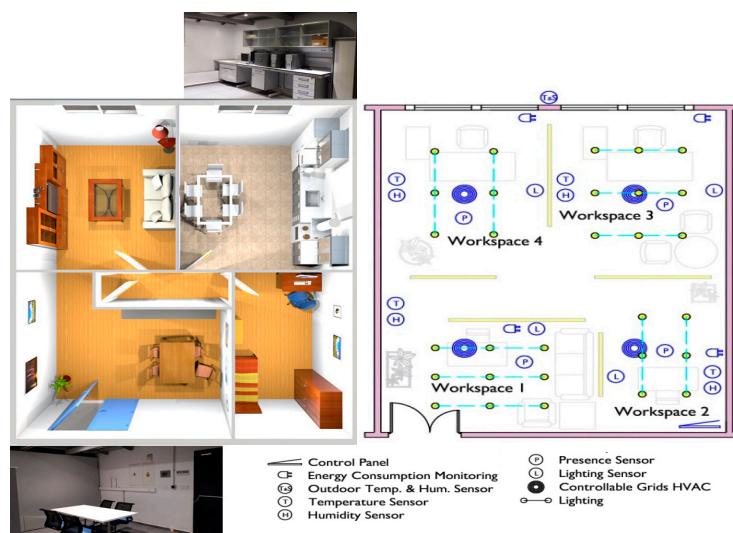


Figure 1. Infographic of a possible setup of the Controlled and Automatized Testing Facility for Human Behavior (CASITA) and distribution of all devices (sensors, actuators and controllers).

2.1. Hardware

The home automation system installed in this reference scenario is composed by Programmable Logic Controllers (PLC), and a Supervisory Control and Data Acquisition (SCADA) system. This system has been given the name Domosec Platform [30]. All the sensors and actuators have been selected in accordance with the principles suggested on [31].

The PLC is able to monitor the sensor status and regulate the infrastructures connected to a platform, while the SCADA system collects data and intercommunicates with the PLC using the actuators. This platform has been designed and developed in-house and more information can be provided on request as it is open-source.

The indoor temperature, humidity, and luminosity are measured in several points of the space. This makes possible to have an idea of how homogeneous the conditions are across the monitored areas. Outdoor conditions are also registered by a weather station located on the top of the building.

Human behavior and presence sensors using passive infrared technology are present. The control access system is based on Radio Frequency Identification (RFID) technology (more details about the device deployment can be found in [6]. Systems and location are exposed in Figure 1.

Due to the importance of outside weather in the studies that are being carried out in CASITA (for example, to measure adaptive thermal comfort), its framework also counts on an ad hoc weather forecast algorithm based on Agencia Estatal de Meteorología (AEMET), the Spanish Meteorology Agency [32], but post-processed further to improve accuracy. This will be explained further in the following sections.

Regarding the actuators deployed in CASITA, there are two Heating, Ventilating and Air Conditioning (HVAC) systems installed in the ceiling that consist on an electric air-to-air heat pump (TOSHIBA RAV-SM803AT-E and 2xTOSHIBA SM806BT-E, Toshiba Carrier Corporation, Tokyo, Japan). Therefore, the indoor temperature and humidity can be modified in CASITA at the user's will. The system has two levels of air velocity (fan power) and a thermostatic proportional control. The primary energy of the system is electricity.

Lighting is provided via light-emitting diodes (LED) placed in the ceiling in accordance with current Spanish regulations. However, they are easy to move as the ceiling space is formed by removable panels. All lighting can also be controlled via the SCADA using the internet. A schema of the hardware and communication architecture in CASITA is shown in Figure 2, and the connection of all this equipment can be seen in Figure 3.

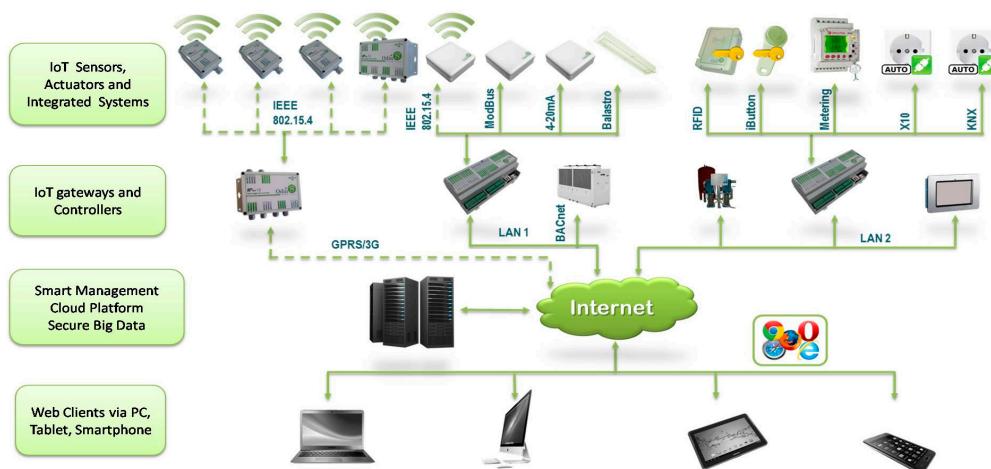


Figure 2. Hardware and communication architecture in CASITA.

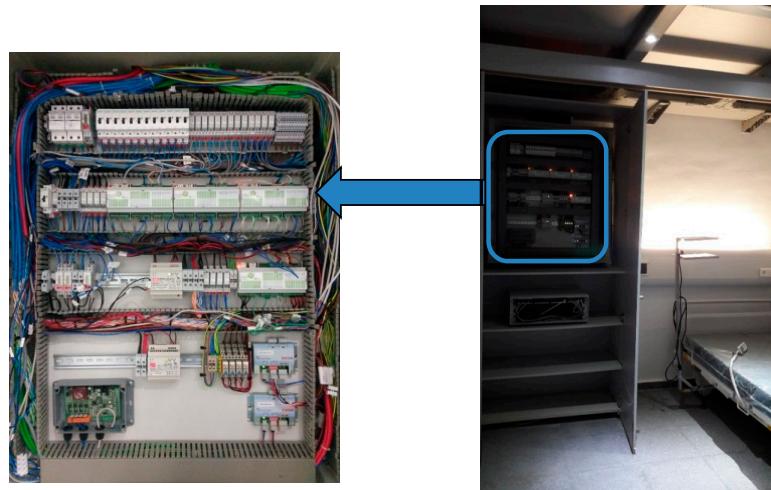


Figure 3. Wiring of data loggers, Supervisory Control and Data Acquisition (SCADA), and other devices in the main wiring cabinet of CASITA.

The electrical consumption of lights, HVAC units, and other electrical appliances are continuously being monitored and collected in the SCADA of the platform. Sensors can report at any given sampling period, which is at least 1 min long, but some of them are able to report at higher frequencies, e.g., high-frequency reporting of electrical grid to verify harmonics. Next table (Table 1) summarizes measured features and actuators that can be found in CASITA.

Table 1. Description of the sensors and actuators available in CASITA.

Features	Sensor Deployments Allow Measurement of a Wide Set of Data
Weather data	Temperature and humidity.
Weather forecast	Up to 4 days.
Indoor conditions	In four different locations, temperature and humidity.
Occupancy and activity	A control access system in the test lab entrance and volumetric detectors in each room let predict in an accurate way the tracking of human presence.
Energy consumption:	For this purpose, and to monitor each component separately, non-intrusive load monitoring techniques have been considered [33]. We distinguish:
Electrical devices	Computers and other appliance are monitored.
Lighting	Differentiating each room.
Heating, Ventilation, and Air Conditioning (HVAC)	Each air-conditioned machine is quantified but is much bigger than the previous consumptions, which makes it energetically undesirable.
Actuators	It is Possible to Modify the Test Lab Features, Comfort and Energy Consumption, Adapting the Next Actuators
Access	Test lab can be completely locked, rendering it impossible to enter.
Control of the energy supplies	The plugs can be disabled completely.
Control of the HVAC machines	It is possible to force a shutdown or a start. The temperature set point and fan velocity mode can be chosen.
Ventilation grilles	Each air supply duct ends in a motorized ventilation grille (one per room), which can be opened or closed depending on the nature of its use in the area.

With this infrastructure, different choices can be combined in order to reach a goal of sufficient comfort, reduced energy consumption, combination thereof, or other objectives.

2.2. Software: The PROPHET Package

The PROPHET package is an utility to model time series and that serves as the perfect soft counterpart of the infrastructure shown in this paper. PROPHET is an R library that has been recently developed and seems to give promising results in other disciplines [34–39]. It is a modular regression model with interpretable parameters that can be intuitively adjusted with domain knowledge about the time series [40].

PROPHET conducts an automatic procedure for forecasting time-series data. The implemented algorithm uses Stan modelling language (allows to share the same core procedure between Python and R implementations) for optimization in order to fit a non-linear additive model and generate uncertainty intervals.

The additive regression model has four main components: a piecewise linear or logistic growth curve trend. Prophet detects changes in trends by selecting changepoints from the data, a yearly seasonal component modeled using Fourier series, a weekly seasonal component using dummy variables, and a user-provided list of important holidays. PROPHET is robust enough to address missing data, shifts in the trend, and typically handles outliers well.

It allows the prediction of a horizon of observations for a given time series that fulfills some characteristics that are common to the time series generated by human actions, where factors such as holidays could be known in advance.

In order to create the model, a decomposable time series model with three main model components will be used: trend, seasonality, and holidays. This is shown in Equation (1),

$$y(t) = g(t) + s(t) + h(t) + \varepsilon_t, \quad (1)$$

where:

- $y(t)$: time series of interest.
- $g(t)$: represents non-periodic components (using piecewise linear or logistic growth curve trend). PROPHET implements two trend models that cover many applications: a saturating growth model and a piecewise linear model with automatic change point selection.
- $s(t)$: trend factor that represents periodic changes. Time series often have multi-period seasonality as a result of the human behaviors they represent. To fit and forecast these effects, we must specify seasonality models that are periodic functions of t . This part relies on Fourier series to provide a flexible model of periodic effects.
- $h(t)$: effects of holidays (a list provided by the user). Holidays and events provide large, somewhat predictable shocks to many time series and often do not follow a periodic pattern, so their effects are not well modeled by a smooth cycle.
- ε_t : error which will be assumed to follow a normal distribution.

This formulation is similar to a generalized additive model (GAM), a class of regression models with non-linear smoothers applied to the regressors. This approach has the advantage in that it decomposes easily and accommodates new components as necessary; for instance, when a new source of seasonality is identified. Thus, PROPHET frames the forecasting problem as a curve-fitting exercise which differs from the traditional models used for time series that account for the temporal dependence structure in the data: ARIMA. This formulation provides several functional advantages with respect to ARIMA formulations: flexibility regarding seasonality with multiple periods, measurements do not need to be regularly spaced and missing values are handled, fitting is very fast, and the parameters of the forecasting model are easily interpretable [41].

3. Commissioning and Example of Data Analysis

For the commissioning of CASITA, we developed a test that involves the use of all of the main systems (sensors, meters, and actuators) that are found in CASITA. With this test, we verified the

validity of the installations. We also made a valid test of a software package that has not been previously used for this purpose to forecast energy consumption. To do this, two models were built with the collected data. Their subsequent improvement became a topic of discussion as the inclusion of the weather forecast as a variable or not had to be determined.

The weather forecast was obtained from an official source (AEMET). The experiment consisted of generating simulated data of office use during 60 days from using the actuators, turning on and off equipment, and interfering with the conditioning system. This was done in an emulated manner to test the actuators and remote controllers of CASITA (all this was designed, run, and measured from an office 40 km away) and because it allowed us to access the ground truth. To ensure that other researchers interested in using CASITA would know appropriately what this facility has to offer, all of the data for this commissioning is available upon request.

The experiment was conducted from 10 June to 14 August 2017. In this period, up to 4 workers were working in a normal schedule from 9:00 to 17:00. It is presumed that they developed their usual functions in an office environment, working at their desks, but also sometimes working in pairs or holding meetings all together. We do not consider metabolic activity of the workers or their humidity emission. Some workers had the possibility to work from home, so the number of people at the office fluctuated between 1 and 4 people; at other times, the place was empty (without air conditioning). The occupancy was registered from presence detectors and door-opening sensors, as well as energy consumption and distribution of the operating grilles and HVAC machines that were activated by employees on-demand. All data were collected hourly, even outside of the working schedule (24 h). The operating temperature of the HVAC machines was fixed to 20 °C. Representative variables of this experiment can be seen in Figure 4.

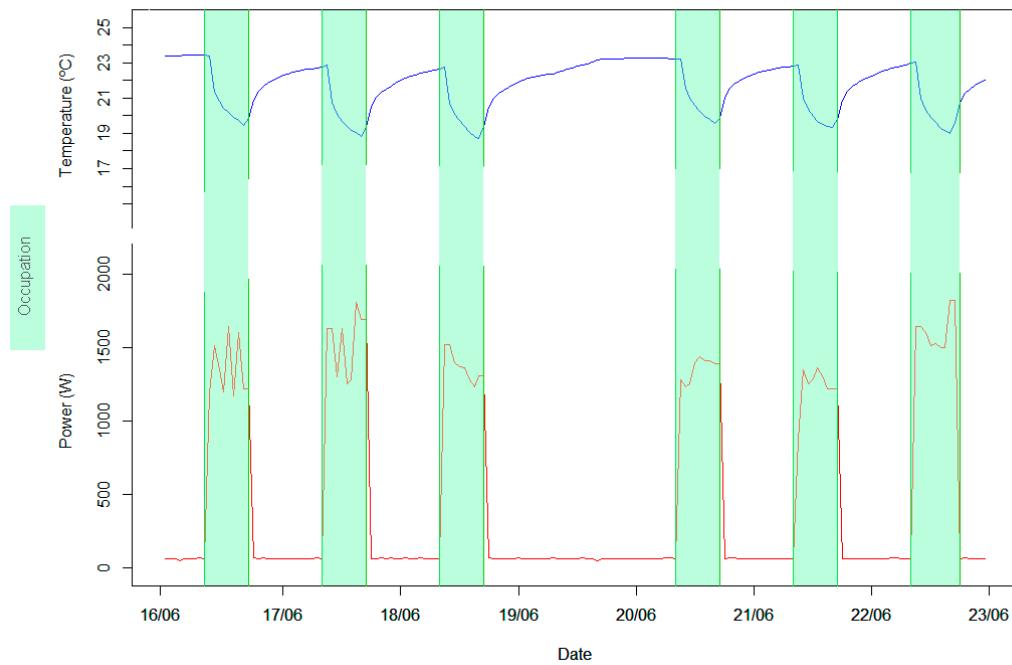


Figure 4. Representation of temperature and power for seven days of data in CASITA.

It is possible to appreciate that during the working time, energy consumption is triggered by the operation of office equipment air conditioning, which lowers the environmental temperature. In the graph, days are differentiated by higher or smaller occupancy and a local bank holiday, where no one was working at CASITA. Once the occupation is zero (in a working day at 17:00), it is easy to

identify the fast rise of the indoor temperature due to the high temperatures outside. When night falls, the indoor temperature changes more slowly due to the lowering of the external temperature.

The aim of this verification is two-fold. First, we will evaluate the commissioning of CASITA, and second, as we have access to the ground truth of the test, we will verify the performance of PROPHET in the field of energy use prediction. We believe that if the results were positive, PROPHET could be used synergistically with CASITA for further research. We have aimed for a data-driven approach that does not take into account the physical properties of the building itself since it has been shown to be appropriate in similar scenarios [42]. Our models are used for predicting a horizon of energy consumption. This makes it different from other approaches [43], whose goal is the punctual prediction of a particular moment.

3.1. Verification of Accessible National Weather Forecasting in CASITA Using PROPHET

To make sure that CASITA offers well-tested weather data and weather forecasts, a stand-alone parallel study was performed.

For the start, it was necessary to study the relationship between the weather forecast obtained from the AEMET web page and the real outdoor conditions of our test lab. If the correlation between them was strong, it would be feasible to anticipate and predict the real outdoor conditions.

In the case of an observed discrepancy between prediction and real data, steps were taken that allowed us to understand that error and to create a correction algorithm that reduces it substantially, adding value to the CASITA research facility. This is seen in Figure 5.

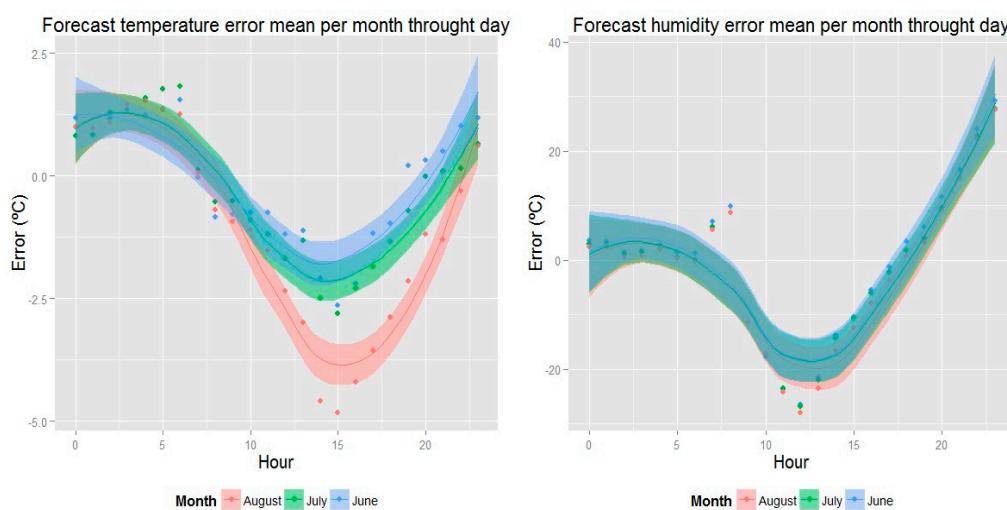


Figure 5. (a) Error mean between weather forecast (temperature) and outdoor temperature. (b) Error mean between weather forecast (humidity) and outdoor humidity.

When one signal was subtracted from another, an error signal was obtained (see Figure 5), in which the mean square error per hour of this signal, organized by month, shows that the value of the discrepancy is predictable; this makes it possible to conclude that the weather forecast always has a similar lack of precision per hour, which can be modeled. We considered this an effect of the geographical surroundings of CASITA that are different to those of the location of the closest weather station of AEMET (Fuente Álamo).

With this implemented, it is easy to introduce this correction into the weather forecast, and assess the achieved improvement that is related to the real outdoor conditions measured. As can be seen in Figure 6, the root-mean-square error (RMSE) has decreased substantially, (especially in August).

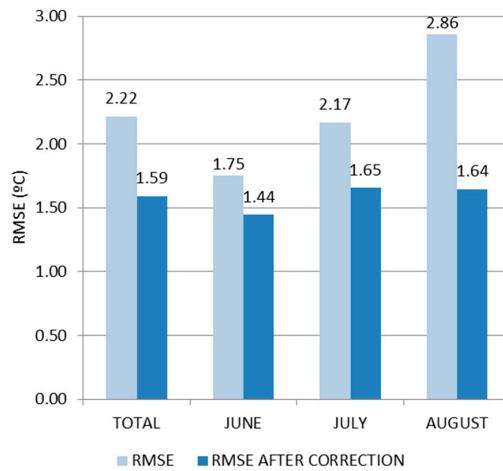


Figure 6. Root-mean-square error (RMSE) achieved after having into account the error mean evolution.

3.2. Validating Influence of the Variables Using PROPHET

A preliminary study of the data was done to ensure that there were no missing values or misleading results that could crash a computer code for data analysis and to perform preliminary sanity checks. The corrplot routine was applied to relevant variables of energy use. And the results can be seen in the following (see Figure 7).



Figure 7. Cross correlation between influential variables in energy consumption. Done with a native routine in R: corrplot.

Once the calculation was done, it was learned that the correlation between conditions out of the building and inside are not very strong, which demonstrates that the conditioning systems work well, and the space is not sensitive to fluctuations on the weather outside. Indoor conditions are clearly affected by the HVAC operation, which is directly in relation to occupancy. In other words, the more

people are working, the harder the air-conditioning devices are working; this is an expected result that demonstrates the validity of the data.

In one part of this preliminary study, we performed a validation exercise that uses the software package PROPHET. It is an open source code that runs in R and performs predictions of time series for many kinds of variables due to its large popularity in other disciplines; we thought it was interesting to test its performance in building physics.

Focusing on the variable that aimed to be forecasted, after studying the results it is possible to conclude that there is a correlation between energy consumption and indoor temperature and humidity. If one sees the results, outdoor temperature is an important variable, instead outdoor humidity is not that significant as one could expect. The same interpretation could be made regarding the weather forecast; temperature is relevant, and humidity is not. Occupation is revealed to be the more influential variable, as the presence of people is an essential requirement to have energy consumption.

In order to make a model of the energy consumption, there are some variables which show this to be influential:

- Occupation;
- Indoor conditions: temperature and humidity;
- Outdoor temperature;
- Forecasted temperature.

The energy consumption in buildings has several characteristics appropriate for the PROPHET algorithm and thus should perform well for energy prediction. These are:

- Strong multiple human-scale seasonality (such as day of the week and the time of year);
- Important holidays that occur at irregular intervals that are known in advance; and
- A certain random component.

Together with previous observations about energy consumption, in our problem the domain knowledge was defined by the inclusion of external regressors that were selected after observing the results:

- Model 1: Forecasting energy consumption in a 24-h predictive horizon.
- Previous energy consumption.
- Previous occupation and future values of this variable with a known pattern and schedule.
- Model 2: Forecasting energy consumption in a 24-h predictive horizon.
- Previous energy consumption.
- Previous occupation and future values of this variable with a known pattern and schedule.
- Outdoor temperature values with temperature predictions filling the time series to be predicted.

These models and their differences are explained in Figure 8. In Model 1 we forecast energy consumption only with previous values of this variable and previous occupation, completed with future occupation. Model 2 introduces, in addition, previous outdoor temperature measurements and the forecast are helped with future temperature values obtained from the national meteorological agency.

It was decided to perform the energy prediction using a sampling period of one hour, this was because that granularity captures most of the dynamics of the building without compromising the volume of data. An extra seasonality component was added that relates to the daily periodic of any energy-related variable linked to human behavior. The implementation was run on the R environment [41].

In order to make a first approximation, a prediction was performed for 12 August at 9:00, (start of working hours) with a predictive horizon of 24 h. In the following paragraphs, we studied two predictive models. These models are compared with the real measured energy consumption.

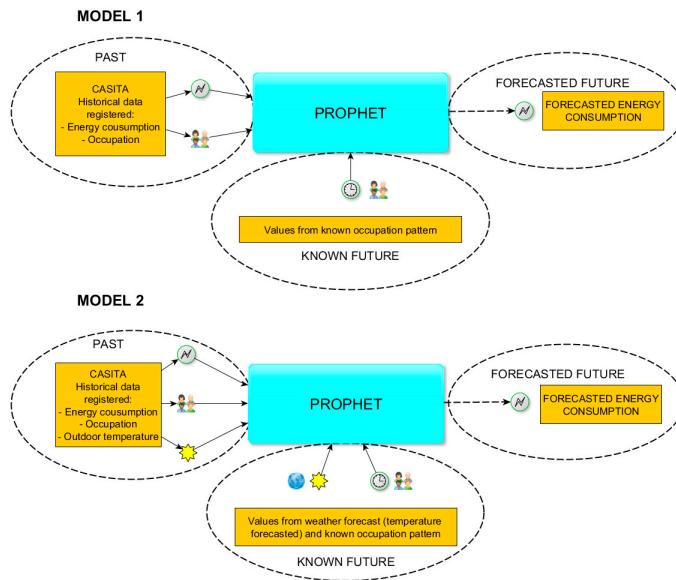


Figure 8. Schema of Model 1 and Model 2.

In the following points, two scenarios were tested, which contrasted with two different situations that provide the two different regressors previously mentioned.

After running the models, it was possible to ask for the next 24-h prediction. Figures 9 and 10 show the 24-h predictions performed with the fitter model (blue line) and the true values (black dots), while the last peak represents the forecasted energy consumption on 12 August, with a previously given occupation schedule.

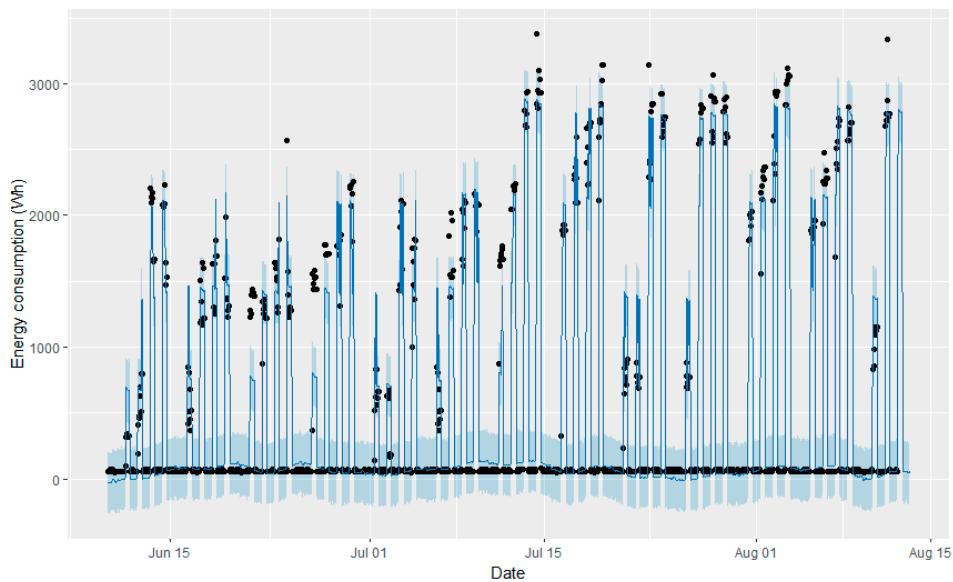


Figure 9. Twenty-four hour predictions performed with the fitter model (blue line) and the true values (black dots) with Model 1.

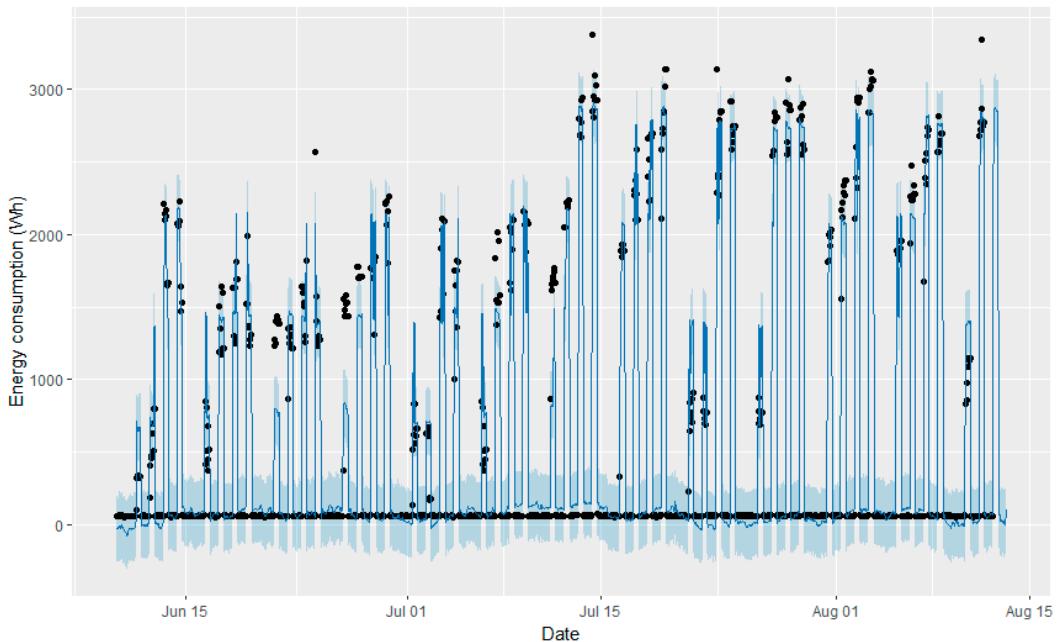


Figure 10. Twenty-four hour predictions performed with the fitter model (blue line) and the true values (black dots) with Model 2.

Although both graphs seem to be very similar, the forecasted hours of the predictions slightly differ for the two models. These results serve as proof that both models have a good approximation to real measurements, but that there are some slight differences. It seems that Model 2 is closer to reality. A comparison between the real measures, and the forecasts with Model 1 and Model 2, is shown in Figure 11.

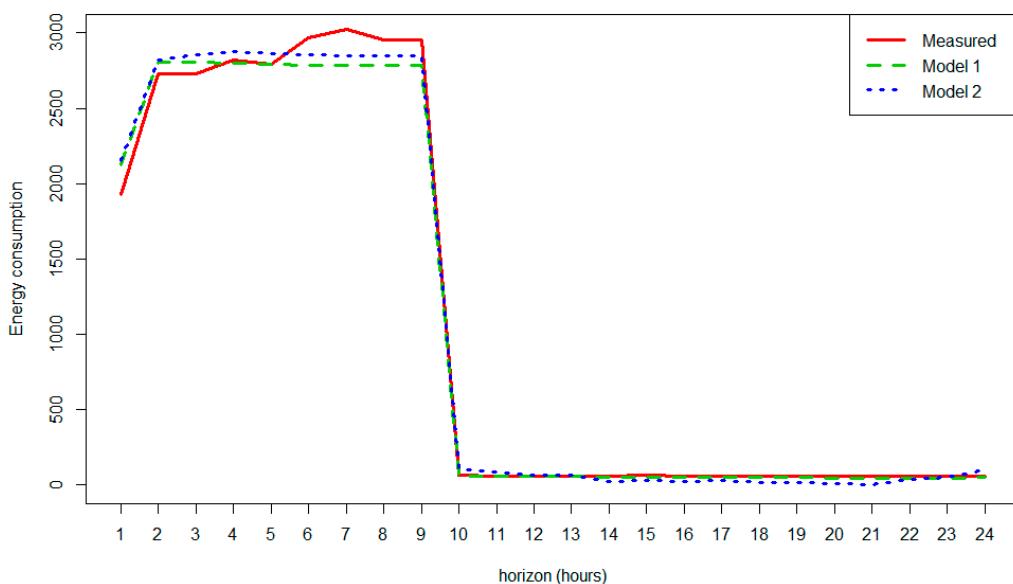


Figure 11. Energy consumption-real measures vs prediction (Model 1/Model 2).

After this example, and in order to evaluate this in a more comprehensive and general way, a cross validation was made to extract some conclusions of the different approaches. The cross validation assumes that these models could be generalized and that their accuracy of predictions estimated. In this iterative process, each hour is predicted using the rest of the available data, obtaining a measure of the error. This is done in Model 1 and Model 2 per hour several times. In addition, a visualisation of a given prediction is shown in Figure 12. It is then possible to estimate the mean absolute error (MAE) for each case under study by measuring the difference between the subsequent real measures and the predicted values. As we can see in Figure 12, combining the MAE per hour makes it possible to see that both have similar behavior. The MAE results in being higher in the working hours and are maximal at 9:00 when the HVAC starts; this is when the variability of the energy consumption is high and rapid. As is easily seen in Figure 12, Model 2 presents a better performance for nearly the entire day. Hence, the MAE is smaller and the accuracy of the model has been improved by adding outdoor temperature and temperature forecast in the prediction phase.

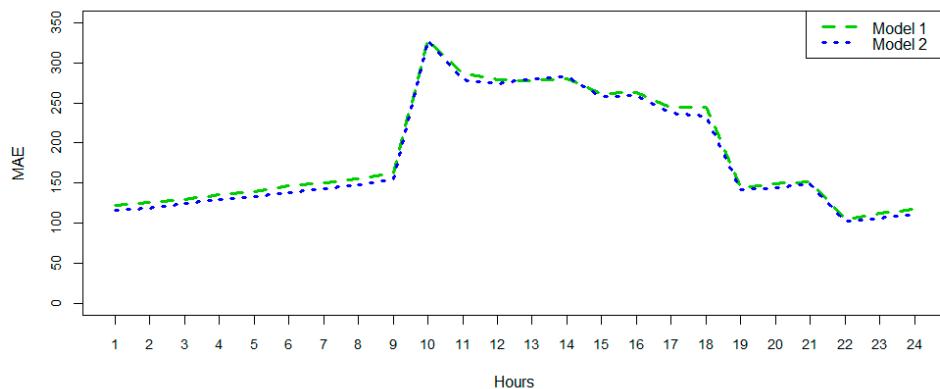


Figure 12. Mean absolute error (Model 1/Model 2).

The improvement achieved can also be shown in terms of RMSE, offering an increase of accuracy. Global reduction of RMSE has been quantified in 4.54%, from the data exposed in Table 2.

Table 2. Evolution of RMSE values over 24 h.

Hour	RMSE Model 1	RMSE Model 2	Improvement	Hour	RMSE Model 1	RMSE Model 2	Improvement
01	192.93	176.80	8.36%	13	378.35	384.33	-1.58%
02	200.24	182.96	8.63%	14	381.95	381.93	0.00%
03	210.39	191.86	8.81%	15	358.22	358.05	0.05%
04	222.05	202.28	8.90%	16	358.66	352.96	1.59%
05	231.73	212.67	8.23%	17	349.19	342.75	1.84%
06	243.96	222.99	8.59%	18	356.19	344.36	3.32%
07	251.60	230.77	8.28%	19	249.11	247.70	0.57%
08	262.76	239.10	9.00%	20	258.60	255.62	1.15%
09	275.33	250.72	8.94%	21	269.52	265.63	1.44%
10	427.56	432.00	-1.04%	22	160.57	149.72	6.75%
11	381.20	377.35	1.01%	23	172.48	159.43	7.57%
12	376.05	374.41	0.43%	24	181.99	167.22	8.12%

4. Conclusions and Future Work

This work describes the commissioning of a new testing facility that has been given the name of Controlled and Automatized Testing Facility for Human Behavior (CASITA). The facility includes a large variety of sensors, meters, and actuators that allow the research to focus on fundamental aspects

of the interactions of humans with built environments. The new contribution is that we have conceived this facility as a pair between the hardware and the software package PROPHET that provides the soft components (algorithms and analysis tools) to make the facility complete.

The first testing of this facility consisted of an occupation experiment that was performed to facilitate the posterior analysis of the software package PROPHET. The results of this software used in publications in other fields convinced us that it could be an excellent addition to CASITA for experiments that involved prediction (as there are many).

Building energy consumption models with new techniques are of considerable interest to the scientific community. Our test experiment was to evaluate the functionalities of CASITA, and to ascertain the improvement of the PROPHET algorithms. Once the correlations were studied, two models were presented. After a brief explanation about a new tool for modeling and forecasting, the PROPHET package of the R software, some parameter settings and a comparison between the models became topics of discussion. The results indicate that introducing outdoor temperature into the model that uses the forecasted temperature provided by AEMET (an official source of weather forecast) improves the accuracy of its predictions.

The variables chosen in this work can be found in any residential or commercial building. As far as a sensor network that would be deployed, the same data can be collected and the models replicated. Therefore, the approach in this paper proposes an improved general model for forecasting energy consumption in buildings. A good approximation to this problem could enable one to plan for energy requirements, achieve energy and economic savings, and contribute to a more effective energy consumption policy.

In essence, the results show that CASITA is an excellent research facility that can be used for the testing of human modeling algorithms, IoT platforms, control strategies, and many more applications. With respect to the prediction algorithms tested here, both models are acceptable and achieve a good level of representativeness.

In addition, it is necessary to note that the algorithms tested have good accuracy but that they have not been compared with other methods, as the main aim of this work was commissioning CASITA and evaluating PROPHET as a side tool for it. We believe that the testing of its suitability was sufficient.

In future works, other scenarios will be tested. Weather forecasting will be added again along with other possible forecasted variables. We also plan to introduce human components into the equation, which would be interesting and will exploit the capabilities of CASITA well. In another vein, the PROPHET package became a tool whose benefits need to be studied further in this and other fields.

Author Contributions: Conceptualization, I.R.-R. and M.Á.Z.; Methodology, I.R.-R. and A.G.V.; Software, A.G.V.; Validation, I.R.-R., A.G.V. and A.P.R.G.; Investigation, I.R.-R., A.G.V., A.P.R.G. and M.Á.Z.; Resources, M.Á.Z.; Writing-Original Draft Preparation, I.R.-R. and M.Á.Z.; Writing-Review and Editing, A.P.R.G.; Supervision, I.R.-R. and A.P.R.G.; Funding Acquisition, M.Á.Z.

Funding: This work has been sponsored by the Spanish Ministry of Economy and Competitiveness through 387 the PERSEIDES (ref. TIN2017-86885-R) and CHIST-ERA (ref. PCIN-2016-010) projects; by MINECO grant BES-2015-071956 and by the European Comission through the H2020-ENTROPY-649849 EU Project.

Acknowledgments: Ramallo-González would like to thank the program Saavedra Fajardo (grant number 220035/SF/16) funded by Consejería de Educación y Universidades of CARM, vía Fundación Séneca-Agencia de Ciencia y Tecnología de la Región de Murcia.

Conflicts of Interest: The authors declare no conflict of interest. The funders had no role in the design of the study; in the collection, analyses, or interpretation of data; in the writing of the manuscript; and in the decision to publish the results.

References

1. Perez-Lombard, L.; Ortiz, J.; Pout, C. A review on buildings energy consumption information. *Energy Build.* **2008**, *40*, 394–398. [[CrossRef](#)]
2. Schnieders, J. CEPHEUS—Measurement Results from More Than 100 Dwelling Units in Passive Houses. Available online: https://www.researchgate.net/publication/237709858_CEPHEUS_-_Measurement_results_from_more_than_100_dwelling_units_in_passive_houses (accessed on 19 July 2018).
3. Schnieders, J.; Hermelink, A. CEPHEUS results: Measurements and occupants' satisfaction provide evidence for Passive Houses being an option for sustainable building. *Energy Policy* **2006**, *34*, 151–171. [[CrossRef](#)]
4. Mogles, N.; Walker, I.; Ramallo-González, A.P.; Lee, J.; Natarajan, S.; Padgett, J.; Gabe-Thomas, E.; Lovett, T.; Ren, G.; Hyniewska, S.; et al. How smart do smart meters need to be? *Build. Environ.* **2017**, *125*, 439–450. [[CrossRef](#)]
5. Terroso-Saenz, F.; González-Vida, A.; Ramallo-González, A.P.; Skarmeta, A.F. An open IoT platform for the management and analysis of energy data. *Future Gener. Comput. Syst.* **2017**. [[CrossRef](#)]
6. Moreno, V.; Úbeda, B.; Skarmeta, A.F.; Zamora, M.A. How can we tackle energy efficiency in IoT basedsmart buildings? *Sensors* **2014**, *14*, 9582–9614. [[CrossRef](#)] [[PubMed](#)]
7. Darby, S. Making it obvious: Designing feedback into energy consumption. In *Energy Efficiency in Household Appliances and Lighting*; Springer: Berlin/Heidelberg, Germany, 2001; pp. 685–696.
8. Karatasou, S.; Laskari, M.; Santamouris, M. Models of behavior change and residential energy use: A review of research directions and findings for behavior-based energy efficiency. *Adv. Build. Energy Res.* **2014**, *8*, 137–147. [[CrossRef](#)]
9. Newton, D.; James, R.; Bartholomew, D. Building energy simulation—A user's perspective. *Energy Build.* **1988**, *10*, 241–247. [[CrossRef](#)]
10. Haldi, F.; Robinson, D. The impact of occupants' behaviour on building energy demand. *J. Build. Perform. Simul.* **2011**, *4*, 323–338. [[CrossRef](#)]
11. Rouleau, J.; Ramallo-González, A.; Gosselin, L. Towards a comprehensive tool to model occupant behaviour for dwellings that combines domestic hot water use with active occupancy. In Proceedings of the 15th IBPSA Conference, San Francisco, CA, USA, 7–9 August 2017.
12. Agarwal, Y.; Balaji, B.; Gupta, R.; Lyles, J.; Wei, M.; Weng, T. Occupancy-driven energy management for smart building automation. In Proceedings of the 2nd ACM Workshop on Embedded Sensing Systems for Energy-Efficiency in Building, Zurich, Switzerland, 2 November 2010; pp. 1–6.
13. Pettersen, T.D. Variation of energy consumption in dwellings due to climate, building and inhabitants. *Energy Build.* **1994**, *21*, 209–218. [[CrossRef](#)]
14. Lindberg, R.; Binamu, A.; Teikari, M. Five-year data of measured weather, energy consumption, and time-dependent temperature variations within different exterior wall structures. *Energy Build.* **2004**, *36*, 495–501. [[CrossRef](#)]
15. Keller, J.; Heiko, A. The influence of information and communication technology (ICT) on future foresight processes—Results from a Delphi survey. *Technol. Forecast. Soc. Chang.* **2014**, *85*, 81–92. [[CrossRef](#)]
16. Weiser, M. The computer for the 21st century. *Sci. Am.* **1991**, *265*, 94–104. [[CrossRef](#)]
17. Atzori, L.; Iera, A.; Morabito, G. The internet of things: A survey. *Comput. Netw.* **2010**, *54*, 2787–2805. [[CrossRef](#)]
18. Perera, C.; Zaslavsky, A.; Christen, P.; Georgakopoulos, D. Sensing as a service model for smart cities supported by internet of things. *Trans. Emerg. Telecommun. Technol.* **2014**, *25*, 81–93. [[CrossRef](#)]
19. European Commission. *Benchmarking Smart Metering Deployment in the EU-27 with a Focus on Electricity*; Publications Office of the European Union: Luxembourg, 2014.
20. Voss, K.; Sartori, I.; Napolitano, A.; Geier, S.; Gonçalves, H.; Hall, M.; Heiselberg, P.; Widén, J.; Candanedo, J.A.; Musall, E. Load matching and grid interaction of net zero energy buildings. In Proceedings of the EUROSUN 2010 International Conference on Solar Heating, Cooling and Buildings, Graz, Austria, 28 September–1 October 2010.
21. Zhao, H.; Magoules, F. A review on the prediction of building energy consumption. *Renew. Sustain. Energy Rev.* **2012**, *16*, 3586–3592. [[CrossRef](#)]
22. Han, D.M.; Lim, J.H. Design and implementation of smart home energy management systems based on zigbee. *IEEE Trans. Consum. Electron.* **2010**, *56*, 1417–1425. [[CrossRef](#)]

23. Oksa, P.; Soini, M.; Sydneimo, L.; Kivikoski, M. Kilavi platform for wireless building automation. *Energy Build.* **2008**, *40*, 1721–1730. [CrossRef]
24. Moreno, M.V.; Zamora, M.A.; Santa, J.; Skarmeta, A.F. An indoor localization mechanism based on RFID and IR data in ambient intelligent environments. In Proceedings of the 2012 Sixth International Conference on Innovative Mobile and Internet Services in Ubiquitous Computing, Palermo, Italy, 4–6 July 2012.
25. Moreno-Cano, M.; Zamora-Izquierdo, M.A.; Santa, J.; Skarmeta, A.F. An indoor localization system based on artificial neural networks and particle filters applied to intelligent buildings. *Neurocomputing* **2013**, *122*, 116–125. [CrossRef]
26. Hernández-Ramos, J.L.; Moreno, M.V.; Bernabé, J.B.; Carrillo, D.G.; Skarmeta, A.F. SAFIR: Secure access framework for IoT-enabled services on smart buildings. *J. Comput. Syst. Sci.* **2015**, *81*, 1452–1463. [CrossRef]
27. Moreno, M.V.; Zamora, M.A.; Skarmeta, A.F. An IoT based framework for user-centric smart building services. *Int. J. Web Grid Serv.* **2015**, *11*, 78–101. [CrossRef]
28. Moreno, V.; Zamora, M.A.; Skarmeta, A.F. A low-cost indoor localization system for energy sustainability in smart buildings. *IEEE Sens. J.* **2016**, *16*, 3246–3262. [CrossRef]
29. Moreno, M.V.; Terroso-Sáenz, F.; González-Vidal, A.; Valdés-Vela, M.; Skarmeta, A.F.; Zamora, M.A.; Chang, V. Applicability of big data techniques to smart cities deployments. *IEEE Trans. Ind. Inform.* **2017**, *13*, 800–809. [CrossRef]
30. Zamora-Izquierdo, M.A.; Santa, J.; Gmez-Skarmeta, A.F. An integral and networked home automation solution for indoor ambient intelligence. *IEEE Pervasive Comput.* **2010**, *9*, 66–77. [CrossRef]
31. Hazas, M.; Friday, A.; Scott, J. Look back before leaping forward: Four decades of domestic energy inquiry. *IEEE Pervasive Comput.* **2011**, *10*, 13–19. [CrossRef]
32. Agencia Estatal de Meteorología-AEMET. Gobierno de España. Available online: <http://www.aemet.es/es/portada> (accessed on 18 July 2018).
33. Zoha, A.; Gluhak, A.; Ali Imran, M.; Rajasegarar, S. Non-intrusive load monitoring approaches for disaggregated energy sensing: A survey. *Sensors* **2012**, *12*, 16838–16866. [CrossRef] [PubMed]
34. Santander, P.; Elórtegui, C.; González, C.; Allende-Cid, H.; Palma, W. Redes sociales, inteligencia computacional y predicción electoral: El caso de las primarias presidenciales de Chile 2017. *Cuadernos. Info* **2017**, *41*, 41–56. (In Spanish) [CrossRef]
35. Hamadeh, A. Anomaly Detection in a Multivariate DataStream in a Highly Scalable and Fault Tolerant Architecture. Master's Thesis, KTH, Stockholm, Sweden, 2017.
36. Brügner, H. Holt-Winters Traffic Prediction on Aggregated Flow Data. In Proceedings of the Seminars Future Internet (FI) and Innovative Internet Technologies and Mobile Communication (IITM) Focal Topic: Advanced Persistent Threats, Munich, Germany, 24 February–16 August 2017; pp. 25–32.
37. Riihijarvi, J.; Mahonen, P. Machine Learning for Performance Prediction in Mobile Cellular Networks. *IEEE Comput. Intell. Mag.* **2018**, *13*, 51–60. [CrossRef]
38. Alkharif, S.; Lee, K.; Kim, H. Time-Series Analysis for Price Prediction of Opportunistic Cloud Computing Resources, In Proceedings of the 7th International Conference on Emerging Databases, Busan, Korea, 7–9 August 2017.
39. Saad, F.; Mansinghka, V. Temporally-Reweighted Chinese Restaurant Process Mixtures for Clustering, Imputing, and Forecasting Multivariate Time Series. In Proceedings of the International Conference on Artificial Intelligence and Statistics, Playa Blanca, Lanzarote, Spain, 9–11 April 2018; pp. 755–764.
40. Taylor, S.; Letham, B. Prophet: Automatic Forecasting Procedure. Available online: <https://cran.r-project.org/web/packages/prophet/prophet.pdf> (accessed on 25 June 2018).
41. Taylor, S.J.; Letham, B. Forecasting at scale. *Am. Stat.* **2018**, *72*, 37–45. [CrossRef]
42. González-Vidal, A.; Ramallo-González, A.P.; Terroso-Sáenz, F.; Skarmeta, A. Data driven modeling for energy consumption prediction in smart buildings. In Proceedings of the 2017 IEEE International Conference on Big Data, Boston, MA, USA, 11–14 December 2017; pp. 4562–4569.
43. González-Vidal, A.; Moreno-Cano, V.; Terroso-Sáenz, F.; Skarmeta, A.F. Towards energy efficiency smart buildings models based on intelligent data analytics. *Procedia Comput. Sci.* **2016**, *83*, 994–999. [CrossRef]



© 2018 by the authors. Licensee MDPI, Basel, Switzerland. This article is an open access article distributed under the terms and conditions of the Creative Commons Attribution (CC BY) license (<http://creativecommons.org/licenses/by/4.0/>).

4.4 Applicability of Big Data Techniques to Smart Cities Deployments

Title	Applicability of Big Data Techniques to Smart Cities Deployments
Authors	M. Victoria Moreno, Fernando Terroso-Sáenz, Aurora González-Vidal, Mercedes Valdés-Vela Antonio F. Skarmeta, Miguel A. Zamora and Victor Chang
Type	Journal
Journal	IEEE TRANSACTIONS ON INDUSTRIAL INFORMATICS
Impact factor (2018)	7.377
Rank	Q1
Publisher	IEEE
Volume	13
Issue	2
Pages	800 - 809
Year	2017
Month	April
ISSN	1551-3203 (Print), 1941-0050 (Electronic)
DOI	10.1109/TII.2016.2605581
URL	https://ieeexplore.ieee.org/abstract/document/7558230/
State	Published
Author's contribution	The PhD student, Aurora González Vidal, contributed to the methodology, software, analysis and research and writing the publication

Applicability of Big Data Techniques to Smart Cities Deployments

M. Victoria Moreno, Fernando Terroso-Sáenz, Aurora González-Vidal, Mercedes Valdés-Vela,
Antonio F. Skarmeta, Miguel A. Zamora, and Victor Chang

Abstract—This paper presents the main foundations of big data applied to smart cities. A general Internet of Things based architecture is proposed to be applied to different smart cities applications. We describe two scenarios of big data analysis. One of them illustrates some services implemented in the smart campus of the University of Murcia. The second one is focused on a tram service scenario, where thousands of transit-card transactions should be processed. Results obtained from both scenarios show the potential of the applicability of this kind of techniques to provide profitable services of smart cities, such as the management of the energy consumption and comfort in smart buildings, and the detection of travel profiles in smart transport.

Index Terms—Big data, Internet of Things (IoT), predictive models, smart city, transit-card mining.

I. INTRODUCTION

A SMART city emerges when the urban infrastructure is evolved through the information and communication technologies [1]. The paradigm of Internet of Things (IoT) [2] has enabled the emergence of a high number of different communication protocols, which can be used to communicate with commercial devices using different data representations. In this context, an IoT-based platform is necessary to manage all interoperability aspects and enable the integration of optimal artificial intelligence (AI) techniques in order to model contextual relationships.

In urban environments, there is a huge amount of different data sources. Plenty of sensors are distributed around cities, most

Manuscript received March 10, 2016; revised July 21, 2016; accepted August 15, 2016. Date of publication September 1, 2016; date of current version April 18, 2017. This work was supported in part by the Spanish Seneca Foundation by means of the PD Program under Grant 19782/PD/15, in part by the MINECO TIN2014-52099-R Project under Grant BES-2015-071956 and ERDF funds, and in part by the European Commission through the H2020-ENTROPY-649849 and FP7-SMARTIE-609062 EU Projects. Paper no. TII-16-0282.R1. (Corresponding author: M. V. Moreno.)

M. V. Moreno, F. Terroso-Sáenz, A. González-Vidal, M. Valdés-Vela, A. F. Skarmeta, and M. A. Zamora are with the Department of Information and Communications Engineering, University of Murcia, Murcia 30100 Spain (e-mail: mvmoreno@um.es; fterroso@um.es; auroragonzalez2@um.es; mdvaldes@um.es; skarmeta@um.es; mzamora@um.es).

V. Chang is with Xi'an Jiaotong Liverpool University, Suzhou 215123, China (e-mail: ic.victor.chang@gmail.com).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/TII.2016.2605581

of them installed in indoor spaces. This situation has brought new analytics mechanisms and tools that provide insight allowing us to have an effective and collaborative way to operate the machines [3]. Furthermore, there are numerous mobile data sources like smart phones, smart cards, wearable sensors and, in the case of vehicles, on-board sensors. All these sensors provide information that makes possible to detect urban dynamic patterns. Nonetheless, the most existing management systems of cities are not able to utilize fully and effectively this vast amount of data and, as a result, there are large volumes of data that are not exploited. In this direction, many AI techniques in computer science have been introduced to deal with the processing of huge amount of data to extract useful information (or termed by knowledge) from data [4], this trend is known as big data.

This paper is intended to analyze the interest of big data for smart cities. In order to face the above-mentioned aspects we propose a general architecture for smart city applications, which is modeled in four layers with different functionalities. Then, we show some applications of big data analysis in two scenarios, both focusing on sensed data coming from both static and dynamic sources. Among other objectives, the first scenario intends to create a distributed framework to share large volumes of heterogeneous information for their use in smart building applications. In this paper, we focus on presenting the deployments and implementations carried out in smart buildings to achieve energy efficiency. For this, different problems like indoor localization, thermal comfort characterization, and energy consumption modeling have been solved through the application of big data techniques. The second example is centered on the public tram service in the City of Murcia (Spain), looking for giving insight into the great amount of data generated by the service's transit cards. In this scenario, big data techniques are applied to extract mobility patterns in public transport.

Hence, this paper discusses three aspects of nowadays smart cities that need to be solved, and for each one of them we provide some research contributions through the application of convenient big data techniques. These contributions are:

- 1) The design and instantiation of an IoT-based architecture for applications of smart cities.
- 2) The approach of an efficient management of energy in smart buildings.
- 3) The extension of the data analysis for the detection of urban patterns, which can be used to improve public transport applied to the public tram service.

The structure of this paper is as follows: Section II enumerates the challenges that current smart cities still have to face, and proposes a general IoT-based architecture for smart city services, which is modeled in layers. Section III describes a first application of smart city where big data techniques have been applied to get energy efficiency in the buildings of a smart campus. Section IV presents a second smart city application that addresses the urban pattern recognitions in public transport. Section V summarizes the main benefits of applying big data techniques to the two scenarios of smart cities addressed in this paper. Finally, Section VI gives some conclusions and an outlook of future work.

II. IoT-BASED ARCHITECTURE FOR SMART CITIES

In this section, we enumerate the main challenges that most current smart cities still have to face. Then, motivated by these challenges, we make a proposal of a general IoT-based architecture for smart city applications.

A. Challenges of Smart Cities

The global challenges that smart cities still have to face can be summarized in the following way:

- 1) *Sensors integration and abstraction capability*: It provides means to integrate different sensor types in a common platform taking into account the different technologies, legacy systems, and communication protocols with focus on IPv6 solutions.
- 2) *Individual intelligence and local reasoning*: Apart from data fusion, more complex data processing can be implemented by smart objects.
- 3) *Learning and adaptation*: Most of the patterns generated in smart cities are sensitive to contextual changes and are able to learn and adapt themselves to such changes as well as to human dynamicity.
- 4) *Dynamic human centric services*: This work designs and implements smart mobility and smart building services that use the patterns generated to provide customized and efficient services taking into account the dynamicity of the citizens behavior.
- 5) *User privacy and security control mechanisms*: In the context of smart cities, it is important to manage the way the user is able to control its data and how they are exposed to third parties and applications.

B. IoT-Based Architecture

Several layers compound the proposed platform that was created with the goal of serving to many applications of smart cities. Fig. 1 depicts the layered IoT-based architecture, which is detailed below.

1) Technologies Layer: In the basis part of Fig. 1, it is observed that a plethora of sensors and network technologies provide the input attributes using wireless sensor networks, wired sensors, gateways, etc., which can be self-configured and remotely controlled through the Internet. Focusing on our first application that consists of the instantiation of the architecture for building management systems (BMS), in this layer it is

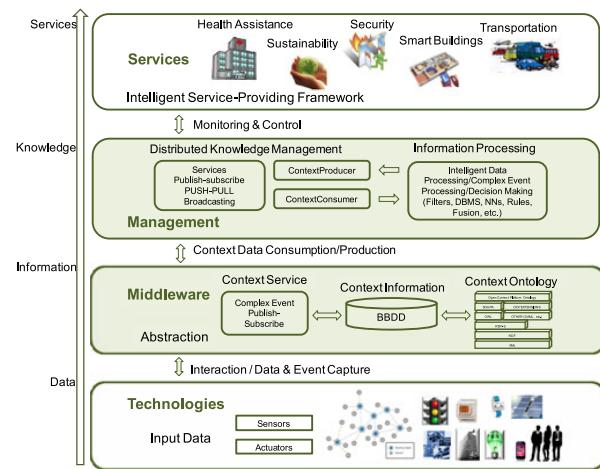


Fig. 1. Layers of the base architecture for smart city services.

gathered information from sensors and actuators deployed in strategic points of the building. But the aforementioned data sources in smart cities are not limited to static devices reporting measurements associated to a particular location, there are also moving ones capable to deliver measurements at different points within a geographical area. This is mainly due to the rapid development of wireless technology, mobile sensor networks and, above all, the advent of smartphones [5]. Although approaches based on mobile-phone sensing require a demanding usage of the communication, location, and other attributes of the smartphone, which can bother some people due to battery draining [6], data captured by static, mobile, and smart-phone sensors can be extended or enriched with the data generated by several social-media channels—like Twitter or Facebook—giving rise to a new generation of *soft sensors* from which we can extract relevant knowledge [7]. As a result, an alternative course of action aims at mining relevant knowledge from users on the basis of nonintrusive ways to obtain data, for example, transit cards in public transport scope.

2) Middleware Layer: The first layer provides a wide variety of data, so a second layer is needed where all collected data from seamless sources are expressed in the same way, this is done in the middleware layer. The context information can be collected in an ontology defined according to the model that represents the knowledge of the specific application domain. Thus, for our energy efficiency semantic model, the devices and building concepts are borrowed by the Smart Appliances REference (SAREF) ontology [8]. The agents representation is made using the DOLCE+DnS Ultralite (DUL) ontology [9], while the observation values of the monitored sensors are represented based on the Semantic Sensor Network (SSN) ontology [10]. However, when it comes to process the incoming data in a real-time manner, it is necessary to use a lightweight representation. As a matter of fact, a sensor-data representation using a simple attribute-value schema for event-based systems is described in [11].

3) Management Layer: After having extracted information from the previous layers, the management layer is in charge of determining decisions bearing in mind the target services provided in smart cities. Different big data analytic techniques can

be used for the intelligent decision making processes. Algorithms like artificial neural networks (ANNs) using backpropagation methods [12] and support vector machines (SVMs) [13] are good to solve nonlinear problems, making them very applicable to build energy prediction issues, ranging from those associated with lighting and heating, ventilation and air condition (HVAC) [14] to the prediction of the heating energy requirements [15]. For optimization problems in BMS addressing energy efficiency, genetic algorithms (GAs) constitute a commonly applied heuristic that can be used in several optimization scenarios such as scheduling cooling operation decisions [16]. Regarding to the smart public transport application, the extraction of users behaviors from transition records have been studied by using different algorithms and techniques like maximum-likelihood estimation [17], probabilistic models [18], conditional random fields [19], graphical information system based processing [20] or database management system (DBMS) based processing [21].

4) Services Layer: Finally, the upper layer (see Fig. 1) shows some examples of smart city services that can be provided following the proposed architecture. Thus, this architecture can be applied to provide applications of smart cities like environmental monitoring, energy efficiency in buildings and public infrastructures [22], environmental monitoring [23], traffic information and public transport, locating citizens, manage emergencies, saving energy, and other services. These actions can either involve citizens or be automatically set.

III. SMART CAMPUS OF THE UNIVERSITY OF MURCIA (UMU)

The UMU has three main campus and several facilities deployed throughout different cities in the Region of Murcia. One of these campuses is currently serving as pilot of two European Projects, the SMARTIE [24] and the ENTROPY [25] project. The goal of this use in the case of smart city is to provide a reference system that is able to manage intelligently the energy use of the most relevant contributor to the energy use at city level, i.e., buildings. The BMS implemented as a part of this smart campus adapts the performance of automated devices through decisions made by the system and the interaction with occupants in order to keep comfort conditions while saving energy. We start by the most representative source of energy consumption at building level: HVAC systems.

A. System Overview

Using a BMS system, it is possible to predict users future behavior from their recorded activities that are measured with sensors. This information allows us to provide convenient environments looking for keeping their comfort while saving energy. The first need for a building to become smart is to know the location of occupants. Once solved the indoor localization problem, it is time to propose a solution to the energy efficiency of buildings associated with the thermal comfort provisioning service related to the HVAC management. For this, energy consumption models of the building need to be generated and used to implement the optimization mechanism able to maximize comfort at the same time that energy consumption is minimized. Therefore,

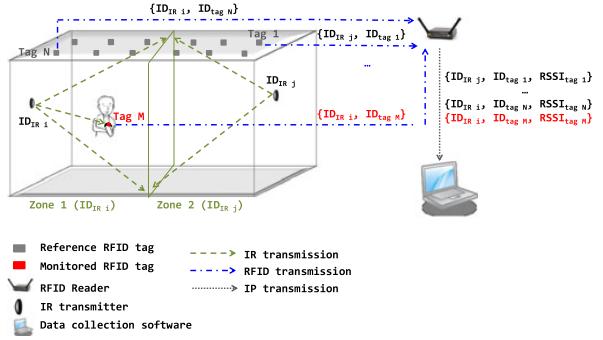


Fig. 2. Localization scenario.

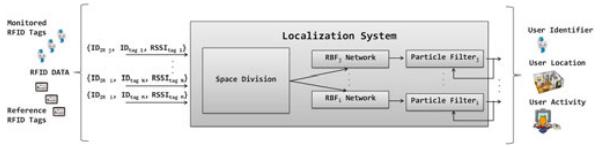


Fig. 3. Data processing for location estimation.

the different problems addressed in this scenario of smart city through the application of big data techniques are as follows:

- 1) indoor localization estimation;
- 2) building energy consumption prediction;
- 3) comfort provisioning and energy saving through an optimization problem.

In the following sections these problems are described with more details, as well as the techniques implemented and the results obtained.

B. Indoor Localization Estimation

Considering the information concerning to the identification and location of the building's occupants, it is necessary to reach the required accuracy in the location in order to provide the indoor services in a comfortable and energy efficient way. Our technological solution to cover the localization needs (i.e., those required by smart buildings to provide occupants with customized comfort services) is based on a single active RFID system and several infrared (IR) transmitters. In Fig. 2, we can observe the data exchange carried out among the different technological devices that compose our localization system.

The final mechanism implemented to solve the indoor localization problem is shown in Fig. 3. In this figure, we can see that the first phase of the mechanism is the space division through the installation of IR devices in the walls of the building area where localization wants to be solved. Therefore, for each space division, there is an IR identifier value (ID_{IR}) associated to this region. For each one of these regions, we implement a regression method based on radial basis functions (RBF) networks. The RBF estimates user positions given different RFID tags situated in the roof. Then, after the position estimation using the RBF network, a particle filter (PF) is applied as a monitoring technique, which takes into account previous user position data for estimating future states according to the current system model.

The PF used in this work is slightly different from its generic definition (which can be found in [26]). The main difference of our filter is in the correction stage. In this stage, the generic definition of the PF applies the resampling using the sequential importance sampling (SIS) algorithm [26] to carry out the filtering of such particles which minimize the deviation of their predicted trajectory. In our implementation, in addition to apply the SIS algorithm to correct the particles positions, we also use in this step the information about the specific IR region at a given instant of time to benefit those particles which fall inside this area. Therefore, before applying the SSI algorithm, we filter according to the coverage area of the IR transmitter identified by the monitoring RFID tag. The main advantage of this constraint is the faster convergence of the filter, because extra information is available to carry out the correction stage of the filter.

C. Building Energy Consumption Prediction

The energy performance model of our BMS is based on the *CEN Standard EN 15251* [27]. This standard proposes the criteria of design for any BMS. It establishes and defines the main input parameters for estimating building energy requirements and evaluating the indoor environment conditions. The inputs considered to solve our problem are the data coming from the RFID cards of users, the user interaction with the building automation system through the control panels or the web access, environmental parameters coming from temperature, humidity, and lighting sensors installed in outdoor and indoor spaces, the consumption energy sensed by the energy meters installed in the building, and the generated energy sensed by the energy meters installed in the solar panels deployed in our testbed. After collecting the data it is mandatory to continue with their cleaning, preprocessing, visualization, and correlation calculation in order to find determining features, which can be used to generate optimal energy consumption models of buildings (management layer of the architecture presented in Section II). Over the input set, we perform the standardization and reduction of data dimensionality using principal components analysis (PCA) [28], identifying the directions in which the observations of each parameter mostly vary.

Regarding the big data techniques that have been already applied successfully to generate energy consumption models of buildings in different scenarios (as such we mentioned in the management layer of the architecture presented in Section II), we propose to evaluate the performance of multilayer perceptron (MLP), Bayesian regularized neural network (BRNN) [29], SVM [30] and Gaussian processes with RBF Kernel [31]. They were selected because of the good performance that all of them have already showed when they are applied to building modeling. All these regression techniques are implemented following a model-free approach, which is based on selecting—for a specific building—the optimal input set and technique, i.e., such input set and technique that provides the most accurate predictive results in a test dataset. In order to implement this free-model approach, we use the R[32] package named CARET [33] to train the energy consumption predictive algorithms, looking for the optimal configuration of their hyperparameters (more infor-

mation can be found in [34]). The selected metric to evaluate the models generated for each technique using test sets is the well-known root mean square error (RMSE), whose formulation appears in (1). This metric shows the error by means of the quantity of kWh that we deviate when predicting

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (1)$$

But in order to get a better understanding of the uncertainty of the model, we also show its coefficient of variation (CVRMSE). This coefficient is the RMSE divided by the mean of the output variable (energy consumption) for the test set [see (2)], giving us a percentage of error adjusted to the data, not just a number in general terms:

$$\text{CVRMSE} = \frac{\text{RMSE}}{\bar{y}}. \quad (2)$$

D. Optimization Problem

Once the building energy consumption is modeled, we focus on the optimization of the HVAC operation trying to keep comfort conditions at the same time at which energy consumption restrictions are considered. As starting point, we establish the comfort extremes considering location type, user activity, and date [35]. Understanding the building thermal and energetic profiles allows us to quantify the effects of particular heating/cooling set point decisions. To derive a heating or cooling schedule, it is necessary to formulate the target outcome. In our buildings, it is possible to perform the following:

- 1) optimizing the indoor temperature during occupation, i.e., minimize the building temperature deviations from a target temperature;
- 2) minimizing daily energy consumption; or
- 3) optimizing a weighted mixture of the criteria, a so-called multiobjective optimization problem.

The definition of building temperature deviation influences the results strongly: Taking the minimum building temperature will result in higher set point choices and higher energy use than using, for instance, the average of indoor temperatures. Constraints on maximum acceptable deviation from target comfort levels or an energy budget can be taken into account to ensure required performance. In our optimization problem, we apply GA using the implementation provided by R (the “genalg” package [36]), to provide schedules for heating/cooling set points using the predictive building models (comfort and energy consumption models).

E. Evaluation and Results

1) Scenario of Experimentation: The reference building where our BMS for energy efficiency is deployed is the Technology Transfer Centre (TTC) of the UMU.¹ Every room of this building is automated through a home automation module unit. It permits us to consider a granularity at room level to carry out the experiments.

¹www.um.es/otri/?opc=cttfuentealamo

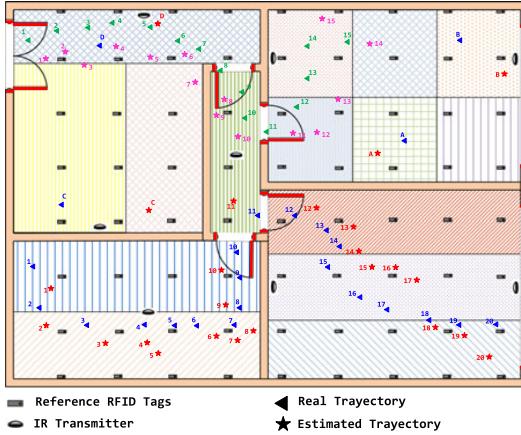


Fig. 4. Tracking processes with a reference tag distribution of $1 \text{ m} \times 1 \text{ m}$.

2) Evaluation. Indoor Localization Mechanism: Different tracking processes are carried out in the environments considered in our tests (the TTC building), implementing our PF. In Fig. 4, examples of some tracking processes are shown considering transition between different spaces of the TTC. For these paths, our system was configured to acquire data in every $T = 10 \text{ s}$. Taking into account the target location areas involved in comfort provisioning (lighting and thermal comfort, represented in different colors), and the real and estimated location data provided by our mechanism.

Thus, with a $1 \text{ m} \times 1 \text{ m}$ distribution of reference RFID tags placed on the roof of the test room, a 65% success percentage in localization is obtained having an error lower than 1 m. A total of 98% of cases have as much as 2.5 m of error. Therefore, it can be safely said that our localization system is able to track users with a sufficient level of accuracy and precision for the location requirements associated with the comfort and energy management in buildings. More details about this indoor localization system can be found in [37].

3) Evaluation. Energy Consumption Prediction: Fig. 5(a) shows the correlation heatmap between the electrical consumption of the TTC building and the outdoor environmental conditions. It is observed that energy consumption correlates significantly ($\alpha = 0.95$) and positively with temperature, radiation, wind speed variables, vapor pressure deficit, and dew point; and negatively with wind direction and humidity variables. This means that we can use these variables safely as inputs of the energy consumption model of our reference building, because they have a clear impact in the energy consumption. Otherwise, precipitation is so unusual that they do not have an association with the output.

Also, a logic differentiation between temporal situations has been considered in order to label behavior. Situation 1: holidays and weekends; situation 2: regular mornings; and situation 3: regular afternoons. The nonparametric Kruskall–Wallis test shows that the energy consumption differs significantly between situations [$H(2) = 547.7, p < 0.01$]. Also, the post-hoc pairwise comparisons corrected with Holm's method retrieve a

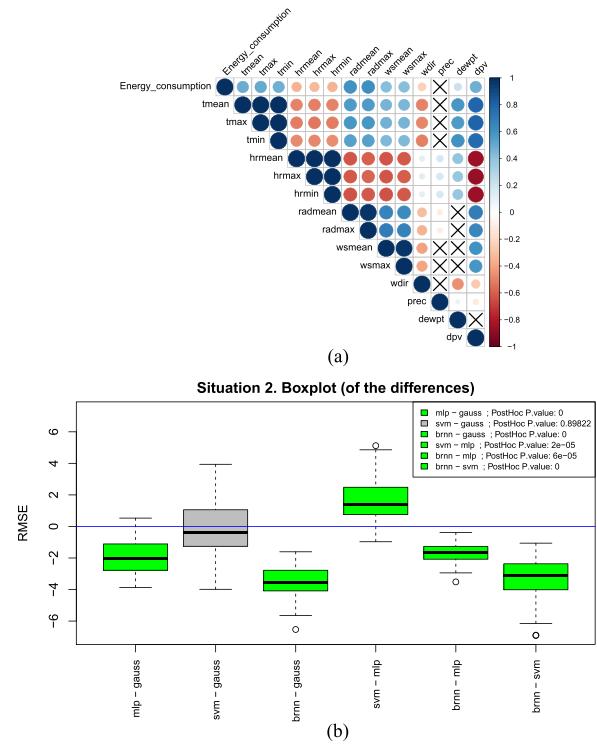


Fig. 5. Modeling results. (a) Correlation heatmap between consumption and outdoor environmental conditions. (b) Boxplots comparing models pairwise (situation 2).

p -value smaller than 0.01, supporting the decision of creating three different models [38]. Thus, for each of the three situations identified for the TTC building, we have evaluated not only the punctual value of RMSE, but also we have validated whether one learning algorithm outperforms statistically significantly the others using the nonparametric Friedman test [39] with the corresponding post-hoc tests for comparison.

Let x_i^j be the i th performance RMSE of the j th algorithm, for this building, we have used five-times ten-fold cross validation, so $i \in \{1, 2, \dots, 50\}$ and four techniques, so $j \in \{1, 2, 3, 4\}$. For every situation, we find significant differences ($\alpha = 0.99$) between every pair of algorithms, except for SVM and Gauss RBF ($p > 0.01$), as shown in Fig. 5(b) for the particular case of situation 2.

The three models have in common that BRNN yields a better result than the other tested techniques, based on the RMSE metric. Thus, BRNN is able to generate a model with a very low mean error of 25.17 kWh—which only represents 7.55% of the sample (this is the most accurate result) in terms of the CVRMSE. And for the worst case, BRNN provides a mean error of 43.76 kWh—which represents 10.29% of the sample in the reference TTC building—that is acceptable enough considering that our final aim is to save energy.

4) Evaluation. Optimization Mechanism: To evaluate our GA-based optimization strategy, controlled experiments were carried out in the TTC building with different occupant's

behaviors. The results show that we can accomplish energy savings between 15% and 31%. Trying to validate the applicability of our proposal, we have also made experiments in a different scenario with limited monitoring and automation technologies, achieving energy saving of about 23%.

IV. PUBLIC TRAM SERVICE OF MURCIA CITY

The second scenario is focused on the information analysis related to use of the tram service of the Region of Murcia [40]. In this case, the main goal was to perform a profiling process of the trips carried out by the users of such public service. For that aim, a fuzzy clustering algorithm is used to automatically extract tram user's profiles. Bearing in mind the architecture introduced in Section II, this system is enclosed in the management layer. The main tasks needed to reach the goal are explained in the following sections.

A. Generation of the Trip Dataset

According to the tram experts, information relevant to trip profiling must include data about time (in terms of day of the week and time of the day), origin and destination stations, and approximate age of the traveler. This information is being continuously recorded in different databases of the tram service. Nevertheless, certain operations of joining, transformation, and preprocessing (discretization and numerization) have been performed in order to compile all this information into a set of tuples susceptible of feeding the subsequent fuzzy clustering algorithm. The two most remarkable operations are the following.

On the one hand, according to the infrastructure of the tram service, users only need to swipe the smart card when they get into the tram. Hence, the recorded data only comprises transactions at the origin of each users trip so it can be regarded as incomplete. In order to deal with this incompleteness, a well-known solution is the *trip-chaining method*, which focus on recovering the origin and destination of the trips. In this case, such a method is based on the assumption that a traveler who takes the tram at an origin station, OS, ended their previous trip on that station OS. Due to the event-based nature of the card records, the complex event processing (CEP) paradigm [11] was adopted to come up with a palette of event-condition-action rules to uncover the trips. While the condition part of the rules performs a match between consecutive records of the same traveler following the aforementioned trip-chaining method, the action part generates a new trip tuple (comprising the origin and destination stations) in case the condition is fulfilled.

On the other hand, as clustering techniques are based on distance calculations among data, a set of numbered (and ordered) geographical areas, each one covering some close stations are identified by the tram experts. Then, instead of having nominal values for origin and destination features these numbered areas make it easier to calculate the distance about tuples in the clustering process.

In summary, the tuples composing the dataset for the subsequent clustering task are composed by the following attributes: $tt_e:\{travellerAge, dayOfTheWeek, hourOfDay, originArea, destArea\}$.

Algorithm 1: Cluster-based Trip profiling process.

```

Input:  $TT$ : dataset of raw trip tuples.
Output:  $P_{TT}$ : Traveller profiles extracted from  $TT$ .
1 if  $t_{now} - t_{prev} > tp_{max} \vee |TT| - |TT_{prev}| > nt_{max}$ 
then
2    $TT_e \leftarrow \text{preProcessing}(TT)$ 
3   foreach  $c \in \{2, \dots, c_{max}\}$  do
4      $clust_c = \text{GKCM}(TT_e, c)$ 
5     if  $clust_c.r_{cs} < r_{cs}^{min}$  then
6        $r_{cs}^{min} \leftarrow clust_c.r_{cs}$ 
7        $P_{TT} \leftarrow clust_c.\text{centroids}$ 
8    $t_{prev} \leftarrow t_{now}$ 
9    $TT_{prev} \leftarrow TT$ 
10  return  $P_{TT}$ 
```

B. Trip Profiling

Clustering mechanisms are suitable when it comes to find out the most representative trips profiles. For that aim, the Gustafson–Kessel clustering method (GKCM) has been chosen since it is able to identify arbitrarily oriented ellipsoidal fuzzy clusters unlike, for instance, the fuzzy C means clustering method, which impose spherical shapes to the data clusters. After the clustering task the identified prototypes (centroids) will summarize the whole dataset of trips. GKCM requires to be supplied with the quantity of potential clusters (c). This is an important parameter since it determines the ability of the potential centroids to represent the real underlying structure of the data.

Therefore, several GKCM executions were performed with different values of c and the *goodness* of the different identified set of clusters was measured. One of the most used measurements is the one proposed in [41] and denoted here as r_{cs} . This magnitude quantifying both the total compactness within clusters and the total separation among them being the greater the better.

Once the number of clusters c has been decided on the basis of r_{cs} , GKCM is executed in order to find the c profiles that best represent the trip dataset. Nevertheless, when exceed a time tp_{max} or a number of trips nt_{max} the algorithm is recomputed in order to detect new profiles which could rise up.

C. Evaluation and Results

The subject of evaluation is the tram service of the region of Murcia (Spain), which includes 18 km railway and 28 stations (see Fig. 6). Fig. 7 depicts the set of predefined geographical areas used in the experiment.

The evaluated dataset contained 37 8719 trips from 2 3400 users in November, 2013. For our experiment, the system was able to uncover 11 0697 trips. Expert knowledge was used to define the types of days and times of the day used in the aforementioned data preprocessing step as [Monday–Thursday, Friday, Saturday, Sunday] and [0–6, 6–10, 10–12, 12–16, 16–20, 20–00]. As a result, a generated TT_e dataset was split up into four different subsets based on the fact that traveler profiles depend



Fig. 6. Line map of the tram service in Murcia.

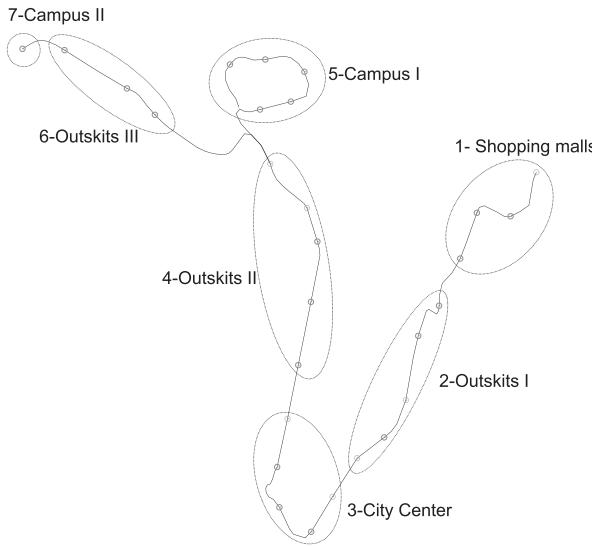


Fig. 7. Geographical regions for the numerization of tuples' station fields.

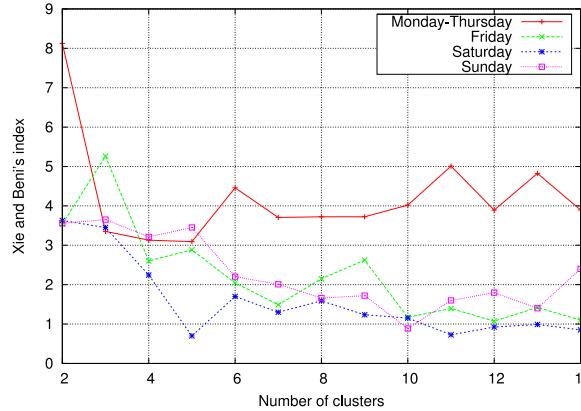


Fig. 8. Cluster validation rate for different cluster partitions.

on the day of the week (regarding, for example, differences of traffic flow between regular workdays and weekends). Next, the GKCM was launched with each of these subsets with different number of clusters.

In Fig. 8, the cluster validation ratio r_{cs} is shown for every TT_e subset, being the lower value the better. As it can be observed,

TABLE I
MONDAY–THURSDAY TRIPS' PROFILES

Profile	Age	Origin Area	Dest. Area	Time of the day
P1	23.37	City Center	Campus I	0–6
P2	25.74	City Center	Campus I	6–10
P3	28.22	Outskirts II	City Center	12–16
P4	32.77	Outskirts I	Outskirts II	6–10
P5	22.20	Campus I	City Center	16–20

while the optimal cluster partition is reached at $c = 5$ for the Monday–Thursday subset, for the remaining subsets minima r_{cs} values are reached at higher number of clusters c . In other words, a higher number of profiles is needed to represent the weekend trips. This is reasonable given that most people postpone leisure activities to the weekend and given that there exist quite a variety of leisure activities that can be done at different hours of the day.

As Table I shows, GKCM extracts five profiles for Monday–Thursday trips. Profiles 1 and 2 correspond to young people travelling in the morning to go toward one of the university zones from the station close the inner city. Besides, profile 5 represents a kind of traveler going back home from the university from 4 to 8 p.m. Finally, profiles 3 and 4 correspond to middle-young age people (28–33 years) that take the tram around the outskirts and city center environments. These could reflect people going from residential areas.

Finally, the heatmap shown in Fig. 9 represents the membership of the Monday–Thursday trips to the defined profiles. If we interpret this plot as a time-framed sequence, a great amount of the traffic focuses on the right side of the line, which connects the city center and the university areas. Nevertheless, such load is more spread along the whole line during the evening.

V. DISCUSSION

In this paper, we propose a general IoT-based architecture, which can be implemented for different applications of smart cities. This architecture is modeled in four layers, being the third one—the management layer—the layer where big data techniques are implemented to provide different services from those offered in the corresponding service layer (last layer).

The big data paradigm can be understood through the lens of 7 Vs [42] (challenges). Regarding the application of different big data techniques to the specific scenarios of smart cities presented in this paper, we have overcame the challenge of *velocity* by collecting hourly data in the smart building application (consumption of energy, outdoor environmental conditions) and even in sorter intervals of time for the public transport application (many people validates their transit cards within seconds). Although we have not tackled *volatility*, it is clearly a goal when looking for the real-time smart city because behavioral scenarios like ours change depending on many social aspects. The *veracity* of the data is guaranteed by the exhaustive preprocessing steps included in the modeling process. We have extracted *value*, making sense of the wide mentioned *variety* of data, and with the described analysis and techniques, we have *validated*

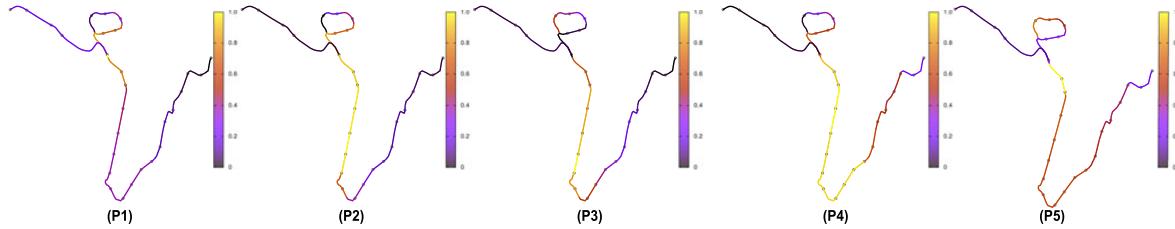


Fig. 9. Tram-line heatmap of the five profiles for Monday–Thursday trips.

TABLE II
MAIN FEATURES OF THE TWO ARCHITECTURE INSTANTIATIONS

Smart City Application	Data	Information	Knowledge	Services
Smart Campus	IR Sensors, RFID tags, Environmental Sensors, Weather Station, Presence Sensors, Energy Consumption Meters, Weather Forecast	Data Transformation through SAREF ontology [9], DUL ontology [10] and SSN ontology [11]	Data Modeling, Predictive Regression (RBFs, SVM, ANN, RF, ARIMA), Tracking algorithm (PFs), Optimization Mechanism (GA)	Indoor localization, Building energy consumption prediction, Energy saving through the HVAC operation optimization
Public Tram Service	Mobile Sensors, Smart Cards	CEP-based filtering, Event Processing in Action [12]	Fuzzy Clustering	Infrastructure monitoring, Mobility patterns.

their usability for solving different problems of smart cities with high accuracy.

In both applications tackled in this paper, the huge *volume* of historical data is being stored using a NoSQL database. At the moment, the storage system is been adapted so as to be compliant with the FI-WARE architecture,² which intends to ease the development of novel applications based on the Future Internet. In particular, the Orion Context Broker³ and the COMET⁴ modules are used in order to store in a NoSQL repository the historical data comprising the measurements from the different data sources.

On the whole, both instantiations of the architecture described above are summarized in Table II. In the next sections, we summarize the main benefits obtained after applying the most suitable big data techniques to the two scenarios of smart cities addressed in this paper.

A. Benefits of Big Data Applications in Smart Buildings for Energy Efficiency

Here, we summarize the main findings extracted from all the experiments and analysis carried out during the application of big data techniques to the smart campus of the UMU.

- 1) *The resolution of the indoor localization problem:* Applying regression techniques based on RBFs and a tracking algorithm applying PFs to data coming from RFID and IR sensors installed in buildings, it was possible to solve the indoor localization problem with a mean accuracy of 1.5 m. Then, indoor localization data can be used to provide customized services in buildings.

²<https://www.fiware.org> [Available Feb. 2016]

³<http://catalogue.fiware.org/enablers/publishsubscribe-context-broker-orion-context-broker> [Available Feb. 2016]

⁴<https://github.com/telefonicaid/fiware-sth-comet>. [Available Feb. 2016]

2) *The resolution of the building energy consumption estimation:* Applying PCA and BRNN techniques to data related to outdoor environmental conditions and energy consumption of buildings, it was possible to generate energy consumption predictive models of buildings with a very low mean error of 43.76 kWh—which only represents the 10.29% of the sample—in the worst case. Then, energy consumption predictions can be used to design the optimal strategies to save energy in buildings.

3) *The resolution of the optimization problem related to the maximization of thermal comfort and minimization of energy consumption in buildings:* Applying optimization methods based on GAs to optimize the energy consumption of buildings, meanwhile, comfort conditions are satisfied, and after including user localization data and user comfort preference prediction, it was possible to get energy savings in heating of about 23% compared with the energy consumption in a previous month without any energy BMS.

B. Benefits of Big Data Applications in Urban Pattern Recognition to Improve Public Tram Service

After applying big data techniques to the urban pattern extraction in the public tram service, all the results from the experiments allowed the service staff to draw up quite interesting conclusions. These are summarized below:

- 1) *Regarding the resolution of the trip extraction:* The formal discovery of the stations' load in terms of trips' origin and destination would allow the service provider and the city council to better plan the whole public transport service of the city. In this way, the more important stations might be considered as “hub” points where commuters can easily transfer from tram to other kinds of transport. Moreover,

such an information could be also useful so as to forecast future infrastructure needs in each part of the tram line (e.g., location and number of places of new parking lots for bicycles close to tram stations).

- 2) *Concerning the resolution of the urban profiles generation:* Experiments pointed out the importance of undergraduates as tram users. Hence, most of the traffic load concentrated in the line between the city center and the campuses. This was really helpful in order to design promotional campaigns for these type of travelers. Moreover, results also confirmed that the line segment toward the shopping mall areas was underused. Thus, campaigns to promote the use of the tram to go shopping was also considered.

VI. CONCLUSION AND FUTURE WORK

This paper displays the benefits of applying big data techniques over data originated by IoT-based devices deployed in smart cities. A general architecture modeled in four layers is proposed to be applied in smart city applications considering big data issues. As part of this overview, a differentiation between static and mobile data sources is made, proposing suitable techniques for them to extract relevant knowledge from their data. Then, we describe two big data applications for smart city services. Specifically, the services of energy efficiency and comfort management in the buildings of a smart campus, and the public transport service of a city. In the first scenario of smart city, we have demonstrated that after applying appropriate big data techniques to problems like indoor localization, energy consumption modeling, and optimization, we are able to provide mean energy savings of 23% per month, while indoor comfort is ensured. Regarding to the urban pattern recognition carried out using data related to the public tram service of the city of Murcia, experiments were performed to confirm that the proposed patterns ended up being of great interest for the service provider in order to better understand how travelers make use of the transportation system. This was fairly useful in order to come up with better planning protocols and more tempting promotional campaigns.

The ongoing work is focused on the inclusion of people behavior during the operational loop of this kind of systems for smart cities. Thus, for the case of smart building applications, users will be encouraged to participate in an active way through their engagement to save energy. On the other hand, in the case of the public tram service, data coming from crowdsensing initiatives will be integrated to improve the estimation of the urban mobility patterns.

REFERENCES

- [1] N. Komninos, "Intelligent cities: Variable geometries of spatial intelligence," *Intell. Build. Int.*, vol. 3, no. 3, pp. 172–188, 2011.
- [2] L. Atzori, A. Iera, and G. Morabito, "The Internet of Things: A survey," *Comput. Netw.*, vol. 54, no. 15, pp. 2787–2805, 2010.
- [3] L. Da Xu, W. He, and S. Li, "Internet of Things in industries: A survey," *IEEE Trans. Ind. Informat.*, vol. 10, no. 4, pp. 2233–2243, Nov. 2014.
- [4] R. Iqbal, F. Doctor, B. More, S. Mahmud, and U. Yousuf, "Big data analytics: Computational intelligence techniques and application areas," *Int. J. Inf. Manage.*, pp. 10–15, in press doi: <http://dx.doi.org/10.1016/j.ijinfomgt.2016.05.020>.
- [5] Z. Yan and D. Chakraborty, *Semantics in Mobile Sensing (Synthesis Lectures on the Semantic Web: Theory and Technology)*, vol. 4, 1st ed. San Rafael, CA, USA: Morgan Claypool, 2014, pp. 1–143.
- [6] A. Carroll and G. Heiser, "An analysis of power consumption in a smartphone," in *USENIX Annu. Tech. Conf.*, 2010, pp. 1–14.
- [7] M. Kantardzic, *Data Mining: Concepts, Models, Methods, and Algorithms*. Hoboken, NJ, USA: Wiley, 2011.
- [8] L. Daniele, F. den Hartog, and J. Roes, "Created in close interaction with the industry: The smart appliances reference (SAREF) ontology," in *7th Int. Formal Ontologies Meet Ind.*, 2015, pp. 100–112.
- [9] K. Janowicz and M. Compton, "The stimulus-sensor-observation ontology design pattern and its integration into the semantic sensor network ontology," in *Proc. 3rd Int. Conf. Semantic Sensor Netw.*, 2010, vol. 668, pp. 64–78. [Online]. Available: CEUR-WS.org
- [10] M. Compton *et al.*, "The SSN ontology of the W3C semantic sensor network incubator group," *Web Semantics: Sci., Serv. Agents World Wide Web*, vol. 17, pp. 25–32, 2012.
- [11] O. Etzion and P. Niblett, *Event Processing in Action*, 1st ed. Greenwich, CT, USA: Manning Publ. Co., 2010.
- [12] T. Maniak, C. Jayne, R. Iqbal, and F. Doctor, "Automated intelligent system for sound signalling device quality assurance," *Inf. Sci.*, vol. 294, pp. 600–611, 2015.
- [13] A. H. Neto and F. A. S. Fiorelli, "Comparison between detailed model simulation and artificial neural network for forecasting building energy consumption," *Energy Build.*, vol. 40, no. 12, pp. 2169–2176, 2008.
- [14] H.-X. Zhao and F. Magoulès, "A review on the prediction of building energy consumption," *Renew. Sustainable Energy Rev.*, vol. 16, no. 6, pp. 3586–3592, 2012.
- [15] B. B. Ekici and U. T. Aksoy, "Prediction of building energy consumption by using artificial neural networks," *Adv. Eng. Softw.*, vol. 40, no. 5, pp. 356–362, 2009.
- [16] F. Ascione, N. Bianco, C. De Stasio, G. M. Mauro, and G. P. Vanoli, "Simulation-based model predictive control by the multi-objective optimization of building energy performance and thermal comfort," *Energy Build.*, vol. 111, pp. 131–144, 2016.
- [17] W. Wang, J. P. Attanucci, and N. H. Wilson, "Bus passenger origin-destination estimation and related analyses using automated data collection systems," *J. Public Transp.*, vol. 14, no. 4, pp. 131–150, 2011.
- [18] M.-P. Pelletier, M. Trépanier, and C. Morency, "Smart card data use in public transit: A literature review," *Transp. Res. Part C: Emerging Technol.*, vol. 19, no. 4, pp. 557–568, 2011.
- [19] N. Yuan, Y. Wang, F. Zhang, X. Xie, and G. Sun, "Reconstructing individual mobility from smart card transactions: A space alignment approach," in *2013 IEEE 13th Int. Data Mining*, Dec. 2013, pp. 877–886.
- [20] M. A. Munizaga and C. Palma, "Estimation of a disaggregate multimodal public transport origin-destination matrix from passive smartcard data from Santiago, Chile," *Transp. Res. Part C: Emerging Technol.*, vol. 24, pp. 9–18, 2012.
- [21] J. P. Attanucci and N. H. Wilson, "Bus passenger origin-destination estimation and related analyses using automated data collection systems," *J. Public Transp.*, vol. 14, no. 4, pp. 131–150, 2011.
- [22] P. Palensky and D. Dietrich, "Demand side management: Demand response, intelligent energy systems, and smart loads," *IEEE Trans. Ind. Informat.*, vol. 7, no. 3, pp. 381–388, Sep. 2011.
- [23] S. Fang *et al.*, "An integrated system for regional environmental monitoring and management based on internet of things," *IEEE Trans. Ind. Informat.*, vol. 10, no. 2, pp. 1596–1605, May 2014.
- [24] EU Smartie Consortium, EU Smartie Project, 2013–2016. [Online]. Available: <http://www.smartie-project.eu>
- [25] EU Entropy Consortium, EU Entropy Project, 2015–2018. [Online]. Available: <http://entropy-project.eu/>
- [26] A. Haug, *A tutorial on Bayesian estimation and tracking techniques applicable to nonlinear and non-Gaussian processes*. MITRE Corporation, McLean, VA, USA, 2005.
- [27] *Indoor Environmental Input Parameters for Design and Assessment of Energy Performance of Buildings Addressing Indoor Air Quality, Thermal Environment, Lighting and Acoustics*, EN Standard 15251, 2007.
- [28] H. Abdi and L. J. Williams, "Principal component analysis," *Wiley Interdisciplinary Rev.: Comput. Statist.*, vol. 2, no. 4, pp. 433–459, 2010.
- [29] L. Hawarah, S. Ploix, and M. Jacomino, "User Behavior Prediction in Energy Consumption in Housing Using Bayesian Networks," in *Artificial Intelligence and Soft Computing*. Berlin, Germany: Springer, 2010, pp. 372–379.

- [30] Y. Fu, Z. Li, H. Zhang, and P. Xu, "Using support vector machine to predict next day electricity load of public buildings with sub-metering devices," *Procedia Eng.*, vol. 121, pp. 1016–1022, 2015.
- [31] M. Alamaniotis, D. Bargiolas, and L. H. Tsoukalas, "Towards smart energy systems: Application of kernel machine regression for medium term electricity load forecasting," *SpringerPlus*, vol. 5, no. 1, pp. 1–15, 2016.
- [32] R Core Team, *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing, 2015. [Online]. Available: <http://www.R-project.org/>
- [33] M. Kuhn, "Building predictive models in R using the caret package," *J. Stat. Softw.*, vol. 28, no. 5, pp. 1–26, 2008.
- [34] A. González-Vidal, V. Moreno-Cano, F. Terroso-Sáenz, and A. F. Skarmeta, "Towards energy efficiency smart buildings models based on intelligent data analytics," *Procedia Comput. Sci.*, vol. 83, pp. 994–999, 2016.
- [35] J. A. Orosa, "A new modelling methodology to control HVAC systems," *Expert Syst. Appl.*, vol. 38, no. 4, pp. 4505–4513, 2011.
- [36] E. Willighagen, "Genalg: R based genetic algorithm," *R package v0.1*, vol. 1, 2005.
- [37] M. V. Moreno, M. Zamora-Izquierdo, J. Santa, and A. F. Skarmeta, "An indoor localization system based on artificial neural networks and particle filters applied to intelligent buildings," *Neurocomputing*, vol. 122, pp. 116–125, 2013.
- [38] J. M. Andy Field and Z. F. Niblett, *Discovering Statistics Using R*, 1st ed. Newbury Park, CA, USA: Sage, 2012.
- [39] J. Demšar, "Statistical comparisons of classifiers over multiple data sets," *J. Mach. Learn. Res.*, vol. 7, pp. 1–30, 2006.
- [40] (2016). [Online]. Available: www.tranviademurcia.es
- [41] S. Miyamoto, H. Ichihashi, and K. Honda, "Algorithms for fuzzy clustering," in *Methods in c-Means Clustering With Applications*, J. Kacprzyk, Ed. Berlin, Germany: Springer-Verlag, 2008.
- [42] M. Ali-ud-din Khan, M. F. Uddin, and N. Gupta, "Seven V's of big data understanding big data to extract value," in *2014 Zone 1 Conf. Am. Soc. Eng. Educ.*, 2014, pp. 1–5.



M. Victoria Moreno received the B.S. (Hons.) and M.S. degrees in telecommunications engineering both from the School of Telecommunication Engineering, Cartagena, Spain, in 2006 and 2009, respectively, and the Ph.D degree in computer science from the University of Murcia, Murcia, Spain, in 2014.

She is currently a Postdoctoral Researcher in the Seneca Foundation, Murcia (Spain). Her research interests include data analysis and modeling, and energy efficiency in smart environments.



Fernando Terroso-Sáenz received the Graduate and Master's degrees in computer science from the University of Murcia, Murcia, Spain, in 2006 and 2010, respectively, and the Ph.D. degree in computer science from the Department of Communications and Information Engineering, University of Murcia, in 2013.

Since 2009, he has been a Researcher in Department of Communications and Information Engineering, University of Murcia. His research interests include complex event processing, ubiquitous computing and intelligent transportation systems.



Aurora González-Vidal received the Graduate degree in mathematics from the University of Murcia, Murcia, Spain, in 2014. Since 2015, she has been working toward the Ph.D. degree in computer science.

In 2015, she got a fellowship to work in the Statistical Division of the Research Support Services, where she specialized in statistics and data analysis. Her research interest includes data analytics for energy efficiency.



Mercedes Valdés-Vela received the Computer Engineering degree and the Ph.D. degree (Hons.) in computer science from the University of Murcia, Murcia, Spain, in 1998 and 2003, respectively.

In 2000, she started to work as Research Staff in the Department of Information and Communications Engineering, University of Murcia, where she is currently a Full Time Assistant Professor. Her main research interests include soft computing, complex event processing, and ambient intelligence.



Antonio F. Skarmeta received the B.S. (Hons.) degree in computer science from the University of Murcia, Murcia, Spain, the M.S. degree in computer science from the University of Granada, Granada, Spain, and the Ph.D. degree in computer science from the University of Murcia.

He is currently a Full Professor in the Department of Information and Communications Engineering, University of Murcia. He is involved in numerous projects, both European and National. His research interests include mobile communications, artificial intelligence, and home automation.



Miguel A. Zamora received the M.S. degree in automation and electronics and the Ph.D. degree in computer science both from the University of Murcia, Murcia, Spain, in 1997 and 2003, respectively.

He is currently a Senior Professor in the Department of Information and Communication Engineering, University of Murcia. His research interests include consumer electronics, home and building automation, and sensor fusion.



Victor Chang received the Ph.D. degree in computer science from the University of Southampton, Southampton, U.K., in 2013.

Since June 2016, he has been an Associate Professor (Reader) with Xi'an Jiaotong Liverpool University, Suzhou, China. He helps organizations in achieving good Cloud design, deployment, and services. He is one of the most active practitioners and researchers in Cloud computing, big data, and Internet of Things in the U.K. Dr. Chang has received the European Award on Cloud Migration in 2011, Best Papers in 2012 and 2015, and numerous awards since 2012.

4.5 An open IoT platform for the management and analysis of energy data

Title	An open IoT platform for the management and analysis of energy data
Authors	Fernando Terroso-Saenz, Aurora González-Vidal, Alfonso P. Ramallo-González and Antonio F. Skarmeta
Type	Journal
Journal	Future Generation Computer Systems
Impact factor (2018)	5.768
Rank	Q1
Publisher	ELSEVIER
Volume	92
Pages	1066 - 1079
Year	2019
Month	March
ISNN	0167-739X
DOI	10.1016/j.future.2017.08.046
URL	https://www.sciencedirect.com/science/article/pii/S0167739X17304181?via%3Dihub
State	Published
Author's contribution	The PhD student, Aurora González Vidal, contributed to the methodology, research and writing the publication



An open IoT platform for the management and analysis of energy data

Fernando Terroso-Saenz ^{*}, Aurora González-Vidal, Alfonso P. Ramallo-González,
Antonio F. Skarmeta



Department of Information and Communications Engineering, Computer Science Faculty, University of Murcia, Spain

HIGHLIGHTS

- IoT platform for the management of energy data in buildings.
- Includes several inner features to support data analytics in the energy domain.
- Based on the open IoT initiative FIWARE.
- Evaluated in a real pilot with comprising several buildings.

ARTICLE INFO

Article history:

Received 15 March 2017
Received in revised form 24 May 2017
Accepted 23 August 2017
Available online 6 September 2017

Keywords:

IoT platform
Energy consumption
FIWARE
Data mining

ABSTRACT

Buildings are key players when looking at end-use energy demand. It is for this reason that during the last few years, the Internet of Things (IoT) has been considered as a tool that could bring great opportunities for energy reduction via the accurate monitoring and control of a large variety of energy-related agents in buildings. However, there is a lack of IoT platforms specifically oriented towards the proper processing, management and analysis of such large and diverse data. In this context, we put forward in this paper the IoT Energy Platform (IoTEP) which attempts to provide the first holistic solution for the management of IoT energy data. The platform we show here (that has been based on FIWARE) is suitable to include several functionalities and features that are key when dealing with energy quality insurance and support for data analytics. As part of this work, we have tested the platform IoTEP with a real use case that includes data and information from three buildings totaling hundreds of sensors. The platform has exceeded expectations proving robust, plastic and versatile for the application at hand.

© 2017 Elsevier B.V. All rights reserved.

1. Introduction

Several reports claim that residential and commercial buildings represent around 30%–40% of the overall energy consumption in Europe and in the United States [1,2]. Because of this, buildings are known to be the largest end-use energy contributor followed by transport and industry, and therefore they are a clear target for potentially reducing global energy consumption substantially.

Despite being great consumers, there is some evidence that shows that public and private buildings have not fully exploited all opportunities available to increase their energy efficiency. On the contrary, they suffer from a rather substantial energy waste that is partly due to inefficient heating, cooling, lighting and other power system (equipment) [3], due to bad use of the systems (behavior) [4] and due to poor fabric efficiency [5]. Although the implementations of measurements to improve the first or the third category can be rather expensive, it has been seen that soft

measurements that focus on the change of behavior of buildings' users are cheap, but yet, can contribute greatly to the reduction of energy use [6].

In order to address the aforementioned inefficiencies due to lack of understanding on how the systems should be operated and other behavioral related aspects in the building sector, one could consider the use of Information and Communication Technologies (ICT) and, more specifically, of the Internet of Things (IoT). This new paradigm that also exists at the domestic level could be used as an instrument to make a realization of the so called *Smart Building*. In fact, it is foreseen that from 2 to 3 houses out of 10 will be equipped with up to 500 smart devices in the near future [7].

The installation of smart meters and In Home Energy Displays to make households aware of their energy consumption is not new [8,9]. The adoption of these devices seems to be an opportunity to exploit them for the reduction of energy use when looking at the available scientific literature (will be detailed later). However, one may also think that the technological effort to deploy such systems may be substantial and become a barrier to achieve this level of technification of the buildings. Nevertheless this technification seems to be happening naturally.

* Corresponding author.

E-mail address: fterroso@um.es (F. Terroso-Saenz).

The large amounts of IoT data that will be coming from buildings in the near expected future will have to be analyzed to reveal insights that could help to obtain, expose and understand knowledge from buildings. In turn, this derived knowledge should be able to help to achieve meaningful energy saving strategies and interventions in the targeted buildings [10].

These wealth of information about energy use, offers a great opportunity according to some literature on energy feedback that suggests that intelligent feedback, (that with an extra larger of computation over simple observation) is an effective technique for the reduction of energy demands via behavioral change [11]. Only with a platform capable of making this possible, the implementation of this new paradigm will be successful.

In the IoT ecosystem, several platforms have emerged providing support from the sensorization stage to the stage of management and storage of the data in different forms [12]. In that sense, one of the most large-scale affords is the FIWARE platform, a key initiative of the Future Internet Public–Private Partnership (PPP) to create a well-aligned set of open enablers to receive, process, contextualize and publish IoT data from and for smart cities including from city-wide information to dwelling specific data.¹

Despite all the reasons exposed before, little efforts have been made so far in order to adapt such platforms to building energy management. This energy ecosystem comprises a set of particularities that should be targeted in a specific manner. After analyzing the few examples of studies that have tried to tackle this problem, one can see that it exists a pressing need to apply different data mining techniques in the building energy domain mainly focusing on consumption prediction and pattern discovery or failure tolerance [13]. Thus, IoT energy platforms should include functions for data analysis among their features.

Although giving insight knowledge behind data is an instrumental aspect of the wealth produced by the IoT, existing platforms are still limited when it comes to integrate data processing and analytic techniques suitable for IoT ecosystems [14]. This is a fundamental limitation of the state of the art as it is key to ensure that the platform will work on the new paradigm of providing tailored, real-time energy feedback to people. This also includes features to support the easy extension of platforms to allocate new data mining techniques comprising common steps in the data mining process. Examples of such features are built-in data-cleaning mechanisms for data pre-processing and storage solutions that would facilitate the execution of online and offline data mining algorithms.

All the aforementioned limitations have motivated us to envision, design, develop and validate what we called the IoT Energy Platform (IoTEP). The key strength of IoTEP is that it is, to our knowledge, the first holistic solution to large scale building energy data management from IoT.

Unlike existing IoT platforms, IoTEP is mainly oriented to support and ease the analysis of large amounts of heterogeneous energy data. A simplified overview of the platform IoTEP is shown in Fig. 1 representing its key features.

To begin with, IoTEP has been designed to easily retrieve either the most up-to-date readings of each sensor within a building, or to retrieve the historic data from such sensors. By means of these two types of access, the platform facilitates the application of both online and offline data analyses over the collected data. As we will see on further sections, this functionality is implemented with two FIWARE storage components, the ORION context broker and COMET. For both enablers, a NGSI-based information model has been defined in order to homogenize all the measured energy-related data.

Secondly, a real-time data cleaning module has been designed as a built-in component of IoTEP. With this, sensor readings are filtered by discarding potential outliers before injecting them in the storage components. This ensures a more efficient use of the resources. For this feature, we have followed a Complex Event Processing (CEP) approach that allows the real-time processing of event streams.

In addition to the above mentioned features, the platform includes also a mechanism to detect volatility changes in the incoming energy data. This mechanism intends to perceive meaningful shifts in such data that might need to re-launch the data-mining services that run within the platform.

Finally, IoTEP features a novel mechanism to automatically identify high-level areas in a building with certain energy-related similarities by means of clustering techniques. The benefit of these virtual areas is twofold. Firstly, they provide alternative representations of the energy status of a building beyond its physical structure; and secondly, they can help in the performance of other data mining analyses by reducing redundancies and defining different granularity levels in the captured sensor data.

Summarizing, the platform presented in this paper intends to be the first stage towards the full adaptation of the IoT paradigm in the retrieval, management and, above all, analysis of energy data in buildings. Considering the need of developing tools that are able to provide personalized real time feedback to change behaviors, and with them, have the potential to reduce energy use, IoTEP is intended to become the stepping stone for the development of such tools.

The paper is structured as it follows: Section 2 provides an overview of the state of the art in this research area. Section 3 looks into the IoT energy platform, including its architecture and its functional modules. Section 4 provides an evaluation of some of the features of the platform; and Section 5 concludes the paper with some final remarks and conclusions.

2. Related work

The present work is based upon two different lines of research, the management of energy data and the implementation of IoT platforms. Consequently, an overview of both lines is put forward in this section.

2.1. Energy data management systems

During the last years, some initiatives within the cloud computing domain have been made to intelligently manage energy data of buildings. In that sense, Zhou et al. [15] described a model for big-data energy management ranging from the collection and pre-processing of data to its further analysis and the final exposition to services. However, it only provides a theoretical approach.

From a practical perspective, the Dynamic Demand Response (D^2R) platform [16] makes use of public and private clouds combined with infrastructure and platform as a service for data storage. This platform was extended with *Cryptonite*, a repository to store sensitive Smart Grid data [17]. Then, different classes of data-driven forecasting models were generated on top of the whole platform with the purpose of carrying out energy prediction among others.

ElasticStream also provides a prototype solution for energy data management and analysis. In this case, the proposed mechanism transfers energy data to a cloud platform for further analysis on the basis of rate changes in the input data streams [18]. Moreover, Vastardis et al. [19] described a centralized architecture to monitor energy consumption in houses including features of pattern-matching related to the behavioral habits of the target users.

In the work of Ozadowicz [20], the authors propose different approaches to calculate the power demand related to energy

¹ <https://catalogue.fiware.org/>.

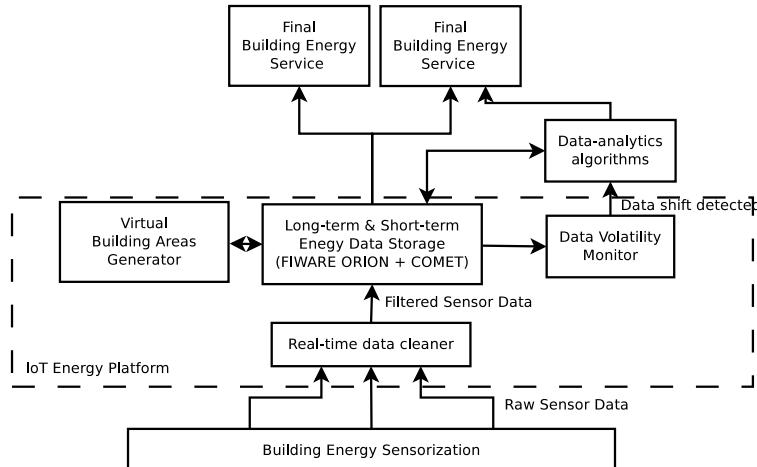


Fig. 1. Conceptual view of the IoT Energy Platform (IoTEP).

consumption using time-driven and event-driven mechanisms for Building Automation and Control Systems. Their Building Energy Management Systems (BEMS) implementation is realized with an IoT platform, introduced by Echelon Corp that includes chips, stacks, communication, application interfaces (API) and management software. Their approaches to calculate the energy demand are based in time (fixed or sliding length of the time windows with the possibility of overlapping) and in events (occupancy).

The MultiAgent System (MAS) named SAVES (Sustainable multiAgent systems for optimizing Variable objectives including Energy and Satisfaction) defined in [21] is used in [22] regarding actual occupant preferences and schedules, actual energy consumption and loss data measured from a real test bed building at the University of Southern California in order to predict energy consumption at different levels (frequency of prediction and device aggregation).

Other works provide energy data management solutions without focusing on analytic aspects. This is the case of the Virtual SCADA architecture for cloud computing (VS-Cloud) that encompasses Cloud Computing for energy data storage [23]. VS-Cloud mainly focuses on the orchestration of components in Smart Grids and the safety storage of sensitive data executed actions, incidents or alarms. Therefore, its domain of application is more related to risk management.

Similarly, the work in [24] proposes an automation platform for energy monitoring. However, such platform does not provide any particular feature to support energy data analytics as it focuses more on the definition of control strategies for energy saving.

Unlike the aforementioned initiatives, our work provides a holistic energy data management and analysis solution. Our platform also follows an open approach by relying on the well-established FIWARE initiative. In that sense, the present work includes explicit features like data volatility monitoring and outliers detection to ease the deployment of data mining algorithms and other services over of the stored data.

FIWARE brings other advantages with respect to previous solutions: firstly, the whole platform orchestration is done by means of lightweight RESTful APIs, that facilitate its further extension; and secondly, the definition of an information model compliant with NGSI standard allows to come up with a homogeneous view of the energy-related data within a building. This feature is key to exploit the potential of gathering energy data. What we propose here is not only an archive of data, but a comprehensive flexible and powerful tool that will serve as the breeding ground for the

creation of context-aware tailored energy feedback platforms that could be realized at a scale never considered before, even reaching national levels.

2.2. IoT platforms

The Internet of Things paradigm is the second pillar of this initiative. All the literature indicates that small devices connected to the internet in buildings will be the norm in the near future. With the right algorithms and communication mechanisms, this situation will enable the monitoring and characterization of energy behaviors and energy consumption in buildings.

The need of effective instantiation of IoT under realistic conditions has generated a varied ecosystem of methodologies and tools taking the form of integrated IoT platforms. In that sense, it is possible to find several surveys in the literature that review existing proprietary and open-source platforms in the IoT ecosystem [12,14,25]. Other important aspects like data ownership, security and privacy [26] or data storage [25] have been also deeply studied in the IoT domain. The reader is referred to this sources to expand on the state of the art.

According to such reviews, some relevant IoT platforms follow a similar open-source and centralized approach along with heterogeneous sensor support like IoTEP. This is the case of Nimbitts² that provides an open source Java library for developing Java, Web and Android solutions to connect to a Nimbitts Server. This back-end part enables simple processing of the collected data based on rules. However, it does not comprise any advanced data-analytics support. ThingSpeak³ features the acquisition, visualization and analysis of data but this is done by means of the proprietary Matlab tool, what may make more difficult the popularization of the platform.

One feature frequently neglected by existing IoT platforms is the support of built-in data mining features able to generate new useful knowledge from the collected and stored data [14]. In real IoT deployments, this processing and analysis task has been frequently done by third-party services. However, integrating certain data mining functionalities as built-in features of platforms would provide a great benefit in a wide range of domains, for example: quick statistics, easy to generate digests or sanity checks. In

² <https://www.nimbitts.com/>.

³ <https://thingspeak.com/>.

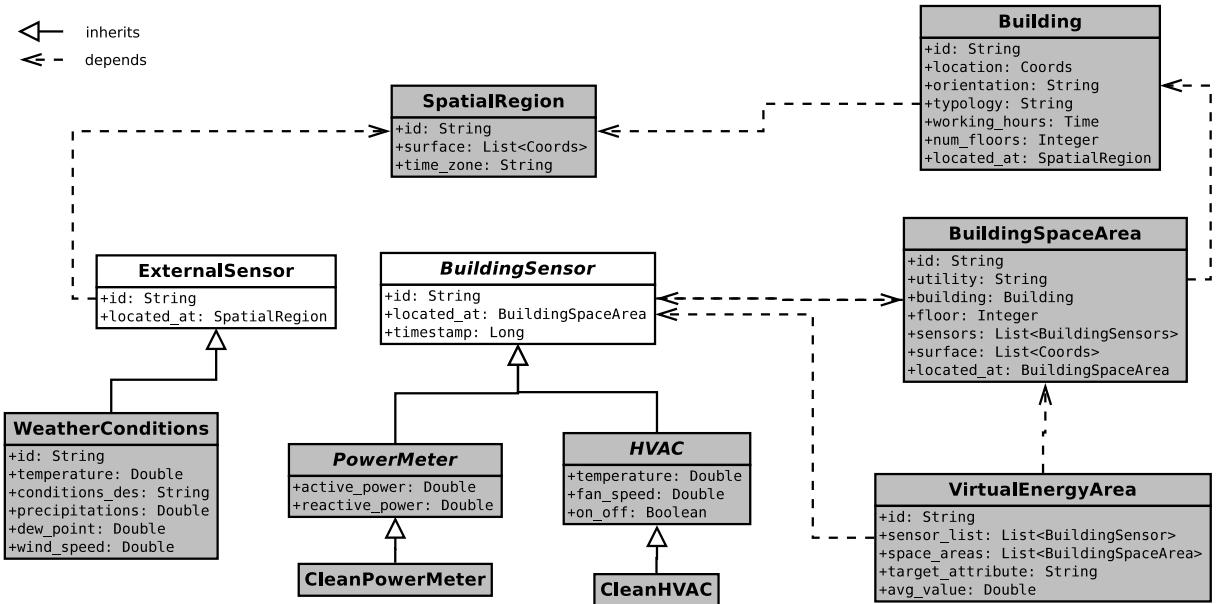


Fig. 2. IoTEP information model.

that sense, only a few IoT platforms actually include native data-analytics features. As a matter of fact, SensorCloud⁴ enables a simple interface for common operations like smoothing, filtering and interpolation whereas GroveStreams⁵ provides some real-time data analytics mechanisms. However, none of them support sensor heterogeneity nor follow an open source approach like IoTEP does.

As for the energy ecosystem, several research lines have already stated the feasibility and suitability of data analysis in order to increase energy awareness within a building [13]. In that sense, IoTEP provides one of the first steps towards such a data-mining enrichment by providing several features fully focused on easing the analysis of IoT energy data namely, real-time data cleaning, data volatility detection and data reduction procedures.

Finally, our work is enclosed within the FIWARE architecture. The high-level goal of this architecture is to build the Core Platform of the Future Internet, introducing an innovative infrastructure for cost-effective creation and delivery of versatile digital services, providing high QoS and security guarantees. In that sense, FILAB [27] conforms live instances of generic enablers, available to developers for free experimentation within this technology.

Some initiatives have started to profit from FIWARE in several domains. One of the most ambitious works is the application on [28] which established a world-wide semantic interoperability solution combining the NGSI, which is part of the core of the FIWARE initiative, and oneM2M context interfaces. Apart from that, [29] demonstrated the suitability of the FIWARE paradigm to compose Future-Internet applications by means of the integration of generic enablers. In a similar manner, [30] put forward a semantic mechanism to integrate data from different types of devices by also using FIWARE components. Finally, in a more functional domain, [31] made use of certain enablers, like ORION context broker, to create a cloud-based gesture recognition application. Also, [32] describes a sensor management for seaports based on the FIWARE platform. It is therefore possible to say that our work

seems to be one of the first efforts to make use of FIWARE enablers in the building energy domain, and furthermore in the energy domain in general.

3. IoT Energy Platform (IoTEP)

This section explains in detail the proposed IoTEP solution. Since the management of the energy data is its key feature, we firstly describe the information model used to define all the data within the IoTEP ecosystem; next, we put forward the specific architecture of the platform that deals with the energy data according to the model.

3.1. Information model

One of the first steps towards the realization of IoTEP was to define a common information model for the whole platform. Such a model must be compliant with the NGSI information model commonly accepted in the FIWARE ecosystem, what facilitates interconnection with other models and other users. This information model follows an entity–attribute approach where entities represent real or virtual elements of interest. Each entity has a type what allows to define type-based hierarchies. In this way, an entity has its own defined attributes and the inherited ones from its ancestors. The IoTEP information model is depicted in Fig. 2. The model design follows the UML class notation with two types of relationships, inheritance and dependence. Each of them is represented by a different arrow in the figure. Whilst inheritance indicates that the child element comprises all the attributes of its parent element, the dependence relationship indicates that an instance of the element at the arrow's origin contains an attribute referencing at one or more instances of the element at the arrow's destination.

Focusing on the content of the model, one can find among its components three key elements related to the energy ecosystem of a building by means of NGSI entities.

To begin with, the entity **building** models the target building. Several operational and architectonic details of the building are included as attributes on this entity. Examples of information in

⁴ <http://www.sensorcloud.com/>.

⁵ <https://grovestreams.com/>.

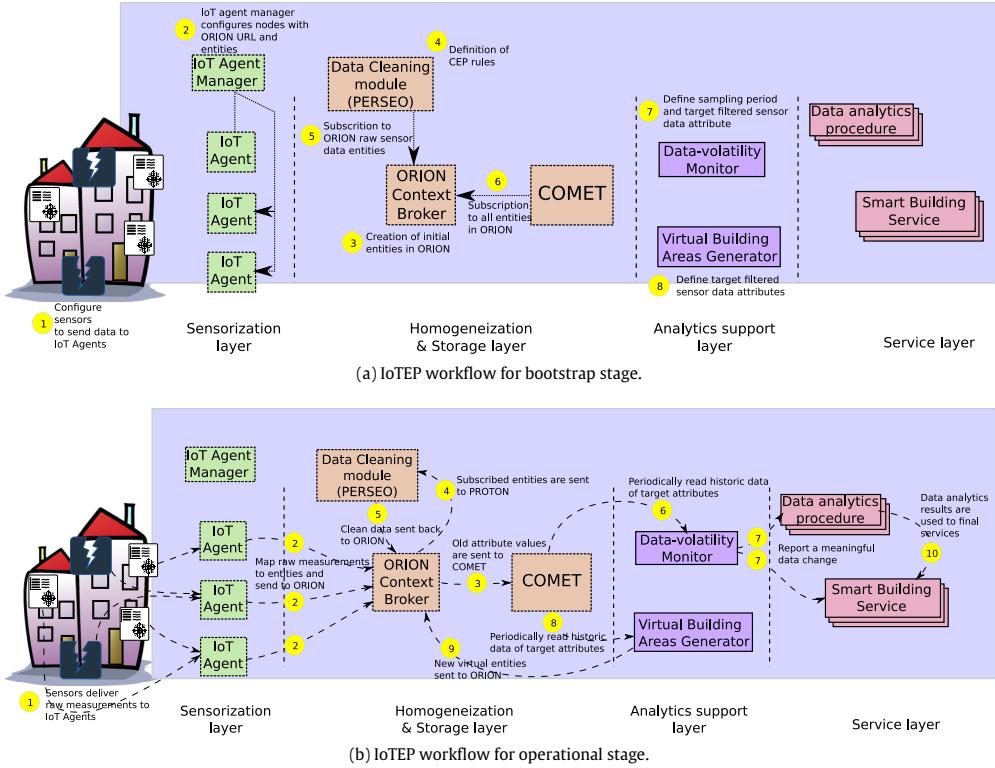


Fig. 3. Platform general workflow.

this section are: opening hours or building use (e.g., company headquarters, university faculty, etc.) but also physical relevant attributes such as fabrics, windows, orientation, and so forth. Moreover, the *spatial region* entity defines the geographic region containing the building. This entity would help to link together buildings located in similar geographic regions that, as a consequence, might share certain energy-related characteristics. The inner structure of a building is represented with the *building space area* entity. This entity gathers the different spatial areas within a building (e.g., classrooms, corridors, halls, landings, etc.). Furthermore, a recursive structure of these areas can be made with their *located at* attribute to represent, for example, that a classroom is inside a *teaching zone*.

This way of introducing data about the buildings and the spaces will make the communication between a Building Information Modeling (BIM) platforms and the IoTEP platform straight forward, what would facilitate the transfer of information among members of a given team.

The second group of entities refers to the energy sensors deployed in the building and the data they collect. This is modeled by means of the *building sensor*, *power meter* and *hvac* entities. Each entity includes the set of attributes monitored by the corresponding energy sensor along with other metadata (e.g., location of the sensor or timestamp of each observation). The *clean* version of these entities refer to the sensor data generated after the data filtering process as described in Section 3.2.2.

The third group of entities focus on representing sensors that are not necessarily within the infrastructure of the building but that may provide useful when collecting energy data. This is the case, for example, of weather stations reporting conditions of the building site. As Fig. 2 shows, this is defined by means of the *external sensor* and *weather conditions* entities.

Finally, only the entities in gray in Fig. 2 have instances stored in ORION and COMET as we will see later.

3.2. Platform architecture

The proposed IoTEP has been structured in four different layers in an incremental approach (this is shown in Fig. 3). In the upcoming sections, a detailed description of each layer is given.

3.2.1. Sensorization layer

This layer is in charge of connecting physical devices or actuators that are going to provide energy data to the platform. Once this is done, it maps the collected data to the NGSI entities of the information model (described in the previous section) and sends the mapped information to the upper homogenization and storage layer.

For the realization of this layer, we have made use of the FIWARE IoT Agent enabler [33]. In a nutshell, this enabler allows to automatically perform the aforementioned data mapping. Different types of this enabler support transport protocols to connect to the physical devices like MQTT,⁶ or Lightweight M2M (LwM2M)⁷.

Consequently, during the bootstrapping phase of the platform, a set of IoT Agents are configured with the NGSI entity type associated to each of its associated sensor by means of the IoT Agent Manager (see Fig. 3(a)). In particular, power meters deployed in the target building are mapped to the *power meter* entity type whereas HVAC devices are mapped to the *hvac* one. Furthermore, we developed an ad-hoc agent to parse the weather conditions coming

⁶ <http://mqtt.org>.

⁷ <http://openmobilealliance.org/iot/lightweight-m2m-lwm2m/>.

from an external third-party weather service to the *weather conditions* entity on a regular basis. During the operational phase (see Fig. 3(b)) each time an IoT Agent receives the raw measurements from a physical device, it *inflates* the entity instance associated to the device in upper layer by means of a RESTfull API, in the homogenization and storage layer (will be described in the next section).

3.2.2. Homogenization and storage layer

In this layer, all the collected energy data from the previous layer is conveniently stored in a uniform solution. This way, this layer addresses the heterogeneity of the incoming energy-related data. Moreover, it contains real time data cleaning stage what ensures the quality of the data collected.

Sensor data repository. Regarding the energy-related data storage, this has been achieved by integrating two FIWARE components.

Firstly, ORION context broker [34] implements a publish-subscribe store providing data access by means of the NGSI-10 API [35]. In IoTEP, this enabler stores the entity instances of the information model. By means of the NGSI update operation, IoT Agents in the sensorization layer update the sensor entities' attributes in real time with the new readings from the devices.

Secondly, the COMET enabler [36] is used for supporting access to historic time series data extending the ORION functionality. In that sense, COMET adheres to the same information model, thus, it does not require any further data harmonization process. It incorporates an ad-hoc API to retrieve raw historical sensor data along with several built-in simple aggregation functions over such data (e.g., provide the sum, min or max of the collected observations for a specific time period).

During the bootstrapping phase of the platform, ORION is initiated with the *static* attributes of the entities in the information model (e.g., 'identifier', 'located at' or 'orientation' attributes) and COMET subscribes in ORION to the *dynamic* attributes of the entities to receive each new value (see Fig. 3(a)).

Sensor data cleaning. Concerning the data quality assurance, we developed a data cleaning module to remove the outliers that might be contained in the raw measurements from the sensors. In that sense, outliers have been reported to be the most prominent quality issue of energy data [37,38].

This module had two key requirements. To begin with, the data cleaning process must be done in a timely manner in order to avoid potential bottlenecks. Furthermore, in an IoT ecosystem we should expect a great variety of data formats and structure. Thus, such data cleaning should be done after data homogenization in order to simplify the overall computational cost of the cleaning stage.

In order to cope with the time-processing constraints, we opted for following the Complex Event Processing (CEP) paradigm to develop a real-time data cleaning module. CEP focuses on timely processing streams of information items, so-called events, by filtering, aggregation or pattern discovery using predefined rules following the event-condition-action paradigm [39]. In the present setting, the incoming events are the readings from the energy sensors, the conditions to be detected are whether a reading should be considered or not an outlier and the action of the final insertion of the cleaned data in the storage structure of the platform.

For the outlier definition, we followed a strategy based on quartiles with fences [40]. In brief, such a strategy extracts the median, the lower Q_1 and upper quartiles Q_3 (aka 25th and 75th percentiles) along with the interquartile range $IQ (= Q_3 - Q_1)$ of the data set under study. On the basis of such statistics, two fences are defined,

- Lower outer fence: $Q_1 - 3 \times IQ$
- Upper outer fence: $Q_3 + 3 \times IQ$

This way, a measurement beyond such fences is considered an *extreme outlier*.

The translation of this strategy to CEP allows to calculate such fences incrementally and update their boundaries each time that a sensor pushes in new data. In particular, two types of CEP rules were defined. The first one comprises the rules in charge of computing for each sensor the aforementioned statistics with respect to each of its parameters. For the sake of clarity, the pseudocode of the CEP rule in charge of calculating the fences for power meter sensors is shown here and it looks as it follows:

```
CONDITION PowerMeter.groupBy(id).within(tintclean) as A
ACTION    new PowerMeterStats(A.id,
           calculateLowerOuterFence(A.active_energy)
           calculateLowerOuterFence(A.reactive_energy),
           calculateUpperOuterFence(A.active_energy)
           calculateUpperOuterFence(A.reactive_energy))
```

where *groupBy* and *within* are two sliding windows. While *groupBy* splits the stream of power-meter data with respect each particular device, *within* defines a time window to retain the last power-meter data generated during the last t_{int}^{clean} time units. After this, the action part of the rule, generates a new power meter stats event comprising the percentiles for each sensor's attribute considering the data included in the time window. It is important to note that this rule would fire each time that new power meter data is injected into the CEP system.

The second set of rules performs the actual extreme outliers detection. Again, there is one rule per sensor type in charge of this task. The pseudocode of the CEP rule to detect the outliers in the power meter data is shown next,

```
CONDITION PowerMeter as A
        AND PowerMeterStats as B
        AND A.id = B.id
        AND A.active_energy ∈ [B.active_energy_lowerFence,
                               B.active_energy_upperFence]
        AND A.reactive_energy ∈ [B.reactive_energy_lowerFence,
                               B.reactive_energy_upperFence]
ACTION    new CleanPowerMeter(A.id, A.timestamp, A.located_at,
                           A.active_energy, A.reactive_energy)
```

Describing it briefly, this rule fires each time that a new power-meter reading is received. The condition part of the rule matches such reading with its associated statistics and checks whether each parameter is contained in its own fences. If that is the case, the reading is considered that has been cleaned. As a result, the action part creates a new clean power meter event with the pre-processed data.

A very similar approach is followed for the HVAC data but, this time, using the thermostat temperature attribute of this type of sensor in order to give rise to clean hvac events.

The implementation of this CEP mechanism has been made with the Perseo FIWARE enabler [41]. This component incorporates a CEP engine and an SQL-based event processing language to define and execute the CEP rules. Furthermore, it leverages the publish-subscription capabilities of ORION. This way, the engine receives each entity instance, which data has been just updated in ORION, as incoming events; and the cleaned events generated by the rules, automatically update their associated entities in ORION (Fig. 3(b)). Hence, during the bootstrapping phase (see Fig. 3(a)) this component is configured with the rules to be executed and the list of entities in ORION to subscribe (in this case, *power meter* and *hvac* entities).

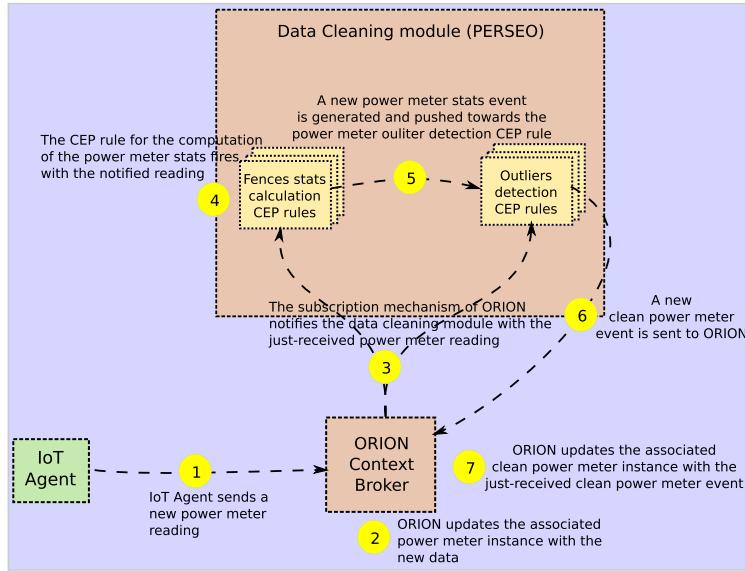


Fig. 4. Workflow of the cleaning of power meter readings.

Finally, Fig. 4 shows an illustrative example of the workflow of the CEP cleaning mechanism and its connection with the sensor data repository. As this figure depicts, each raw sensor reading coming from the IoT Agents is initially stored in ORION by updating its associated *building sensor* instance. In the figure's scenario, a new power-meter reading will update the *power meter* instance representing the sender's sensor (steps 1 and 2 in the figure).

Next, ORION automatically notifies to the data cleaning module the new reading (step 3). This notification fires the two types of CEP rules described before (steps 4 and 5). At the end, the module outcome takes the form of a *clean power meter* event that updates the associated *clean power meter* instance in ORION. This *clean power meter* instance represents the cleaned version of the power meter sensor updated in step 2. Moreover, we should note that all the aforementioned interactions occur following a push-style communication enabling the real-time processing.

3.2.3. Analytics support layer

The third layer of the platform embraces all the functionalities of the platform to provide support for data mining services that can run on top of the platform. In particular, two features have been included in this layer, an energy data volatility detector and a virtual entities generator.

Virtual energy building areas generator (VEBAG). The amount of data that we are able to collect in smart buildings by means of large sensor networks sometimes does not increase the *information volume* because of redundancy. Depending on its nature, this redundancy is treated using different approaches: redundancy detection, data compression, feature extraction, and some others [42].

IoTEP works under the hypothesis that a clever way to reduce the number of variables taking part in the models can not only decrease the computation costs but also increase the accuracy on predictions and classification. In this way, the creation of abstract entities will be justified from the data analytics side, based on the assumption of the existence of this redundancy.

Therefore, the goal of the VEBAG module is the creation of high level entities that preserve as much information as possible in the data set but yet, reducing the volume of it. In this case,

we want to create virtual areas comprising several *building space areas*, finding patterns in the energy-related use and defining these virtual areas according to such information to optimize the content of information.

To do so, we aggregate each attribute per energy device daily. This aggregation can be easily done with the built-in RESTful aggregation functions provided by COMET within the homogenization and storage layer. That way, we can represent each device as a time series having one attribute measurement per day and with this, it is possible to find a clustering algorithm that groups every attribute of the time series finding some distinctions between them, like DBSCAN or longitudinal k-means.

Once every device is assigned to a cluster or virtual area, the generator computes the mean of the elements of each cluster to get an average measurement. Finally, each generated cluster is stored in the storage layer as an instance of the *virtual energy area* entity (see Fig. 3(b)). In that sense, this generator is launched on a regular basis or when certain data shifts are detected in the data by the data volatility monitor (described in the next section). Fig. 5 depicts an illustrative example of this process given the building's floor.

Firstly, Fig. 5(a) shows the distribution of room-based building space areas along with their HVACs. It should be recalled that each of these areas and sensors will be stored as different instances in IoTEP. Furthermore, the figure also shows an example of a possible time-series plot of the regulated temperature for each HVAC for illustration purposes.

Next, Fig. 5(b) shows the *virtual energy areas* generated on the basis of the aforementioned temperature time series. As we can see, the six initial room-based building space areas have been merged into three instances of *virtual energy areas* by grouping together the HVACs with similar time series. This way, rooms 4, 5 and 6 and their associated HVACs have been merged into a single area (*virtual energy area* 3 in the figure).

All in all, the generation of these virtual energy areas enables the platform to provide multiple views of the energy status of a building. In a low-level setting, we can monitor energy parameters from a single-sensor point of view. Over such simple view, we can also extract energy parameters related to a particular building spatial area (e.g., room, corridor and the like) by simple aggregation

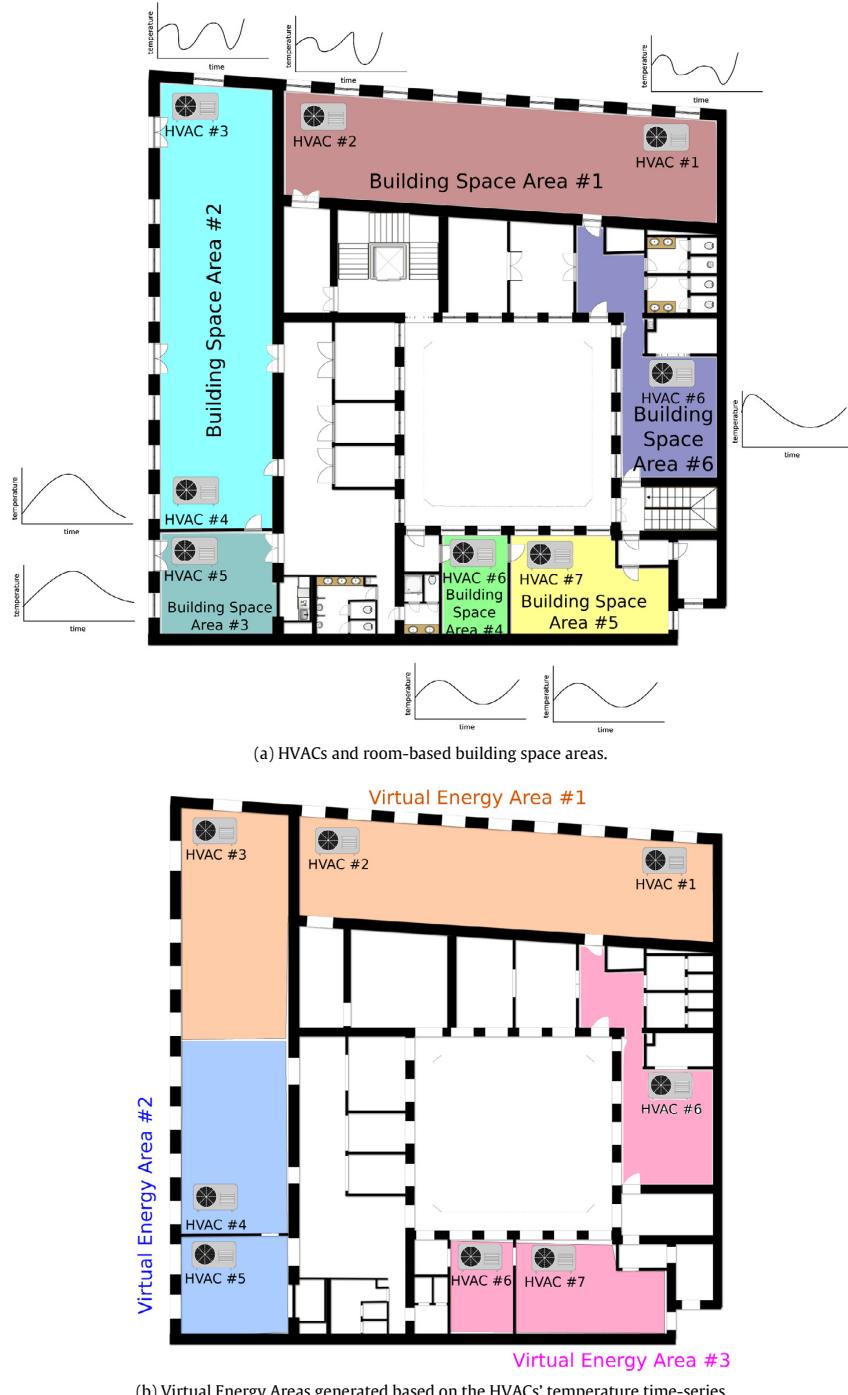


Fig. 5. Example of generation of virtual energy areas considering the HVAC temperature in a building floor.

using the *building spatial area* instances. Finally, *virtual energy area* instances enrich the energy awareness by providing an extra layer of perception that is not constrained by the building architectural structure. This way, it is possible to monitor building areas with similar energy behaviors simultaneously.

Data volatility monitor. In order to come up with real energy-aware services, the monitoring of certain energy parameters of a building becomes paramount. This includes detecting either abnormal energy consumption related to building spaces or an abnormal temperature setting related to HVACs.

For that goal, the data volatility monitor focuses on computing the current rate of change of each energy sensor parameter included in the storage layer. This is done in three steps.

Firstly, we extract the historic data set of the target energy parameter for a particular sensor with respect to a pre-defined time period t_{int}^{vol} from COMET. Then, the average rate of change among pairs of consecutive observations of the attribute is computed. Finally, if such averaged value is substantially different than the historic rate of change of that attribute then an alarm is triggered. For the sake of clarity, the pseudo-code of this process is shown in Algorithm 1.

Algorithm 1: Data volatility calculation.

```

Input: Type, identifier and energy parameter of the monitored sensor
        ( $sensor_{type}$ ,  $sensor_{id}$ ,  $sensor_{attr}$ ), time interval under study ( $t_{int}^{vol}$ )
        and historic rate of change of the considered parameter for the
        target sensor ( $r_{attr}^{sensor}$ ).
Output: Data volatility alarm, if any.
/* Historic data extraction */  

1  $\mathcal{D} \leftarrow \text{get\_COMET\_raw\_historic\_data}(sensor_{type}, sensor_{id}, sensor_{attr}, t_{int}^{vol})$   

   /* Average data-rate change calculation */  

2  $d_{prev} \leftarrow 0$   $r_{avg} \leftarrow 0$   $n \leftarrow 0$   

3 for each  $d \in \mathcal{D}$  do  

4    $r \leftarrow |d - d_{prev}|$   

5    $r_{avg} \leftarrow r_{avg} + \frac{d - r_{avg}}{n}$   

6    $d_{prev} \leftarrow d$   $n \leftarrow n + 1$   

   /* Meaningful data-rate change detection */  

7 if  $r_{avg} >> r_{attr}^{sensor}$  then  

8   return data_volatility_alarm( $sensor_{type}$ ,  $sensor_{id}$ ,  $sensor_{attr}$ ,  $r_{avg}$ )
```

This alarm is received by the final energy services on top of the platform and the VEBAG module. If this module receives a set of consecutive alarms related to the same energy parameter in a short period of time then it might indicate that the energy similarities in between building areas have changed. In order to capture such shift, VEBAG re-launches the clustering process to reconfigure the virtual areas related to such energy factor. In that sense, this monitor is endlessly executed every t_{int}^{vol} time units in order to keep a continuous control over the sensor data streams.

Finally, we would like to notice that this last mechanism along with the CEP data cleaning described in Section 3.2.2 might provide some clues to building operators about data inconsistencies due to sensor interferences. In particular, the data cleaning module can remove readings that are not consistent with the normal operation of a sensor whereas the data volatility mechanism can also detect abnormal disturbances in the data rate change of a sensor reporting that something unusual is happening.

3.2.4. Service layer

Although not that central when considering the architecture of the platform here developed, the Service Layer is the last level of the IoTEP. This layer serves as interface between the IoTEP and the user, that could be anything from a building services manager to the back end of a smartphone application.

At this level, the data analytics procedures can be invoked and their results visualized. Also, smart-building services that may be the norm when the smart-building paradigm is fully established will be nested at this level of the IoTEP platform, and will allow features such as advanced HVAC predictive control, home automation, fuel poverty evaluation, sick building syndrome diagnostics, risk situations for vulnerable people (as in heat waves), smart tariff strategies, and many others.

4. Validation of the platform

In order to test the feasibility of the proposed platform, IoTEP has been instantiated in a real pilot that allowed us to evaluate functionalities of the new platform. Here we provide some details of the evaluation scenario.

4.1. Pilot description

IoTEP was instantiated at the University of Murcia, Spain. During the last three years, this university has carried out an ambitious plan to monitor and control its buildings' infrastructures distributed across the university premisses. The number of buildings monitored and the automated services have increased quickly in the last years, what serves well the purpose of testing the plasticity of the platform presented in this paper. It should be noted that the sensorization of the buildings at the University of Murcia was done independently of this project, so the fact that the platform was able to allocate the data coming from all the sensors was already a proof of its validity.

In this context, IoTEP was used as the main enabler of an energy efficiency campaign at three cases, namely the Faculty of Chemistry and two multi-disciplinary research and technological transfer centers within the university. Details of the three buildings are provided in Table 1.

Lastly, the evaluation of IoTEP covered a three-month winter campaign from 01/10/2016 to 28/02/2017.

Platform configuration. IoTEP was installed in a centralized server with CentOS 6.7 as operating system, 8 GB RAM and 250 GB hard disk. Besides, Table 2 sums up the configuration of the inner parameters of the platform. It should be reminded that t_{int}^{clean} defines the time interval used by the CEP cleaning mechanism to compose the quartile fences (Section 3.2.2) whereas t_{int}^{vol} indicates the length of the time series considered by the data volatility mechanism to infer meaningful data shifts (see Algorithm 1).

Before the deployment of IoTEP in the pilot, a full covering of energy related variables was done in the buildings under study. After preliminary evaluations, it was discovered that there are three families of data that are fundamental to understand the energy behavior of the building users and heat losses of the envelopes. The three families are: building characteristics, energy streams and building state.

The building characteristics are the physical description of the building. Detailed blueprints of the building were obtained from the department of estates of the university together with detailed plans of constructions. This information together with visual inspections carried out by the members of our team have allowed as to have a rather full description of the condition of the building thermal envelope. With this, it was possible to use building physical models to analyze and predict the heat flows of the building and therefore the energy performance of the fabrics.

About the second family, we were able to monitor in real time with a sampling period of 10 min the operation of more than 200 conditioning units in real time. This included the status of the machines (on/off) and the set point temperatures. It was also possible to obtain the technical characteristics of the machines, what together with the rest of the data allowed us to have a rather accurate proxy of specific power consumption in real time. To contextualize this individual power consumption, the total power consumption of the building was also measured.

Finally, it was needed to know what the conditions on the interior of the spaces of the building were. For this, we monitored in real time the temperature of more than 200 spaces. These temperatures are in accordance with the data taken from the conditioning systems what allowed us to create virtual control volumes/zones in which to evaluate energy flows.

Table 1
Use case building characterization.

	Faculty of chemistry (FC)	Technological transfer center (TTC)	Research center (RC)
Location (coords)	38.02, -1.16	37.72, -1.09	38.02, -1.17
Orientation	south-west	south-west	south-west
Surface area	1500 m ²	3323 m ²	1000 m ²
Floors	6	4	2

Table 2
IoTEP parameters setting.

Parameter	Description	Value
t_{int}^{clean}	Time window length for sensor stream fence calculation	30 days
t_{int}^{vol}	Time period for data volatility calculation	2 hours

Table 3
Information model entities distribution per building.

Entity	Number of instances		
	FC	TTC	RC
Spatial region	1	1	1
Building	1	1	1
Building space area	344	16	10
HVAC	239	0	4
Clean HVAC	239	0	4
Power meter	1	13	4
Clean power meter	1	13	4
Weather conditions	1	1	1

The IoTEP was created in such manner that it allows to allocate all this information in two ways: in the form of data stream, and in the form of “static” information. In this way, the description of the building is allocated on the *building* entity previously described. The characteristics of the conditioning system and the data stream can be placed on the *HVAC* and *power meter* entities created for this purpose.

This comprehensive set-up fully monitors the most important energy related aspects of the building, what could be a two-bladed sword. In principle, this allows to do high level reasoning on the data with the high added value that this represents; however, such a large flow of data may render the infrastructure slow and inefficient with such a heterogeneous data. With the solution proposed in this paper we overcome the problems, leading to a platform that, because of the efficient handling of data inherited from FIWARE, allows for the true real time comprehensive data analysis of buildings. With the advantages that this represents.

As a result of this study, Table 3 shows the distribution of instances of the entities of the IoTEP information model stored in ORION per building.

4.2. Pilot objectives

The goal for this testing campaign was to develop a new service able to predict the next-day energy consumption of each of the three buildings, and with this to evaluate the framework we present at all the different levels. However, it should be reminded that this is only an example of the variety of features that could be implemented on IoTEP. The service tested would be instrumental for the department of estates of the university in order to plan energy-saving actions and advanced versions of model predictive control.

As Fig. 6 shows, this service was developed on top of IoTEP i.e. on the service layer shown in Section 3.2.4, by using its functionalities. It was implemented as a web application allowing the control of some of the IoTEP features by the buildings manager to carry on decisions according to data analysis results. Consequently, this application acts as a dashboard that allows users to control

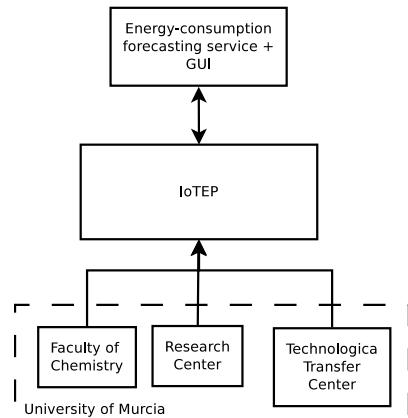


Fig. 6. Representation of the IoTEP pilot evaluation.

the platform and access the aforementioned energy consumption service (see Fig. 7).

In terms of access of the inner features of IoTEP the application includes the following actions,

- Firstly, it is possible to visualize the most recent readings of the HVAC devices per each room of the building. For this feature, the application makes use of the ORION component of the platform.
- Secondly, it is also possible to visualize the HVAC data given a time range defined by the user. For this purpose, the application leverages the raw historic data extraction method of COMET.
- Moreover, this dashboard also allows to control and visualize the results of the *virtual energy areas* generation of the platform (VEBAG module). In that sense, the user can also select the clustering method, and the number of clusters will be selected automatically by the Calinski-Harabasz index.

Finally, the energy consumption prediction service was also integrated in this application. On this way, building managers have full control over all the data analytic process starting from data visualization, aggregation and clustering to the final energy prediction procedure. This integration allows to perform such prediction for several granularity levels targeting from single devices, space areas or *virtual energy areas*. This multi-faced prediction is a key innovation aspect of the application.

For the evaluation of the platform, we studied the suitability and feasibility of the multi-layered view of the energy-related information proposed by IoTEP by means of the *virtual energy areas* generation. Additionally, we also studied the accuracy of the



Fig. 7. IoTEP dashboard and energy consumption prediction service.

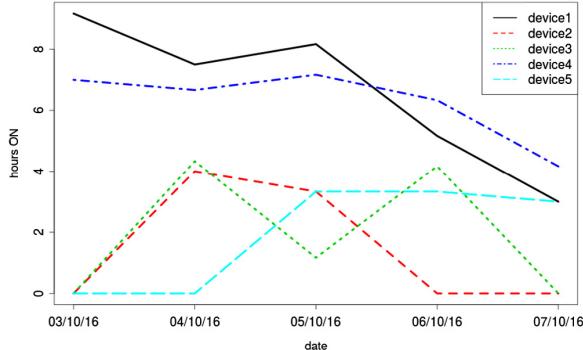


Fig. 8. Time series of 5 HVAC devices.

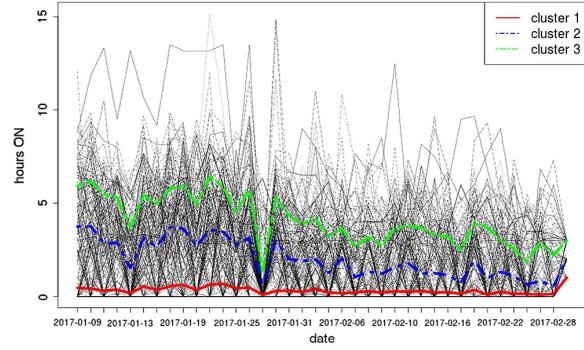


Fig. 9. Cluster evolutions.

energy prediction service when such areas are included as the target entities.

For the generation of these areas, the daily aggregation of data made by the VABAG module was based on counting the hours that each device is tuned on during the day (24 h). As an example, the number of hours that five devices were on during five days is shown in Fig. 8

For the clustering of such aggregated data, we relied on the k-means algorithm [43], but as mentioned before, more algorithms can be used for this purpose. We arbitrarily selected 3 clusters, but a different number can be selected if needed. In Fig. 9 we show the three evolutions of the groups of HVAC within FC that this algorithm identified for working days during the period of study. That way, we found rooms in this building with high use pattern (cluster 3, comprising 47 devices), rooms with little use (cluster 1 with 118 HVACs) and rooms presenting an intermediate frequency of use of the HVAC system (cluster 2 with 74 HVACs). The separation of these clusters could be the first step to an intervention strategy to modify the behavior of big consumers.

In the same way, and looking at the infrastructure level, we represent 239 values taken from the HVAC devices into 3 variables providing a 98.7% reduction of data.

Regarding the energy-prediction service, it makes its prediction according to the previous HVAC grouping within FC. Hence, we compare its performance with the use of the raw data set and in

combination with environmental variables. Being the inputs and outputs of the model identified, we followed the next steps [44]: Being the inputs and outputs of the model identified, we followed the next steps [44]:

1. Standardization of inputs
2. Splitting the data into training (75%) and test set (25%)
3. Validation: 10-fold cross validation and 5 repetitions over the training data set using several models: random forest, artificial neural networks and support vector regression.
4. Evaluation: Using the RMSE metric to evaluate the models and its coefficient of variation for comparison.

The scenarios to compare are based on the different inputs to consider:

- "Hours on" average per cluster of the previous day
- Weather predictions from Weather Underground API.⁸
- Raw HVAC data (every HVAC device daily usage)
- Both average per cluster and weather predictions

As we can see in Table 4, with a really reduced number of inputs (only 3 variables), for every model we obtain very good results compared to the others. That way, the use of clusters for creating

⁸ <https://www.wunderground.com/>.

Table 4

RMSE (and CV-RMSE) of the different models and inputs.

Model	HVAC clusters	Weather	Raw HVAC	Clust + Weath
RF	0.32 (10.53)	0.513 (17.74)	0.358 (11.83)	0.356 (11.76)
SVM	0.316 (11.03)	0.635 (22)	0.446 (14.76)	0.461 (15.23)
BRNN	0.281 (9.48)	0.423 (14.63)	0.347 (11.47)	0.398 (13.15)
# Inputs	3	23	239	26

higher level entities is proved to be useful. Although this is a rather arbitrary method, we prove with this that the platform serves to host algorithms for data analysis and prediction on a very versatile way

Comparative results. In the work [22], CV-RMSE is used in order to validate their results. They are evaluating both aggregated (total) and disaggregated (cooling and ventilating) energy consumption in a daily, weekly and monthly basis. When we compare our results with theirs, we are obtaining 6% less of variance for the RMSE, which is very satisfactory.

In addition, the Recommended Values for Baseline Model from ASHRAE Guideline 14 [45] account for the CV-RMSE smaller than 30% for daily predictions which we reach with ease (our best performance returns a 9.48 %, see Table 4).

To sum up, with this small example we show what can be implemented on the service layer of the IoT EP. With this, we intend to prove how rather complex methods can be implemented on a simple way in our platform. Also, we have shown an example of reducing data volume taking advantage of data redundancy reduction doing clustering. For this specific example we have taken three clusters as an arbitrary number and we have shown that total energy can be predicted with them. This was done as it evaluates all the features of the platform that we show in this paper, but many other applications and examples can be developed following the principles shown in Section 2.

4.3. Lessons learnt

From this first deployment of IoT EP, we can draw up some remarks.

Firstly, the results of the preliminary sensorization study of pilot were easily integrated in the IoT EP information model. This allowed to homogenized all such results in a common format and showed the versatility of the model.

Secondly, the integration of data mining support procedures as part of the platform made possible the easy development of a final service for energy data mining. In that sense, developers only needed to focus on the actual functionality of the service related to the prediction algorithms since other important tasks of the data analysis like data pre-processing or clustering were already provided by the platform.

Finally, the idea of providing a multi-layered view of the energy status of a building by means of clustering techniques has proved its suitability in the energy prediction service in two aspects. From a data-mining point of view, it reduces the redundancy of data and, thus, making up lightweight models. From a more functional point of view, the level of abstraction that the virtual energy areas provide might help building managers to better understand certain energy behaviors within the building.

All in all, this pilot has helped us to confirm that the integration of data analytics support features as part of the IoT platform is currently a key requirement in the energy domain. This enables the development of more sophisticated energy-aware services in a fast-pace process what seem to be the next natural step towards a more efficient energy-literate society.

5. Conclusions

Due to the importance of the building sector in the end-use energy consumption, it becomes a foremost task to achieve meaningful energy savings that will reduce this energy use in reality.

Despite the fact that IoT technologies have been widely used for the realization of the smart building concept, the simple sensorization of buildings is not enough to make a housing stock that consumes fewer energy resources a reality. IoT is also required to properly process, manage and, above all, analyze the energy-related data that would help to develop final energy-aware services targeting the energy efficiency goal.

In this context, several multi-purpose IoT platforms already provide generic solutions to manage IoT data. However, there is a lack of platforms in this field focusing on (1) the household energy domain and (2) providing support for data analytics. As a result, the present work shows an IoT Energy Platform (IoT EP) that covers the two aforementioned needs by following an open approach based on FIWARE enablers. IoT EP provides several functionalities oriented to the data analytics domain like the CEP data cleaning module or the times series storage along with functionalities for the correct energy management like the data volatility monitoring or the virtual energy areas detector that will allow with personalized energy feedback for the improvement of energy behavior.

Lastly, the platform has been instantiated in a real use case having a large energy sensor network. In that sense, one of the key novelties of IoT EP is that the virtual areas detection has proved to be of great help when it comes to develop an end-use energy prediction service over the platform, but many other services could be implemented with trivial computational effort under this paradigm.

Regarding further work, IoT EP has been developed re-using several open source components that are orchestrated following lightweight RESTfull calls what allows other scientists and engineers to contribute to this platform, opening the door to crowd sourced development. Consequently, new modules and enablers can be smoothly integrated in the existing architecture. In that sense, the integration of other types of sensing approaches beyond mote-class sensors, like crowdsensing, it foreseen as future actions in the platform. This would allow to capture and analyze other forms of human behavior also relevant for the building energy domain.

Acknowledgments

This paper has been made possible thanks to the support of the European Commission through the H2020-ENTROPY-649849, the Spanish National Project CI-CYT EDISON (TIN2014-52099-R) and MINECO TIN2014-52099-R project (grant BES-2015-071956) granted by the Ministry of Economy and Competitiveness of Spain (including ERDF support). Ramallo-González would like to thank the program Saavedra Fajardo (grant number 220035/SF/16) funded by Consejería de Educación y Universidades of CARM, via Fundación Séneca-Agencia de Ciencia y Tecnología de la Región de Murcia.

References

- [1] Odyssee Mure project, Energy efficiency trends in buildings in the EU, lessons from the ODYSSEE MURE project, 2012. <http://www.odyssee-mure.eu/publications;br/energy-efficiency-in-buildings.html>.
- [2] Buildings Energy Data Book, Buildings share of U.S. primary energy consumption stats. <http://buildingsdatabook.eren.doe.gov/TableView.aspx?table=1.3>.
- [3] GeSI, GeSI smarter2020 report. <http://gesi.org/SMARTer2020>.
- [4] D. Coley, T. Kershaw, M. Eames, A comparison of structural and behavioural adaptations to future proofing buildings against higher temperatures, *Build. Environ.* 55 (2012) 159–166. <http://dx.doi.org/10.1016/j.buildenv.2011.12.011>. <http://www.sciencedirect.com/science/article/pii/S0360132311004276>.
- [5] Project iNSPiRe - Development of Systemic Packages for Deep Energy Renovation of Residential and Tertiary Buildings including Envelope and Systems, 2014. <http://inspirefp7.eu>.
- [6] M. Vellei, S. Natarajan, B. Biri, J. Padget, I. Walker, The effect of real-time context-aware feedback on occupants heating behaviour and thermal adaptation, *Energy Build.* 123 (2016) 179–191. <http://dx.doi.org/10.1016/j.enbuild.2016.03.045>. <http://www.sciencedirect.com/science/article/pii/S0378778816301992>.
- [7] J.A. Stankovic, Research directions for the internet of things, *IEEE Internet of Things J.* 1 (1) (2014) 3–9. <http://dx.doi.org/10.1109/JIOT.2014.2312291>.
- [8] S. Darby, Smart metering: What potential for householder engagement? *Build. Res. Inf.* 38 (5) (2010) 442–457. <http://dx.doi.org/10.1080/09613218.2010.492660>.
- [9] V. Desley, B. Laurie, M. Peter, The effectiveness of energy feedback for conservation and peak demand: A literature review, *Open J. Energy Effic.* 2013 (2013).
- [10] M.V. Moreno, A.F. Skarmeta, L. Dufour, D. Genoud, A.J. Jara, Exploiting IoT-based sensed data in smart buildings to model its energy consumption, 2015 IEEE International Conference on Communications, ICC, 2015, pp. 698–703. <http://dx.doi.org/10.1109/ICC.2015.7248403>.
- [11] N. Simcock, S. MacGregor, P. Catney, A. Dobson, M. Ormerod, Z. Robinson, S. Ross, S. Royston, S.M. Hall, Factors influencing perceptions of domestic energy information: Content, source and process, *Energy Policy* 65 (2014) 455–464. <http://dx.doi.org/10.1016/j.enpol.2013.10.038>. <http://www.sciencedirect.com/science/article/pii/S0301421513010604>.
- [12] M. Zdravković, M. Trajanović, J. Sarraipa, R. Jardim-Gonçalves, M. Lemoche, A. Aubry, H. Panetto, Survey of Internet-of-Things platforms, in: 6th International Conference on Information Society and Technology, ICIST 2016, Vol 1, ISBN: 978-86-85525-18-6, 2016, pp. 216–220, Kopaonik, Serbia. <https://hal.archives-ouvertes.fr/hal-01298141>.
- [13] M. Molina-Solana, M. Ros, M.D. Ruiz, J. Gmez-Romero, M. Martin-Bautista, Data science for building energy management: A review, *Renew. Sustain. Energy Rev.* 70 (2017) 598–609. <http://dx.doi.org/10.1016/j.rser.2016.11.132>. <http://www.sciencedirect.com/science/article/pii/S1364032116308814>.
- [14] J. Mineraud, O. Mazhelis, X. Su, S. Tarkoma, A gap analysis of internet-of-things platforms, *Comput. Commun.* 89 (2016) 5–16.
- [15] K. Zhou, C. Fu, S. Yang, Big data driven smart energy management: From big data to big insights, *Renew. Sustain. Energy Rev.* 56 (2016) 215–225. <http://dx.doi.org/10.1016/j.rser.2015.11.050>. <http://www.sciencedirect.com/science/article/pii/S1364032115013179>.
- [16] Y. Simmhan, S. Aman, A. Kumbhare, R. Liu, S. Stevens, Q. Zhou, V. Prasanna, Cloud-based software platform for big data analytics in smart grids, *Comput. Sci. Eng.* 15 (4) (2013) 38–47. <http://dx.doi.org/10.1109/MCSE.2013.39>.
- [17] A. Kumbhare, Y. Simmhan, V. Prasanna, Cryptonite: A secure and performant data repository on public clouds, 2012 IEEE Fifth International Conference on Cloud Computing, 2012, pp. 510–517. <http://dx.doi.org/10.1109/CLOUD.2012.109>.
- [18] A. Ishii, T. Suzumura, Elastic stream computing with clouds, 2011 IEEE 4th International Conference on Cloud Computing, 2011, pp. 195–202. <http://dx.doi.org/10.1109/CLOUD.2011.11>.
- [19] N. Vastardis, M. Kampouridis, K. Yang, A user behaviour-driven smart-home gateway for energy management, *J. Ambient Intell. Smart Environ.* 8 (6) (2016) 583–602.
- [20] A. Ożadowicz, J. Grela, An event-driven building energy management system enabling active demand side management, in: Event-Based Control, Communication, and Signal Processing, EBCCSP, 2016 Second International Conference on, IEEE, 2016, pp. 1–8.
- [21] L. Klein, J.-y. Kwak, G. Kavulya, F. Jazizadeh, B. Becerik-Gerber, P. Varakantham, M. Tambe, Coordinating occupant behavior for building energy and comfort management using multi-agent systems, *Autom. Constr.* 22 (2012) 525–536.
- [22] N. Li, J.-y. Kwak, B. Becerik-Gerber, M. Tambe, Predicting HVAC energy consumption in commercial buildings using multiagent systems, Proceedings of the 30th International Symposium on Automation and Robotics in Construction and Mining, ISARC, 2013.
- [23] C. Alcaraz, I. Agudo, D. Nunez, J. Lopez, Managing incidents in smart grids a la cloud, 2011 IEEE Third International Conference on Cloud Computing Technology and Science, 2011, pp. 527–531. <http://dx.doi.org/10.1109/CloudCom.2011.79>.
- [24] M.V. Moreno, B. beda, A.F. Skarmeta, M.A. Zamora, How can we tackle energy efficiency in iot basedsmart buildings? *Sensors* 14 (6) (2014) 9582–9614. <http://dx.doi.org/10.3390/s140609582>. <http://www.mdpi.com/1424-8220/14/6/9582>.
- [25] A. Botta, W. de Donato, V. Persico, A. Pescap, Integration of cloud computing and Internet of Things: A survey, *Future Gener. Comput. Syst.* 56 (2016) 684–700. <http://dx.doi.org/10.1016/j.future.2015.09.021>. <http://www.sciencedirect.com/science/article/pii/S0167739X15003015>.
- [26] J. Zhou, Z. Cao, X. Dong, A.V. Vasilakos, Security and privacy for cloud-based IoT: Challenges, *IEEE Commun. Mag.* 55 (1) (2017) 26–33. <http://dx.doi.org/10.1109/MCOM.2017.1600363CM>.
- [27] T. Zahariadis, A. Papadakis, F. Alvarez, J. Gonzalez, F. Lopez, F. Facca, Y. Al-Hazmi, FIWARE Lab: Managing resources and services in a cloud federation supporting future internet applications, IEEE/ACM 7th International Conference on Utility and Cloud Computing, 2014, pp. 792–799. <http://dx.doi.org/10.1109/UCC.2014.129>.
- [28] E. Kovacs, M. Bauer, J. Kim, J. Yun, F. Le Gall, M. Zhao, Standards-based worldwide semantic interoperability for IoT, *IEEE Commun. Mag.* 41 (2016).
- [29] S. Sotiriadis, E.G.M. Petrakis, S. Covaci, P. Zampognaro, E. Georga, C. Thuemmler, An architecture for designing Future Internet (FI) applications in sensitive domains: Expressing the software to data paradigm by utilizing hybrid cloud technology, 13th IEEE International Conference on BioInformatics and Bio-Engineering 2013, pp. 1–6. <http://dx.doi.org/10.1109/BIBE.2013.6701578>.
- [30] F. Ramapany, F.G. Marquez, J. Soriano, T. Elsalih, Handling smart environment devices, data and services at the semantic level with the FI-WARE core platform, 2014 IEEE International Conference on Big Data, Big Data, 2014, pp. 14–20. <http://dx.doi.org/10.1109/BigData.2014.7004417>.
- [31] A. Preventis, K. Stravoskoufou, S. Sotiriadis, E.G.M. Petrakis, Personalized motion sensor driven gesture recognition in the FIWARE cloud platform, 14th International Symposium on Parallel and Distributed Computing, 2015, pp. 19–26. <http://dx.doi.org/10.1109/ISPDC.2015.10>.
- [32] P. Fernández, J.M. Santana, S. Ortega, A. Trujillo, J.P. Surez, C. Domnguez, J. Santana, A. Snchez, Smartport: A platform for sensor data monitoring in a seaport based on fiware, *Sensors* 16 (3) (2016). <http://dx.doi.org/10.3390/s16030417>. <http://www.mdpi.com/1424-8220/16/3/417>.
- [33] Telefonica I+D, IoT Agent documentation, 2017. <http://fiware-iot-stack.readthedocs.io>.
- [34] Telefonica I+D, ORION context broker documentation, 2017. <http://fiware-orion.readthedocs.io>.
- [35] Open Mobile Alliance (OMA) Specification, NGSI Context Management, 2010. http://www.openmobilealliance.org/release/NGSI/V1_0_20100803-C/OMA-TS-NGSI_Context_Management-V1_0-20100803-C.pdf.
- [36] Telefonica I+D, COMET documentation, 2017. <http://fiware-sth-comet.readthedocs.io>.
- [37] W. Chen, K. Zhou, S. Yang, C. Wu, Data quality of electricity consumption data in a smart grid environment, *Renew. Sustain. Energy Rev.* 75 (2017) 98–105. <http://dx.doi.org/10.1016/j.rser.2016.10.054>. <http://www.sciencedirect.com/science/article/pii/S1364032116307109>.
- [38] A. Ramallo-González, New method to reconstruct building environmental data, in: Building Simulation International Conference BS2015, University of Bath, 2015.
- [39] O. Etzion, P. Niblett, Event processing in action, first ed., Manning Publications Co., Greenwich, CT, USA, 2010.
- [40] NIST/SEMAtech, e-Handbook of Statistical Methods, 2012. <http://www.itl.nist.gov/div898/handbook/>.
- [41] Telefonica I+D, PERSEO official repository, 2017. <https://github.com/telefonica/perseo-fe..>
- [42] H. Hu, Y. Wen, T.S. Chua, X. Li, Toward scalable systems for big data analytics: A technology tutorial, *IEEE Access* 2 (2014) 652–687. <http://dx.doi.org/10.1109/ACCESS.2014.2332453>.
- [43] C. Genolini, B. Falissard, Kml: k-means for longitudinal data, *Comput. Stat.* 25 (2) (2010) 317–328. <http://dx.doi.org/10.1007/s00180-009-0178-4>.
- [44] A. Gonzlez-Vidal, V. Moreno-Cano, F. Terroso-Senz, A.F. Skarmeta, Towards energy efficiency smart buildings models based on intelligent data analytics, *Procedia Computer Science* 83 (2016) 994–999. <http://dx.doi.org/10.1016/j.procs.2016.04.213>. The 7th International Conference on Ambient Systems, Networks and Technologies, ANT 2016 / The 6th International Conference on Sustainable Energy Information Technology, SEIT-2016 / Affiliated Workshops. <http://www.sciencedirect.com/science/article/pii/S1877050916302460>.
- [45] ASHRAE, Measurement of Energy and Demand Savings, ASHRAE, 2002.



Fernando Terroso-Sáenz graduated from the University of Murcia with a degree in Computer science in 2006. He also received the master's degree in Computer Science at the same university in 2010. Since 2009, he has been working as a grant student in the Department of Information Engineering and Communications of the University of Murcia where he has published several papers in national and international conference proceedings. His research interests include complex event processing, ubiquitous computing and fuzzy modeling.



Ramallo-González completed his Ph.D. in Building Physics at the University of Exeter with a scholarship from the Wates Foundation. He has worked as post-doctoral researcher on two EPSRC funded projects in the department of Architecture and Civil Engineer of the University of Bath. Currently he is a Savedra-Fajardo Research Fellow in the Faculty of Computer Science at the University of Murcia, and PI of the project ThermaSim.



Aurora Gonzalez Vidal graduated in Mathematics from the University of Murcia in 2014. In 2015 she got a fellowship to work in the Statistical Division of the Research Support Service, where she specialized in Statistics and Data Analysis. In 2015, she started her Ph.D. studies in Computer Science, focusing her research on Data Analysis for Energy Efficiency and studied a Master in Big Data. Her research covers machine learning, data mining, and time series segmentation.



Antonio F. Gómez-Skarmeta received the MS degree in Computer Science from the University of Granada and BS (Hons.) and the Ph.D. degree in Computer Science from the University of Murcia. He is a Full Professor in the same Department and University. He has worked on different research projects at regional, national and especially at the European level in areas related to advanced services like multicast, multihoming, security and adaptive multimedia applications in IP and NGN networks.

4.6 Providing Personalized Energy Management and Awareness Services for Energy Efficiency in Smart Buildings

Title	Providing Personalized Energy Management and Awareness Services for Energy Efficiency in Smart Buildings
Authors	Eleni Fotopoulou, Anastasios Zafeiropoulos, Fernando Terroso-Sáenz, Umutcan Simsek Aurora González-Vidal, George Tsiolis, Panagiotis Gouvas Paris Liapis, Anna Fensel and Antonio Skarmeta
Type	Journal
Journal	Sensors
Impact factor (2018)	3.031
Rank	Q1
Publisher	MDPI
Volume	17
Issue	9
Year	2017
Month	September
ISNN	1424-8220
DOI	10.3390/s17092054
URL	https://www.ncbi.nlm.nih.gov/pubmed/28880227
State	Published
Author's contribution	The PhD student, Aurora González Vidal, contributed to the development of the set of data mining mechanisms, to the problem contextualisation and the experiments

Article

Providing Personalized Energy Management and Awareness Services for Energy Efficiency in Smart Buildings

Eleni Fotopoulou ¹, Anastasios Zafeiropoulos ^{1,*} , Fernando Terroso-Sáenz ² , Umutcan Şimşek ³ , Aurora González-Vidal ² , George Tsilios ¹, Panagiotis Gouvas ¹, Paris Liapis ¹, Anna Fensel ³ and Antonio Skarmeta ²

¹ Ubitech Ltd. Research and Development Department, 15231 Athens, Greece; efotopoulou@ubitech.eu (E.F.); gtsilios@ubitech.eu (G.T.); pgouvas@ubitech.eu (P.G.); liapis.paris@ubitech.eu (P.L.)

² Departamento de Ingeniería de la Información y las Comunicaciones, Facultad de Informática, Universidad de Murcia, 30003 Murcia, Spain; fterroso@um.es (F.T.-S.); aurora.gonzalez2@um.es (A.G.-V.); skarmeta@um.es (A.S.)

³ Semantic Technology Institute (STI) Innsbruck, University of Innsbruck, 6020 Innsbruck, Austria; umutcan.simsek@sti2.at (U.S.); anna.fensel@sti2.at (A.F.)

* Correspondence: azafeiropoulos@ubitech.eu; Tel.: +30-216-500-0507

Received: 19 July 2017; Accepted: 5 September 2017; Published: 7 September 2017

Abstract: Considering that the largest part of end-use energy consumption worldwide is associated with the buildings sector, there is an inherent need for the conceptualization, specification, implementation, and instantiation of novel solutions in smart buildings, able to achieve significant reductions in energy consumption through the adoption of energy efficient techniques and the active engagement of the occupants. Towards the design of such solutions, the identification of the main energy consuming factors, trends, and patterns, along with the appropriate modeling and understanding of the occupants' behavior and the potential for the adoption of environmentally-friendly lifestyle changes have to be realized. In the current article, an innovative energy-aware information technology (IT) ecosystem is presented, aiming to support the design and development of novel personalized energy management and awareness services that can lead to occupants' behavioral change towards actions that can have a positive impact on energy efficiency. Novel information and communication technologies (ICT) are exploited towards this direction, related mainly to the evolution of the Internet of Things (IoT), data modeling, management and fusion, big data analytics, and personalized recommendation mechanisms. The combination of such technologies has resulted in an open and extensible architectural approach able to exploit in a homogeneous, efficient and scalable way the vast amount of energy, environmental, and behavioral data collected in energy efficiency campaigns and lead to the design of energy management and awareness services targeted to the occupants' lifestyles. The overall layered architectural approach is detailed, including design and instantiation aspects based on the selection of set of available technologies and tools. Initial results from the usage of the proposed energy aware IT ecosystem in a pilot site at the University of Murcia are presented along with a set of identified open issues for future research.

Keywords: energy efficiency; behavioral change; personalized recommendations; energy analytics; behavioral analytics; big data analytics; Internet of Things (IoT); Drools; rules management system; semantic reasoning

1. Introduction

Energy consumption in residential and commercial buildings is estimated to account for around 40% of total energy consumption, making the need for promoting solutions that can potentially lead to significant reductions compelling. As stated by the U.S. Energy Information Administration, in 2015, about 40% of total U.S. energy consumption was consumed in residential and commercial buildings [1], while a similar percentage is reported by the European Commission for the overall consumption of the buildings sector in the EU [2].

The design and adoption of novel information and communication technologies (ICT) towards achieving higher levels of energy efficiency in the buildings sector is considered promising, as stated in the Global e-Sustainability Initiative SMARTer2030 report [3]. ICT has the potential to enable a 20% reduction of global CO₂ equivalent emissions by 2030, holding emissions at 2015 levels [3]. The application of ICT-enabled solutions is going to provide residents with greater insight and control, and an enhanced living experience whilst saving energy and resources. However, the application of novel ICT technologies for energy efficiency has also to rely on people adjusting their energy consumption behavior. As stated in the report of European Environment Agency [4], up to 20% of energy savings can be achieved through different measures targeting consumer behavior.

In this article, ENTROPY, an energy-aware IT ecosystem is detailed. ENTROPY aims to support energy efficiency in the buildings sector through behavioral change of the occupants with regards to their daily energy consumption patterns [5]. The main distinguishing characteristic of ENTROPY is that it exploits the advantages provided by a set of novel ICT technologies for enabling the design, development and provision of personalized energy management and awareness services in smart buildings. The philosophy of the proposed approach is based on the provision of personalized services that can lead to behavioral change through energy consumption awareness and motives provided to occupants based on their behavioral profile.

The main adopted ICT technologies include Internet of Things (IoT), information fusion, semantic web, rule-based recommendations, big data mining, and analysis mechanisms. Novel IoT node configuration, networking, and efficient data aggregation mechanisms, including mobile crowd-sensing mechanisms, are applied for the interconnection of any type of sensor (e.g., low-cost sensors such as Arduino and Raspberry Pi), a collection of data, and the application of data quality enhancement mechanisms (e.g., removal of outliers, fix missing values in time-series data). Information processing, semantic mapping, and fusion mechanisms are applied for representing the collected data in a unified way, boosting in this way their exploitability and interoperability with existing services, as well as their interlinking with available data. Recommendation mechanisms are applied for real-time reasoning over the available data and provision of suggestions for personalized actions that can lead to improving energy efficiency through behavioral change. Big data mining and analysis mechanisms are also supported for producing behavioral and energy consumption analytics, targeting at providing advanced insights and increasing the energy-awareness level of end users.

All the aforementioned technologies are supported through an integrated IT ecosystem that comprises the basis for the consumption of existing services, as well as the design and development of further energy management and awareness services, personalized mobile applications, and serious games.

The structure of the paper is as follows: In Section 2, the overall architectural approach for the design of the energy-aware IT ecosystem is provided, including subsections for the description of the IoT nodes management and aggregation mechanisms, the description of the two designed semantic models for representing energy management and occupants' behavioral concepts, and the description of the set of services provided to end users, namely, the personalized recommendations services, the data mining and analysis services, and the set of APIs for the development of personalized mobile applications and serious games. Following this, in Section 3, initial results based on the deployment of the proposed IT ecosystem at the University of Murcia in Spain and the realization of an energy

efficiency campaign are presented, while Section 4 provides a set of conclusions and identifies open issues for future research.

2. Energy-Aware IT Ecosystem Architectural Approach

Prior to delving into the description of the ENTROPY energy-aware IT ecosystem architectural approach, the type of users that interact with the ecosystem along with a concise overview of an indicative workflow for realizing an energy efficiency campaign is provided. Two types of users are considered in ENTROPY, namely, campaign managers and end users. Campaign managers are responsible for setting up an energy efficiency campaign and may consist of smart buildings administrators, energy efficiency experts, data scientists and behavioral scientists. The combination of knowledge, with regard to energy efficiency, data science, and behavioral aspects, is considered necessary towards the setup of sensor data monitoring, data analysis, and personalized recommendation delivery processes considering the infrastructure in the smart buildings and the type of users engaged in the campaign. End users regard the set of users participating in the campaign and they may consist of citizens, students, academic personnel, employees in enterprises, etc.

The basic steps followed by campaign managers for initiating and running an energy efficiency campaign are depicted in Figure 1. Upon registration to the ENTROPY ecosystem, a campaign manager obtains access to the ENTROPY services and is able to define a set of buildings taking place in the campaign, along with their division in subareas. For each area or subarea, a set of characteristics related to the surface, the working hours, the capacity in number of people, the location, the energy class of the building, etc., are provided. The next steps regarding the assignment of sensors per area or subarea and the configuration of the set of sensor data monitoring streams that have to be activated. Based on these, a set of data queries can be designed through a query design editor, the results of which are being used as input for data mining and analysis processes. The latter also have to be specified including information regarding the algorithm to be executed and the type of input and output data. Following this, the campaign can be initiated through the activation of the sensor data monitoring streams and the initiation of interaction with the end users. Continuous monitoring, evaluation and undertaking of corrective actions can be realized by the campaign administrator. It should be noted that the aforementioned steps do not need to be followed in a strict sequential way, depending on the specificities of each campaign.

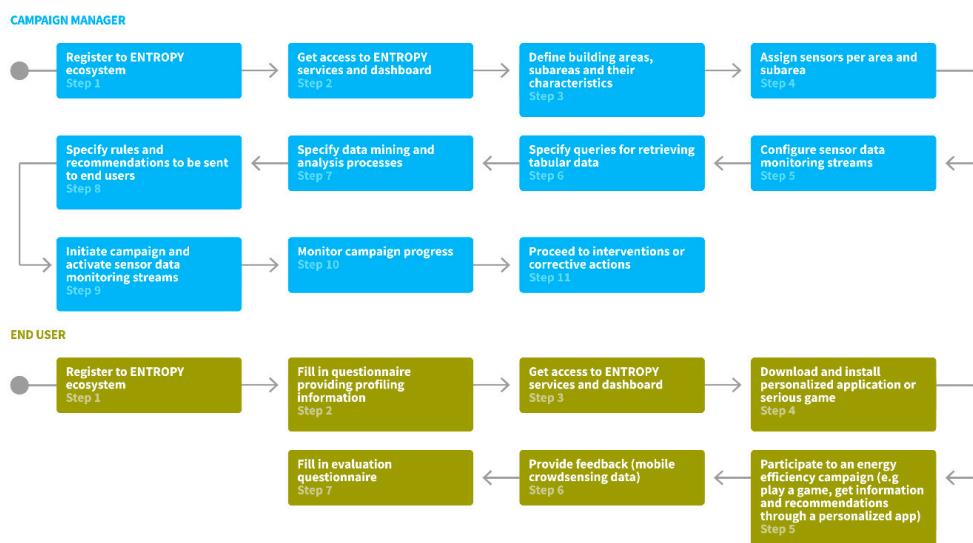


Figure 1. ENTROPY users and basic platform usage workflow.

The basic steps followed by end users for participating at an energy efficiency campaign are also depicted in Figure 1. Upon registration to the ENTROPY ecosystem, the end user has to fill in a questionnaire targeted at providing a user profile with regards to the type of employee personality, work engagement, energy conservation habits, and game interaction preferences. Following, the end user gets access to the ENTROPY services and is able to install and run ENTROPY mobile applications and serious games and, thus, participate to an energy efficiency campaign. Through the ENTROPY applications, the end user gets information regarding energy consumption, as well as environmental parameters in the areas that he has activities, receives personalized recommendations and requests for action while he is also able to provide feedback to the ENTROPY platform (e.g., information regarding malfunctioning of equipment). At the end of an energy efficiency campaign, the end user is requested to fill in an evaluation questionnaire, targeting at measuring perception of behavioral change, as well as any changes with regards to their gaming profile.

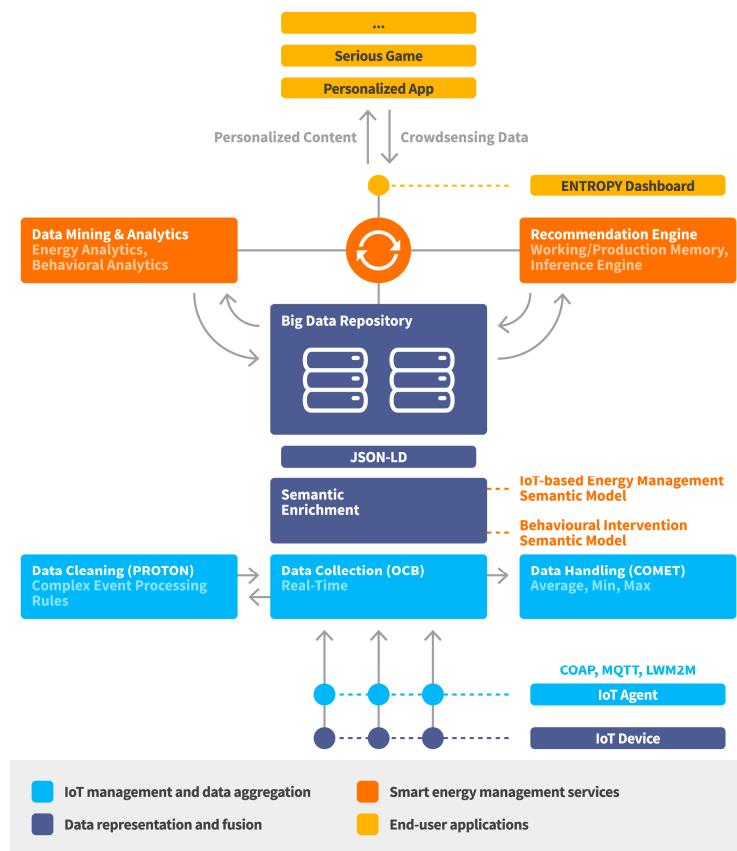


Figure 2. Energy-aware IT ecosystem architectural approach.

A high-level view of the ENTROPY energy-aware IT ecosystem architectural approach is provided at Figure 2. As depicted, a layered architecture is followed with discrete layers for IoT management and data aggregation, data representation and fusion, smart energy management services and end user applications. The IoT management and data aggregation layer is responsible for IoT nodes registration, management and data aggregation and cleaning functionalities at the edge part of the infrastructure. The data representation and fusion layer is responsible for representing the collected data based on a set of defined semantic models as well as supporting a set of data fusion mechanisms over active data streams. The smart energy management services layer is responsible for providing advanced analytics

and recommendations to end users, as well as incorporating learning techniques for continuously exploiting the produced output by each service. The end user applications layer is responsible for the design of personalized mobile applications and web-based serious games able to take advantage of the set of services provided by the lower layers. Following this, detailed information is provided for the designed and implemented mechanisms per layer.

2.1. Internet of Things Node Management and Data Aggregation

The mechanisms designed for the IoT management and data aggregation layer follow an edge computing approach. Edge computing facilitates the processing of information, where required, in the logical extremes of a network, improving in this way the performance and efficiency of applications in terms of usage of resources. It should be noted that the design of energy and information management systems is considered one of the main application areas that combine IoT and edge computing technologies [6,7]. The set of mechanisms support the easy registration, configuration and lightweight management of the infrastructure sensors deployed in the target buildings and a set of data aggregation, pre-processing and cleaning functionalities. The design and development of such mechanisms is based on the adoption and extension of a set of cloud-based open enablers, provided by the European platform for Future Internet FIWARE [8]. These enablers are orchestrated together by means of lightweight RESTful Application Programming Interfaces (APIs) according to the Open Mobile Alliance Next Generation Service Interface (NGSI) 9–10 standard [9].

In the provided approach, a set of different enablers provided by the FIWARE platform may be used, given that we are based on decoupled and self-contained modules. Data access and processing mechanisms can be designed in a future-proof way, given that the NGSI standard intends to provide a uniform cross-domain interface for advanced data access and processing. Since FIWARE enablers are compliant with such an interface, this facilitates the interoperability of FIWARE solutions with other architectures avoiding a silo-effect and making FIWARE an open solution that can be easily adopted by private and public stakeholders.

It should be noted that a set of initiatives and approaches are using FIWARE enablers in various domains. Indicatively, the Global Services Mobile Alliance (GSMA) has defined a generalized architecture for the delivery of “Internet of Things” “Big Data” services to support an ecosystem of third-party application developers [10]. Within this architectural approach, the FIWARE NGSIv2 interface has been specified as the recommended standard for certain interfaces, while a set of FIWARE enablers are used for supporting specific functionalities, including the data and control broker. In [11], a system architecture for achieving world-wide semantic interoperability solution is presented, combining the NGSI, which is part of the core of the FIWARE initiative, and oneM2M context interfaces. In [12], a semantic mechanism to integrate data from different types of devices by using FIWARE components is presented, while, in a more functional domain, and [13] made use of certain enablers, like the ORION context broker, to create a cloud-based gesture recognition application. Furthermore, in [14] it is described a sensor management system for seaports based on the FIWARE platform, while in [15] a novel patient monitoring system based on FIWARE enablers is proposed. Following the existing works on a set of diverse domains, the proposed work in this manuscript is one of the first efforts to make use of FIWARE enablers in the energy management in smart buildings domain and, thus, in the energy domain in general.

Regarding the usage of the FIWARE enablers in the ENTROPY ecosystem, the first step regards the registration of the sensor nodes and the collection of sensor data in real-time. The FIWARE enabler called IoT Agent is used for this purpose [16]. The IoT Agent acts as a gateway for hardware devices. It supports a set of communication protocols (e.g., Constrained Application Protocol (COAP), MQ Telemetry Transport (MQTT), Lightweight Machine-to-Machine (LWM2M)) for establishing connectivity with the sensor nodes and retrieving data in real-time.

The way that this data is stored and managed is tackled by another FIWARE enabler called Orion Context Broker (OCB) [17]. OCB supports the creation of real or virtual elements of interest

by using the term “entities”. Each entity is considered as a virtual sensor node that can obtain data from infrastructure sensor nodes. In the ENTROPY ecosystem, an information model comprising one entity per type of infrastructure sensor has been defined, facilitating the collection of data for all the registered sensor nodes in a homogeneous way. For instance, for collecting energy consumption data, an entity type named “energy_sensor” representing an energy meter installed in the considered building is created. Each entity includes the set of attributes monitored by the infrastructure sensor nodes along with metadata regarding the location of the sensor and the timestamp of each observation.

Following, a sensor data cleaning process takes place for improving the overall data quality through the usage of the FIWARE Complex Event Processing (CEP) enabler called PROTON [18]. CEP focuses on timely processing streams of information items, so-called events, like filtering or aggregation by means of predefined rules following the event-condition-action paradigm [19]. A filtering mechanism is implemented that discards extreme outliers of the different attributes that a sensor measurement contains, avoiding the further transmission of erroneous data. Specifically, a first set of CEP rules focus on calculating certain statistical features of each sensor’s attribute stream (e.g., first, third, and inter-quartile values) over a particular time window. Next, a second set of CEP rules is applied for each new sensor observation, discarding it in case it is considered as an outlier based on the previously defined statistical features.

Given the existence of high quality data, the COMET FIWARE enabler is used for supporting access to historic time series data [20]. COMET adheres to the same information models as the OCB enabler that gets the real-time data, thus, it does not require any further data harmonization process. It incorporates several built-in simple aggregation functions over the historic sensor data (e.g., provide sum, minimum, or maximum values). Access to such data is considered very helpful for realizing comparisons, providing input to data mining and analysis processes as well as describing rules that can lead to personalized recommendations.

2.2. Semantic Representation Models and Data Fusion

Upon making the collected data available through the IoT management and aggregation layer, sensor data streams towards the ENTROPY platform can be activated. Sensor data streams may regard real-time data or aggregated data. For each of the activated data streams, the collected data is mapped to the ENTROPY semantic models [21,22], as detailed in the following subsections, and then stored in the big data repository that is based on MongoDB. The semantic enrichment module is integrated with the ENTROPY platform and operates as an intermediate layer between data coming from FIWARE and the ENTROPY platform (user interface and REST-API) and the MongoDB. Upon the activation of a new sensor data stream, the end user is responsible to denote the mapping between the monitored sensor metric with the relevant parameter in the semantic model, as depicted in Figure 3. Following this, the collected data is stored based on the semantic model denoted parameter, supporting the unified access to the collected data.

The semantic models provide a set of advantages in terms of management and exploitation of the collected data. Through the representation of the main concepts and their relations and the mapping of the data into these concepts, unified data representation and data access mechanisms are designed and implemented. Exchange and reuse of data are also facilitated, especially following the evolving open and linked data principles. By exploiting the plethora of existing linked data tools, interlinking of data collected via the sensor data streams with available open or privately-owned data can be realized. Within ENTROPY, interconnection with the LinDA workbench [23] is realized, that is a complete open-source package of enterprise linked data tools to quickly map and publish your data in the linked data format, interlink them with other public or private data, analyze them, and create visualizations. Such functionalities for data interconnection, exchange and reuse are considered crucial for energy data in the buildings sector for realizing comparisons among data collected via similar campaigns in different regions, as well as setting up target values for energy efficiency. Interconnection of energy data with other environmental monitoring parameters provided as open data by meteorological authorities

or socioeconomic data made available by national or international statistics authorities can be also achieved (e.g., similar to the study realized in [24]), leading to advanced analysis and insights without requiring too much data management effort on behalf of the data scientists. Furthermore, through the semantically-represented data, reasoning mechanisms may be applied by taking advantage of the denoted semantics and leading to advanced insights or recommendations, as they are detailed in Section 2.3.1. Finally, linked data analytics can be produced considering the representation of a data mining and analysis process in the semantic models and the interlinking of the input and output datasets of an analysis. The enriched datasets can be exploited in a twofold way. On one hand, they can be used by data scientists for analysis results comparison purposes, while on the other hand, they can be used for defining set of rules for producing recommendations considering the analysis results.

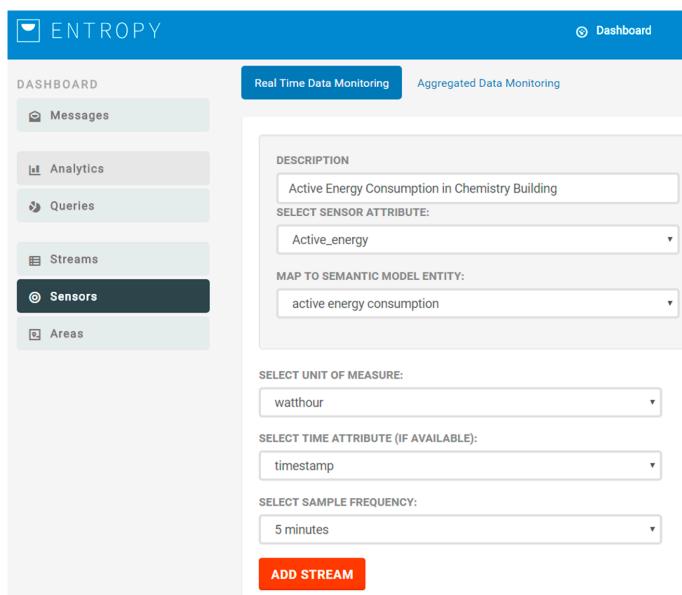


Figure 3. Sensor data stream configuration.

As already mentioned, data is stored in a MongoDB database that is a NoSQL database. Such a choice is made mainly by taking into account the supported load balancing and sharding characteristics. In order to support the storage of data without losing their expressivity in terms of their mapping to the semantic models and in parallel ensure high-performance characteristics during data management and reasoning processes, data is stored in JSON-LD format that stands for JavaScript Object Notation for Linked Data. JSON-LD is a method of encoding linked data using JSON. It is considered as an ideal data format for programming environments, REST web services, and unstructured databases, such as MongoDB. Using MongoDB and JSON-LD together is considered optimal in cases that combination of efficient representation schemes along with efficient data retrieval mechanisms has to be realized. Actually, JSON-LD was created for developers who are working with data and it showcases the power of linked data without having to go through the somewhat steep learning curve that the semantic web usually has. JSON-LD facilitates publishing data through APIs while it splits the representation layer (HTML) from the semantic layer (JSON-LD), characteristics that are not supported through other formats (e.g., Resource Description Framework—RDF). Finally, the representation of the data in JSON-LD format enables the design and implementation of reasoning mechanisms that overcome the performance limitations of ontology-based reasoning through the processing of SPARQL rules. Such an approach has been adopted in the presented IT ecosystem, through the adoption of Drools for semantic reasoning purposes over JSON-LD-represented data.

2.2.1. IoT-Based Energy Management Semantic Model

The IoT-based Energy Management semantic model (IoT-Energy) [21,22] aims to represent the set of concepts related to the support of energy efficiency in smart buildings (Figure 4). It includes conceptualization of the buildings, their structure, the deployed sensor networking infrastructure, the activation of sets of sensor data streams, as well as the realization of analysis over the collected sensor data. IoT-Energy inherits and builds upon well-known ontologies, such as the Friends of a Friend (Foaf), the Smart Appliances REference (SAREF), the Semantic Sensor Network (SSN), and the Linked Data Analytics Ontology (LDAO) [22,23].

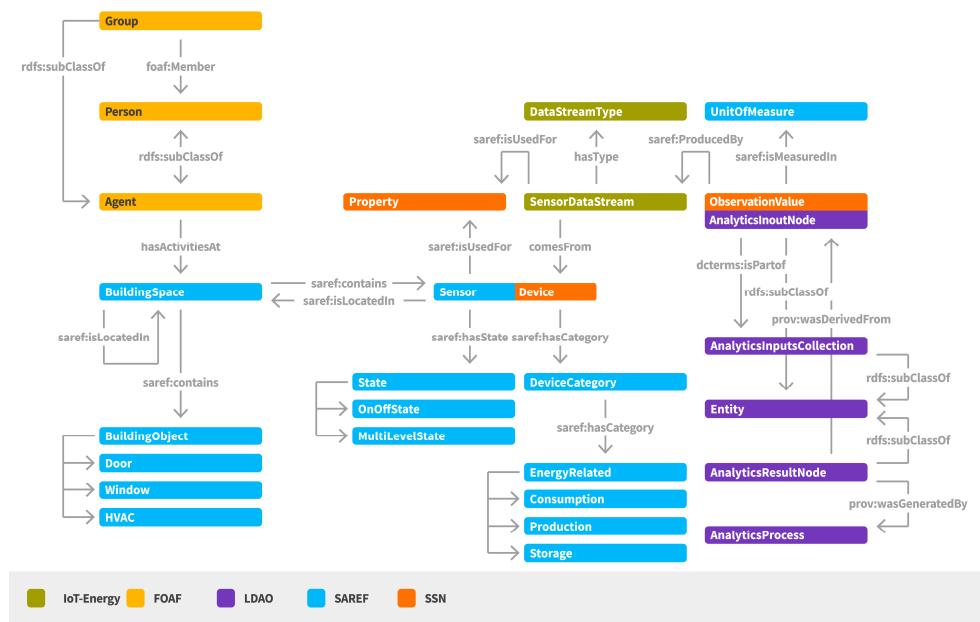


Figure 4. IoT-based energy management semantic model.

A main entity regards the *BuildingSpace* that represents any space in the building. A *BuildingSpace* may be located in another *BuildingSpace*, supporting in this way a hierarchy of building spaces. A set of *Persons* or *Groups* may have activities at a *BuildingSpace*, information that is highly helpful for providing personalized content and recommendations to the associated persons per location.

A *BuildingSpace* contains *BuildingObjects* and *Sensors/Devices*. By *BuildingObjects* we refer to objects that exist within the building space (e.g., door, window, projector, heating, ventilation and air conditioning (HVAC) device) and may be used for realizing an action (e.g., in case of an object of type “window”, send a recommendation for closing the window). By *Sensor/Device* we refer to any sensor node able to provide data upon getting measurements for a specific parameter, denoted as *ObservedProperty* in the semantic model. Each *Sensor/Device* has a category type that in case of energy related sensors can be related with the monitoring of consumption, production or storage of energy.

Another basic concept that is introduced by IoT-Energy is the sensor data stream, denoted as *DataStream*. A *DataStream* generates *ObservationValues* for a specific *ObservedProperty* of a *Device/Sensor*. Different types of *DataStreams* may be activated for providing real-time or aggregated data.

For supporting the representation of data mining and analysis processes and the relations among the provided input and output data, concepts from the Linked Data Analytics Ontology (LDAO) are inherited [23]. Each *ObservationValue* is considered as an *AnalyticInputNode*. A set of *AnalyticInputNodes* are used as an *AnalyticInputCollection* for the realization of an *AnalyticProcess* and the production of *AnalyticResultNodes*.

2.2.2. Behavioral Intervention Semantic Model

The Behavioral Intervention Semantic Model (EBIO) [21,22] aims to represent a set of concepts related to the behavioral profile of occupants in smart buildings and, thus, to facilitate the categorization of users in specific profiles and the provision of personalized content and recommendations for achieving behavioral change (Figure 5). The main concepts represented in EBIO regard the *Agent* and the *Recommendation*.

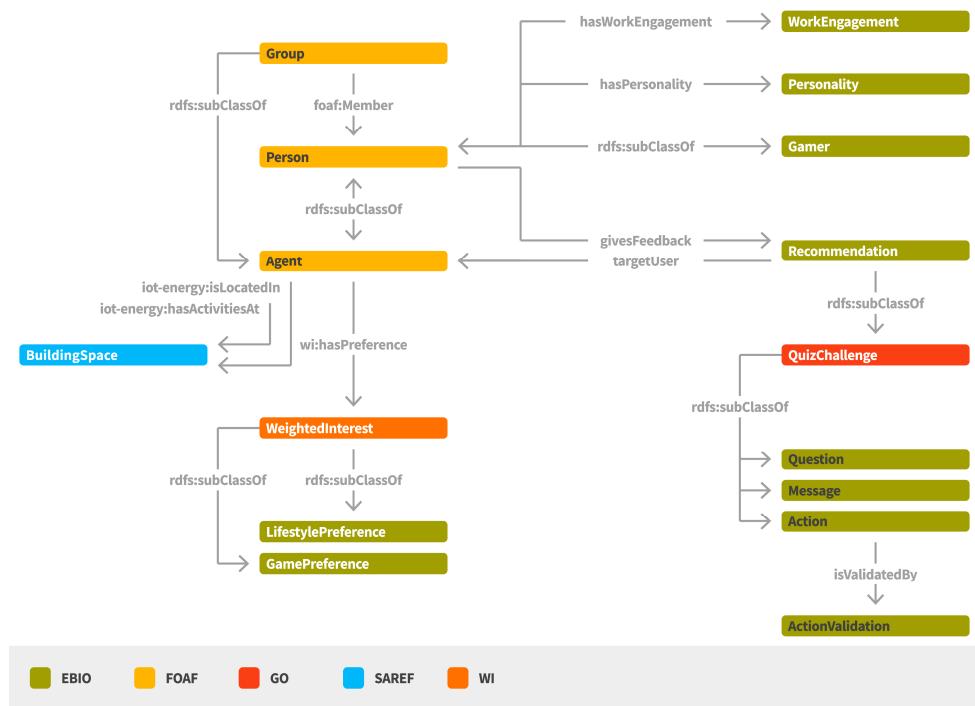


Figure 5. Behavioral intervention semantic model.

An *Agent* can be a *Person* or a *Group* where personalized recommendations can be sent. A *Person* has a “*Personality*” profile for denoting personality traits (e.g., extraversion, agreeableness, conscientiousness, emotional stability, openness to experiences). *WorkEngagement* characteristics can be also associated with a *Person*, mainly by providing indications with regards to the positive work-related state of fulfilment that is characterized by vigor, dedication, and absorption. A *Person* can be also be classified with regards to its gaming preferences, classification that can be proven very useful for providing the suitable content and application interaction mode (e.g., socializer, free spirit, achiever, disruptor). A *Person* may also have a set of interests denoted as *WeightedInterests* that may regard, among others, its game preferences (e.g., rewards, badges, points, levels) or lifestyle preferences.

Different types of *Recommendations* can be provided to *Persons* targeting at their behavioral change. Such *Recommendations* can have the form of a *Message*, a *QuizChallenge*, an *Action* or a *Question*. An *Action* is associated with an activity or a series of activities whose result contributes to elimination of a certain energy waste cause. A *Question* can be posed to a group of people, leading to the collection of crowd-sensing feedback (e.g., information regarding their comfort level). It should be noted that in the context of ENTROPY, the *Group* concept represents a group of users that share a common characteristic (e.g., tendency to sacrifice comfort for energy efficiency). A user can give a positive or negative *Feedback* to a *Recommendation* that is later utilized for the generation of new personalized recommendations.

2.3. Intelligent Energy Management and Awareness Services

Based on the semantically-mapped storage of the collected data in the ENTROPY big data repository, and through the definition of a set of REST APIs, various services can be designed and provided through the ENTROPY platform. Such services include the recommendation engine for providing personalized recommendations to end users, as well as the data mining and analysis mechanisms for providing behavioral and energy analytics. It should be noted that these mechanisms work in a complementary fashion, since produced output from an analysis process can trigger the provision of a new recommendation. Similarly, the feedback provided by end users based on the consumption of recommendations can lead to analysis and classification of end users in specific personality or gamer types.

2.3.1. Recommendation Engine

The recommendation engine is responsible for providing context-aware and personalized recommendations taking into account the occupants' behavioral profiles. It is implemented based on Drools, a rules-based management system [25]. It consists of the working memory, where facts are introduced based on the provided data, the production memory, where the set of defined rules are available, and an inference engine that supports reasoning and conflict resolution, as well as triggering of the appropriate recommendations (see Figure 6).

Triggering of recommendations follows a continuous match-resolve-act approach. Specifically, the match phase regards the mapping of the set of applied rules which are satisfied based on the available data, the resolve phase regards the process of conflict resolution, if any, among the satisfied rules, while the act phase regards the triggering of the recommendations towards the group of the target users. Rules are mostly related with the identification of context change, especially with regards to the location of the end users or changes in the observed values in the activated data streams.

A rule consists of a condition element and a recommendation template in the action part, which connects a context change with specific target user group criteria. When a rule is fired due to a context change (e.g., when average CO₂ measurement within an hour exceeds the defined threshold), the recommendation engine selects the set of target users based on the defined user attribute filters (e.g., players who have activities at a certain location, users that are classified as highly responsive at the proposed actions through the personalized recommendations, users that satisfy specific behavioral criteria) and creates a personalized recommendation for each of them by using the defined recommendation template. Following this, the set of recommendations are published in a publish/subscribe framework and made available for consumption by the set of personalized applications and serious games.

A produced recommendation contains the target user, the related content, the measurement attributes that are involved in the creation of the recommendation, the possible reward for the completion of the recommendation, as well as the validation method for it. The attributes involved in the creation of the recommendation is provided for gamification purposes, since different measurements may work towards earning different rewards (e.g., the points earned from completing a task regarding CO₂ may have an impact on earning a so called "Refresher Badge"). The rewards are registered to a user upon the completion and validation of a recommendation, which differs per type of recommendation. For instance, an action may be validated by checking the status of the sensors on the involved building objects (e.g., a window), while the validation of a quiz is done inherently by answering all the questions. A set of indicative recommendations are provided in Figure 7.

It should be noted that reasoning mechanisms are applied based on the designed ENTROPY semantic models. An implementation of a reasoning business logic is realized based on Drools, mainly for exploiting the high performance and scalability characteristics of Drools-based systems, compared to ontology-based reasoning [26], as well as the separation of the rules-definition business logic from the core ENTROPY functionalities. Campaign managers are able to define and update or extend the set of rules applied for inference purposes in a dynamic way, compared to static approaches realized

in ontology-based reasoning, while rules declaration can be realized in a non-technical way in order to be understandable to people that are not domain experts [27]. In this way, both performance and rules definition complexity aspects are tackled.

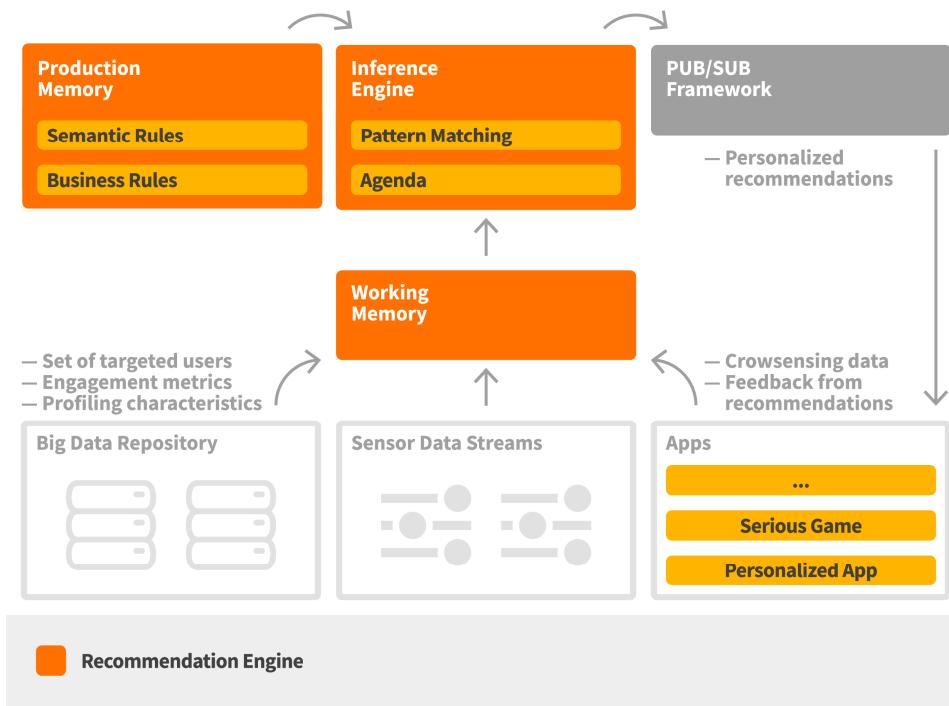


Figure 6. Recommendation engine workflow.

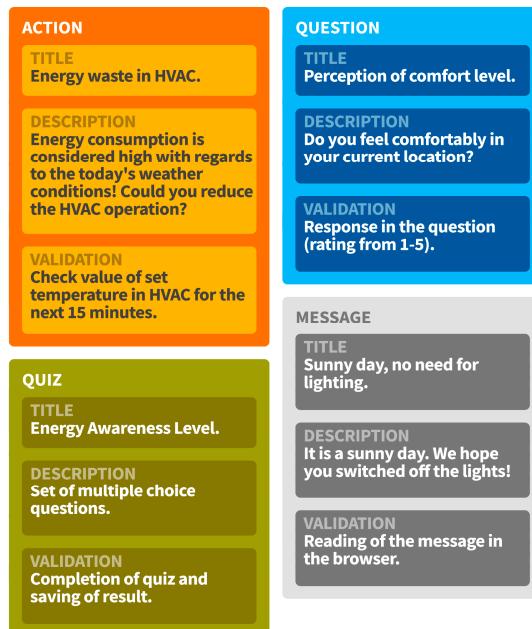


Figure 7. Indicative recommendations.

However, given that an OWL implementation offers advantages in terms of rules expressiveness as well as exploitation of the semantic properties of a semantic model, the development of a module is realized, denoting set of rules in Drools native rule language (DRL) able to support the set of semantic properties. It should be noted that the examination of such a solution was also suggested by existing performance evaluation studies comparing Drools and ontology-based reasoning [28]. In more detail, dynamic production of new knowledge is supported through the execution of specific rules that are responsible to produce reasoning on the fly. These rules are able to perform class transitivity (i.e., if a base-class belongs to a class that has also a parental class, then the new knowledge is the fact that the base-class belongs also to the parental class), supertype inheritance (i.e., if there is an instance of class that has a parental class then the new knowledge is the fact that a replica of the same instance instantiates the parental class), property transitivity and sub-property transitivity. The inference business logic is implemented using the first-order logic capabilities of the Drools engine instead of relying on third-party reasoners, enhancing a significant amount of the overall system efficiency.

Drools-based reasoning mechanisms are also considering a set of axioms that are denoted throughout the usage of the platform for declaring advanced relationships that are observed. For instance, given that by the interpretation of a first set of questionnaires results, it is denoted that humanitarian-oriented persons show a preference for badges and roles in the used mobile applications, we can define a class axiom in OWL Manchester syntax such as “(hasPreference value Badge and hasPreference value Role) subClassOf Humanitarian”. Such an axiom is included in the real-time reasoning process, leading to production of the relevant knowledge each time the axiom description is validated.

2.3.2. Data Mining and Analysis Services

Another service that is implemented and provided within the ENTROPY platform regards the support of a set of big data mining and analysis techniques towards the extraction of energy and behavioral analytics. Insights provided with regards to the energy usage in smart buildings, as well as the behavioral characteristics of the occupants, may lead on one hand on increase of their energy awareness and on the other hand on targeted recommendations for reducing energy consumption.

The supported set of analytics processes concerns descriptive, predictive, classification, clustering, and prescriptive analytics [5,29]. Descriptive analytics are providing summary information regarding the energy usage, as well as other environmental or behavioral attributes. Predictive analytics are providing estimates for usage of energy the upcoming period, as well as examining the relationship among energy consumption and set of parameters, such as average temperature, heating or cooling degree days, day of the week, etc. The considered algorithms include linear regression, multiple linear regression, support vector regression, and principal component analysis. Classification and clustering analytics are applied for identifying or classifying collective behaviors among the involved users. Based on the identification of groups, targeted interventions may be planned, while the produced groups may be also considered as input towards a group-aware forecasting analysis. The considered algorithms include artificial neural networks, Bayesian regularized neural networks, random forest, k-means, density-based spatial clustering, and hierarchical clustering. Prescriptive analytics are applied for combining analytics results with automation solutions considering the interplay among energy efficiency and comfort level of occupants.

The workflow followed for the support of data mining and analysis techniques is depicted in Figure 8. An analysis process is based on the selection of an analysis template and the selection of the queries to be executed for providing the input datasets (training and/or evaluation datasets). Each analysis template represents a specific algorithm and provides to the user the flexibility to adjust the relevant configuration parameters. Such parameters include input parameters for the algorithm along with their description and their default value, as well as output parameters along with their type (text, image, data, html). An indicative analysis template for the calculation of heating or cooling degree days per day [30] for a monthly period is depicted in Figure 9. A set of analysis templates can be made

available and be used for initiating an analysis. It should be noted that an analysis process is also associated with a set of execution parameters that denote whether an analysis should be realized in a manual or automated way, as well as the periodicity factor for the latter case.

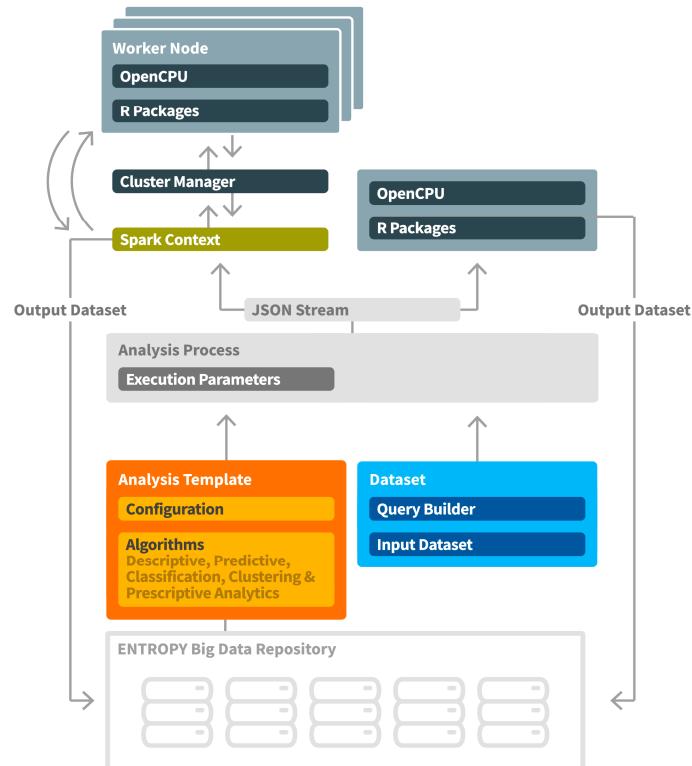


Figure 8. Data mining and analysis workflow.

The design of queries for obtaining the input datasets for the analysis is based on the development of a query builder over MongoDB, facilitating end users to easily prepare their input datasets. Two categories of queries are supported, namely queries for fetching data collected by sensor data streams (e.g., energy consumption, humidity, and indoor temperature data per hour for a specific room) and queries for fetching data related to the set of users participating at the energy efficiency campaign (e.g., a set of users with an educational level relevant to a Master's degree). An indicative query for getting the average power consumption and the external temperature per hour is depicted in Figure 10. Upon the execution of the queries, streams of the input training or evaluation datasets are provided to the analysis toolkits.

In ENTROPY, the R Project for Statistical Computing [31], and the Apache Spark fast and general engine for large-scale data processing [32] are used for this purpose. Depending on the analysis needs in terms of big data management and performance aspects, the optimal tool per case may be selected. Interconnection of the ENTROPY components with the analysis toolkits is based on the OpenCPU system for embedded scientific computing that provides a reliable and interoperable HTTP API for data analysis based on R. In the case of large-scale data processing and the need for a big data analysis framework, the Apache Spark engine is used, where the analysis process is realized in a set of worker nodes, each one of which is hosting an Apache Spark OpenCPU Executor [33]. The set of worker nodes are formulating a cluster orchestrated by a cluster manager.

Upon the realization of an analysis, the produced results (output dataset) are also made available through a set of URLs providing access to the set of results, as defined in the output parameters

of the analysis template. It should be noted that analysis results are also semantically mapped to the ENTROPY semantic models, based on the adoption of the LDAO ontology, as mentioned in the previous section.

The screenshot shows the ENTROPY web application interface. The left sidebar has a dark theme with light-colored buttons for 'Analytics' (selected), 'Messages', 'Queries', 'Streams', 'Sensors', and 'Areas'. The main content area has a blue header bar with the ENTROPY logo and navigation links for 'Dashboard', 'Account', 'Admin', and 'Logout (entropy)'. Below the header, the URL is /ANALYTIC / ALGORITHM and the page title is 'Analytic'. A sub-header says 'Define a Analytic Algorithm Template'. The form fields include:

- NAME:** Cooling Degree Days
- BASE URL:** http://192.168.3.6
- EXECUTION URL:** http://192.168.3.6/ocpu/library/coolingDegreeDays/R/coolingDegreeDaysPlot
- INPUT PARAMETERS:** A table with columns Name, Description, Type, and DefaultValue. It contains two rows:

Name	Description	Type	DefaultValue
base_temp	base temperature	string	21
surface	surface	string	10
- DATA PARAMETERS:** A table with columns Name, Description, Type, and DefaultValue. It contains one row:

Name	Description	Type	DefaultValue
datastream	datastream	query	http://entropy.eu/projects.net/api/v1/query/execut query/595a4a33aafe05530ec1198c
- RESULTS:** A table with columns Description, Type, and Url. It contains one row:

Description	Type	Url
cooling degree days	text	coolingdeggredays.html

Figure 9. Indicative algorithm analysis template.

The screenshot shows the ENTROPY web application interface. The left sidebar has a dark theme with light-colored buttons for 'Messages' (selected), 'Analytics', 'Queries' (selected), 'Streams', 'Sensors', and 'Areas'. The main content area has a blue header bar with the ENTROPY logo and navigation links for 'Dashboard', 'Account', 'Admin', and 'Logout (entropy)'. Below the header, the URL is /QUERIES / EDIT and the page title is 'Queries'. A sub-header says 'Edit Query'. The form fields include:

- NAME:** SampleQuery
- Selected Collection:** SENSORSTREAM
- AND:** A rule editor with two conditions:
 - ObservationValue isProducedBy : equal ExtTempHessoAvgHourDegree
 - ObservationValue isProducedBy : equal AvgPowerlGHourDegreeDays
- Buttons:** UPDATE QUERY, Run Query, Reset

Below the query editor is a table with columns ExtTempHessoAvgHourDegree, AvgPowerlGHourDegreeDays, and inDateTime. The data is:

ExtTempHessoAvgHourDegree	AvgPowerlGHourDegreeDays	inDateTime
13.32	205058.66666666666	2017-08-29T07:00:00.000Z
13.00	190807.75	2017-08-29T06:00:00.000Z
12.60	187019.07692307694	2017-08-29T05:00:00.000Z
12.55	187034.08	2017-08-29T04:00:00.000Z
10.02	188220	2017-08-29T03:00:00.000Z
10.10	186874.76363636364	2017-08-29T02:00:00.000Z
10.01	186811.06666666666	2017-08-29T01:00:00.000Z
11.00	186772	2017-08-29T00:00:00.000Z

Figure 10. Indicative query design through the query builder.

At the current phase, a set of initial algorithms are considered, however, the overall implementation facilitates the incremental addition of further analysis mechanisms.

2.3.3. Personalized Applications and Serious Games Development

In addition to the set of intelligent energy management and awareness services supported by the ENTROPY IT ecosystem, the development of mobile applications is facilitated. Given the unified representation of data through the semantic models independently of the underlying sensor infrastructure, as well as the design and implementation of set of REST APIs for accessing and storing data to the big data repository, personalized applications and serious games development is enabled, while their applicability may regard various smart building cases. Such APIs include, among others, the provision of information for the available building spaces and their energy consumption profiles, the activated sensor data streams, latest data per sensor data stream, the recommendations provided per user along with the collected feedback, the set of actions that may be requested to be realized by an end user, the execution of queries, user demographic data, functionalities for user registration, authentication and login in the IT ecosystem, initialization and update of the user profile per application, as well as the retrieval of the top users per application. An overview of the developed APIs along with the associated input and output parameters is provided at Figure 11. In each case, the set of GET or POST body parameters, the headers and the type of the response are specified, as detailed in [34]. The aforementioned functionalities can be proven beneficial towards the development of smart applications that can combine energy and behavioral data, along with the services provided by the ENTROPY IT ecosystem, and lead to the improvement of energy efficiency in smart buildings.

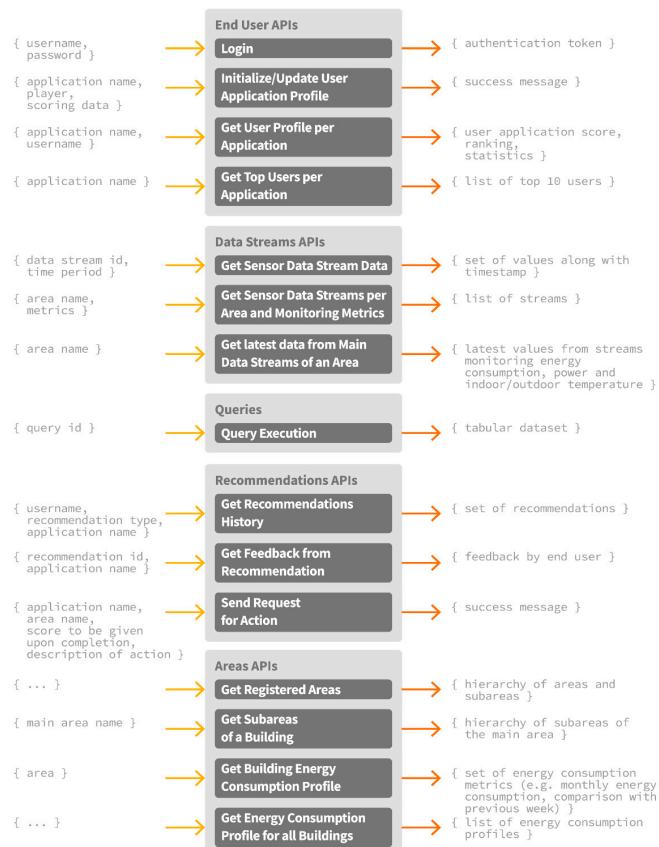


Figure 11. ENTROPY API overview.

3. Improving Energy Efficiency at the University of Murcia

One of the pilot cases where the ENTROPY IT ecosystem is instantiated is at the University of Murcia in Spain. The pilot regards an energy efficiency campaign at three cases, namely the Faculty of Chemistry and two multi-disciplinary research centers. In each of these cases, a set of infrastructure sensors and actuators have been installed to capture energy-related, as well as environmental data. The three cases comprise of 25 energy smart meters and 190 HVACs monitoring energy consumption and temperature in various building spaces, ranging from a room or corridor level to the overall building.

The targeted users regard students, professors and the administrative staff. While students' behavioral changes are tackled on the shared labs as a group, every professor has their own office, and so personal actions might be carried out. Contrary to the other two groups, the administrative staff is required to register when they enter or leave their offices by means of personal smart cards providing in this manner data regarding their presence.

Following some preliminary results of the aforementioned campaign for the building of the Faculty of Chemistry are detailed. These results regard the realization of a set of data mining and analysis processes with a twofold objective. On one hand, targeting at the prediction of the energy consumption for the following day and on the other hand targeting at the grouping of the set of considered building spaces based on the usage of the HVAC equipment during the day.

The first part of the analysis is realized per day at evening time and is based on weather forecasting data provided by the Weather Underground API service [35] and energy consumption data of the overall building. The output data regards energy consumption for the following day. The applied algorithms are random forest (RF), support vector regression (SVR), and Bayesian regularized neural network (BRNN), providing the root-mean-square error (RMSE) of the prediction, which varies between 0.4 and 0.6 for the tested algorithms, where $RMSE = \sqrt{\left(\frac{\sum(y_i - \bar{y}_i)^2}{n}\right)}$. In Figure 12, the application of the built models to the test dataset is shown for a set of dates, where random forest provides the best results. When normalizing the RMSE by the mean of the real tested values, we obtain the coefficient of variation (CVRMSE). CVRMSE is used to avoid ambiguity when comparing models. In this case it varies between 9.5% (RF) and 12.7% (BRNN), meaning good predictive results. By having accurate predictions regarding the energy consumption, optimal planning of usage of energy can be achieved.

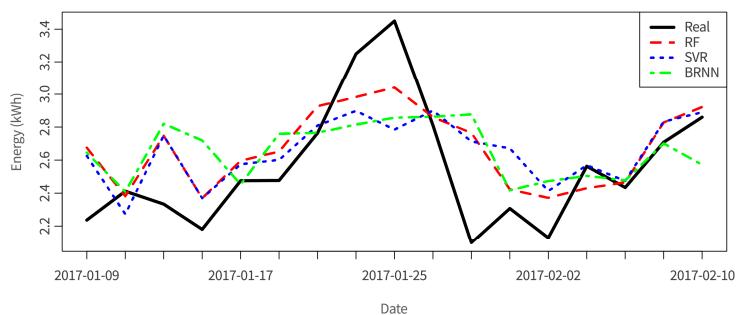


Figure 12. Analysis results: energy consumption prediction based on environmental variables.

In the second part of the analysis, information regarding the usage of the HVACs in a set of building spaces is collected and used for clustering purposes. Based on the cluster where a building space is assigned, targeted recommendations to the users that have activities in this building space may be provided. The information used for clustering purposes regards the state of the HVAC devices (on/off), the indoor and outdoor temperature, as well as the target temperature set. An indicative graph of the outdoor temperature for a weekly period is depicted in Figure 13, as it is produced by

the ENTROPY platform, while an indicative figure of the configuration provided for a registration of building spaces and subspaces is also depicted in Figure 14.

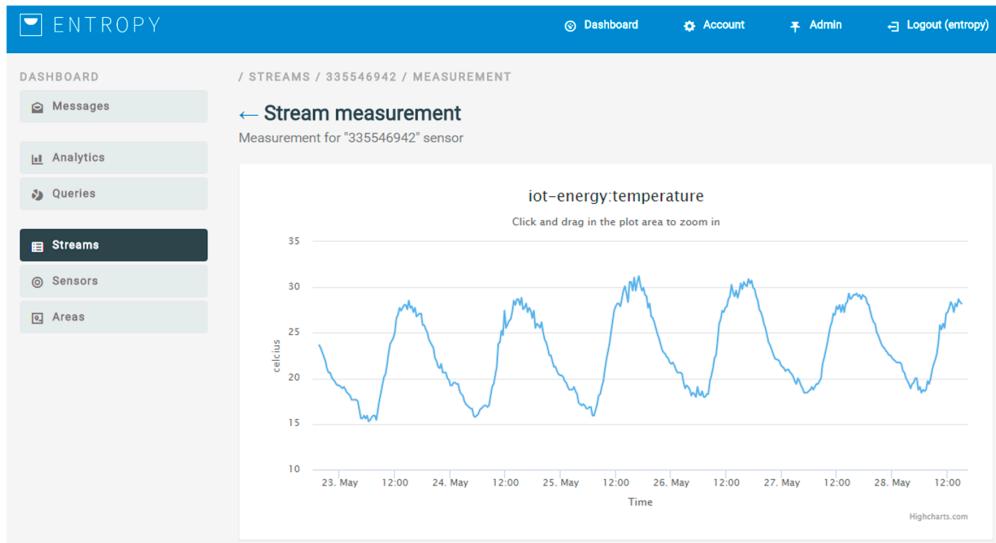


Figure 13. Outdoor temperature at the University of Murcia Campus.

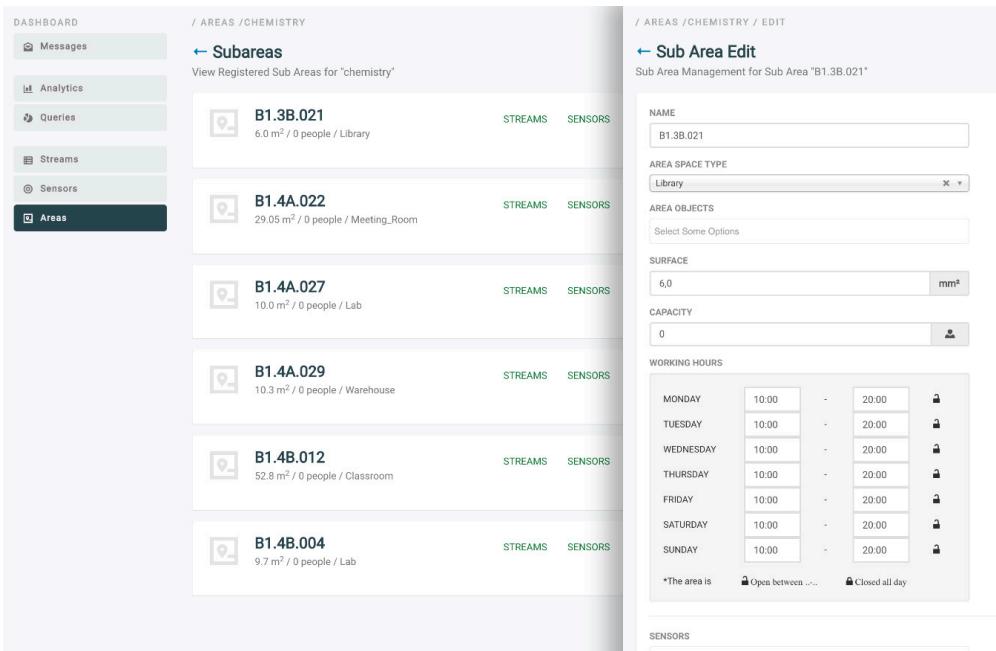


Figure 14. Building space and subspaces registration at the University of Murcia Campus.

The set of algorithms applied for clustering purposes are hierarchical clustering, longitudinal k-means and density-based spatial clustering of applications with noise. In Figure 15, the three trajectories of the HVAC groups that hierarchical clustering identifies for working days of January and February 2017 are colored. Each black line corresponds to the usage graph (percentage of daily active

time period) of a single HVAC device through both months. Clusters 1, 2, and 3 regard building spaces with low, intermediate, and high usage patterns, accordingly. Each cluster trajectory line corresponds to the mean daily value of the set of building spaces that belong to the cluster. These clusters can be introduced in the energy consumption model as input variables, since the energy consumption is linearly dependent to the HVAC usage.

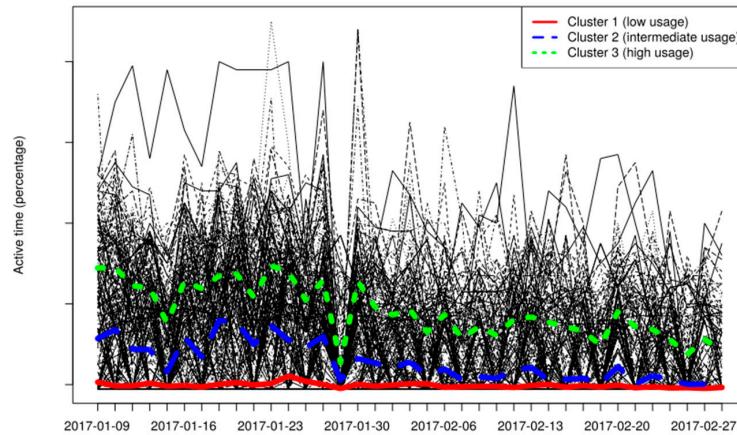


Figure 15. Analysis results: Clustering of building spaces according to their HVAC usage through time (Cluster 1/2/3: red/blue/green line denoting trajectory of set of building spaces with low/intermediate/high usage patterns, accordingly).

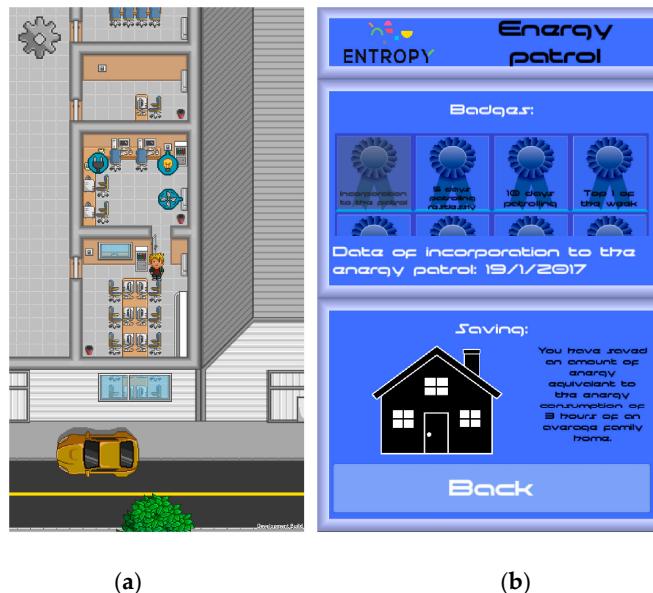


Figure 16. Screenshots from interactive games: (a) My Green Avatar; and (b) The Energy Patrol.

For interacting with end users participating in the energy efficiency campaign and providing personalized recommendations, two games have been developed by exploiting the ENTROPY IT ecosystem. Indicative screenshots from these games are depicted in Figure 16. The first game (Figure 16a)—called “My Green Avatar”—allows users within the same building to report their actions to achieve energy savings (e.g., switch-off HVACs, turn off appliances or lights, and the

like) by means of their own virtual avatar. Then, such actions are confirmed by analyzing the data streams of the related actuators. The second game (Figure 16b)—called “The Energy Patrol”—provides recommendations to users to undertake specific actions to improve the energy savings of their building like turning off the lights in potentially empty rooms or adjust the temperature of the HVACs to an efficient setting. Such recommendations are provided to end users based on a set of rules defined in the ENTROPY recommendation engine and are consumed in real-time through the deployed publish/subscribe framework. Validation of the successful realization of the proposed actions is also supported, exploiting the mechanisms provided by the recommendation engine.

4. Conclusions

In the current article, a novel IT ecosystem is presented that aims to improve energy efficiency in smart buildings through behavioral change of the occupants, based on the exploitation of emerging ICT technologies.

A set of innovative characteristics of the ENTROPY IT ecosystem are detailed. The set of open and extensible mechanisms for sensors registration, configuration, and sensor data management at the edge or the core part of the infrastructure facilitates the easy adoption of the overall solution and its instantiation in diverse and heterogeneous infrastructure cases. The representation of data based on the specification of energy and behavioral semantic models facilitate the unified access to them by numerous applications and services as well as their interconnection with available open and linked data for further processing. The set of data analysis and recommendation services, designed in a way that they collaborate among each other can lead to targeted recommendations, energy and behavioral analytics, actions and decisions with direct impact on behavioral change of occupants and, thus, energy efficiency increase. Finally, the set of REST APIs provided for consuming the set of available services can boost the design and development of personalized applications and serious games targeted to specific type of buildings with minimal effort.

The detailed IT ecosystem can be applied in diverse cases with minimal configuration effort, the supported energy management and awareness services can be easily consumed while the design and development of further services and mobile applications is highly facilitated through the exploitation of the unified way of representation of the collected data.

Building upon the presented energy-aware IT ecosystem and taking into account the set of initial results produced, several open issues and ideas for extensions are identified. Based on the existing implementation of data mining and analysis services, the design and implementation of a set of data mining and analysis processes stemming from various tools can be realized, leading to advance insights through the processing of energy data. Such tools include the R statistics project, SparkR as a light-weight frontend to use Apache Spark from R, as well as analysis software implemented via other tools (e.g., Python scripts), exploiting the interfaces provided through OpenCPU. With regards to the recommendation engine, extensive performance evaluation results for the usage of Drools for semantic reasoning purposes taking into account the number of introduced rules and the volume of the processed data can be realized, leading to meaningful and exploitable insights for adopting such a solution in energy management solutions as well as other domains. In parallel, collection of feedback for the provided semantic models can lead to extensions or minor modifications of them aiming to serve a wider community and improve data interoperability aspects. Further extensions in the FIWARE enablers can be also implemented and proven beneficial in order to introduce advanced complex event processing rules for improving data quality prior to transmitting and processing them at the centralized infrastructure. Finally, realization of a set of energy efficiency campaigns, part of which are already planned to be realized within the ENTROPY H2020 project, and evaluation of the potential for reducing energy consumption through the usage of the ENTROPY IT ecosystem services has to be achieved combined with a set of dissemination activities for adoption of the ecosystem by a wider community.

Supplementary Materials: The following are available online at <http://www.mdpi.com/1424-8220/17/9/2054/s1>; Video S1. Introduction to ENTROPY; Video S2. Walkthrough in the ENTROPY platform.

Acknowledgments: This work is supported by the European Commission Research Programs through the Entropy Project under Contract H2020-649849.

Author Contributions: Eleni Fotopoulou, Anastasios Zafeiropoulos, Fernando Terroso-Sáenz, Umutcan Şimşek, and Antonio Skarmeta were mainly involved in the design of the ENTROPY architectural approach, as well as the implementation of the ENTROPY platform; Aurora González-Vidal developed a set of data mining mechanisms; Panagiotis Gouvas developed a set of semantic reasoning mechanisms based on Drools; George Tsilios designed and implemented the ENTROPY user interface; and Paris Liapis and Anna Fensel contributed to the definition of the ENTROPY semantic models and the design of the recommendation engine.

Conflicts of Interest: The authors declare no conflict of interest.

References

1. U.S. Energy Information Administration. Monthly Energy Review. Available online: <https://www.eia.gov/totalenergy/data/monthly/#consumption> (accessed on 31 August 2017).
2. European Commission. Energy Efficiency of Buildings. Available online: <https://ec.europa.eu/energy/en/topics/energy-efficiency/buildings> (accessed on 31 August 2017).
3. Global e-Sustainability Initiative (GeSI). *Global e-Sustainability Initiative (GeSI) SMARTer2030 Report, SMARTer2030, ICT Solutions for 21st Century Challenges*; Global e-Sustainability Initiative (GeSI): Brussels, Belgium, 2015.
4. European Environment Agency. *European Environment Agency Technical Report No 5/2013, Achieving Energy Efficiency through Behavior Change: What Does it Take?* European Environment Agency: Paris, France, 2013.
5. ENTROPY H2020 Project. Available online: <http://entropy-project.eu/> (accessed on 31 August 2017).
6. Baccarelli, E.; Naranjo, P.G.V.; Scarpiniti, M.; Shojafar, M.; Abawajy, J.H. Fog of Everything: Energy-Efficient Networked Computing Architectures, Research Challenges, and a Case Study. *IEEE Access* **2017**, *5*, 9882–9910. [[CrossRef](#)]
7. Naranjo, P.G.V.; Shojafar, M.; Vaca-Cardenas, L.; Canali, C.; Lancellotti, R.; Baccarelli, E. Big Data over SmartGrid—A Fog Computing Perspective. In Proceedings of the 24th International Conference on Software, Telecommunications and Computer Networks (SoftCOM 2016), Split, Croatia, 22–24 September 2016; pp. 1–6.
8. FIWARE Platform Official Website. Available online: <https://www.fiware.org/> (accessed on 31 August 2017).
9. Open Mobile Alliance Next Generation Services Interface (NGSI) Specification. Available online: <http://www.openmobilealliance.org/release/NGSI/> (accessed on 31 August 2017).
10. GSM Association Official Document CLP.25. IoT Big Data Framework Architecture. Available online: <https://www.gsma.com/iot/wp-content/uploads/2016/11/CLP.25-v1.0.pdf> (accessed on 31 August 2017).
11. Kovacs, E.; Bauer, M.; Kim, J.; Yun, J.; le Gall, F.; Zhao, M. Standards-Based Worldwide Semantic Interoperability for IoT. *IEEE Commun. Mag.* **2016**, *54*, 40–46. [[CrossRef](#)]
12. Ramparany, F.; Marquez, F.G.; Soriano, J.; Elsaleh, T. Handling smart environment devices, data and services at the semantic level with the FI-WARE core platform. In Proceedings of the 2014 IEEE International Conference on Big Data (Big Data), Washington, DC, USA, 27–30 October 2014; pp. 14–20.
13. Preventis, A.; Stravoskouflos, K.; Sotiriadis, S.; Petrkakis, E.G.M. Personalized Motion Sensor Driven Gesture Recognition in the FIWARE Cloud Platform. In Proceedings of the 14th International Symposium on Parallel and Distributed Computing, Limassol, Cyprus, 29 June–2 July 2015; pp. 19–26.
14. Fernandez, P.; Santana, J.M.; Ortega, S.; Trujillo, A.; Suárez, J.P.; Domínguez, C.; Santana, J.; Sánchez, A. SmartPort: A Platform for Sensor Data Monitoring in a Seaport Based on FIWARE. *Sensors* **2016**, *16*, 417. [[CrossRef](#)] [[PubMed](#)]
15. Fazio, M.; Celesti, A.; Márquez, F.G.; Glikson, A.; Villari, M. Exploiting the FIWARE Cloud Platform to Develop a Remote Patient Monitoring System. In Proceedings of the 2015 IEEE Symposium on Computers and Communication (ISCC), Larnaca, Cyprus, 6–9 July 2015.
16. FIWARE IoT Agent Provision API Documentation. Available online: <http://docs.telefonicaiotiotagents.apiry.io/#> (accessed on 31 August 2017).
17. FIWARE Orion Context Broker. Available online: <https://github.com/telefonicaid/fiware-orion> (accessed on 31 August 2017).

18. IBM Proactive Technology Online. Available online: <https://github.com/ishkin/Proton> (accessed on 31 August 2017).
19. Etzion, O.; Niblett, P. *Event Processing in Action*; Manning Publications: Greenwich, CT, USA, August 2010.
20. FIWARE Short Time Historic (STH) (Comet). Available online: <https://github.com/telefonicaid/fiware-sth-comet> (accessed on 31 August 2017).
21. ENTROPY Semantic Models. Available online: <http://vocab.sti2.at/entropy/> (accessed on 31 August 2017).
22. Şimşek, U.; Fensel, A.; Zafeiropoulos, A.; Fotopoulou, E.; Liapis, P.; Bouras, T.; Saenz, F.T.; Gómez, A.F.S. A semantic approach towards implementing energy efficient lifestyles through behavioural change. In Proceedings of the 12th International Conference on Semantic Systems (SEMANTiCS 2016), Leipzig, Niedersachsen, Germany, 12–15 September 2016; pp. 173–176.
23. Fotopoulou, E.; Zafeiropoulos, A.; Papaspyros, D.; Hasapis, P.; Tsoliis, G.; Bouras, T.; Mouzakitis, S.; Zanetti, N. Linked Data Analytics in Interdisciplinary Studies: The Health Impact of Air Pollution in Urban Areas. *IEEE Access* **2016**, *4*, 149–164. [[CrossRef](#)]
24. Pooranian, Z.; Nikmehr, N.; Najafi-Ravadanegh, S.; Mahdin, H.; Abawajy, J. Economical and Environmental Operation of Smart Networked Microgrids under Uncertainties Using NSGA-II. In Proceedings of the 24th International Conference on Software, Telecommunications and Computer Networks (SoftCOM), Split, Croatia, 22–24 September 2016; pp. 1–6.
25. Drools Business Rules Management System. Available online: <https://www.drools.org/> (accessed on 31 August 2017).
26. Ordóñez, A.; Eraso, L.; Ordóñez, H.; Merchan, L. Comparing Drools and Ontology Reasoning Approaches for Automated Monitoring in Telecommunication Processes. *Proced. Comput. Sci.* **2016**, *95*, 353–360. [[CrossRef](#)]
27. JBoss Rules Manual. Available online: https://access.redhat.com/documentation/en-US/JBoss_Enterprise_SOA_Platform/4.2/html/JBoss_Rules_Manual/sect-JBoss_Rules_Reference_Manual-Why_use_a_Rule_Engine.html (accessed on 31 August 2017).
28. Hille, P.V.; Jacques, J.; Taillard, J.; Rosier, A.; Delerue, D.; Burgun, A.; Dameron, O. Comparing Drools and Ontology Reasoning Approaches for Telecardiology Decision Support. *Stud. Health Technol. Inf.* **2012**, *180*, 300–304.
29. González-Vidal, A.; Moreno-Cano, V.; Terroso-Sáenz, F.; Skarmeta, A.F. Towards Energy Efficiency Smart Buildings Models Based on Intelligent Data Analytics. *Proced. Comput. Sci.* **2016**, *83*, 994–999. [[CrossRef](#)]
30. American Society of Heating, Refrigerating and Air-Conditioning Engineers (ASHRAE). Chapter 14, Climatic Design Information. In *2017 ASHRAE Handbook—Fundamentals*; ASHRAE Handbook Series; American Society of Heating, Refrigerating and Air-Conditioning Engineers: Atlanta, GA, USA, 2017; ISBN 978-1939200587.
31. R Project for Statistical Computing. Available online: <https://www.r-project.org/> (accessed on 31 August 2017).
32. Apache Spark. Available online: <http://spark.apache.org/> (accessed on 31 August 2017).
33. Apache Spark OpenCPU Executor (ROSE). Available online: <https://github.com/onetapbeyond/opencpu-spark-executor> (accessed on 31 August 2017).
34. ENTROPY H2020 Project Deliverable 2.2. Design of Wireless Sensor Networking and Data Fusion Mechanisms. Available online: <http://entropy-project.eu/wp-content/uploads/2017/05/ENTROPY-D2.2-Design-of-Wireless-Sensor-Networking-and-Data-Fusion-Mechanisms.pdf> (accessed on 31 August 2017).
35. Weather Underground API Service. Available online: <https://www.wunderground.com/weather/api/> (accessed on 31 August 2017).



© 2017 by the authors. Licensee MDPI, Basel, Switzerland. This article is an open access article distributed under the terms and conditions of the Creative Commons Attribution (CC BY) license (<http://creativecommons.org/licenses/by/4.0/>).

BIBLIOGRAPHY

Publications

- [1] Aurora Gonzalez-Vidal, Payam Barnaghi, and Antonio F Skarmeta. Beats: Blocks of eigenvalues algorithm for time series segmentation. *IEEE Transactions on Knowledge and Data Engineering*, 30(11):2051–2064, 2018.
- [2] Aurora Gonzalez-Vidal, Fernando Jimenez, and Antonio F Skarmeta. A methodology for energy multivariate time series forecasting in smart buildings based on feature selection. *Energy and Buildings*, 2019.
- [3] Ignacio Rodríguez-Rodríguez, Aurora González Vidal, Alfonso Ramallo González, and Miguel Zamora. Commissioning of the controlled and automatized testing facility for human behavior and control (casita). *Sensors*, 18(9):2829, 2018.
- [4] M Victoria Moreno, Fernando Terroso-Sáenz, Aurora González-Vidal, Mercedes Valdés-Vela, Antonio F Skarmeta, Miguel A Zamora, and Victor Chang. Applicability of big data techniques to smart cities deployments. *IEEE Transactions on Industrial Informatics*, 13(2):800–809, 2017.
- [5] Fernando Terroso-Saenz, Aurora González-Vidal, Alfonso P Ramallo-González, and Antonio F Skarmeta. An open iot platform for the management and analysis of energy data. *Future Generation Computer Systems*, 2017.
- [6] Eleni Fotopoulou, Anastasios Zafeiropoulos, Fernando Terroso-Sáenz, Umutcan Şimşek, Aurora González-Vidal, George Tsiolis, Panagiotis Gouvas, Paris Liapis, Anna Fensel, and Antonio Skarmeta. Providing personalized energy management and awareness services for energy efficiency in smart buildings. *Sensors*, 17(9):2054, 2017.
- [7] Aurora González-Vidal, Victoria Moreno-Cano, Fernando Terroso-Sáenz, and Antonio F Skarmeta. Towards energy efficiency smart buildings models based on intelligent data analytics. *Procedia Computer Science*, 83:994–999, 2016.

- [8] Aurora González-Vidal, Alfonso P Ramallo-González, and Antonio Skarmeta. Empirical study of massive set-point behavioral data: Towards a cloud-based artificial intelligence that democratizes thermostats. In *2018 IEEE International Conference on Smart Computing (SMARTCOMP)*, pages 211–218. IEEE, 2018.
- [9] Aurora González-Vidal, Alfonso P Ramallo-González, Fernando Terroso-Sáenz, and Antonio Skarmeta. Data driven modeling for energy consumption prediction in smart buildings. In *2017 IEEE International Conference on Big Data (Big Data)*, pages 4562–4569. IEEE, 2017.
- [10] Tomás Vantuch, Aurora González Vidal, Alfonso P Ramallo-González, Antonio F Skarmeta, and Stanislav Misák. Machine learning based electric load forecasting for short and long-term period. In *2018 IEEE 4th World Forum on Internet of Things (WF-IoT)*, pages 511–516. IEEE, 2018.
- [11] Aurora González-Vidal, Victoria Moreno, and Antonio F. Skarmeta. Model-free approach based on iot data analytics for energy efficiency in smart environments. In *2016 European Conference on Networks and Communications EuCNC*, 2016.
- [12] M Victoria Moreno, Fernando Terroso-Sáenz, Aurora González-Vidal, and Antonio F Skarmeta. Data analytics in smart buildings. *Building Blocks for IoT Analytics*, page 167, 2016.
- [13] Fernando Terroso-Sáenz, Jesús Cuenca-Jara, Aurora González-Vidal, and Antonio F Skarmeta. Human mobility prediction based on social media with complex event processing. *International Journal of Distributed Sensor Networks*, 12(9):5836392, 2016.
- [14] Fernando Terroso-Sáenz, Aurora González-Vidal, and Antonio F. Skarmeta. Towards anticipate detection of complex event processing rules with probabilistic modelling. *International Journal of Design & Nature and Ecodynamics*, 11(3):275–283, 2016.
- [15] Fernando Terroso-Sáenz, Mercedes Valdes-Vela, Aurora González-Vidal, and Antonio F Skarmeta. Human mobility modelling based on dense transit areas detection with opportunistic sensing. *Mobile Information Systems*, 2016, 2016.
- [16] Jesus Cuenca-Jara, Fernando Terroso-Sáenz, Mercedes Valdes-Vela, Aurora González-Vidal, and Antonio F Skarmeta. Human mobility analysis based on social media and fuzzy clustering. In *2017 Global Internet of Things Summit (GIoTS)*, pages 1–6. IEEE, 2017.
- [17] Fernando Terroso-Sáenz, Victoria Moreno, Aurora González-Vidal, Miguel Ángel Zamora-Izquierdo, and Antonio F. Skarmeta. Internet de las cosas y gamificación aplicados a eficiencia energética en edificios. In *III Congreso EECN Edificios Energía Casi Nula 2016*, pages 25–30, 2016.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [18] Eleni Fotopoulou, Anastasios Zafeiropoulos, Fernando Terroso, Aurora Gonzalez, Antonio Skarmeta, Umutcan Şimşek, and Anna Fensel. Data aggregation, fusion and recommendations for strengthening citizens energy-aware behavioural profiles. In *2017 Global Internet of Things Summit (GIoTS)*, pages 1–6. IEEE, 2017.
-

References

- [19] Accenture Strategy. # smarter2030: Ict solutions for 21st century challenges. *The Global eSustainability Initiative (GeSI), Brussels, Brussels-Capital Region, Belgium, Tech. Rep*, 2015.
- [20] Evangelos Theodoridis, Georgios Mylonas, and Ioannis Chatzigiannakis. Developing an iot smart city framework. In *IISA 2013*, pages 1–6. IEEE, 2013.
- [21] Dave Evans. The internet of things: How the next evolution of the internet is changing everything. *CISCO white paper*, 1(2011):1–11, 2011.
- [22] Rajendra K Pachauri, Myles R Allen, Vicente R Barros, John Broome, Wolfgang Cramer, Renate Christ, John A Church, Leon Clarke, Qin Dahe, Purnamita Dasgupta, et al. *Climate change 2014: synthesis report. Contribution of Working Groups I, II and III to the fifth assessment report of the Intergovernmental Panel on Climate Change*. IPCC, 2014.
- [23] Jane Marceau. Innovation in the city and innovative cities introduction, 2008.
- [24] María Victoria Moreno Cano, José Santa, Miguel Angel Zamora, and Antonio F Skarmeta Gómez. Context-aware energy efficiency in smart buildings. In *Ubiquitous Computing and Ambient Intelligence. Context-Awareness and Context-Driven Interaction*, pages 1–8. Springer, 2013.
- [25] EC. Benchmarking smart metering deployment in the eu-27 with a focus on electricity. *Report from the Commission*, 2014.
- [26] N. Mogles, I. Walker, A. P. Ramallo-González, S. Lee, J. and Natarajan, J. Padget, E. Gabe-Thomas, T. Lovett, S. Ren, G. and Hyniewska, E. O'Neill, R. Hourizi, and D. Coley. How smart do smart meters need to be? *Building and Environment*, (125):439–450, 2017.
- [27] Raghunath Nambiar, Rajesh Shroff, and Shane Handy. Smart cities: Challenges and opportunities. In *2018 10th International Conference on Communication Systems & Networks (COMSNETS)*, pages 243–250. IEEE, 2018.
- [28] Luis Pérez-Lombard, José Ortiz, and Christine Pout. A review on buildings energy consumption information. *Energy and buildings*, 40(3):394–398, 2008.

- [29] Monthly energy review. *U.S. Energy Information Administration*, September, 2017.
- [30] CABA. Intelligent buildings and the impact of the internet of things. *Landmark Research Project. Executive Summary*, pages 1–18, 2017.
- [31] Elham Delzendeh, Song Wu, Angela Lee, and Ying Zhou. The impact of occupants' behaviours on building energy analysis: A research review. *Renewable and Sustainable Energy Reviews*, 80:1061–1071, 2017.
- [32] Rishee K Jain, Rimas Gulbinas, John E Taylor, and Patricia J Culligan. Can social influence drive energy savings? detecting the impact of social influence on the energy consumption behavior of networked users exposed to normative eco-feedback. *Energy and Buildings*, 66:119–127, 2013.
- [33] Anca-Diana Barbu, Nigel Griffiths, and Gareth Morton. Achieving energy efficiency through behaviour change: what does it take. *European Environment Agency (EEA), Copenhagen*, 2013.
- [34] Maarten De Groote, Jonathan Volt, and Frances Bean. Is europe ready for the smart buildings revolution? *Building Performance Institute Europe. Retrieved July*, 7:2017, 2017.
- [35] Jennifer King and Christopher Perry. *Smart buildings: Using smart technology to save energy in existing buildings*. Amercian Council for an Energy-Efficient Economy, 2017.
- [36] Doug Laney. 3d data management: Controlling data volume, velocity and variety. *META group research note*, 6(70):1, 2001.
- [37] Muhammad Fahim Uddin, Navarun Gupta, et al. Seven v's of big data understanding big data to extract value. In *Proceedings of the 2014 Zone 1 Conference of the American Society for Engineering Education*, pages 1–5. IEEE, 2014.
- [38] John Gantz and David Reinsel. The digital universe in 2020: Big data, bigger digital shadows, and biggest growth in the far east. *IDC iView: IDC Analyze the future*, 2007:1–16, 2012.
- [39] David L Hall and James Llinas. An introduction to multisensor data fusion. *Proceedings of the IEEE*, 85(1):6–23, 1997.
- [40] Meisong Wang, Charith Perera, Prem Prakash Jayaraman, Miranda Zhang, Peter Strazdins, RK Shyamsundar, and Rajiv Ranjan. City data fusion: Sensor data fusion in the internet of things. *International Journal of Distributed Systems and Technologies (IJDST)*, 7(1):15–36, 2016.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [41] Carnot Institutes. White paper: Smart networked objects and internet of things. *Information Communication Technologies and Micro Nano Technologies alliance, White Paper*; (Jan. 2011), 2011.
- [42] Dan Feldman, Melanie Schmidt, and Christian Sohler. Turning big data into tiny data: Constant-size coresets for k-means, pca and projective clustering. In *Proceedings of the twenty-fourth annual ACM-SIAM symposium on Discrete algorithms*, pages 1434–1453. Society for Industrial and Applied Mathematics, 2013.
- [43] Marvin Weinstein, F Meirer, A Hume, Ph Sciau, G Shaked, R Hofstetter, Erez Persi, A Mehta, and David Horn. Analyzing big data with dynamic quantum clustering. *arXiv preprint arXiv:1310.2700*, 2013.
- [44] Muhammad Habib ur Rehman, Chee Sun Liew, Assad Abbas, Prem Prakash Jayaraman, Teh Ying Wah, and Samee U Khan. Big data reduction methods: a survey. *Data Science and Engineering*, 1(4):265–284, 2016.
- [45] Rubén García-Pajares, José M Benítez, and GI Sainz-Palmero. Frasel: a consensus of feature ranking methods for time series modelling. *Soft Computing*, 17(8):1489–1510, 2013.
- [46] Selwyn Piramuthu. Evaluating feature selection methods for learning in data mining applications. *European journal of operational research*, 156(2):483–494, 2004.
- [47] Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine learning research*, 7(Jan):1–30, 2006.
- [48] H Burak Gunay, Weiming Shen, and Guy Newsham. Data analytics to improve building performance: A critical review. *Automation in Construction*, 97:96–109, 2019.
- [49] Marti Frank, Hannah Friedman, Kristin Heinemeier, Cory Toole, David Claridge, Natascha Castro, and Philip Haves. State-of-the-art review for commissioning low energy buildings: Existing cost/benefit and persistence methodologies and data, state of development of automated tools and assessment of needs for commissioning zeb. *NISTIR*, 7356:2007, 2007.
- [50] Sool Yeon Cho. *The persistence of savings obtained from commissioning of existing buildings*. PhD thesis, 2008.
- [51] Fateme Fahiman, Sarah M Erfani, Sutharshan Rajasegarar, Marimuthu Palaniswami, and Christopher Leckie. Improving load forecasting based on deep learning and k-shape clustering. In *2017 International Joint Conference on Neural Networks (IJCNN)*, pages 4134–4141. IEEE, 2017.

- [52] Christos N Schizas, Soteris A Kalogirou, and Costas Neocleous. Artificial neural networks in modelling the heat-up response of a solar steam generating plant. 1996.
- [53] SA Kalogirou, CC Neocleous, and CN Schizas. Building heating load estimation using artificial neural networks. In *Proceedings of the 17th international conference on Parallel architectures and compilation techniques*, volume 8, page 14, 1997.
- [54] Soteris A Kalogirou. Applications of artificial neural-networks for energy systems. *Applied Energy*, 67(1):17–35, 2000.
- [55] D MacKay. Bayesian non-linear modeling for the 1993 energy prediction competition. *Maximum Entropy and Bayesian Methods*, pages 221–234, 1993.
- [56] Fazli Wahid and Do-Hyeun Kim. Prediction methodology of energy consumption based on random forest classifier in korean residential apartments. 2015.
- [57] Yangyang Fu, Zhengwei Li, Hao Zhang, and Peng Xu. Using support vector machine to predict next day electricity load of public buildings with sub-metering devices. *Procedia Engineering*, 121:1016–1022, 2015.
- [58] Christopher KI Williams and David Barber. Bayesian classification with gaussian processes. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 20(12):1342–1351, 1998.
- [59] Douglas J Leith, Martin Heidl, and John V Ringwood. Gaussian process prior models for electrical load forecasting. *Probabilistic Methods Applied to Power Systems*, pages 112–117, 2004.
- [60] Sunil Mamidi, Yu-Han Chang, and Rajiv Maheswaran. Improving building energy efficiency with a network of sensing, learning and prediction agents. In *Proceedings of the 11th International Conference on Autonomous Agents and Multiagent Systems-Volume 1*, pages 45–52. International Foundation for Autonomous Agents and Multiagent Systems, 2012.
- [61] Minoru Kawashima, Charles E Dorgan, and John W Mitchell. Hourly thermal load prediction for the next 24 hours by arima, ewma, lr and an artificial neural network. Technical report, American Society of Heating, Refrigerating and Air-Conditioning Engineers . . . , 1995.
- [62] X. Yan. *Linear Regression Analysis: Theory and Computing*. World Scientific Publishing Company Pte Limited, 2009.
- [63] Coskun Hamzacebi and Huseyin Avni Es. Forecasting the annual electricity consumption of turkey using an optimized grey model. *Energy*, 70:165–171, 2014.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [64] ASHRAE Fundamentals Handbook. American society of heating, refrigerating and air-conditioning engineers. *Inc.: Atlanta, GA, USA*, 2017.
- [65] G Burnand. The study of the thermal behaviour of structures by electrical analogy. *British Journal of Applied Physics*, 3(2):50, 1952.
- [66] Peder Bacher and Henrik Madsen. Identifying suitable models for the heat dynamics of buildings. *Energy and Buildings*, 43(7):1511–1522, 2011.
- [67] Alfonso P Ramallo-González, Matthew Brown, Elizabeth Gabe-Thomas, Tom Lovett, and David A Coley. The reliability of inverse modelling for the wide scale characterization of the thermal properties of buildings. *Journal of Building Performance Simulation*, pages 1–19, 2017.
- [68] Gianluca Serale, Massimo Fiorentini, Alfonso Capozzoli, Daniele Bernardini, and Alberto Bemporad. Model predictive control (mpc) for enhancing building and hvac system energy efficiency: Problem formulation, applications and opportunities. *Energies*, 11(3):631, 2018.
- [69] Sama Aghniaey, Thomas M Lawrence, Javad Mohammadpour, WenZhan Song, Richard T Watson, and Marie C Boudreau. Optimizing thermal comfort considerations with electrical demand response program implementation. *Building Services Engineering Research and Technology*, 39(2):219–231, 2018.
- [70] Yuehong Lu, Shengwei Wang, Yongjun Sun, and Chengchu Yan. Optimal scheduling of buildings with energy generation and thermal energy storage under dynamic electricity pricing using mixed-integer nonlinear programming. *Applied Energy*, 147:49–58, 2015.
- [71] Peter Palensky and Dietmar Dietrich. Demand side management: Demand response, intelligent energy systems, and smart loads. *Industrial Informatics, IEEE Transactions on*, 7(3):381–388, 2011.
- [72] Yang-Seon Kim and Jelena Srebric. Improvement of building energy simulation accuracy with occupancy schedules derived from hourly building electricity consumption. *Ashrae Transactions*, 121(1):353–361, 2015.
- [73] Dean Abbott. *Applied predictive analytics: principles and techniques for the professional data analyst*. John Wiley & Sons, 2014.
- [74] Eamonn J Keogh and Michael J Pazzani. An enhanced representation of time series which allows fast and accurate classification, clustering and relevance feedback. *KDD*, 98:239–243, 1998.
- [75] Eamonn Keogh, Selina Chu, David Hart, and Michael Pazzani. Segmenting time series: A survey and novel approach. *Data mining in time series databases*, 57:1–22, 2004.

- [76] Hafzullah Aksoy, Abdullah Gedikli, N Erdem Unal, and Athanasios Kehagias. Fast segmentation algorithms for long hydrometeorological time series. *Hydrological Processes*, 22(23):4600–4608, 2008.
- [77] Xiaopeng Xi, Eamonn Keogh, Christian Shelton, Li Wei, and Chotirat Ann Ratanamahatana. Fast time series classification using numerosity reduction. In *Proceedings of the 23rd international conference on Machine learning*, pages 1033–1040. ACM, 2006.
- [78] Xiaoyue Wang, Abdullah Mueen, Hui Ding, Goce Trajcevski, Peter Scheuermann, and Eamonn Keogh. Experimental comparison of representation methods and distance measures for time series data. *Data Mining and Knowledge Discovery*, pages 1–35, 2013.
- [79] Yuelong Zhu, De Wu, and Shijin Li. A piecewise linear representation method of time series based on feature points. In *International Conference on Knowledge-Based and Intelligent Information and Engineering Systems*, pages 1066–1072. Springer, 2007.
- [80] Eamonn J Keogh and Michael J Pazzani. A simple dimensionality reduction technique for fast similarity search in large time series databases. In *Pacific-Asia conference on knowledge discovery and data mining*, pages 122–133. Springer, 2000.
- [81] Ngoc Thanh Nguyen, Bogdan Trawiński, Radosław Katarzyniak, and Geun-Sik Jo. *Advanced methods for computational collective intelligence*, volume 457. Springer, 2012.
- [82] Houtao Deng, George Runger, Eugene Tuv, and Martyanov Vladimir. A time series forest for classification and feature extraction. *Information Sciences*, 239:142–153, 2013.
- [83] Mustafa Gokce Baydogan and George Runger. Time series representation and similarity based on local autopatterns. *Data Mining and Knowledge Discovery*, 30(2):476–509, 2016.
- [84] Jessica Lin, Eamonn Keogh, Li Wei, and Stefano Lonardi. Experiencing sax: a novel symbolic representation of time series. *Data Mining and knowledge discovery*, 15(2):107–144, 2007.
- [85] Jessica Lin, Eamonn Keogh, Li Wei, and Stefano Lonardi. Experiencing sax: a novel symbolic representation of time series. *Data Mining and knowledge discovery*, 15(2):107–144, 2007.
- [86] Anthony Bagnall, Jason Lines, Jon Hills, and Aaron Bostrom. Time-series classification with cote: the collective of transformation-based ensembles. *IEEE Transactions on Knowledge and Data Engineering*, 27(9):2522–2535, 2015.
- [87] Ron Kohavi and George H John. Wrappers for feature subset selection. *Artificial intelligence*, 97(1-2):273–324, 1997.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [88] Nathalie Japkowicz and Mohak Shah. *Evaluating learning algorithms: a classification perspective*. Cambridge University Press, 2011.
- [89] Asha Gowda Karegowda, AS Manjunath, and MA Jayaram. Comparative study of attribute selection using gain ratio and correlation based feature selection. *International Journal of Information Technology and Knowledge Management*, 2(2):271–277, 2010.
- [90] Amir Ahmad and Lipika Dey. A feature selection technique for classificatory analysis. *Pattern Recognition Letters*, 26(1):43–56, 2005.
- [91] StevenL. Salzberg. C4.5: Programs for machine learning by j. ross quinlan. morgan kaufmann publishers, inc., 1993. *Machine Learning*, 16(3):235–240, 1994.
- [92] Hai-Xiang Zhao and Frédéric Magoulès. Feature selection for predicting building energy consumption based on statistical learning method. *Journal of Algorithms & Computational Technology*, 6(1):59–77, 2012.
- [93] Oscar Utterbäck. Feature selection methods with applications in electrical load forecasting. *Master's Theses in Mathematical Sciences*, 2017.
- [94] Sven F Crone and Nikolaos Kourentzes. Feature selection for time series prediction—a combined filter and wrapper approach for neural networks. *Neurocomputing*, 73(10–12):1923–1936, 2010.
- [95] Youqiang Sun, Jiuyong Li, Jixue Liu, Christopher Chow, Bingyu Sun, and Rujing Wang. Using causal discovery for feature selection in multivariate numerical time series. *Machine Learning*, 101(1-3):377–395, 2015.
- [96] Shohei Hido and Tetsuro Morimura. Temporal feature selection for time-series prediction. In *2012 21st International Conference on Pattern Recognition (ICPR 2012)*, pages 3557–3560. IEEE, 2012.
- [97] O García-Hinde, Vanessa Gómez-Verdejo, Manel Martínez-Ramón, Carlos Casanova-Mateo, J Sanz-Justo, Silvia Jiménez-Fernández, and Sancho Salcedo-Sanz. Feature selection in solar radiation prediction using bootstrapped svrs. In *Evolutionary Computation (CEC), 2016 IEEE Congress on*, pages 3638–3645. IEEE, 2016.
- [98] Daniel O'Leary and Joel Kubby. Feature selection and ann solar power prediction. *Journal of Renewable Energy*, 2017, 2017.
- [99] Cong Feng, Mingjian Cui, Bri-Mathias Hodge, and Jie Zhang. A data-driven multi-model methodology with deep feature selection for short-term wind forecasting. *Applied Energy*, 190:1245–1257, 2017.

- [100] Rubén García Pajares, Jose M Benítez, and Gregorio Sáinz Palmero. Feature selection for time series forecasting: A case study. In *Hybrid Intelligent Systems, 2008. HIS'08. Eighth International Conference on*, pages 555–560. IEEE, 2008.
- [101] Eija Ferreira. Model selection in time series machine learning applications. *PhD Thesis. University of Oulu Graduate School. University of Oul*, 2015.
- [102] Ling Tang, Chenghao Wang, and Shuai Wang. Energy time series data analysis based on a novel integrated data characteristic testing approach. *Procedia Computer Science*, 17:759–769, 2013.
- [103] Irena Koprinska, Mashud Rana, and Vassilios G Agelidis. Correlation and instance based feature selection for electricity load forecasting. *Knowledge-Based Systems*, 82:29–40, 2015.
- [104] Francisco Martínez-Álvarez, Alicia Troncoso, Gualberto Asencio-Cortés, and José C Riquelme. A survey on data mining techniques applied to electricity-related time series forecasting. *Energies*, 8(11):13162–13193, 2015.
- [105] Cyril Goutte. Note on free lunches and cross-validation. *Neural Computation*, 9(6):1245–1249, 1997.
- [106] Richard De Dear and Gail Schiller Brager. The adaptive model of thermal comfort and energy conservation in the built environment. *International journal of biometeorology*, 45(2):100–108, 2001.
- [107] Anna Carolina Menezes, Andrew Cripps, Dino Bouchlaghem, and Richard Buswell. Predicted vs. actual energy performance of non-domestic buildings: Using post-occupancy evaluation data to reduce the performance gap. *Applied energy*, 97:355–364, 2012.
- [108] Ian Richardson, Murray Thomson, and David Infield. A high-resolution domestic building occupancy model for energy demand simulations. *Energy and buildings*, 40(8):1560–1566, 2008.
- [109] Jean Rouleau, Alfonso Ramallo-González, and Louis Gosselin. Towards a comprehensive tool to model occupant behaviour for dwellings that combines domestic hot water use with active occupancy. In *Proceedings of the 15th IBPSA Conference, San Francisco, CA, USA*, pages 7–9, 2017.
- [110] Nataliya Mogles, Ian Walker, Alfonso P Ramallo-González, JeeHang Lee, Sukumar Nataraajan, Julian Padget, Elizabeth Gabe-Thomas, Tom Lovett, Gang Ren, Sylwia Hyniewska, et al. How smart do smart meters need to be? *Building and Environment*, 125:439–450, 2017.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [111] H Burak Gunay, William O'Brien, Ian Beausoleil-Morrison, and Jayson Bursill. Development and implementation of a thermostat learning algorithm. *Science and Technology for the Built Environment*, 24(1):43–56, 2018.
- [112] John Goins and Mithra Moezzi. Linking occupant complaints to building performance. *Building Research & Information*, 41(3):361–372, 2013.
- [113] Frédéric Haldi and Darren Robinson. On the behaviour and adaptation of office occupants. *Building and environment*, 43(12):2163–2177, 2008.
- [114] H Burak Gunay, William O'Brien, Ian Beausoleil-Morrison, and Sara Gilani. Development and implementation of an adaptive lighting and blinds control algorithm. *Building and Environment*, 113:185–199, 2017.
- [115] Sarah Darby. Smart metering: what potential for householder engagement? *Building Research & Information*, 38(5):442–457, 2010.
- [116] Tom Hargreaves, Michael Nye, and Jacqueline Burgess. Making energy visible: A qualitative field study of how householders interact with feedback from smart energy monitors. *Energy policy*, 38(10):6111–6119, 2010.
- [117] Directive (eu) 2018/844 of the european parliament and of the council amending directive 2010/31/eu on the energy performance of buildings and directive 2012/27/eu on energy efficiency. *Official Journal of the European Union*, L156:75–91, 2018.
- [118] Maitreyee Dey, Manik Gupta, Mikdam Turkey, and Sandra Dudley. Unsupervised learning techniques for hvac terminal unit behaviour analysis. In *2017 IEEE SmartWorld, Ubiquitous Intelligence & Computing, Advanced & Trusted Computed, Scalable Computing & Communications, Cloud & Big Data Computing, Internet of People and Smart City Innovation (SmartWorld/SCALCOM/UIC/ATC/CBDCom/IOP/SCI)*, pages 1–7. IEEE, 2017.
- [119] Maitreyee Dey, Soumya Prakash Rana, and Sandra Dudley. Smart building creation in large scale hvac environments through automated fault detection and diagnosis. *Future Generation Computer Systems*, 2018.
- [120] Michelle Shipworth, Steven K Firth, Michael I Gentry, Andrew J Wright, David T Shipworth, and Kevin J Lomas. Central heating thermostat settings and timing: building demographics. *Building Research & Information*, 38(1):50–69, 2010.
- [121] Rune Vinther Andersen, Bjarne W Olesen, and Jørn Toftum. Modelling occupants' heating set-point preferences. In *Building Simulation 2011: 12th Conference of International Building Performance Simulation Association*, 2011.

- [122] Yu Zheng and Xiaofang Zhou. *Computing with spatial trajectories*. Springer Science & Business Media, 2011.
- [123] Tian Guo, Zhixian Yan, and Karl Aberer. An adaptive approach for online segmentation of multi-dimensional mobile data. In *Proceedings of the Eleventh ACM International Workshop on Data Engineering for Wireless and Mobile Access*, pages 7–14. ACM, 2012.
- [124] Nicholas Jing Yuan, Yingzi Wang, Fuzheng Zhang, Xing Xie, and Guangzhong Sun. Reconstructing individual mobility from smart card transactions: A space alignment approach. In *2013 IEEE 13th International Conference on Data Mining*, pages 877–886. IEEE, 2013.
- [125] Huiji Gao and Huan Liu. Mining human mobility in location-based social networks. *Synthesis Lectures on Data Mining and Knowledge Discovery*, 7(2):1–115, 2015.
- [126] Miao Lin and Wen-Jing Hsu. Mining gps data for mobility patterns: A survey. *Pervasive and Mobile Computing*, 12:1–16, 2014.
- [127] Luca Pappalardo, Filippo Simini, Salvatore Rinzivillo, Dino Pedreschi, Fosca Giannotti, and Albert-László Barabási. Returners and explorers dichotomy in human mobility. *Nature communications*, 6:8166, 2015.
- [128] Miao Lin, Wen-Jing Hsu, and Zhuso Qi Lee. Predictability of individuals’ mobility with high-resolution positioning data. In *Proceedings of the 2012 ACM conference on ubiquitous computing*, pages 381–390. ACM, 2012.
- [129] Hoyoung Jeung, Man Lung Yiu, Xiaofang Zhou, and Christian S Jensen. Path prediction and predictive range querying in road network databases. *The VLDB Journal*, 19(4):585–602, 2010.
- [130] John Krumm, Robert Gruen, and Daniel Delling. From destination prediction to route prediction. *Journal of Location Based Services*, 7(2):98–120, 2013.
- [131] Disheng Qiu, Paolo Papotti, and Lorenzo Blanco. Future locations prediction with uncertain data. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 417–432. Springer, 2013.
- [132] Quan Yuan, Gao Cong, Zongyang Ma, Aixin Sun, and Nadia Magnenat Thalmann. Who, where, when and what: discover spatio-temporal topics for twitter users. In *Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 605–613. ACM, 2013.
- [133] Anastasios Noulas, Salvatore Scellato, Neal Lathia, and Cecilia Mascolo. Mining user mobility features for next place prediction in location-based services. In *2012 IEEE 12th international conference on data mining*, pages 1038–1043. IEEE, 2012.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [134] Natalia Andrienko, Gennady Andrienko, Georg Fuchs, and Piotr Jankowski. Scalable and privacy-respectful interactive discovery of place semantics from human mobility traces. *Information Visualization*, 15(2):117–153, 2016.
- [135] Vanessa Frias-Martinez and Enrique Frias-Martinez. Spectral clustering for sensing urban land use using twitter activity. *Engineering Applications of Artificial Intelligence*, 35:237–245, 2014.
- [136] Samiul Hasan, Xianyuan Zhan, and Satish V Ukkusuri. Understanding urban human activity and mobility patterns using large-scale location-based data from online social media. In *Proceedings of the 2nd ACM SIGKDD international workshop on urban computing*, page 6. ACM, 2013.
- [137] Takeshi Sakaki, Makoto Okazaki, and Yutaka Matsuo. Tweet analysis for real-time event detection and earthquake reporting system development. *IEEE Transactions on Knowledge and data engineering*, 25(4):919–931, 2012.
- [138] Ruchi Parikh and Kamalakar Karlapalem. Et: events from tweets. In *Proceedings of the 22nd international conference on world wide web*, pages 613–620. ACM, 2013.
- [139] Alan Ritter, Oren Etzioni, Sam Clark, et al. Open domain event extraction from twitter. In *Proceedings of the 18th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 1104–1112. ACM, 2012.
- [140] Calin Railean, Monica Borda, and Alexandra Moraru. Complex event processing in social media. *Acta Technica Napocensis*, 55(3):10, 2014.
- [141] Eleonora D’Andrea, Pietro Ducange, Beatrice Lazzerini, and Francesco Marcelloni. Real-time detection of traffic from twitter stream analysis. *IEEE transactions on intelligent transportation systems*, 16(4):2269–2283, 2015.
- [142] Fei Wu, Zhenhui Li, Wang-Chien Lee, Hongjian Wang, and Zhuojie Huang. Semantic annotation of mobility data using social media. In *Proceedings of the 24th International Conference on World Wide Web*, pages 1253–1263. International World Wide Web Conferences Steering Committee, 2015.
- [143] Bartosz Hawelka, Izabela Sitko, Euro Beinat, Stanislav Sobolevsky, Pavlos Kazakopoulos, and Carlo Ratti. Geo-located twitter as proxy for global mobility patterns. *Cartography and Geographic Information Science*, 41(3):260–271, 2014.
- [144] Mordechai Haklay and Patrick Weber. Openstreetmap: User-generated street maps. *IEEE Pervasive Computing*, 7(4):12–18, 2008.

- [145] Shane B Eisenman, Emiliano Miluzzo, Nicholas D Lane, Ronald A Peterson, Gahng-Seop Ahn, and Andrew T Campbell. Bikenet: A mobile sensing system for cyclist experience mapping. *ACM Transactions on Sensor Networks (TOSN)*, 6(1):6, 2009.
- [146] Linjuan Zhang, Deyun Gao, Weicheng Zhao, and Han-Chieh Chao. A multilevel information fusion approach for road congestion detection in vanets. *Mathematical and Computer Modelling*, 58(5-6):1206–1221, 2013.
- [147] Jiafu Wan, Jianqi Liu, Zehui Shao, Athanasios Vasilakos, Muhammad Imran, and Keliang Zhou. Mobile crowd sensing for traffic prediction in internet of vehicles. *Sensors*, 16(1):88, 2016.
- [148] Ármin Petkovics and Károly Farkas. Efficient event detection in public transport tracking. In *2014 International Conference on Telecommunications and Multimedia (TEMU)*, pages 74–79. IEEE, 2014.
- [149] Shaohan Hu, Lu Su, Hengchang Liu, Hongyan Wang, and Tarek Abdelzaher. Smartroad: a crowd-sourced traffic regulator detection and identification system. In *Proceedings of the 12th international conference on Information processing in sensor networks*, pages 331–332. ACM, 2013.
- [150] Bryce W Sharman and Matthew J Roorda. Analysis of freight global positioning system data: clustering approach for identifying trip destinations. *Transportation Research Record*, 2246(1):83–91, 2011.
- [151] Zhenhui Li, Jiawei Han, Bolin Ding, and Roland Kays. Mining periodic behaviors of object movements for animal and biological sustainability studies. *Data Mining and Knowledge Discovery*, 24(2):355–386, 2012.
- [152] Opher Etzion, Peter Niblett, and David C Luckham. *Event processing in action*. Manning Greenwich, 2011.
- [153] Sebastian Stipkovic, Ralf Bruns, and Jürgen Dunkel. Pervasive computing by mobile complex event processing. In *2013 IEEE 10th International Conference on e-Business Engineering*, pages 318–323. IEEE, 2013.
- [154] Lu-An Tang, Yu Zheng, Jing Yuan, Jiawei Han, Alice Leung, Chih-Chieh Hung, and Wen-Chih Peng. On discovery of traveling companions from streaming trajectories. In *2012 IEEE 28th International Conference on Data Engineering*, pages 186–197. IEEE, 2012.
- [155] Andreas Bauer and Christian Wolff. An event processing approach to text stream analysis: basic principles of event based information filtering. In *Proceedings of the 8th acm international conference on distributed event-based systems*, pages 35–46. ACM, 2014.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [156] Kaile Zhou, Chao Fu, and Shanlin Yang. Big data driven smart energy management: From big data to big insights. *Renewable and Sustainable Energy Reviews*, 56:215–225, 2016.
- [157] Yogesh Simmhan, Saima Aman, Alok Kumbhare, Rongyang Liu, Sam Stevens, Qunzhi Zhou, and Viktor Prasanna. Cloud-based software platform for big data analytics in smart grids. *Computing in Science & Engineering*, 15(4):38, 2013.
- [158] Alok Kumbhare, Yogesh Simmhan, and Viktor Prasanna. Cryptonite: a secure and performant data repository on public clouds. In *2012 IEEE Fifth International Conference on Cloud Computing*, pages 510–517. IEEE, 2012.
- [159] Atsushi Ishii and Toyotaro Suzumura. Elastic stream computing with clouds. In *2011 IEEE 4th International Conference on Cloud Computing*, pages 195–202. IEEE, 2011.
- [160] Laura Klein, Jun-young Kwak, Geoffrey Kavulya, Farrokh Jazizadeh, Burcin Becerik-Gerber, Pradeep Varakantham, and Milind Tambe. Coordinating occupant behavior for building energy and comfort management using multi-agent systems. *Automation in construction*, 22:525–536, 2012.
- [161] Nan Li, Jun-young Kwak, Burcin Becerik-Gerber, and Milind Tambe. Predicting hvac energy consumption in commercial buildings using multiagent systems. In *ISARC. Proceedings of the International Symposium on Automation and Robotics in Construction*, volume 30, page 1. IAARC Publications, 2013.
- [162] Cristina Alcaraz, Isaac Agudo, David Nunez, and Javier Lopez. Managing incidents in smart grids a la cloud. In *2011 IEEE Third International Conference on Cloud Computing Technology and Science*, pages 527–531. IEEE, 2011.
- [163] M Moreno, Benito Úbeda, Antonio Skarmeta, and Miguel Zamora. How can we tackle energy efficiency in iot basedsmart buildings? *Sensors*, 14(6):9582–9614, 2014.
- [164] Julien Mineraud, Oleksiy Mazhelis, Xiang Su, and Sasu Tarkoma. A gap analysis of internet-of-things platforms. *Computer Communications*, 89:5–16, 2016.
- [165] Linda Steg, Goda Perlaviciute, and Ellen van der Werff. Understanding the human dimensions of a sustainable energy transition. *Frontiers in psychology*, 6:805, 2015.
- [166] Sam C Staddon, Chandrika Cycil, Murray Goulden, Caroline Leygue, and Alexa Spence. Intervening to change behaviour and save energy in the workplace: A systematic review of available evidence. *Energy Research & Social Science*, 17:30–51, 2016.
- [167] Kathy Kuntz, Rajan Shukla, and I Bensch. How many points for that? a game-based approach to environmental sustainability. *Proceedings of the American Council for an*

- Energy-Efficient Economy Summer Study on Energy Efficiency in Buildings*, 7:126–137, 2012.
- [168] Brian Orland, Nilam Ram, Dean Lang, Kevin Houser, Nate Kling, and Michael Coccia. Saving energy in an office environment: A serious game intervention. *Energy and Buildings*, 74:43–52, 2014.
 - [169] Graham N Dixon, Mary Beth Deline, Katherine McComas, Lauren Chambliss, and Michael Hoffmann. Using comparative feedback to influence workplace energy conservation: A case study of a university campaign. *Environment and Behavior*, 47(6):667–693, 2015.
 - [170] Andreas Kamilaris, Jodi Neovino, Sekhar Kondepudi, and Balaji Kalluri. A case study on the individual energy use of personal computers in an office setting and assessment of various feedback types toward energy savings. *Energy and Buildings*, 104:73–86, 2015.
 - [171] Stephanie N Timm and Brian M Deal. Effective or ephemeral? the role of energy information dashboards in changing occupant energy behaviors. *Energy Research & Social Science*, 19:11–20, 2016.
 - [172] Liga Poznaka, Ilze Laicane, Dagnija Blumberga, Andra Blumberga, and Marika Rosa. Analysis of electricity user behavior: case study based on results from extended household survey. *Energy Procedia*, 72:79–86, 2015.
 - [173] Brian Thomas and Diane Cook. Activity-aware energy-efficient automation of smart buildings. *Energies*, 9(8):624, 2016.
 - [174] Tianzhen Hong, Sarah C Taylor-Lange, Simona D’Oca, Da Yan, and Stefano P Corgnati. Advances in research and applications of energy-related occupant behavior in buildings. *Energy and buildings*, 116:694–702, 2016.
 - [175] Li Da Xu, Wu He, and Shancang Li. Internet of things in industries: A survey. *IEEE Transactions on industrial informatics*, 10(4):2233–2243, 2014.
 - [176] Ian T Jolliffe and Jorge Cadima. Principal component analysis: a review and recent developments. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 374(2065):20150202, 2016.
 - [177] Zhongde Wang. Fast algorithms for the discrete w transform and for the discrete fourier transform. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 32(4):803–816, 1984.
 - [178] K Ramamohan Rao and Ping Yip. *Discrete cosine transform: algorithms, advantages, applications*. Academic press, 2014.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [179] Gregory K Wallace. The jpeg still picture compression standard. *IEEE transactions on consumer electronics*, 38(1):xviii–xxxiv, 1992.
- [180] Alan C Bovik. *The essential guide to image processing*. Academic Press, 2009.
- [181] Martin Hirzel, Henrique Andrade, Bugra Gedik, Gabriela Jacques-Silva, Rohit Khandekar, Vibhore Kumar, Mark Mendell, Howard Nasgaard, Scott Schneider, Robert Soulé, et al. Ibm streams processing language: Analyzing big data in motion. *IBM Journal of Research and Development*, 57(3/4):7–1, 2013.
- [182] Lei Li, Farzad Noorian, Duncan JM Moss, and Philip HW Leong. Rolling window time series prediction using mapreduce. In *Information Reuse and Integration (IRI), 2014 IEEE 15th International Conference on*, pages 757–764. IEEE, 2014.
- [183] Leo Breiman. Random forests. *Machine learning*, 45(1):5–32, 2001.
- [184] Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian Processes for Machine Learning (Adaptive Computation and Machine Learning)*. The MIT Press, 2005.
- [185] S Taylor and B Letham. Prophet: Automatic forecasting procedure. *R package version 0.2*, 1, 2017.
- [186] C. J. Willmott and K. Matsuura. Advantages of the mean absolute error (MAE) over the root mean square error (RMSE) in assessing average model performance. *Climate Research*, 30(1):79–82, 2005. cited By (since 1996)149.
- [187] T Agami Reddy, Namir F Saman, David E Claridge, Jeff S Haberl, W Dan Turner, and Alan T Chalifoux. Baseline methodology for facility-level monthly energy use-part 1: Theoretical aspects. In *ASHRAE transactions*, pages 336–347. ASHRAE, 1997.
- [188] Franklin L Quilumba, Wei-Jen Lee, Heng Huang, David Y Wang, and Robert L Szabados. Using smart meter data to improve the accuracy of intraday load forecasting considering customer behavior similarities. *IEEE Transactions on Smart Grid*, 6(2):911–918, 2015.
- [189] Cheng Fan, Fu Xiao, and Shengwei Wang. Development of prediction models for next-day building energy consumption and peak power demand using data mining techniques. *Applied Energy*, 127:1–10, 2014.
- [190] Richard E Edwards, Joshua New, and Lynne E Parker. Predicting future hourly residential electrical consumption: A machine learning case study. *Energy and Buildings*, 49:591–603, 2012.
- [191] Jeremy Miles Andy Field and Zoe Field Niblett. *Discovering Statistics Using R*. Sage Publications Ltd, 1st edition, 2012.

- [192] Alfonso P Ramallo-González, Matthew Brown, and David A Coley. Identifying the ideal topology of simple models to represent dwellings.
- [193] DA Coley and JM Penman. Second order system identification in the thermal response of real buildings. paper ii: recursive formulation for on-line building energy management and control. *Building and Environment*, 27(3):269–277, 1992.
- [194] F Jiménez, G Sánchez, JM García, G Sciavicco, and L Miralles. Multi-objective evolutionary feature selection for online sales forecasting. *Neurocomputing*, 2016.
- [195] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice-Hall, 2 edition, 2003.
- [196] Fernando Jiménez, Gracia Sánchez, and José M Juárez. Multi-objective evolutionary algorithms for fuzzy classification in survival prediction. *Artificial intelligence in medicine*, 60(3):197–219, 2014.
- [197] Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and TAMT Meyarivan. A fast and elitist multiobjective genetic algorithm: Nsga-ii. *IEEE transactions on evolutionary computation*, 6(2):182–197, 2002.
- [198] F Jiménez, G Sánchez, JM García, G Sciavicco, and L Miralles. Multi-objective evolutionary feature selection for online sales forecasting. *Neurocomputing*, 234:75–92, 2017.
- [199] Eckart Zitzler, Kalyanmoy Deb, and Lothar Thiele. Comparison of multiobjective evolutionary algorithms: Empirical results. *Evolutionary computation*, 8(2):173–195, 2000.
- [200] E. Zitzler, L. Thiele, M. Laumanns, C.M. Fonseca, and V. Grunert da Fonseca. Performance assessment of multiobjective optimizers: An analysis and review. *IEEE Transactions on Evolutionary Computation*, 7:117–132, 2002.
- [201] J. Novakovic. Toward optimal feature selection using ranking methods and classification algorithms. *Yugoslav Journal of Operations Research*, 21(1), 2016.
- [202] Huan Liu and Rudy Setiono. A probabilistic approach to feature selection - a filter solution. In *Proceedings of the 13th International Conference on Machine Learning (ICML)*, volume 96, pages 319–327, 1996.
- [203] Kenji Kira and Larry A. Rendell. A practical approach to feature selection. In *Proceedings of the Ninth International Workshop on Machine Learning*, ML92, pages 249–256, San Francisco, CA, USA, 1992. Morgan Kaufmann Publishers Inc.
- [204] Hervé Abdi and Lynne J. Williams. Principal component analysis. *Wiley Interdisciplinary Reviews: Computational Statistics*, 2(4):433–459, 2010.

4.6. PROVIDING PERSONALIZED ENERGY MANAGEMENT AND AWARENESS SERVICES FOR ENERGY EFFICIENCY IN SMART BUILDINGS

- [205] Corinna Cortes and Vladimir Vapnik. Support-vector networks. *Machine learning*, 20(3):273–297, 1995.
- [206] Henrik Spiiid. *Multivariate ARIMA and ARIMA-X Analysis*. CRAN, 2017. License GPL-2, Version 2.2, RoxygenNote 5.0.1.
- [207] Li-Yeh Chuang, Chao-Hsuan Ke, and Cheng-Hong Yang. A hybrid both filter and wrapper feature selection method for microarray classification. *CoRR*, abs/1612.08669, 2016.
- [208] Luis Carlos Molina, Lluís Belanche, and Àngela Nebot. Feature selection algorithms: A survey and experimental evaluation. In *Proceedings of the 2002 IEEE International Conference on Data Mining (ICDM 2002), 9-12 December 2002, Maebashi City, Japan*, pages 306–313, 2002.
- [209] J Fergus Nicol and Michael A Humphreys. Adaptive thermal comfort and sustainable thermal standards for buildings. *Energy and buildings*, 34(6):563–572, 2002.
- [210] Balakrishnan Narayanaswamy, Bharathan Balaji, Rajesh Gupta, and Yuvraj Agarwal. Data driven investigation of faults in hvac systems with model, cluster and compare (mcc). In *Proceedings of the 1st ACM Conference on Embedded Systems for Energy-Efficient Buildings*, pages 50–59. ACM, 2014.
- [211] Yu Zheng, Xing Xie, Wei-Ying Ma, et al. Geolife: A collaborative social networking service among user, location and trajectory. *IEEE Data Eng. Bull.*, 33(2):32–39, 2010.
- [212] Elias Frentzos, Kostas Gratsias, and Yannis Theodoridis. Index-based most similar trajectory search. In *2007 IEEE 23rd International Conference on Data Engineering*, pages 816–825. IEEE, 2007.
- [213] Fernando Terroso-Saenz, Mercedes Valdes-Vela, and Antonio F Skarmeta-Gomez. A complex event processing approach to detect abnormal behaviours in the marine environment. *Information Systems Frontiers*, 18(4):765–780, 2016.
- [214] Fernando Terroso-Sáenz, Mercedes Valdés-Vela, Francisco Campuzano, Juan A Botia, and Antonio F Skarmeta-Gómez. A complex event processing approach to perceive the vehicular context. *Information Fusion*, 21:187–209, 2015.
- [215] Robert Babuška. *Fuzzy modeling for control*, volume 12. Springer Science & Business Media, 2012.
- [216] Umutcan Şimşek, Anna Fensel, Anastasios Zafeiropoulos, Eleni Fotopoulou, Paris Liapis, Thanassis Bouras, Fernando Terroso Saenz, and Antonio F Skarmeta Gómez. A semantic approach towards implementing energy efficient lifestyles through behavioural change.

- In *Proceedings of the 12th International Conference on Semantic Systems*, pages 173–176. ACM, 2016.
- [217] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2014.
- [218] Jessica Granderson, Phillip N Price, David Jump, Nathan Addy, and Michael D Sohn. Automated measurement and verification: Performance of public domain whole-building electric baseline models. *Applied Energy*, 144:106–113, 2015.